

# Lecture 2: Operating-System Structures

# COMP362 Operating Systems

## Prof. AJ Bieszczad

# Outline



- System Startup
- Operating System Services
  - For users
  - For resource sharing
- User Operating System Interface
- System Calls
- Types of System Calls
- System Programs
- Operating System Design and Implementation
- Operating System Structure

# OS Startup



- Operating systems are designed to run on a variety of machines
  - Therefore, the system must be configured for each specific computer
  - During installation of the OS, a “system configuration” program obtains information concerning the specific configuration of the hardware system
- Before any operating system can be loaded and started, the computer needs to go through a **bootstrapping** process; **boot** for short.
  - When system is powered up, execution starts at a fixed memory location in NVRAM (non-volatile RAM) that holds code provided by the computer vendor called **firmware**:
    - **BIOS (Basic Input-Output System)** on older computers, or
    - **UEFI (Unified Extensible Firmware Interface)** intended to replace BIOS on modern computers.
  - After power-on self-tests (**POST**), the code that is called **bootstrap** (or **boot**) **loader** tests and initializes the hardware and then following a **boot sequence** loads an operating system
    - boot sequence is configured in BIOS or UEFI
    - all devices in the boot sequence are examined in sequence until a valid operating system is found and loaded.
  - On some computers, boot sequence can be overridden during the boot by invoking a **boot menu**, and then selecting a device to boot from disregarding the boot sequence
    - for example, it can be done on Dell computers by pressing the F12 key (as used in the lab)

# OS Startup (cont.)

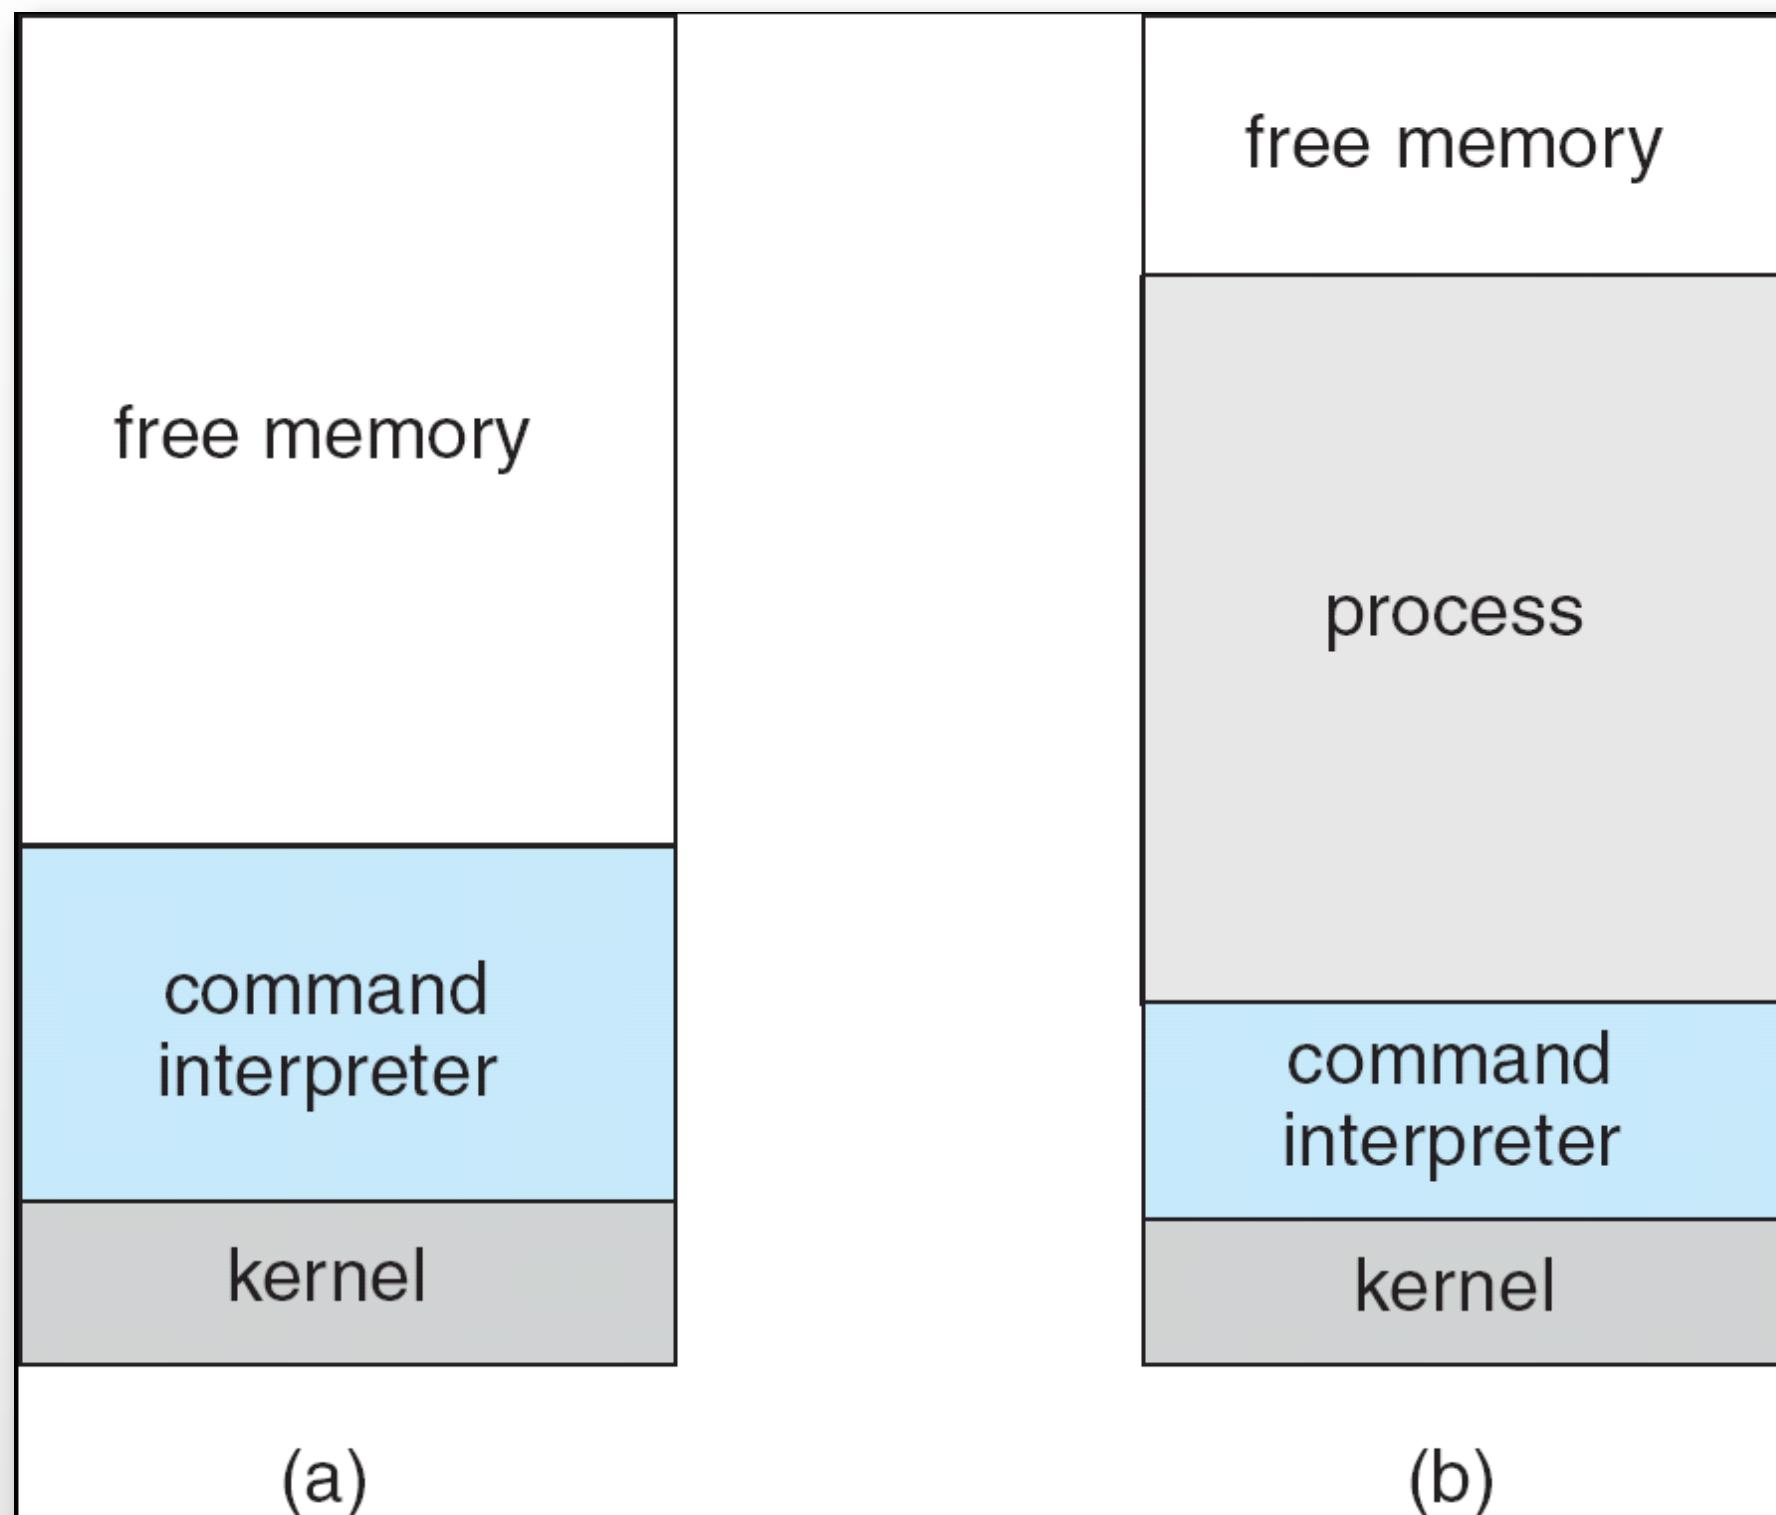


- The firmware on UEFI-based system supports loading an operating systems from any location on any boot device.
- In contrast, on older BIOS-based systems, the operating system is always loaded from the pre-assigned **master boot sector (MBR)** of the selected boot device
  - however, UEFI is sometimes configured to behave like BIOS for compatibility reasons.
- Instead of loading an operating system directly, some computers may be configured to load an intermediary software, so-called **OS loader**, that may allow users to choose from several operating systems present on the boot device
  - for example, on **dual-boot computers** you can choose between two operating systems; say between Windows and Ubuntu, or between macOS and Windows.
- LILO, GRUB, Bootcamp, are examples of such OS loaders
  - options may be set by the administrator through a configuration file that the loader reads before asking the user for selection
  - the configuration file has specifics such as the default OS that is loaded after some timeout, order of the OSes, location of specific loaders on the disk, etc.
  - For more details on Linux default OS Loader GRUB see: <https://help.ubuntu.com/community/Grub2>

# OS Memory Layout

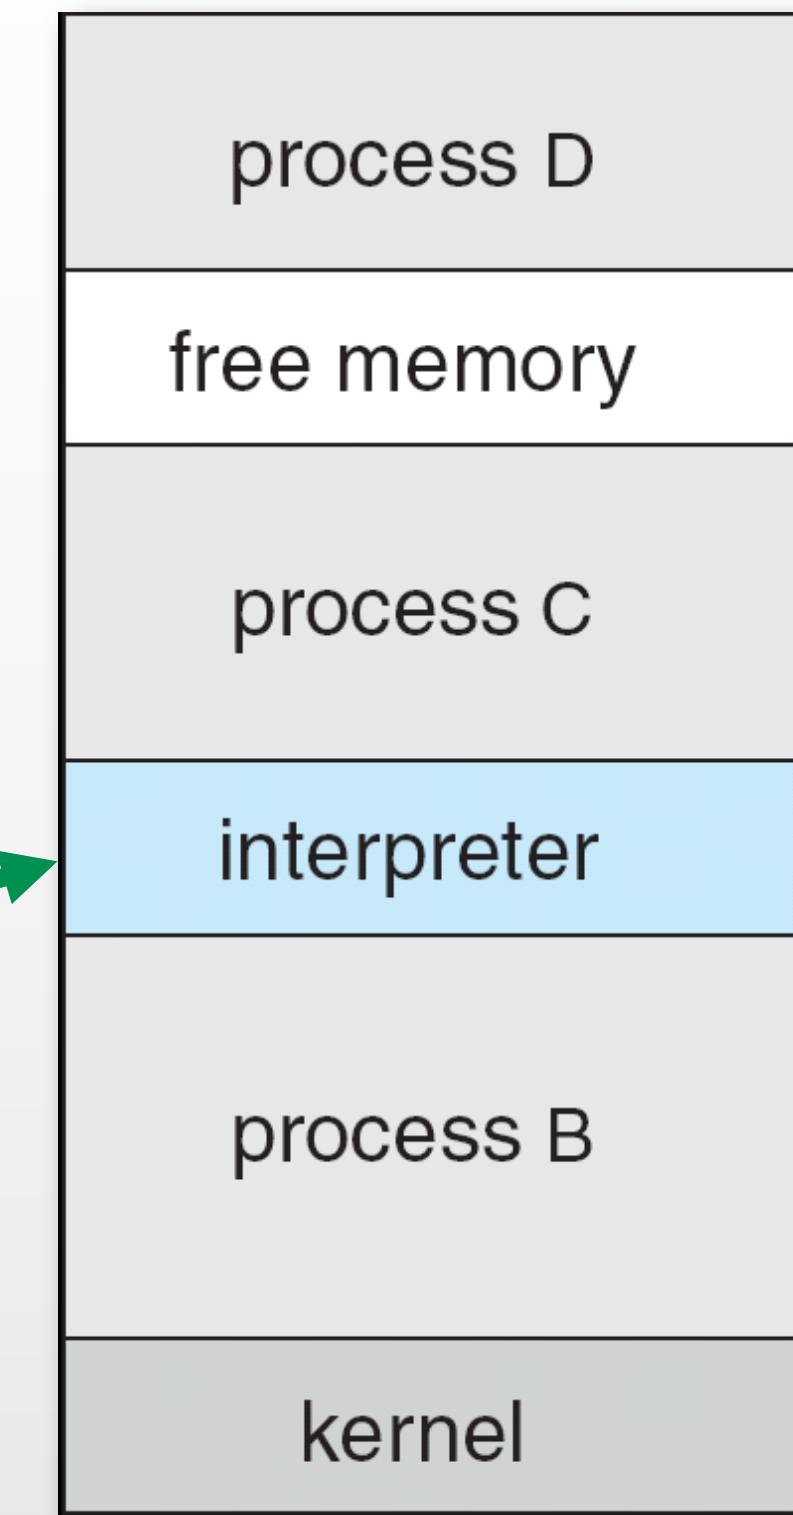


MS-DOS execution

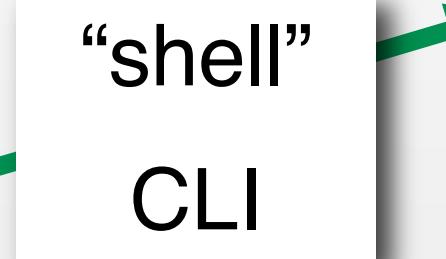


(a) At system startup (b) running a program

Linux execution

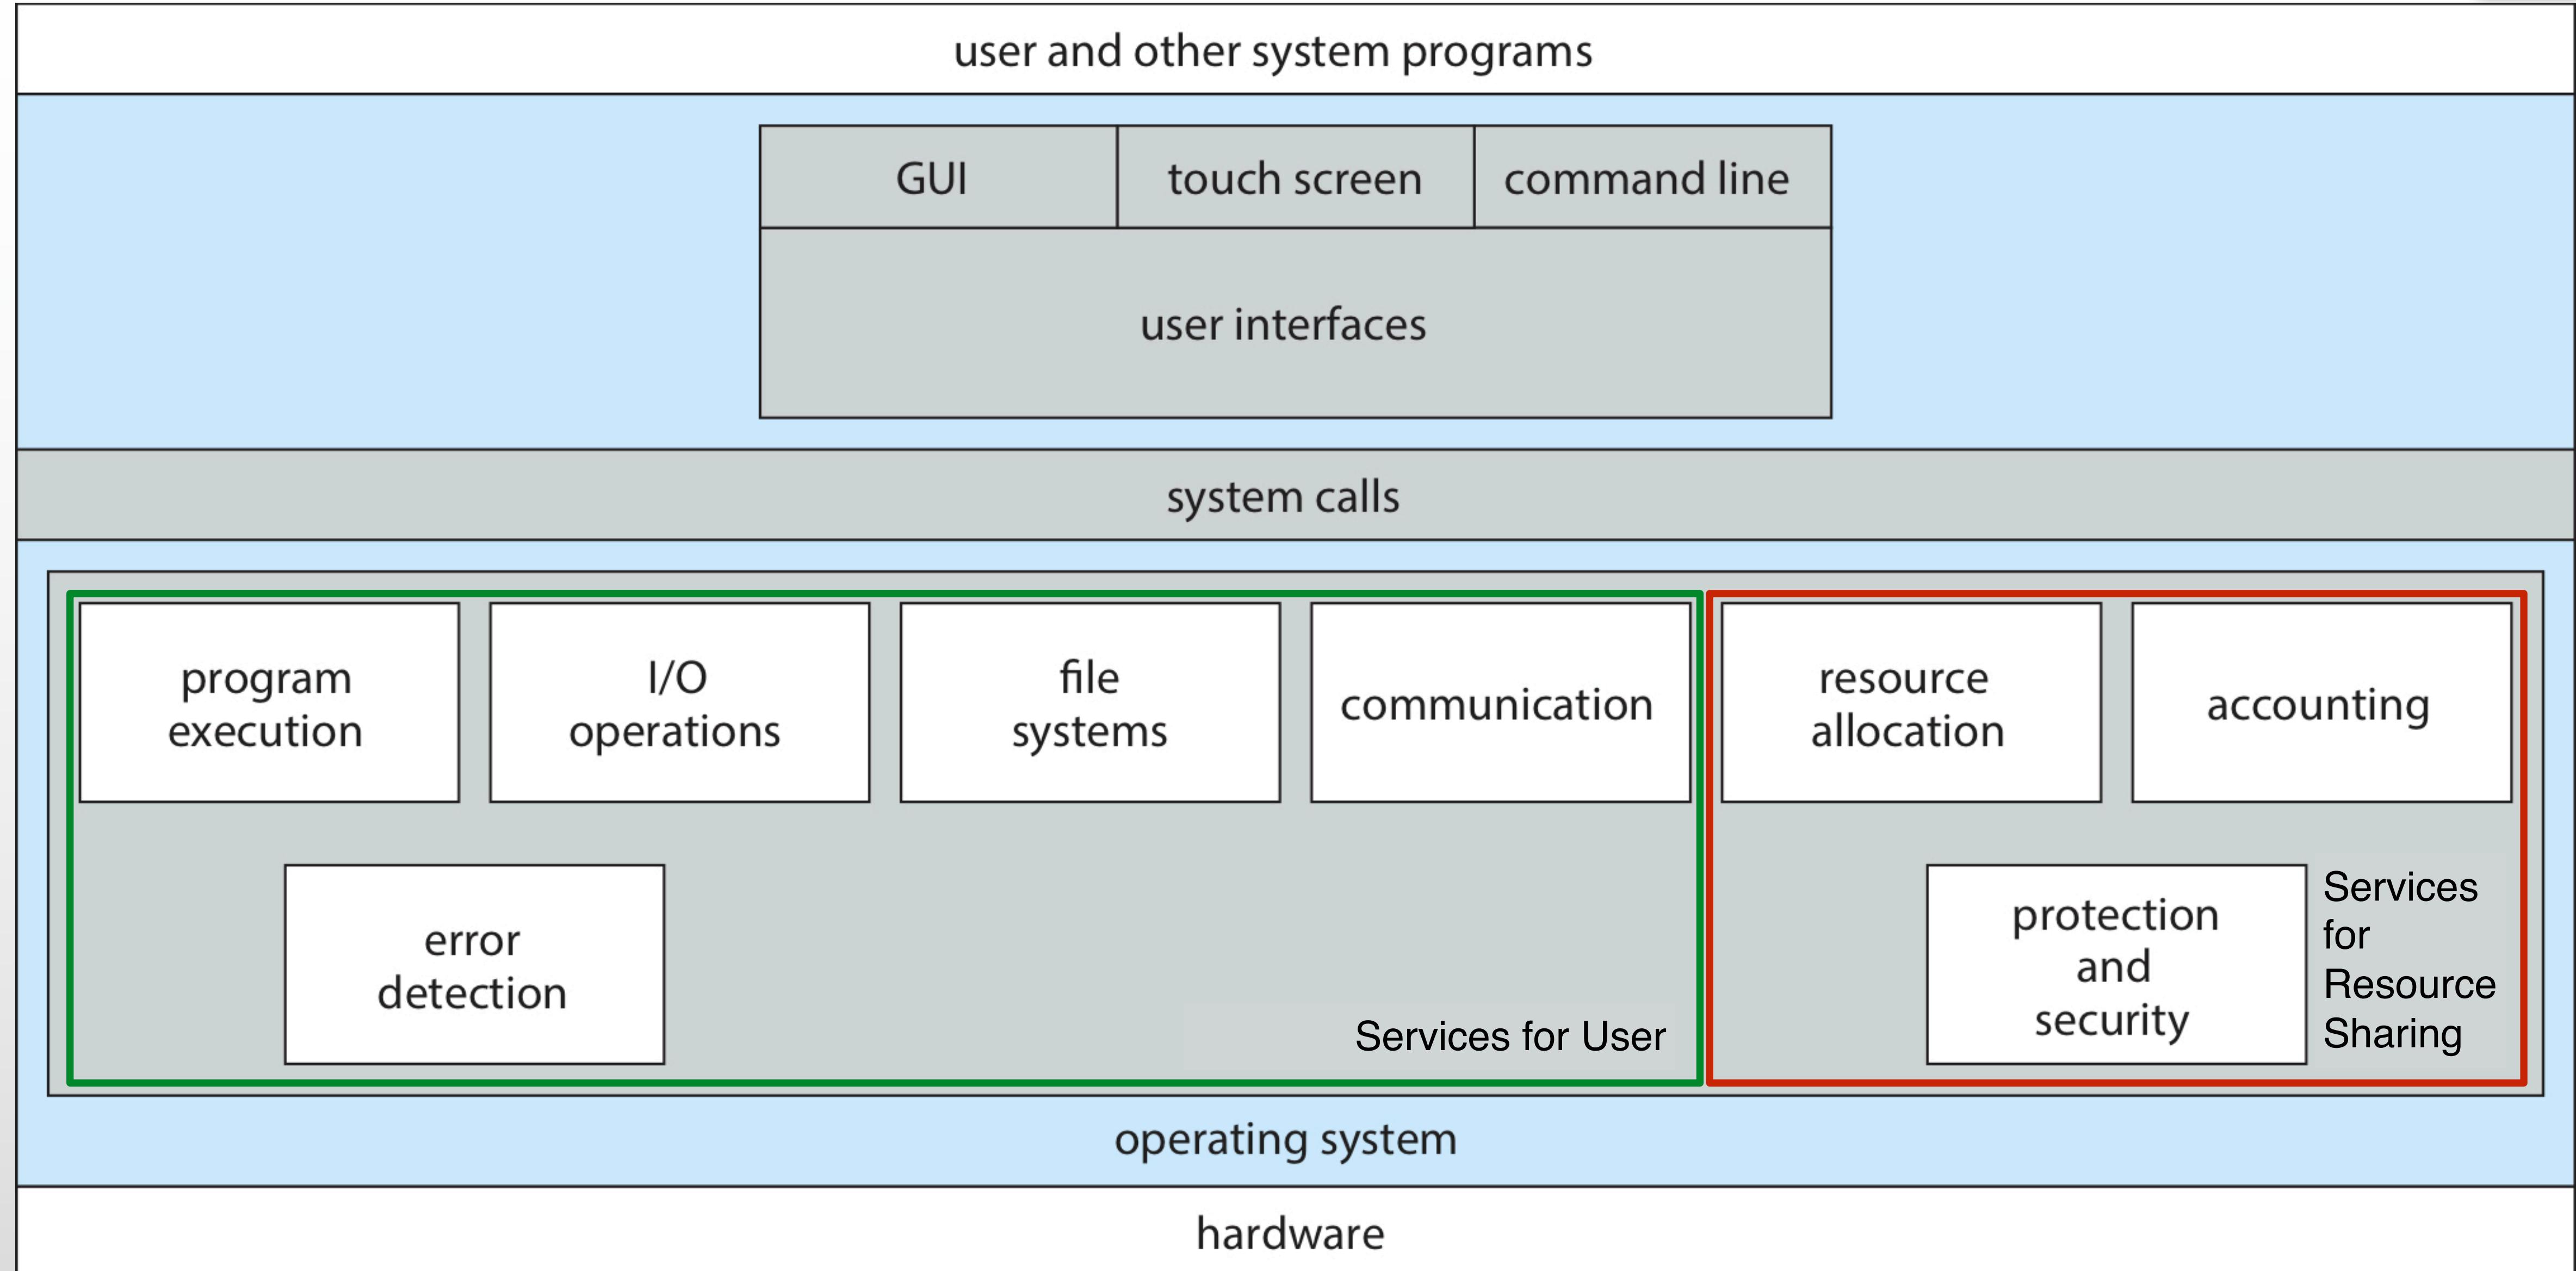


“shell”  
CLI



**Modern OSes have  
flexible organization of  
memory**

# OS Services



# OS Services for the User



- **Program execution** - The system must be able to load a program into memory and to run that program, end execution, either normally or abnormally (indicating error)
- **I/O operations** - A running program may require I/O, which may involve a file or an I/O device.
- **File-system manipulation** - The file system is of particular interest. Obviously, programs need to read and write files and directories, create and delete them, search them, list file information, permission management.
- **Communication** – Processes may exchange information, on the same computer or between computers over a network
  - Communications may be via shared memory or through message passing (packets moved by the OS)
- **Error detection** – OS needs to be constantly aware of possible errors
  - May occur in the CPU and memory hardware, in I/O devices, in user program
  - For each type of error, OS should take the appropriate action to ensure correct and consistent computing
  - Debugging facilities can greatly enhance the user's and programmer's abilities to efficiently use the system

# OS Services for Resource Sharing



## ● **Resource allocation**

- When multiple users or multiple jobs run concurrently, resources must be allocated to each of them
  - Some resources such as CPU cycles, main memory, and file storage need customized allocation algorithms (e.g., dynamically depending on the state)
  - Others such as I/O devices may use generalized request and release algorithms

## ● **Accounting/Logging**

- To keep track of which users use how much and what kinds of computer resources

## ● **Protection and security**

- The owners of information stored in a multiuser or networked computer system may want to control use of that information, concurrent processes should not interfere with each other

### ● **Protection is a mechanism**

- i.e., involves ensuring that all access to system resources is controlled; for example, from erroneous code like runaway (“floating”) pointers

### ● **Security is a policy**

- i.e., access requires user authentication to defending resources from invalid access attempts; for example, from unauthorized access

# OS User Interfaces



- All operating systems come with a **user interface (UI)**
  - Varies between Command-Line (**CLI**), Graphics User Interface (**GUI**), Batch
  - A UI is not necessarily considered part of the operating system
- **CLI (command-line interface)** — character-based command entry
  - Can be implemented in kernel or by systems program
  - Often multiple flavors implemented – **shells**
    - We use **bash** shell, but there are many others; e.g., sh, csh, zsh, ksh, tcsh
  - Primarily fetches a command from user and executes it
    - Some commands are built-in (i.e., compiled with other shell parts); others may be just names of programs to execute
    - The latter - **utilities** - are convenient: adding new features doesn't require shell modification
- **GUI (graphical user interface)** - user-friendly desktop metaphor interface
  - Used to be mouse, keyboard, and monitor, but touch screen interface is increasingly popular
  - Icons represent files, programs, actions, etc
  - Mouse button clicks, gestures, over objects in the interface cause various actions like providing information, options, execution functionality, opening a file, directory, etc.
  - Invented at Xerox PARC (Palo Alto Research Center) — just like Smalltalk, one of Java precursors
  - Many systems now include both CLI and GUI interfaces
    - MS Windows is a GUI with CLI “command” shell (also “powershell” starting with Win7)
    - Apple macOS has “Cocoa” GUI interface with UNIX kernel underneath (Darwin) and shells available in Terminal application (bash is the default)
    - Linux has Gnome or KDE (or any of many, many other) GUIs; it also has CLI in Terminal application (with bash as the default)
    - Some OSes start to be exclusively gesture-based (like iOS); others implement multiple paradigms (like Windows 10, for example)

# System Calls



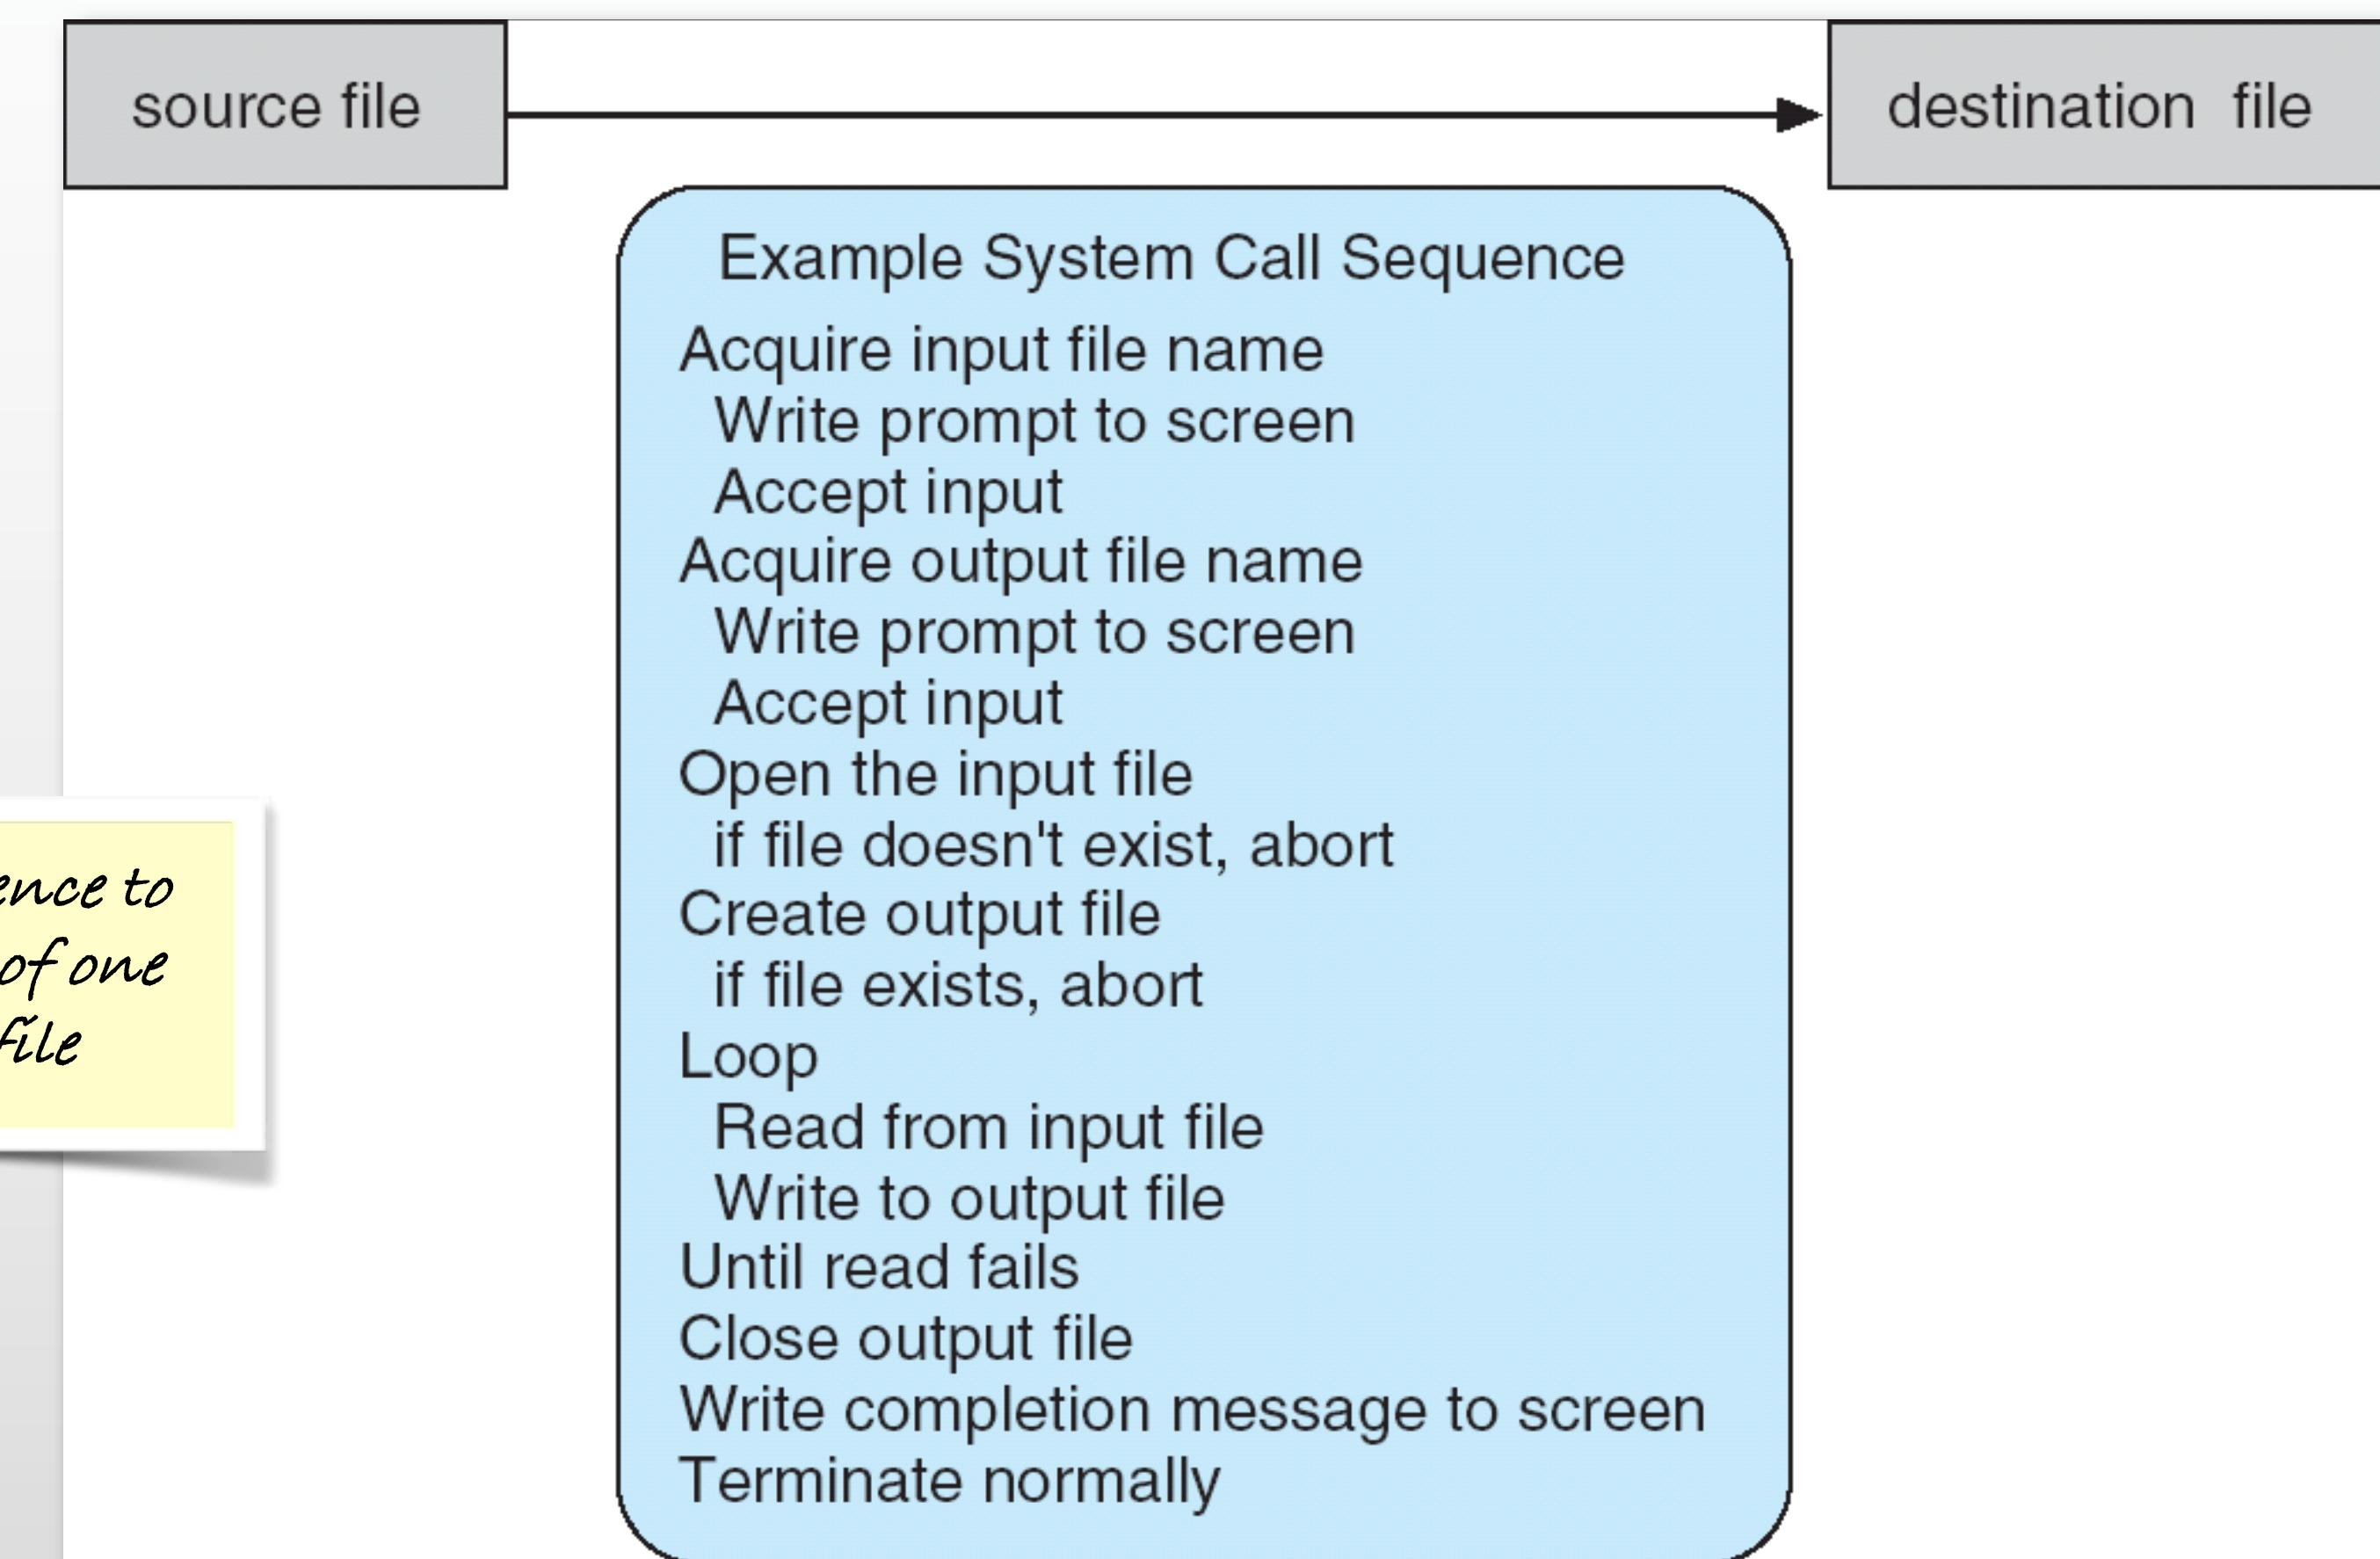
- Programming interface to the services provided by the OS
  - Usually very topical in functionality
    - for example, copying a file may involve many system calls:

See:

\$ man 2 syscall  
\$ man 2 syscalls

```
$ mycopy
< in.txt
> out.txt
$
```

*System call sequence to copy the contents of one file to another file*



# System Call APIs



- Usually accessed by programs via a high-level **Application Program Interface (API)** rather than direct system call
  - \$ man 2 syscall
- Often, but not always, the name of the wrapper function is the same as the name of the system call that it invokes.
- A number of system calls may be combined as needed to implement a complex function
- Error handling added through perror()

```
$ man perror  
$ sudo apt install moreutils  
$ man errno  
$ errno -l
```

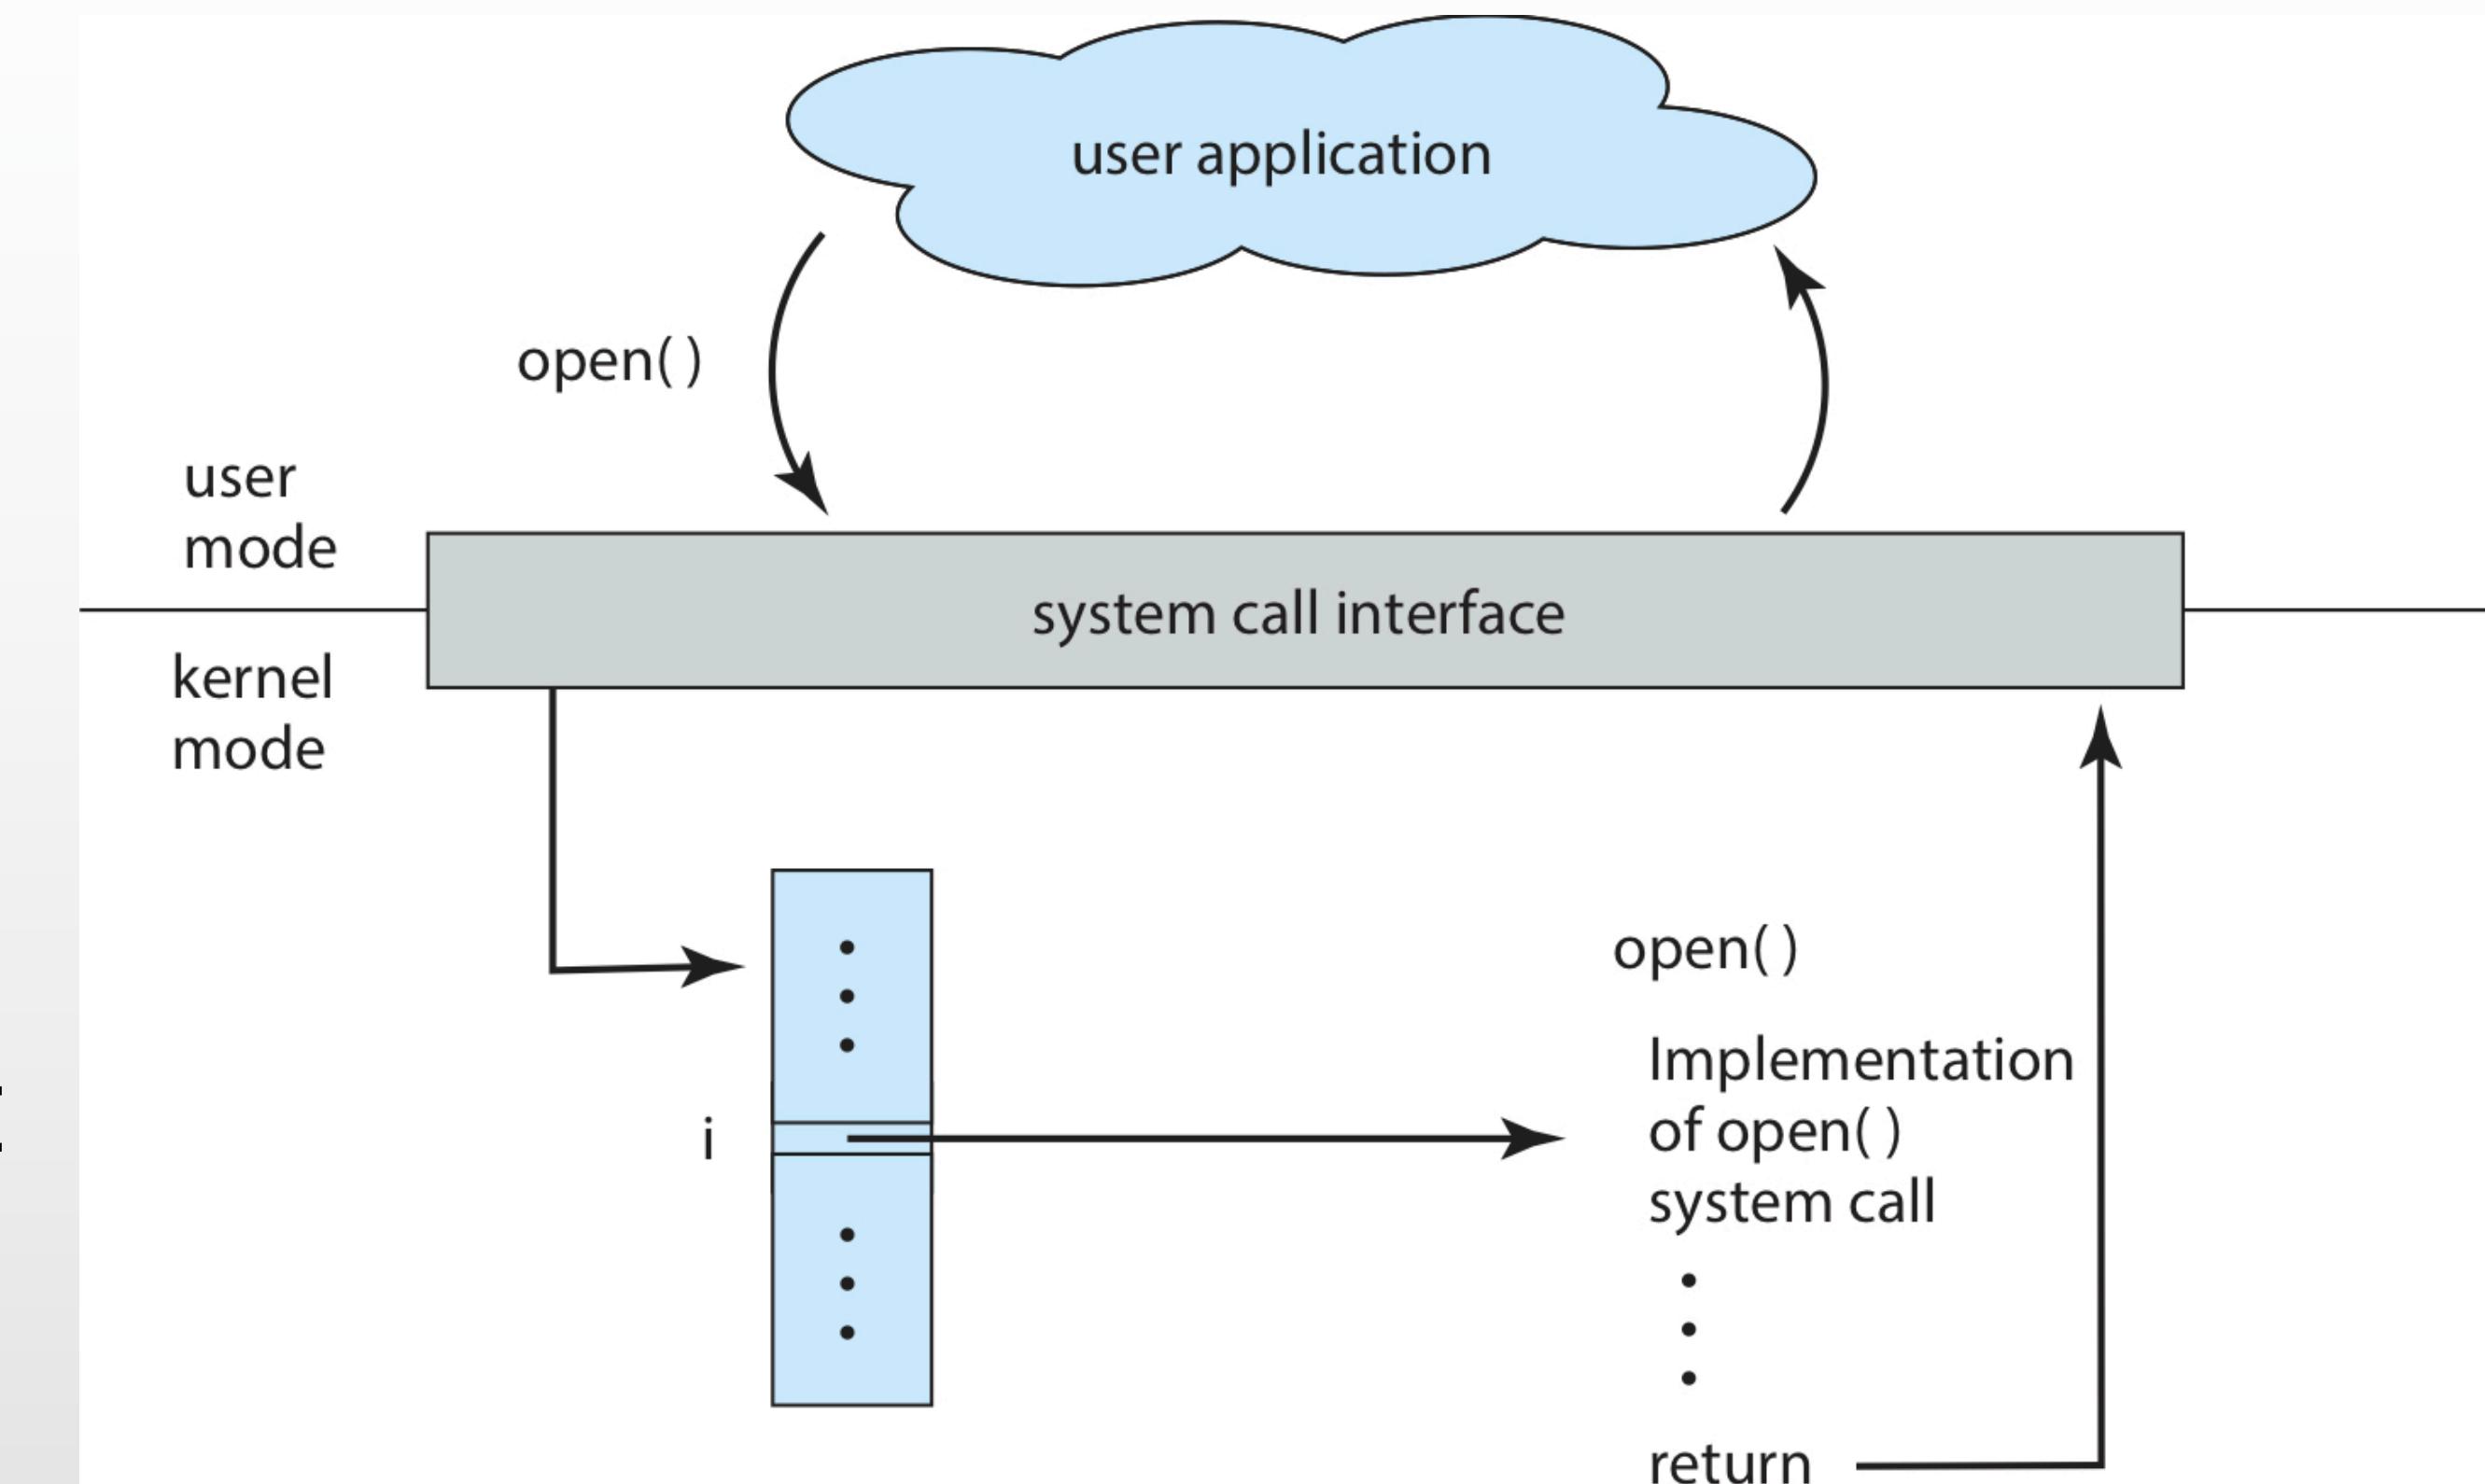
- Most known OS APIs are:
  - POSIX API for POSIX-based systems including virtually all versions of UNIX, Linux, and macOS
    - \$ sudo apt install manpages-posix-dev
  - Win32 API for Windows
    - “Win32” also applies to 64-bit
- System calls can also be invoked directly through a low-level interface
  - implemented in assembler
  - architecture dependent

```
#define _GNU_SOURCE  
#include <unistd.h>  
#include <sys/syscall.h>  
#include <sys/types.h>  
#include <signal.h>  
#include <stdio.h>  
#include <pthread.h>  
  
int  
main(int argc, char *argv[])  
{  
    pid_t tid;  
  
    // NOTE: gettid() wrapper IS NOT IMPLEMENTED IN C LIBRARY  
  
    // get thread id using a direct system call  
  
    tid = syscall(SYS_gettid); // same as tid = syscall(224);  
  
    printf("TID from a direct system call: %ld\n", (long int) tid);  
  
    // send a hungup (SIGHUP) signal to the thread  
    tid = syscall(SYS_tgkill, getpid(), tid, SIGHUP);  
}
```

# System Call Implementation



- Typically, there is a number associated with each system call
  - system-call interface maintains a table of references to system call handlers indexed according to these numbers
- The system call interface invokes intended system call in OS kernel and returns status of the system call and any return values
- The caller needs know nothing about how the system call is implemented
  - just needs to obey the API and understand what OS will do as a result of the call

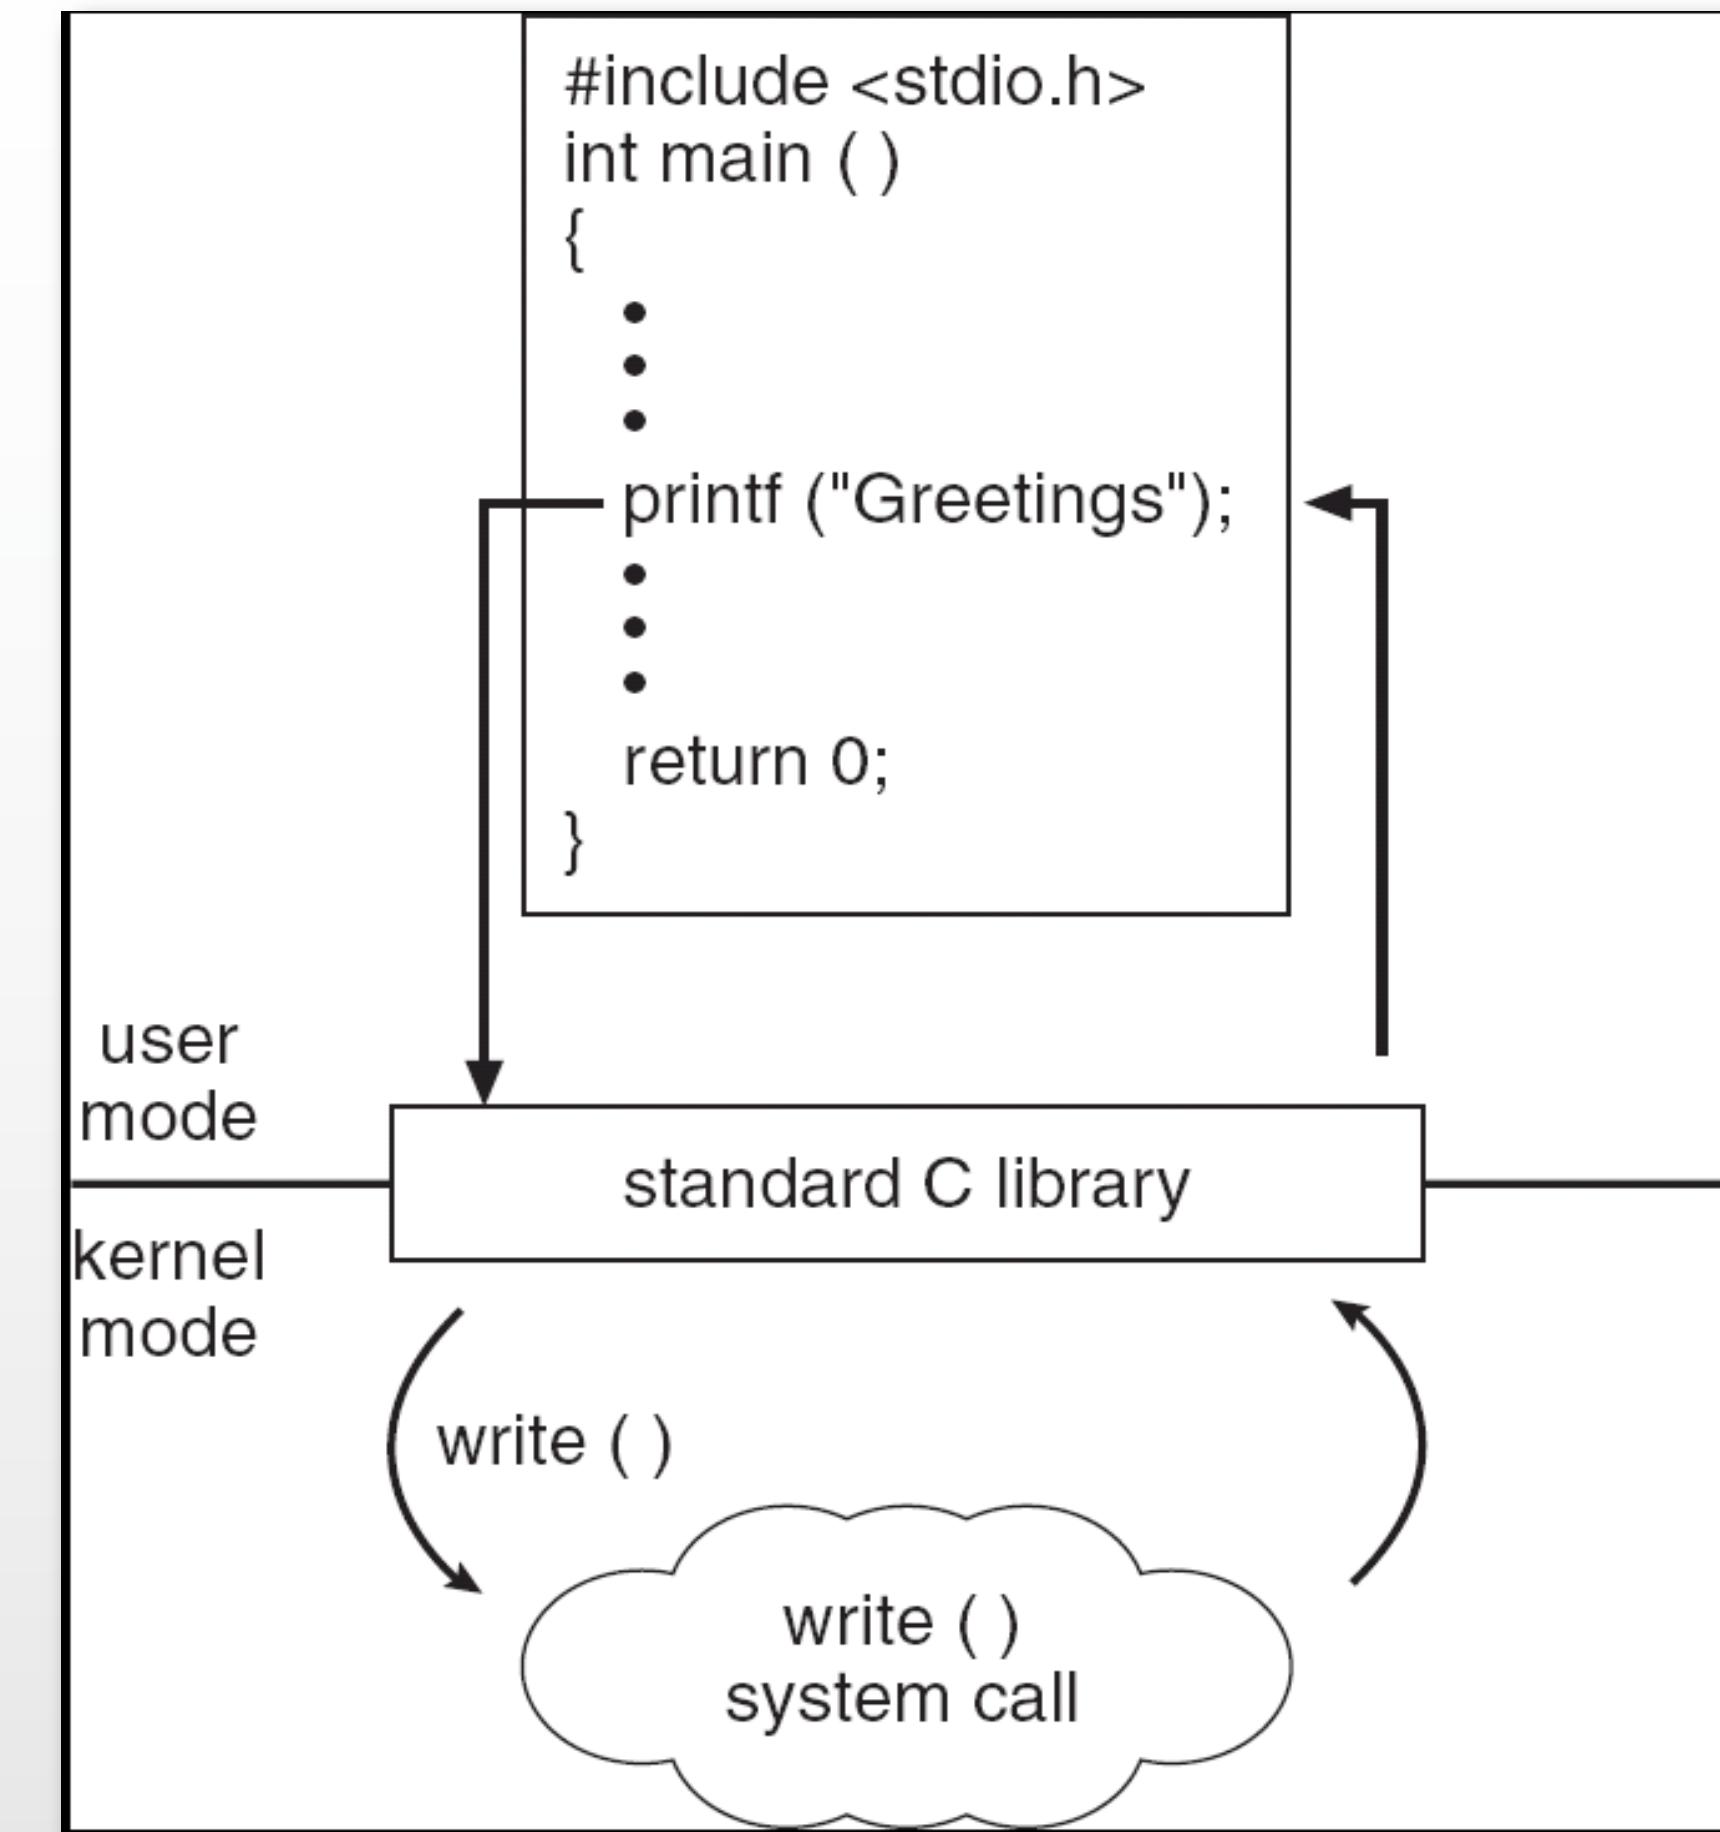


# Standard C Library Example



- Most details of OS interface hidden from programmer by the API
- managed by run-time support library
- set of functions built into libraries included with the compiler and embedded into every program

The library that is called `libc` is added automatically to every C program



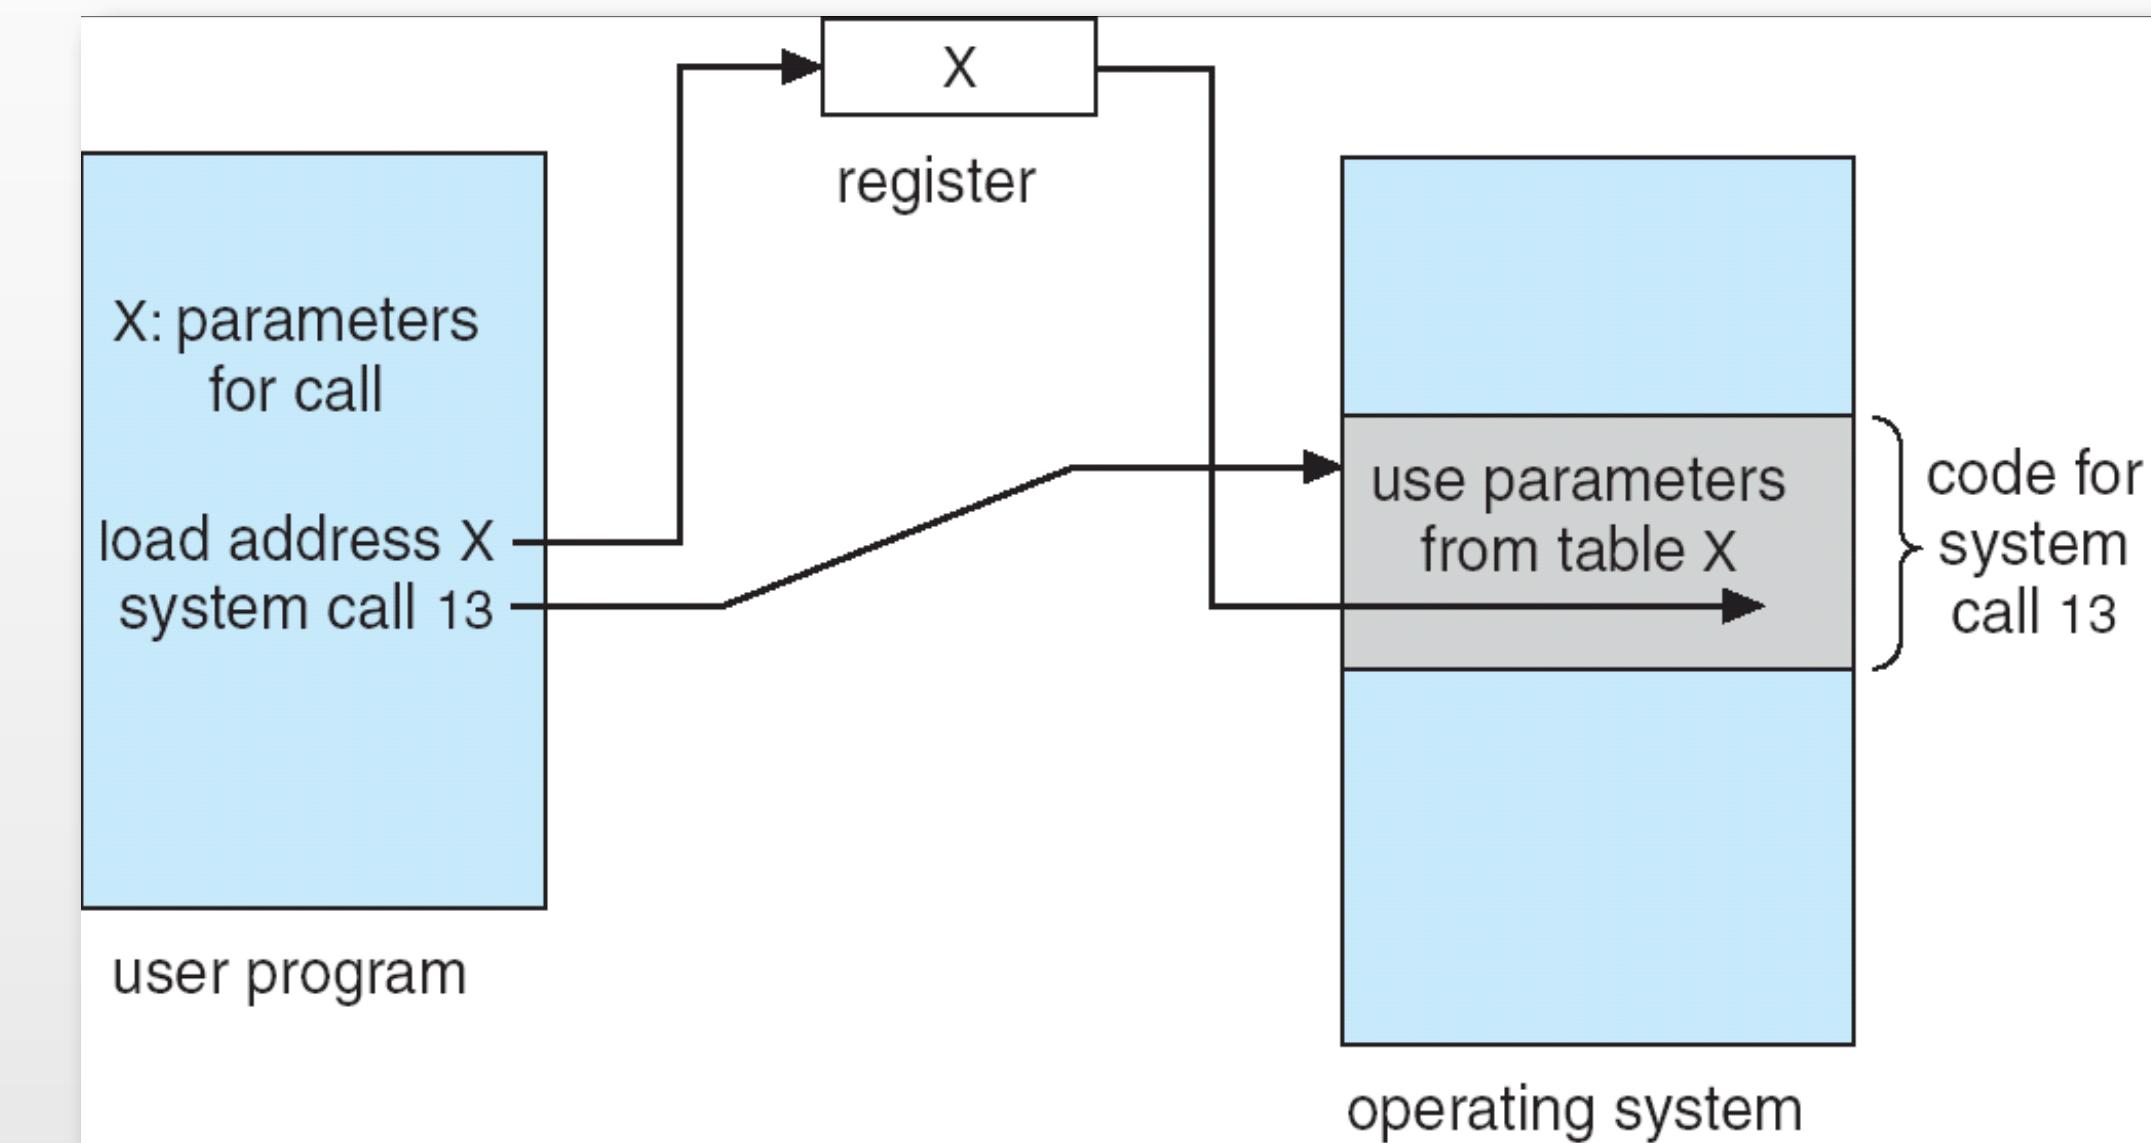
- C program invoking `printf()` library call, which calls `write()` system call

# System Call Parameter Passing



- As in other type of programming, usually more information is required than simply the identity of desired system call
- Exact type and amount of information vary according to OS and call
- There are three general methods used to pass parameters to the OS
  - NOTE: This is a choice that the OS designers make, and not the users
    - The developers of the APIs have to stick to the OS convention
- Simplest: pass the **parameters in registers** (e.g., MS-DOS)
  - In some cases, may be more parameters than registers
- Parameters placed, or pushed, onto the stack** by the program and popped off the stack by the operating system (e.g., NetBSD)
  - This is just like a regular function call

- Parameters stored in a block, or table, in memory**, and address of the block passed as a parameter in a register (e.g., Linux, Solaris)



- Block and stack methods do not limit the number or length of parameters being passed

# Types of System Calls



- Process control
- File management
- Device management
- Information maintenance
- Communication
- As we continue with the course, we will be discovering some of the APIs to these system calls at two levels:
  - What they are and how we are using them?
  - How are they (or rather: should be) implemented?

# System Programs



- Most users' view of the operating system is defined by UI functions (e.g., GUI clicks, CLI commands) and by **OS utilities**, and not by the actual system calls
  - In fact, most OS utilities are **system programs** rather than shell commands
    - open a man page for bash to see which commands are listed there; anything else called from the shell is a utility
- \$ **man bash**
- System programs provide a convenient environment for program development and execution.
  - They can be divided into utilities for:
    - File manipulation
    - Status information
    - Programming language support
    - Program loading and execution
    - Communication
    - Common application programs (e.g., date, time, text processing, pattern matching, searching, calculators, formatters, etc.)
    - Some applications may become recognized as system programs as they gain popularity and are included with the distribution of the OS
      - e.g., many GNU utility programs

# System Programs (cont.)



- Some are simply user interfaces to system calls; others are considerably more complex
- **File management**
  - Create, delete, copy, rename, print, dump, list, and generally manipulate files and directories
- **Status information**
  - Some ask the system for info - date, time, amount of available memory, disk space, number of users
  - Others provide detailed performance, logging, and debugging information
  - Typically, these programs format and print the output to the terminal or other output devices
    - the text can be piped into another program rather than sent to the terminal
  - Some systems implement a system database used to store and retrieve configuration information; for example, Registry in Windows
  - Others, provide utilities to manipulate per-program configuration files; for example, plists in macOS

# System Programs (cont.)

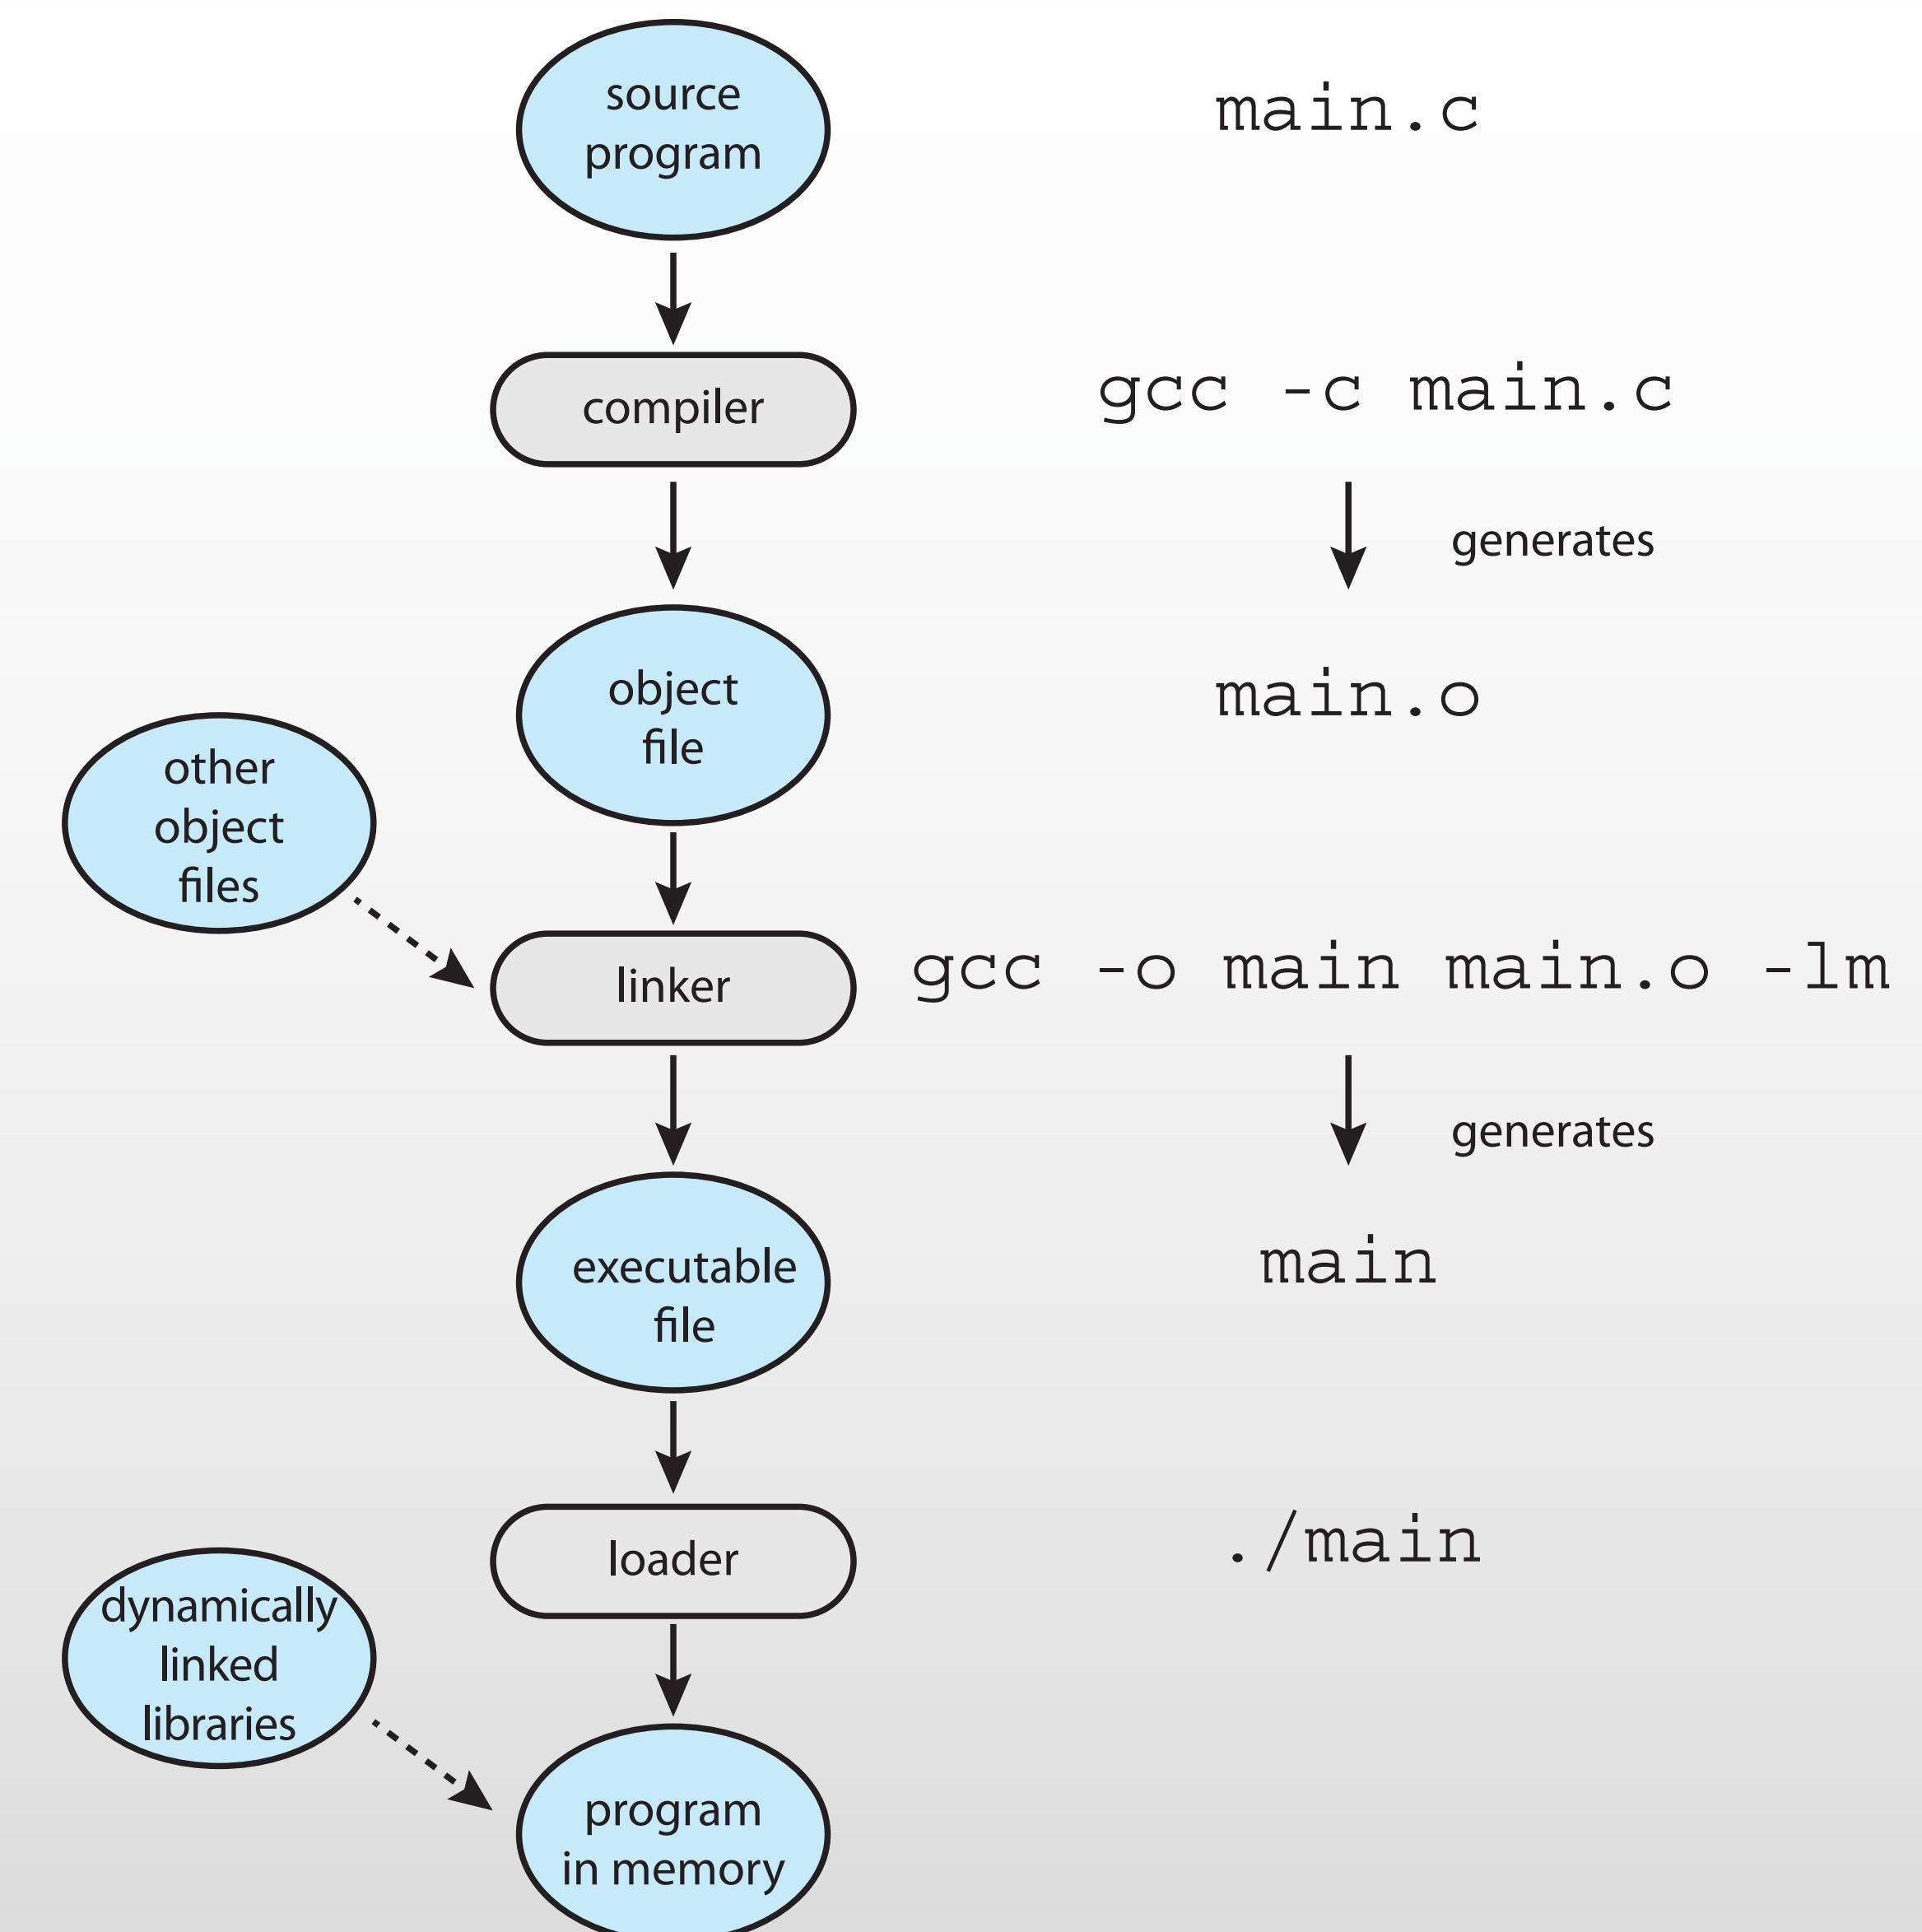


- **File modification**
  - Text editors to create and modify files
  - Special commands to search contents of files or perform transformations of the text
- **Programming-language support**
  - Compilers, assemblers, debuggers and interpreters sometimes provided
- **Program loading and execution**
  - Absolute loaders, relocatable loaders, linkage editors, and overlay-loaders, debugging systems for higher-level and machine language
- **Communication**
  - Provide the mechanism for creating virtual connections among processes, users, and computer systems
  - Allow users to send messages to one another's screens, browse web pages, send electronic-mail messages, log in remotely, transfer files from one machine to another

# Program Building, Loading, and Execution



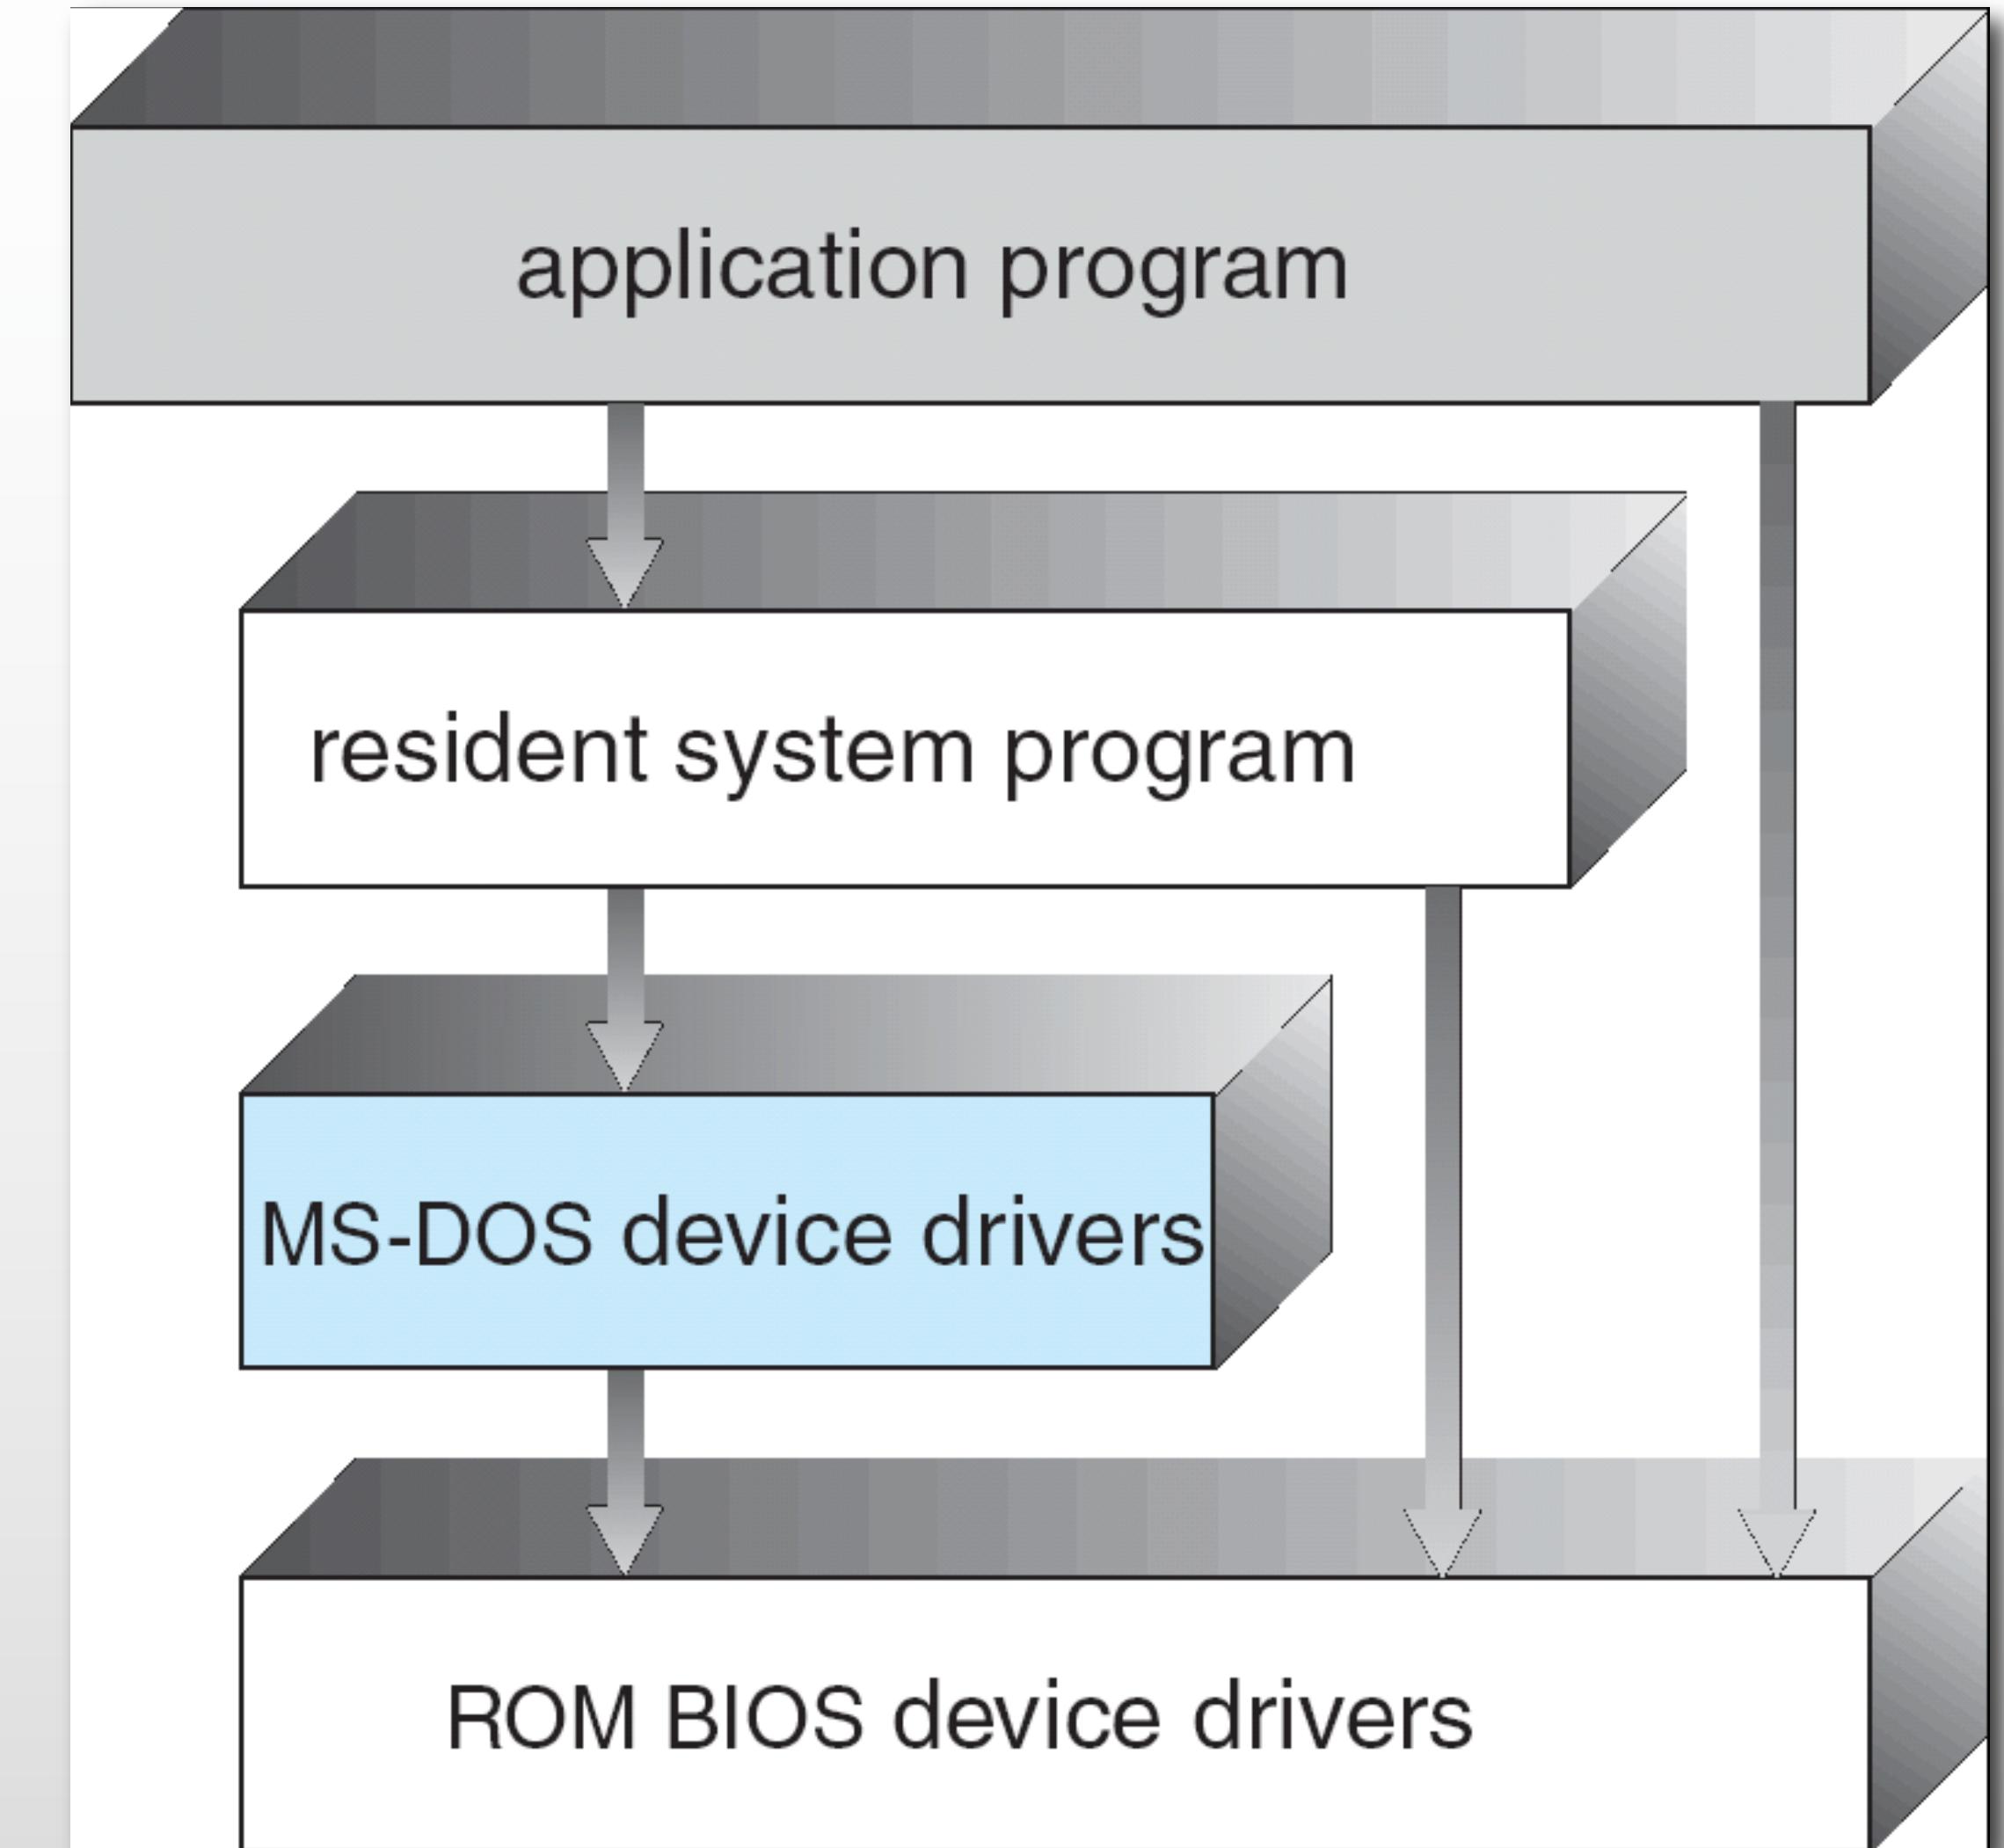
- Recall from COMP232



# OS Structures: Simple Structure



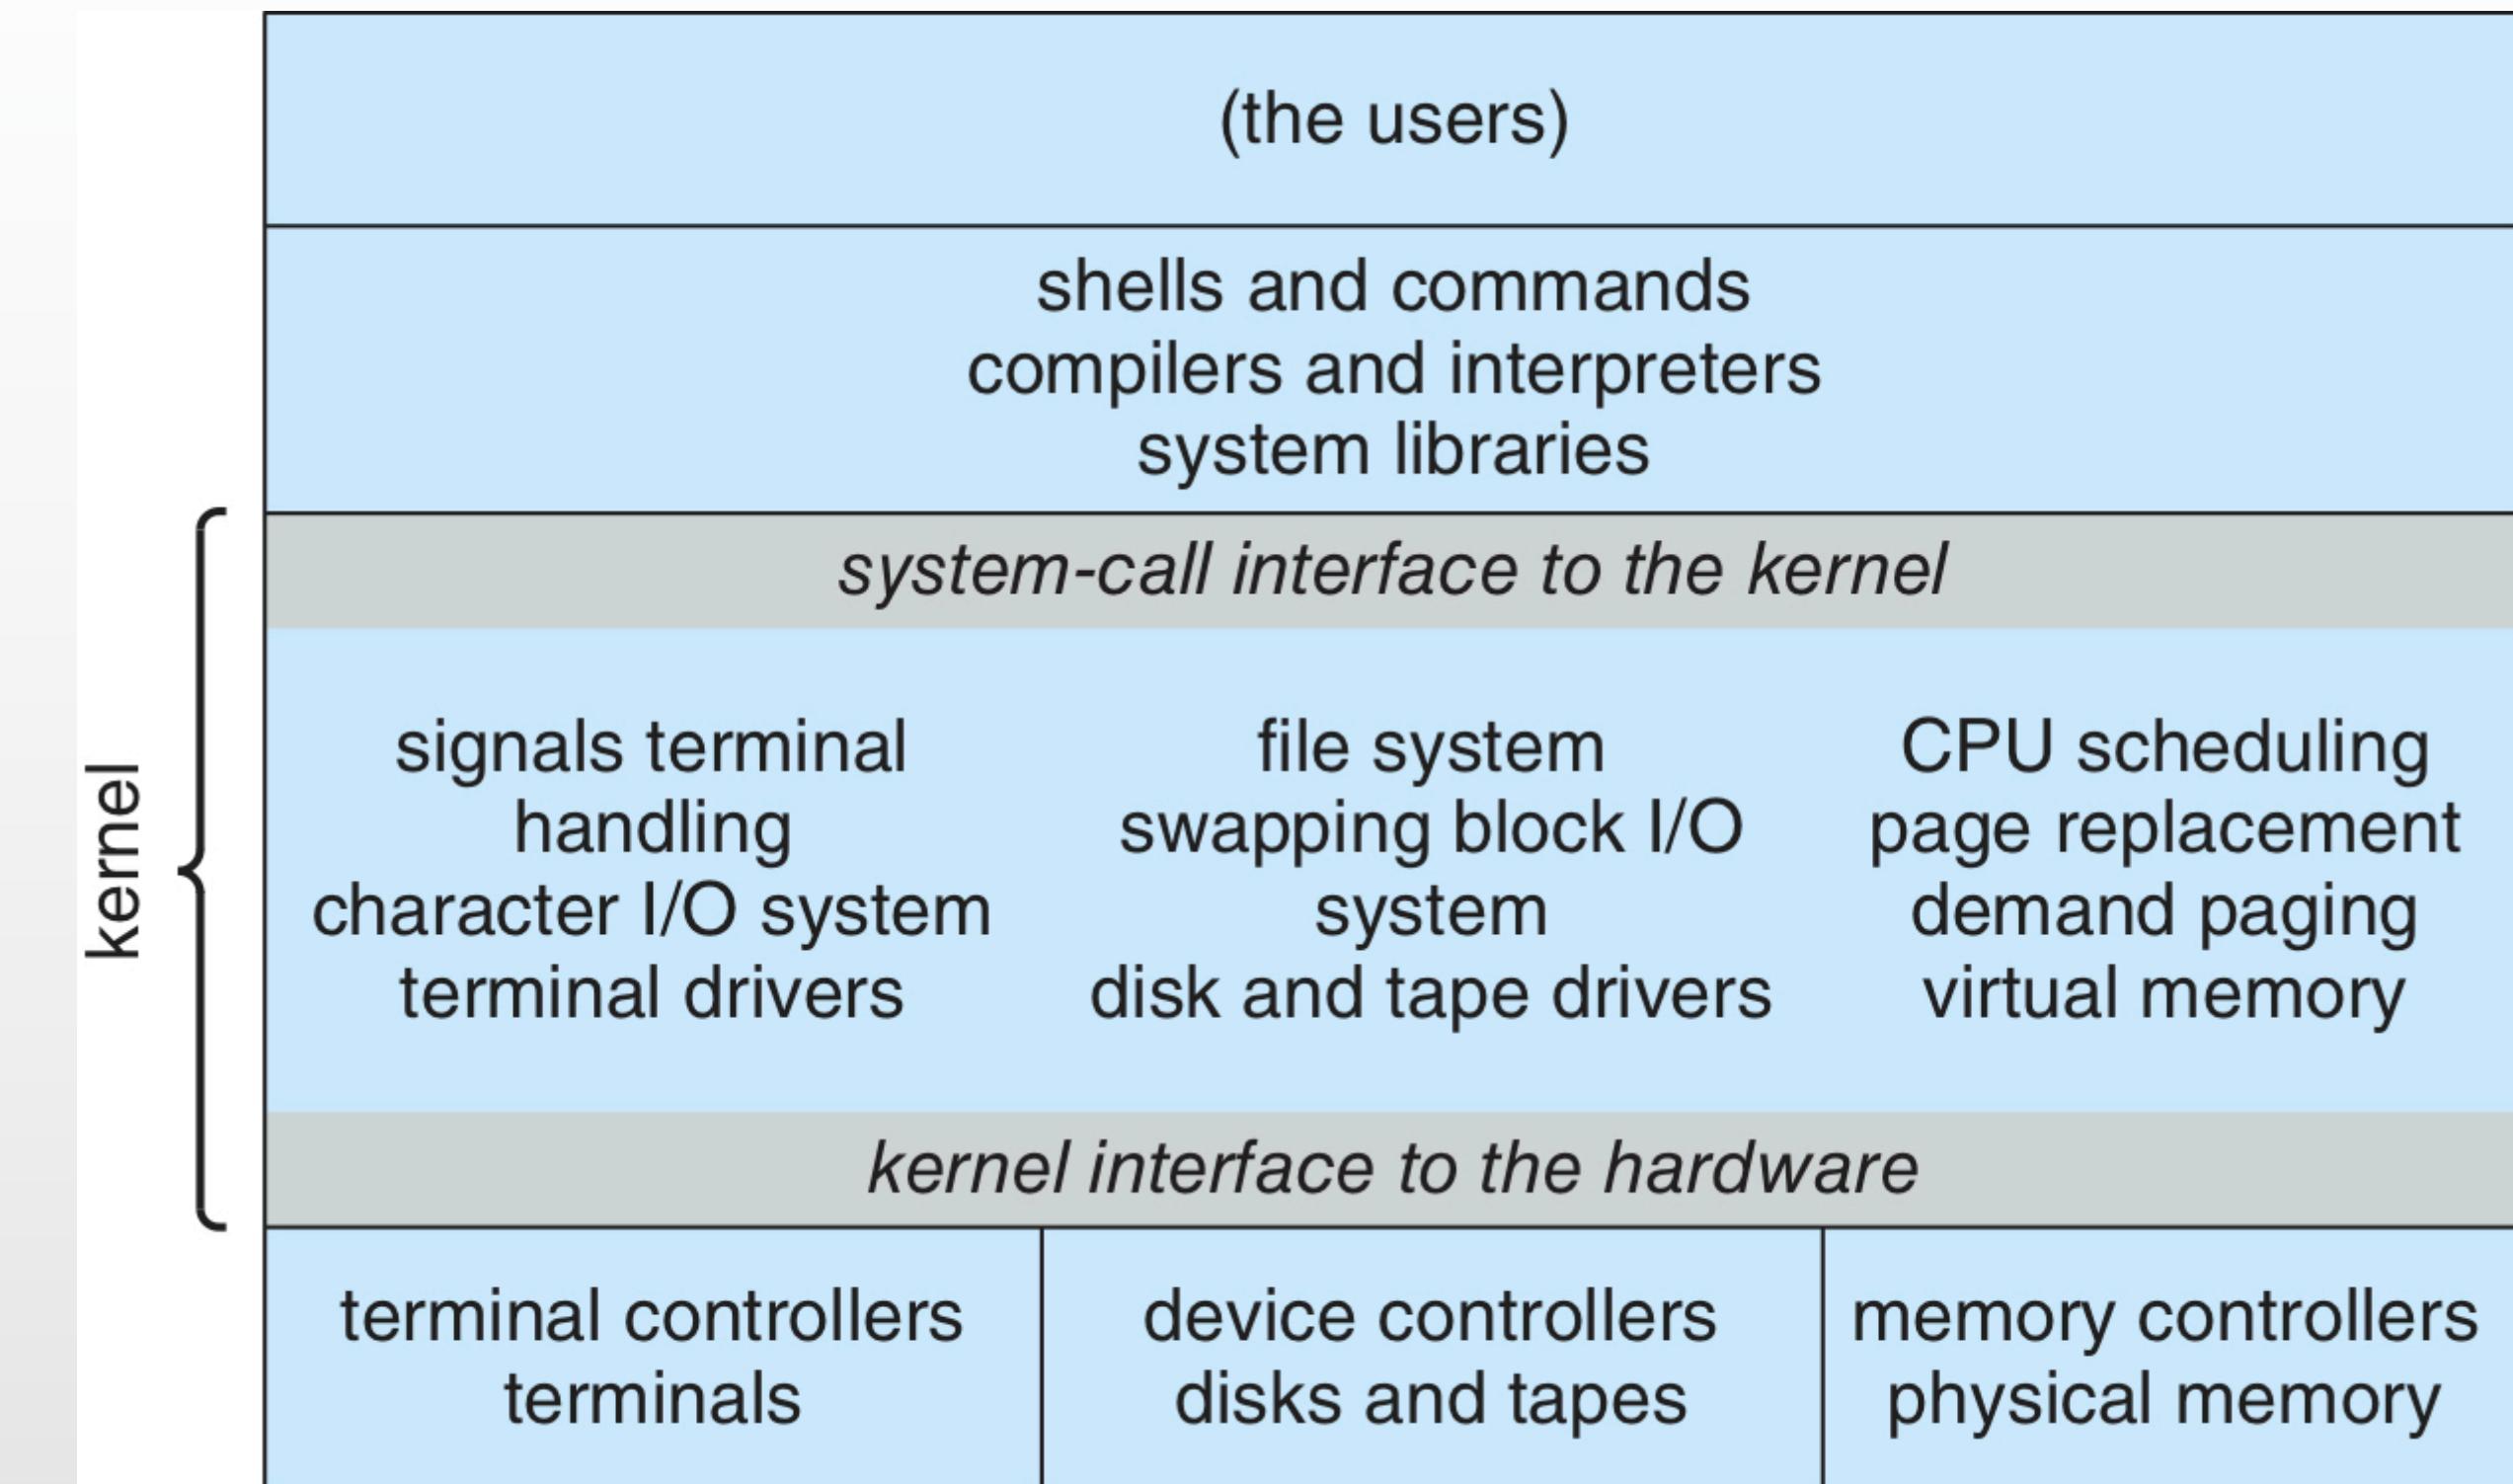
- Early OSes (e.g., MS-DOS) written to provide the most functionality in the least space
  - Monolithic (i.e., not divided into modules)
  - Although they had some structure, the interfaces and levels of functionality were not well separated
    - A lot of problems!
- Famous funny quotes:
  - Mr. Gates (former CEO of Microsoft):  
*"Nobody ever will need more memory!"*
    - (about the 640 kBytes allocation to programs; that is **kilo**!)
  - A bit better than Mr. Watson's (at the time, the CEO of IBM): *"I think there is a world market for maybe five computers."*



# Structure of Original UNIX



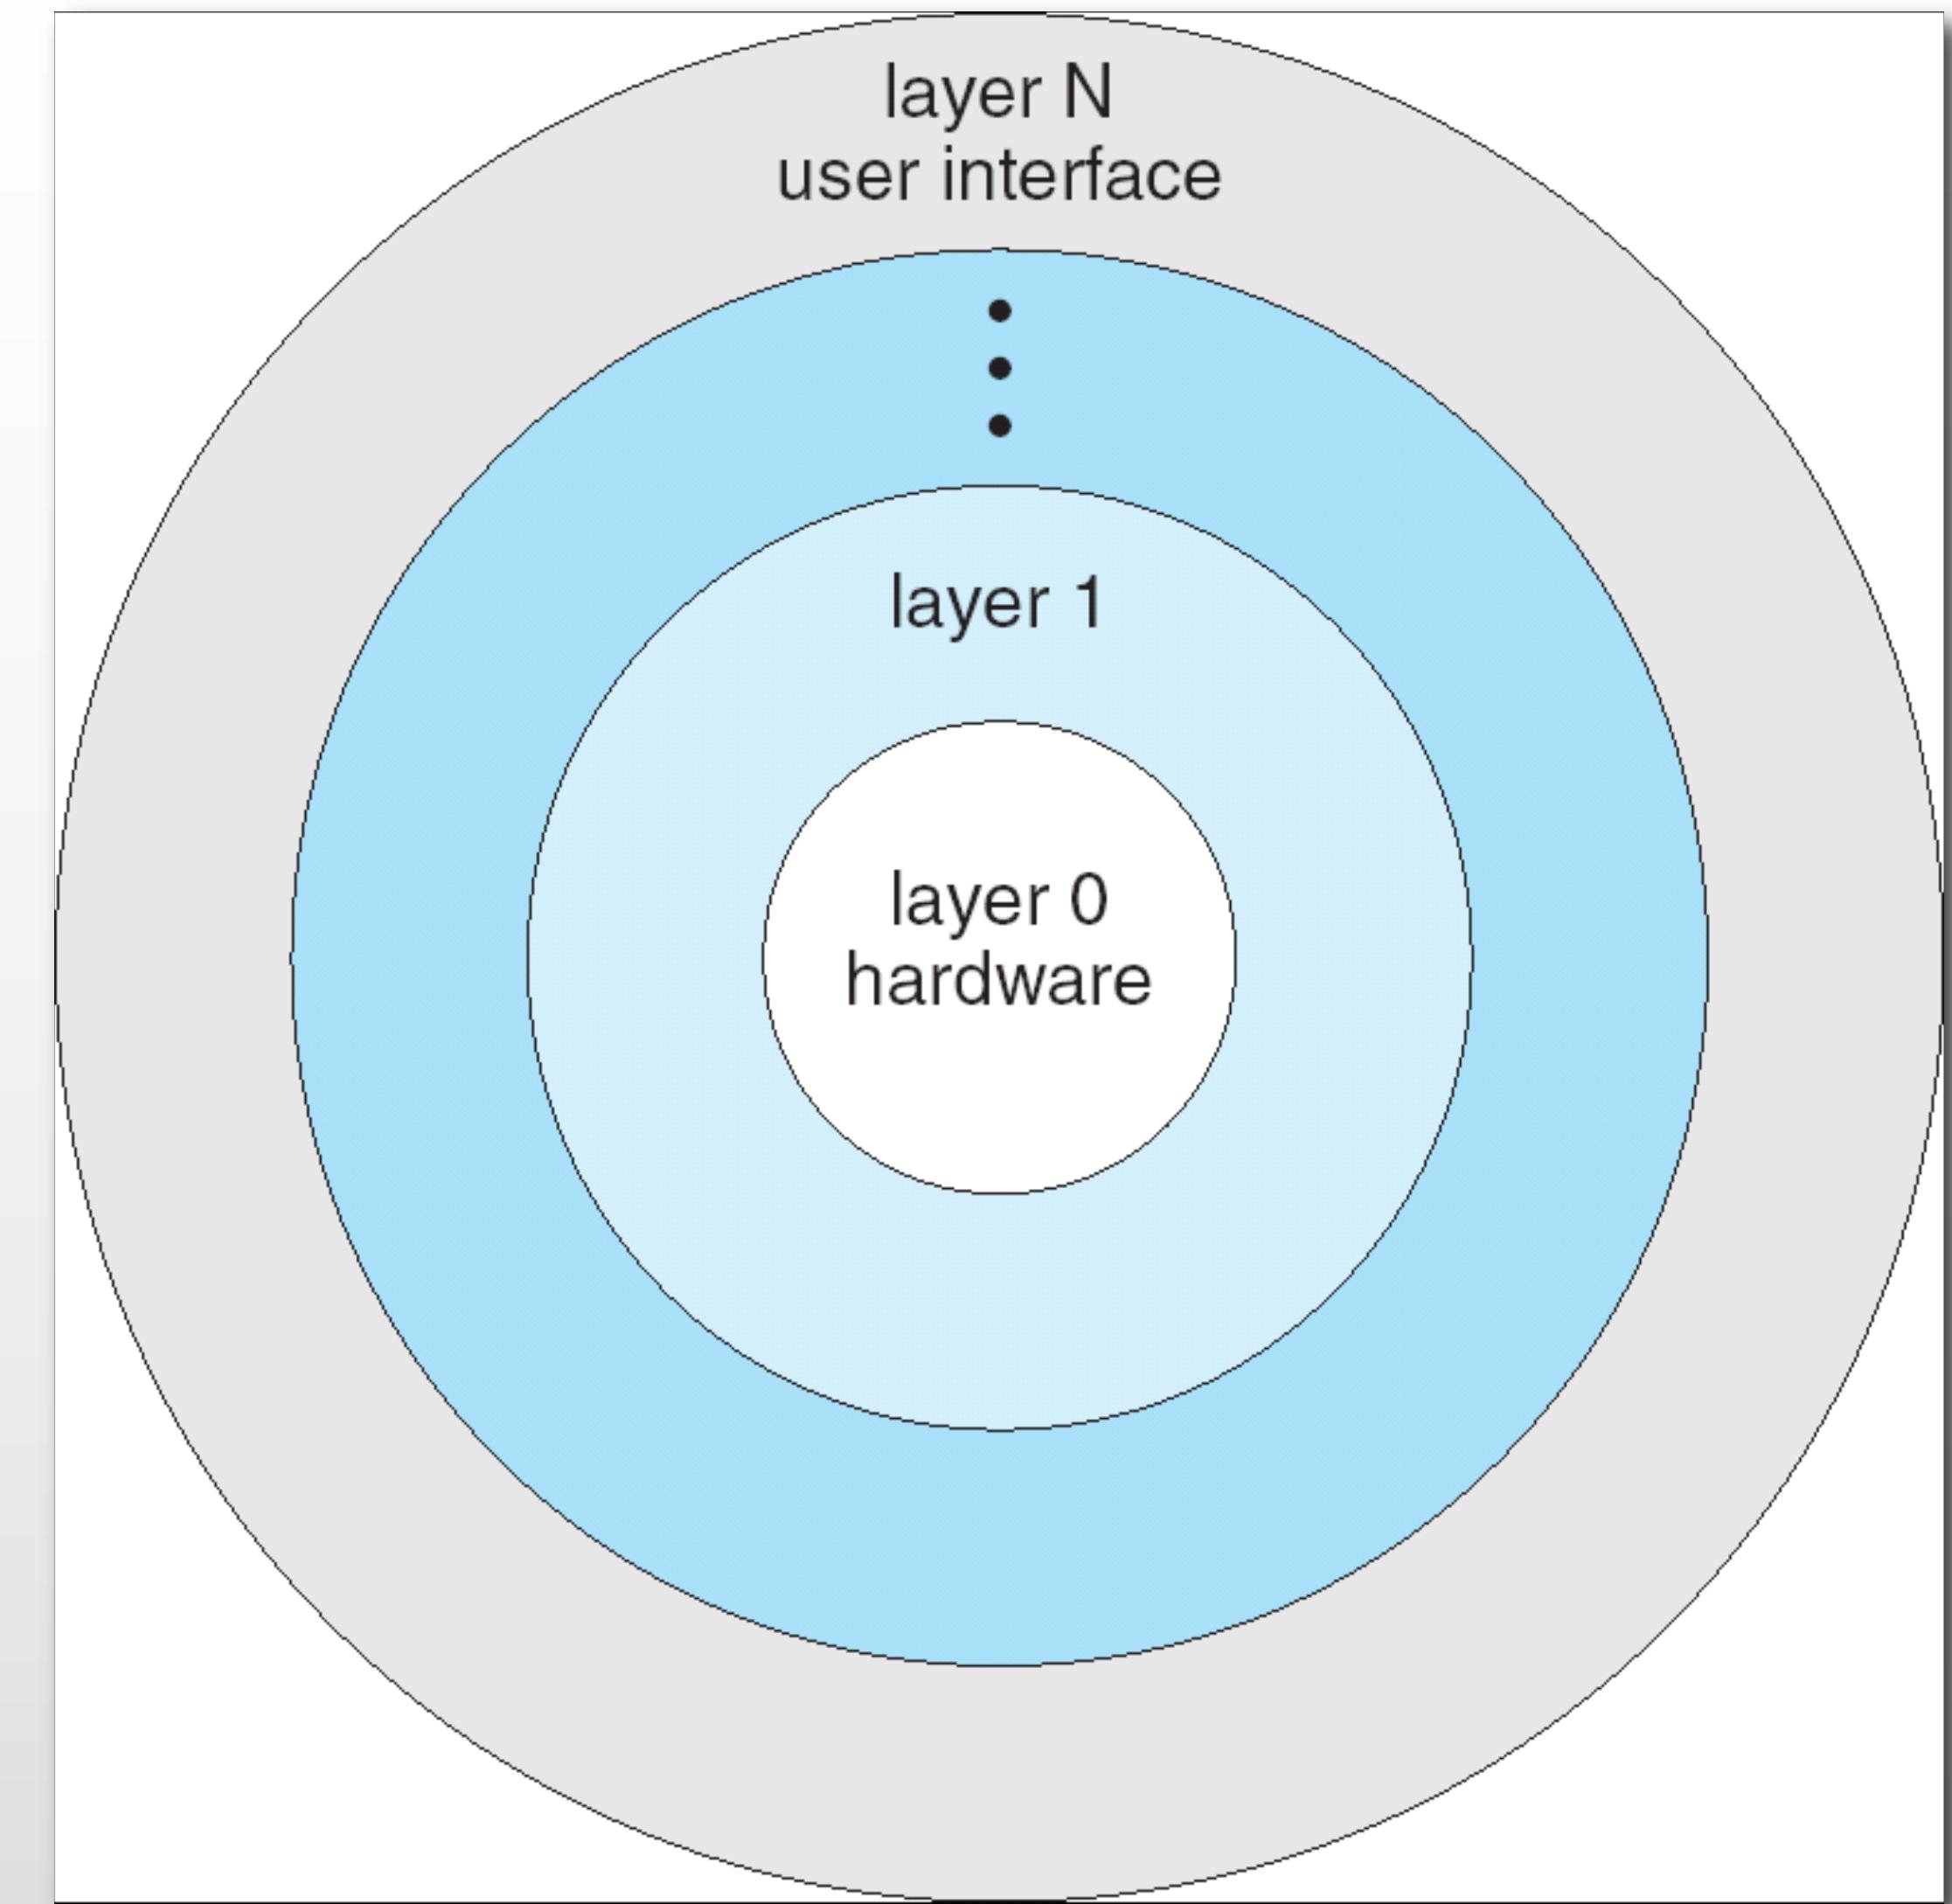
- UNIX – limited by hardware functionality, the original UNIX operating system had limited structuring
  - a lot of this original design is still well visible in modern implementations
- The UNIX OS is layered
  - it consists of two separable parts:
    - systems programs
    - the kernel
  - consists of everything below the system-call interface and above the physical hardware
- Problem
  - kernel provides the file system, CPU scheduling, memory management, and other operating-system functions
  - many functions for one level



# OS Structures: Layered Approach



- Inspired by the extremely successful multi-layer network communication stack
- The operating system is divided into several layers (levels), each built on top of lower layers. The bottom layer (layer 0), is the hardware; the highest (layer N) is the user interface.
- With modularity, layers are selected such that each uses functions (operations) and services of only the lower-level layers
- Nice idea, but serious pragmatic problems
  - tricky to decide what goes into which layer
  - efficiency

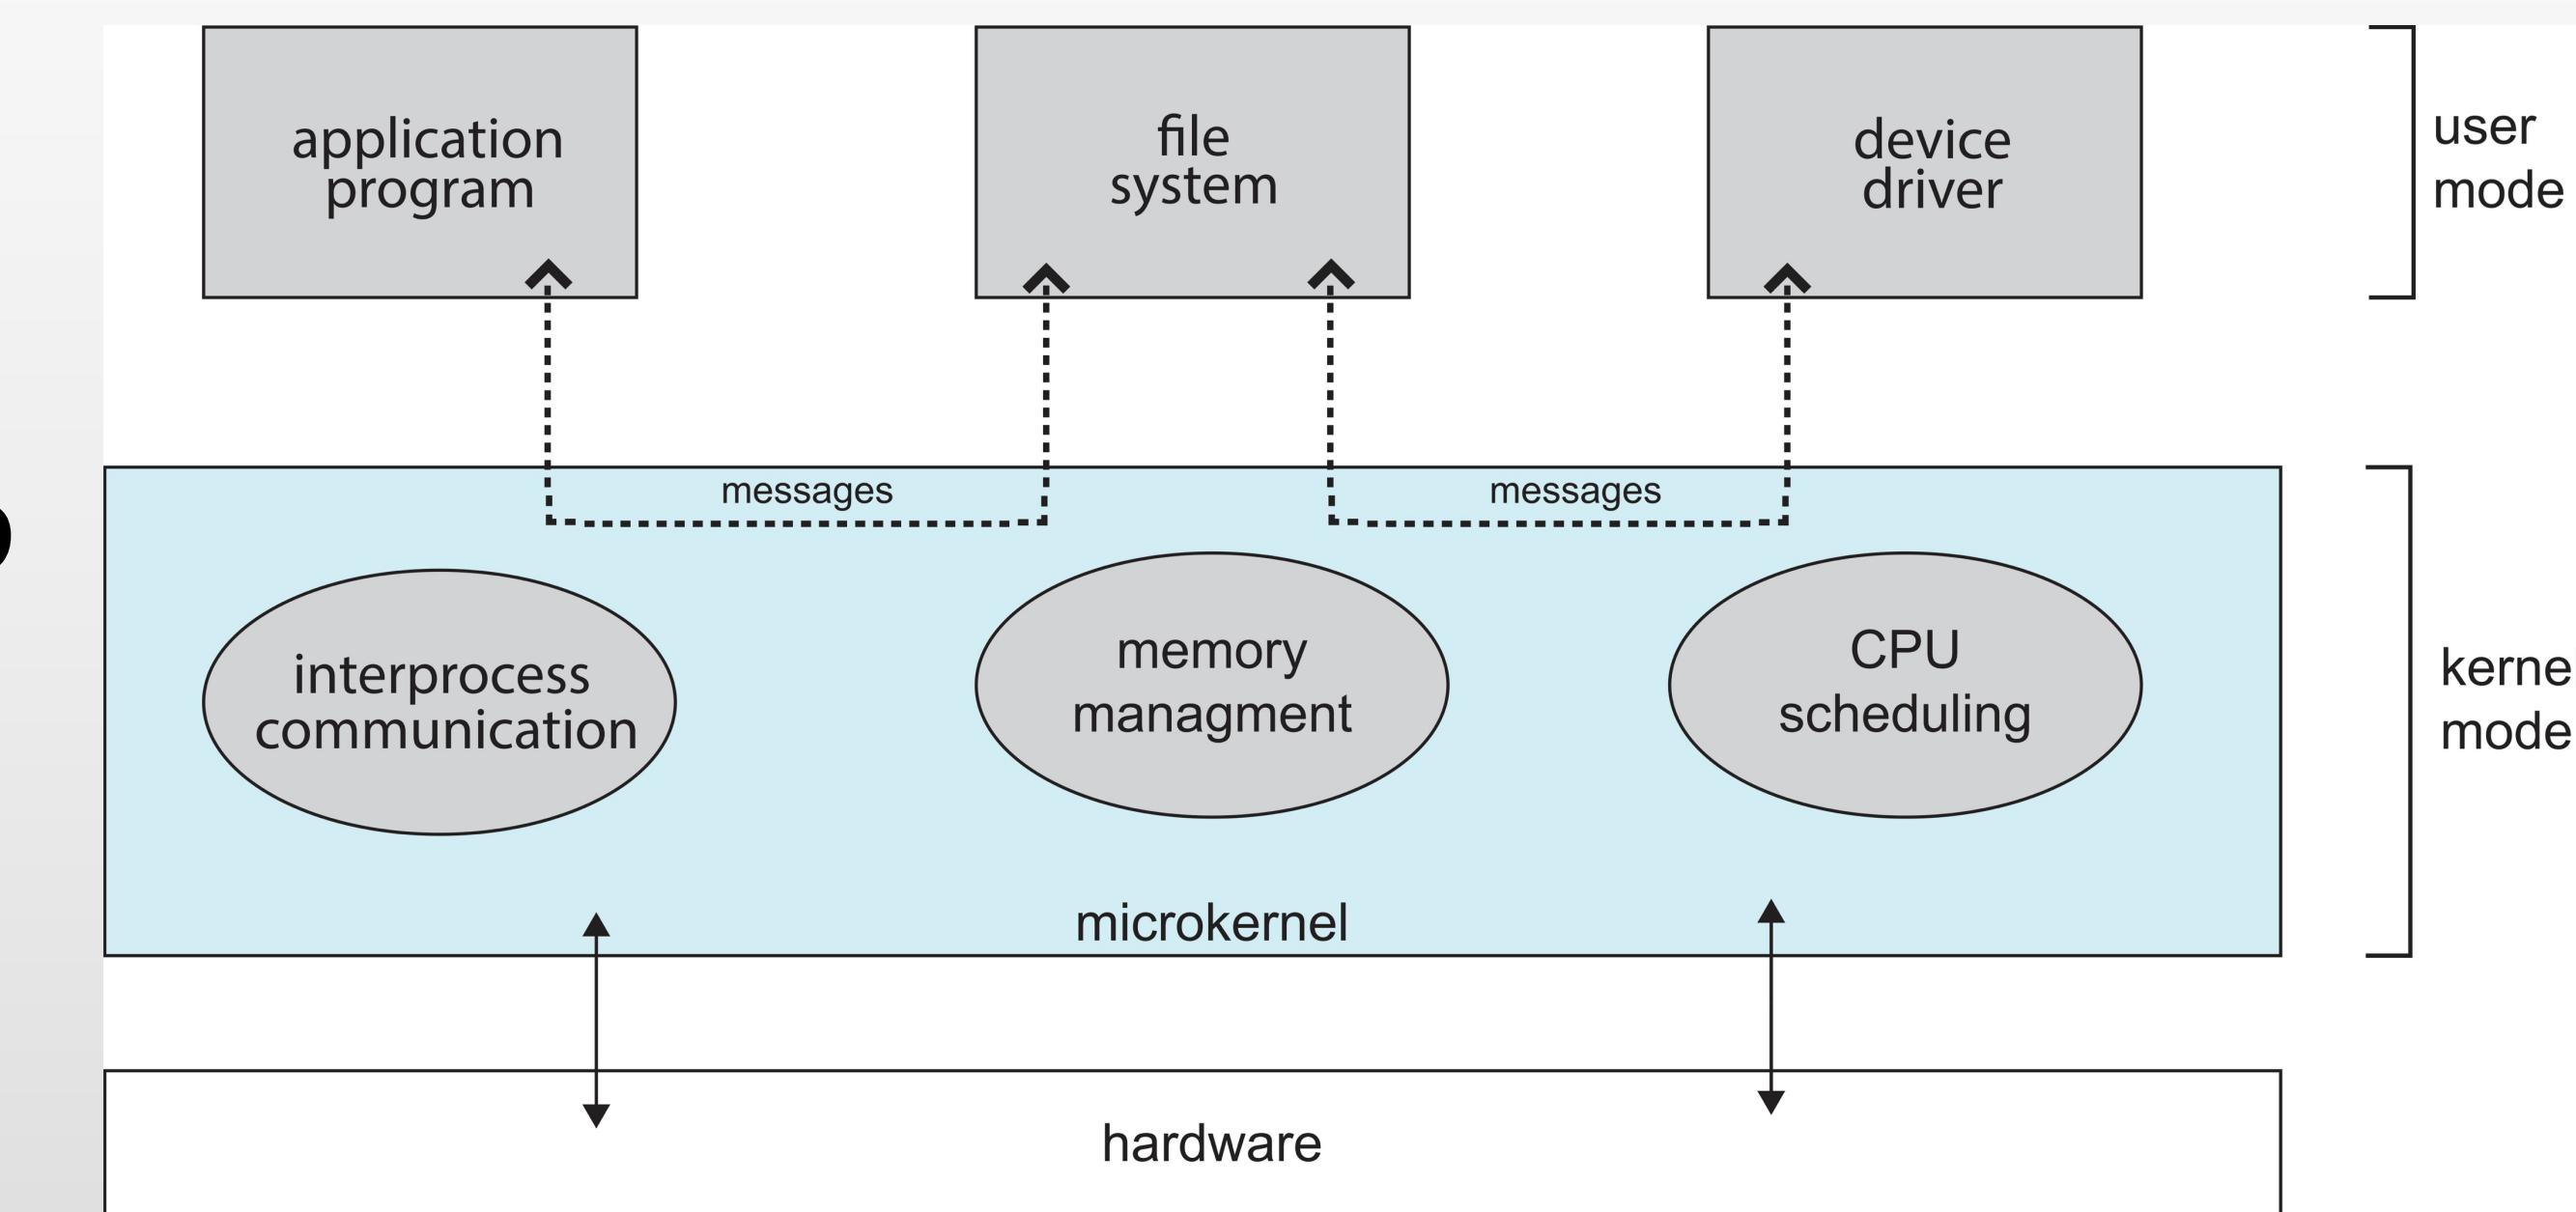


# Microkernel System Structure



- Moves as much from the kernel into “user” space as possible
  - What’s left: **microkernel**
- Communication takes place between user modules using message passing
- Benefits:
  - Easier to extend a microkernel
  - Easier to port the operating system to new architectures
  - More reliable (less code is running in kernel mode)
  - More secure

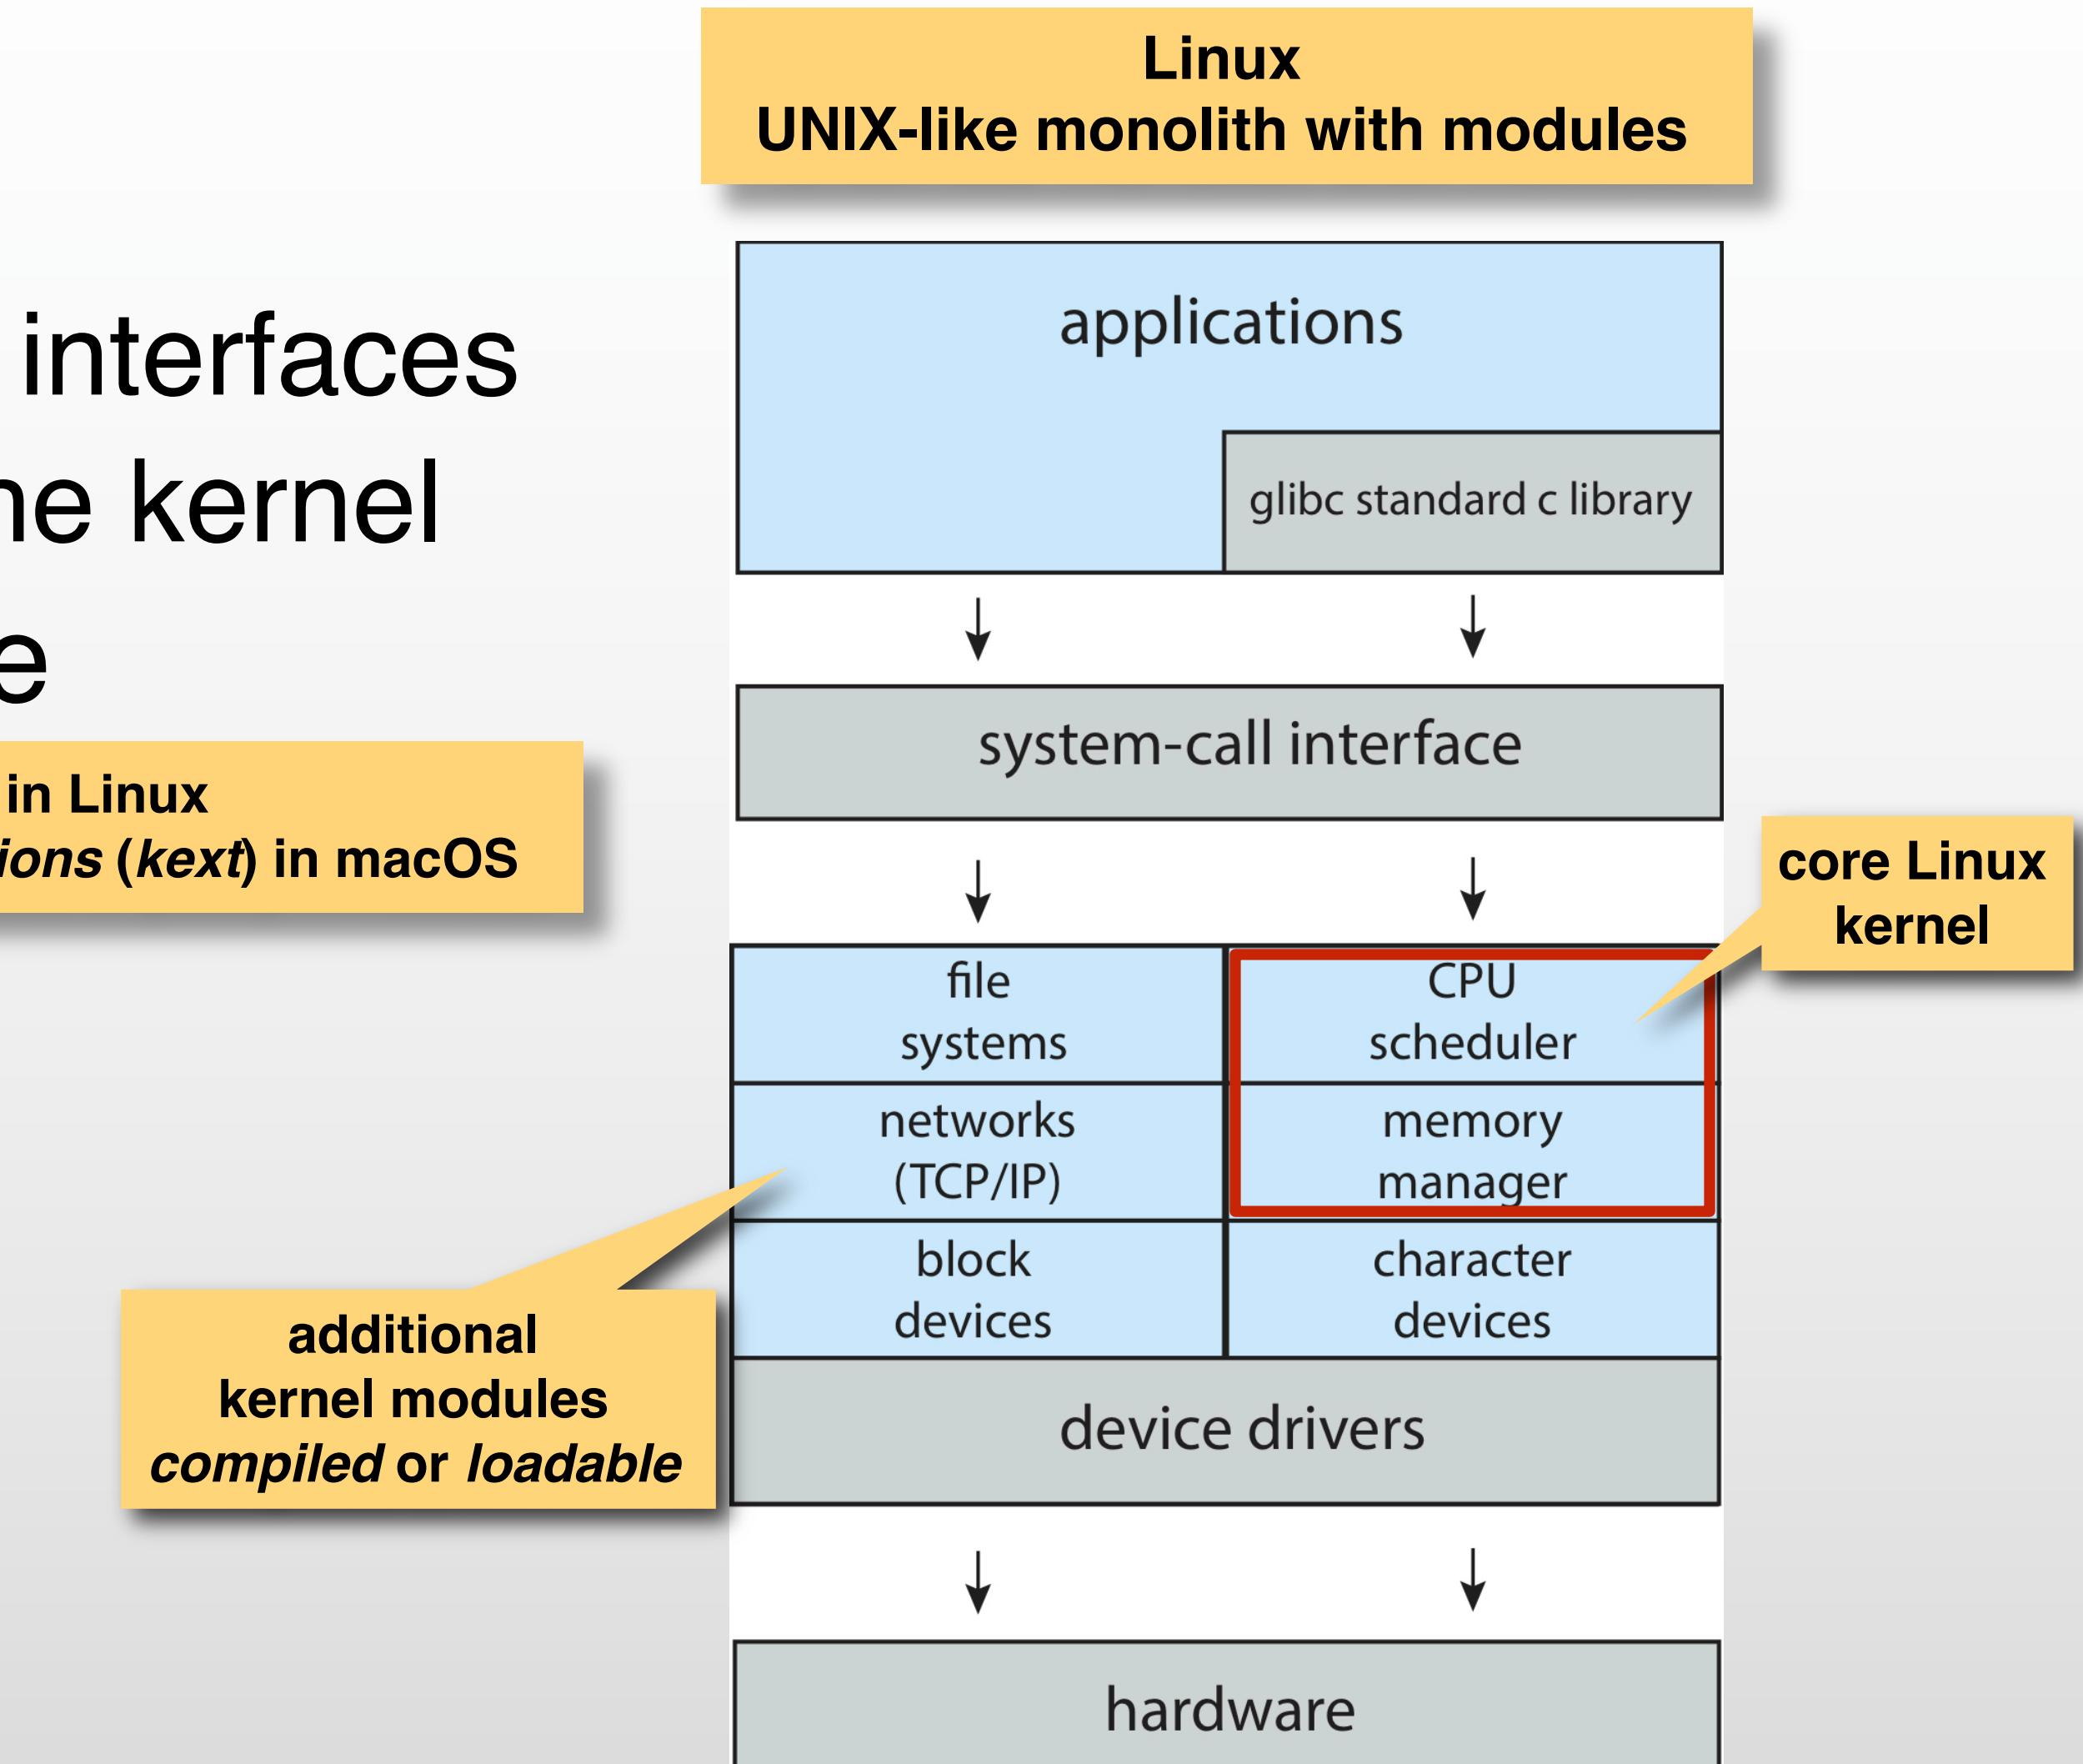
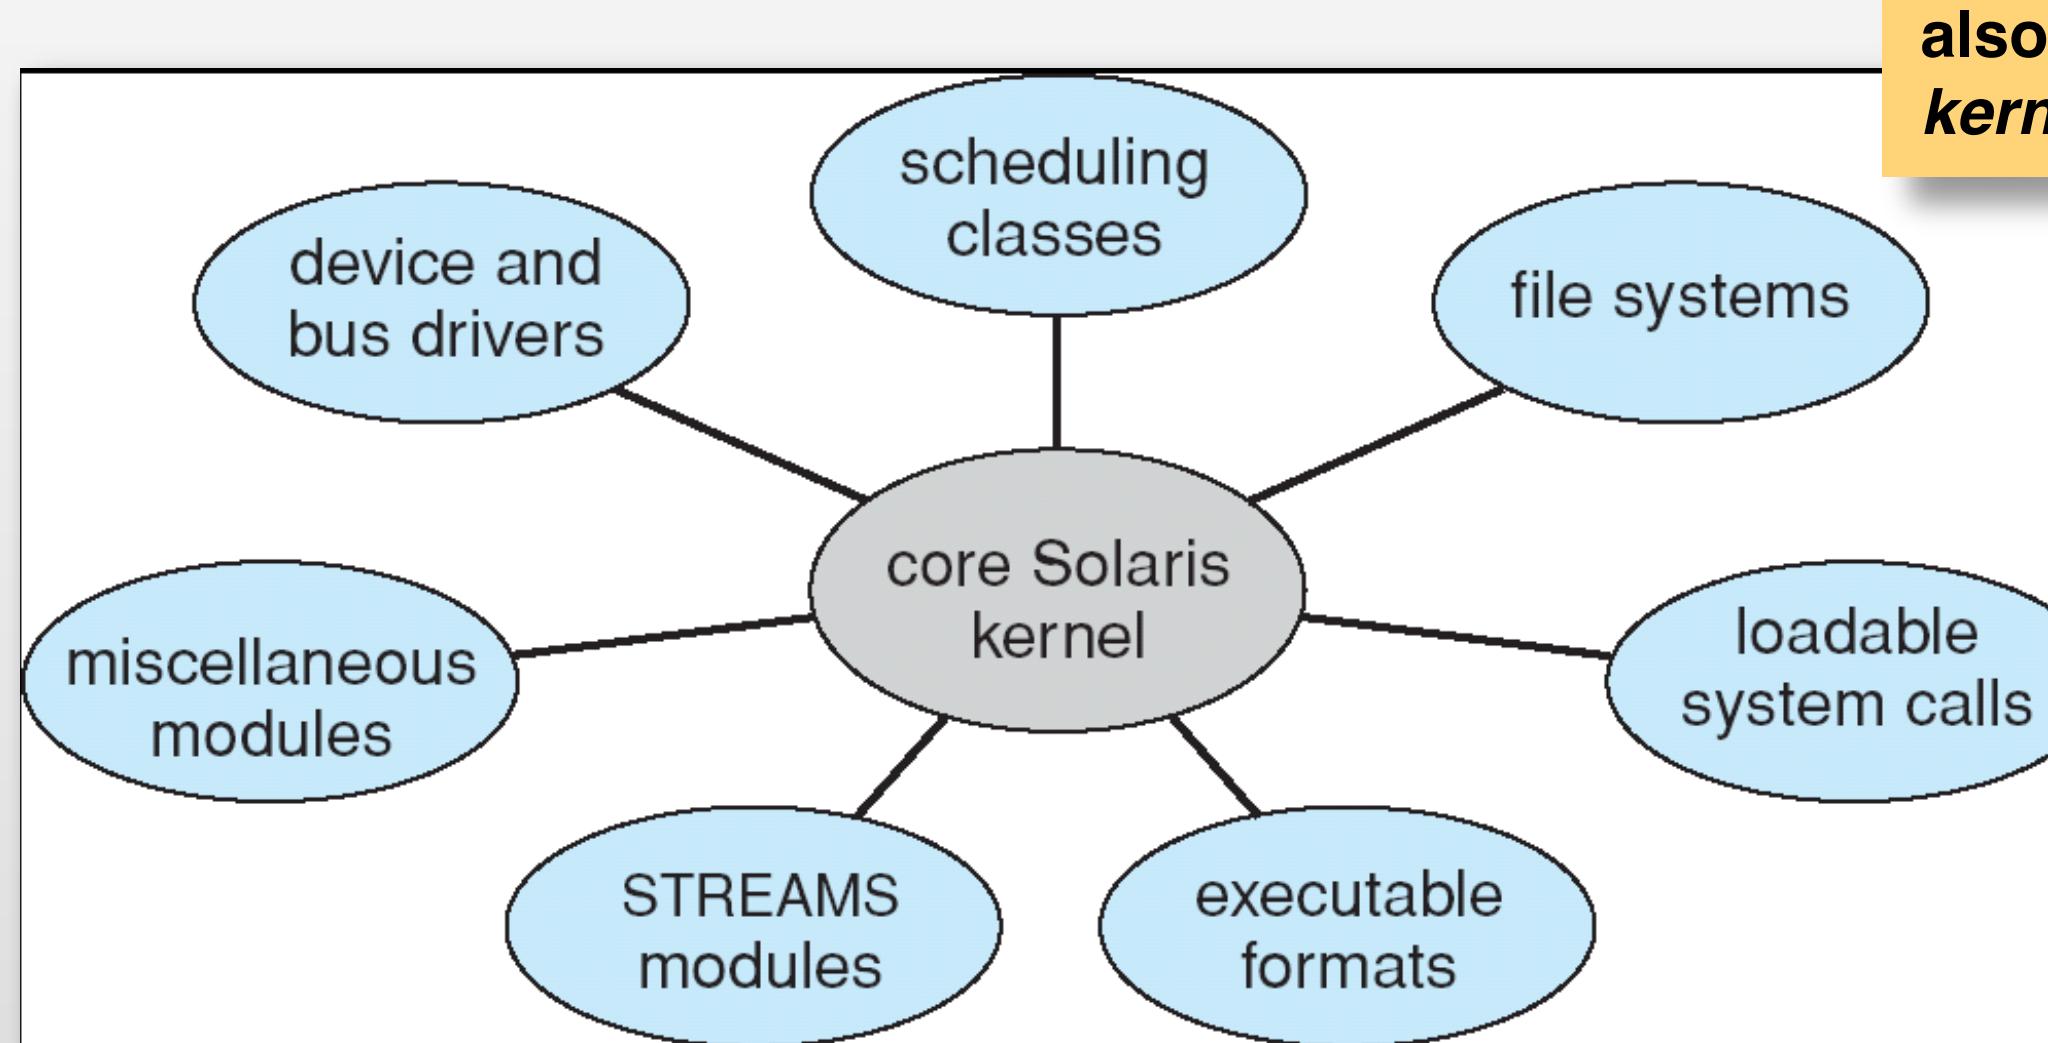
- Detriments:
  - Performance overhead of user space to kernel space communication
  - Often, pragmatic compromises are used to improve efficiency



# Modular OS Architecture



- Most modern operating systems implement kernel modules
  - Object-oriented approach
  - Each core component is separate
  - Each talks to the others over known interfaces
  - Each is loadable as needed within the kernel
- Overall, like layers but more flexible

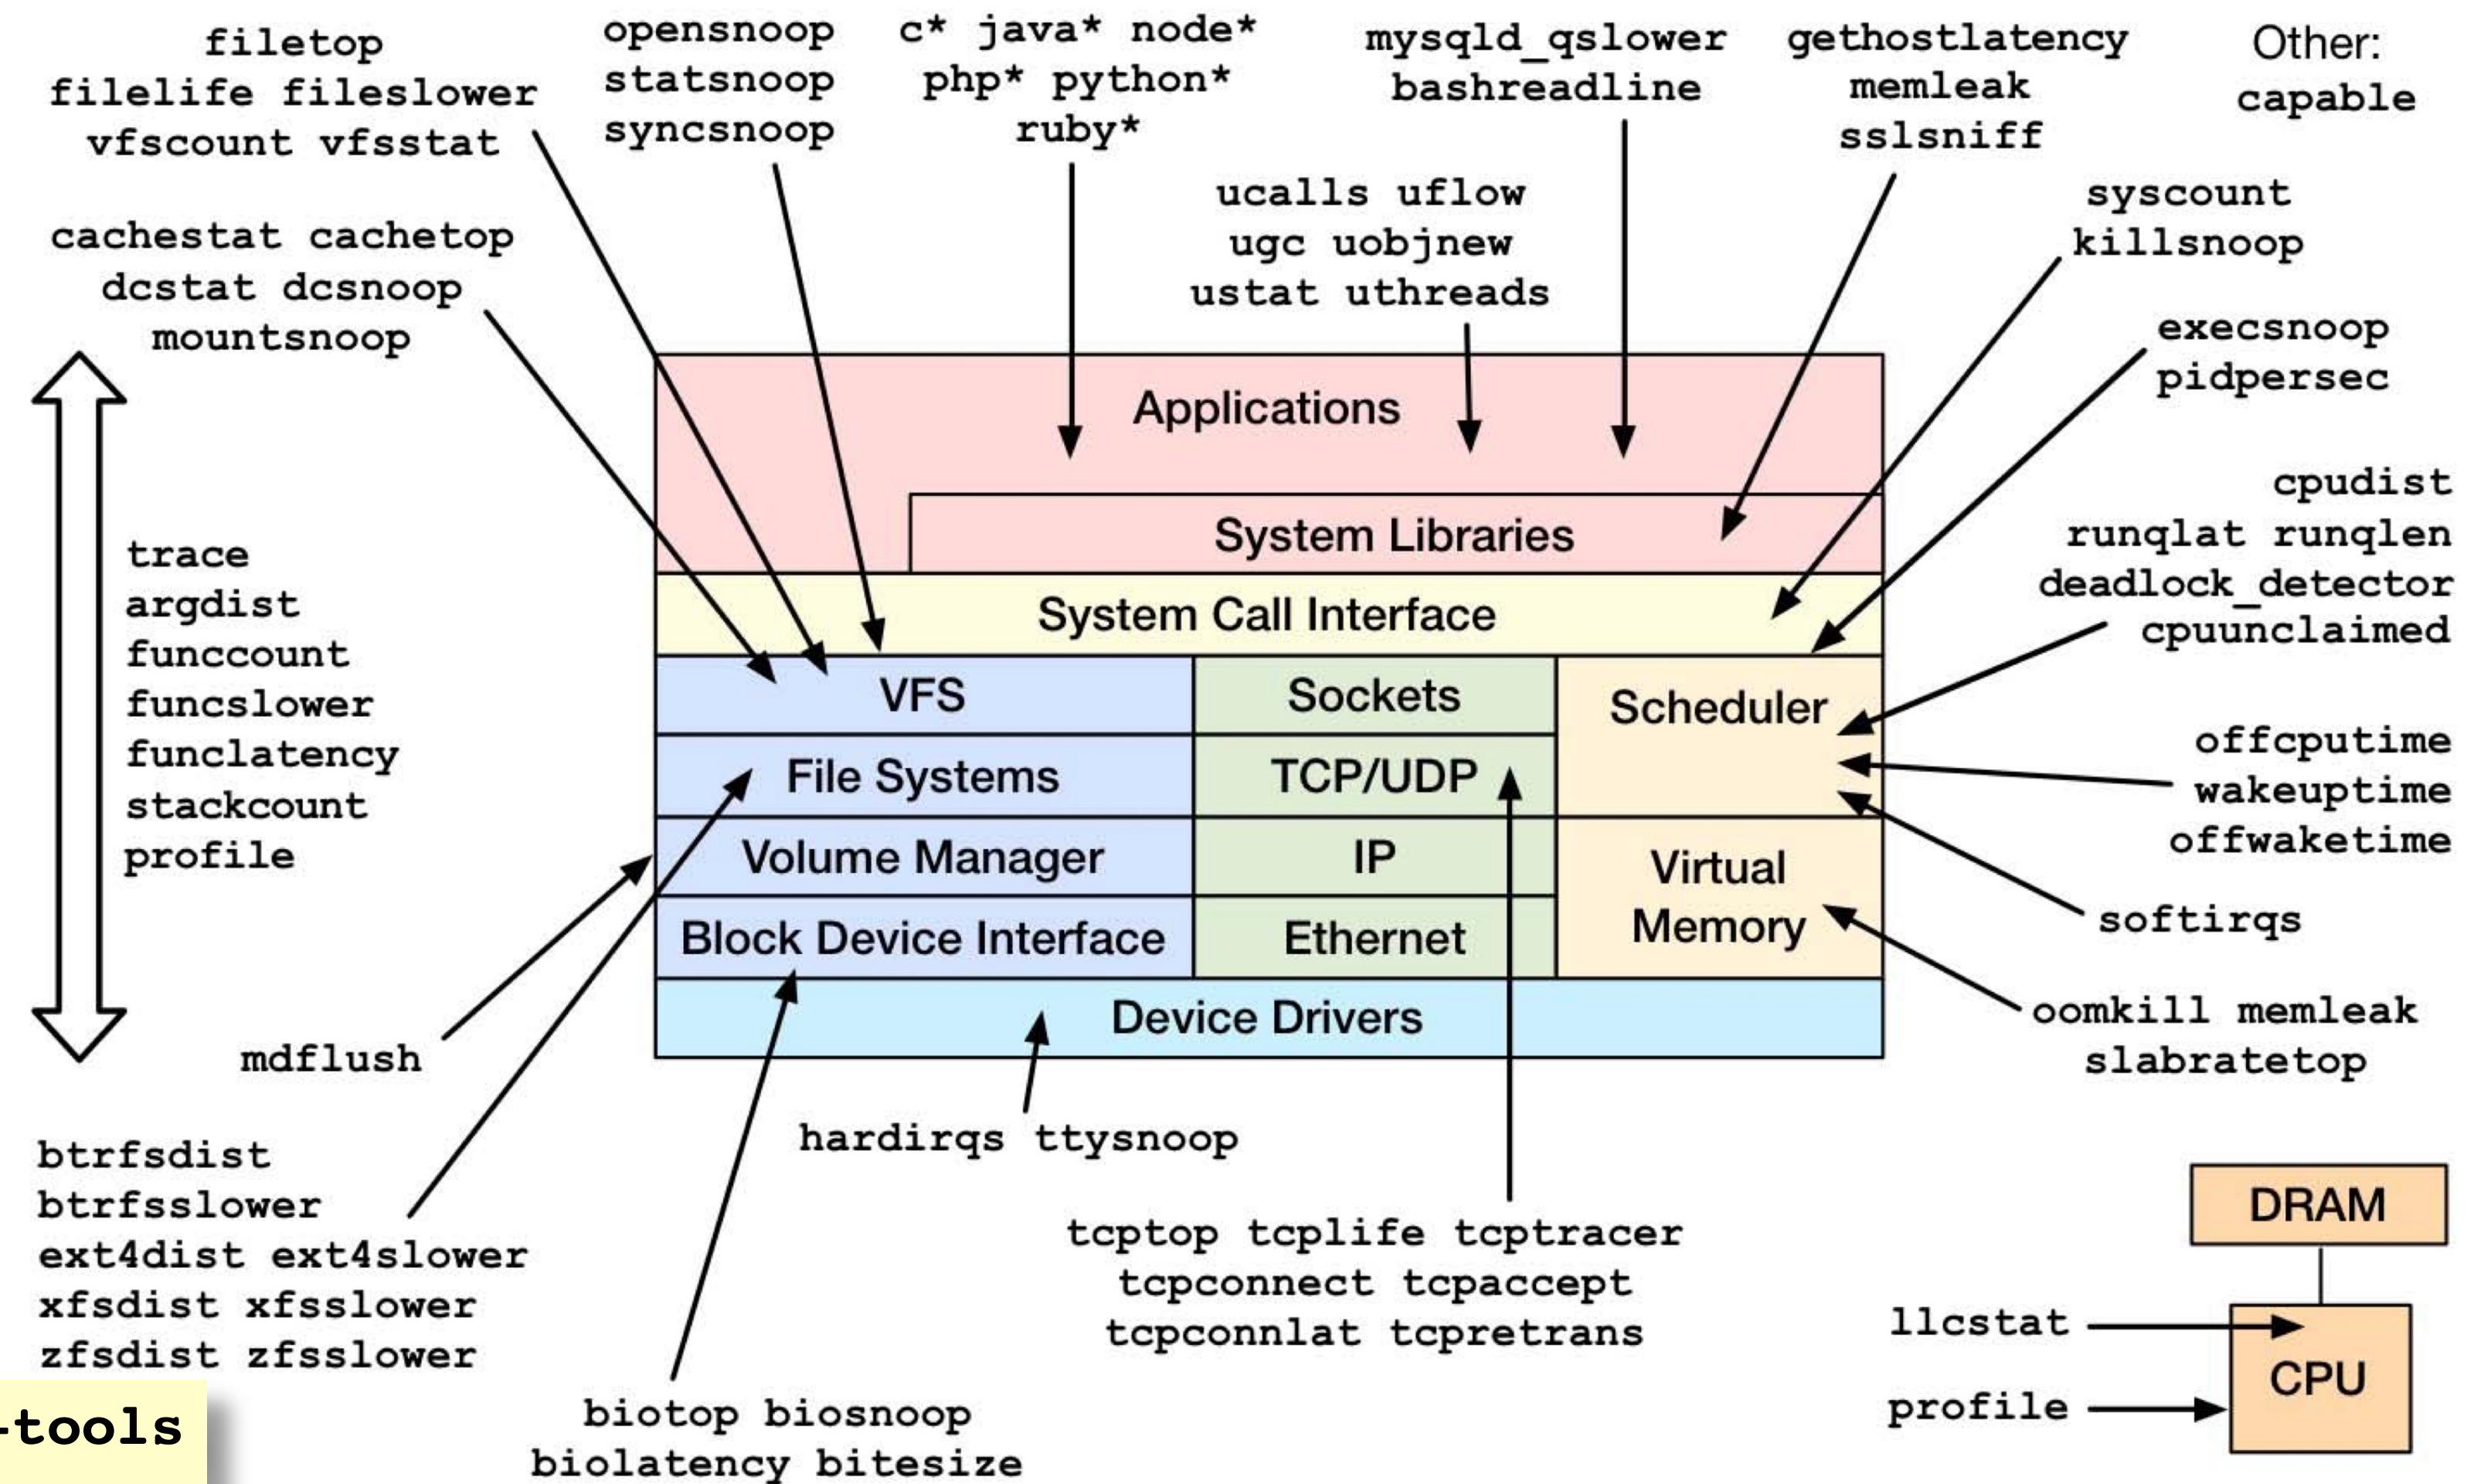


# OS Debugging

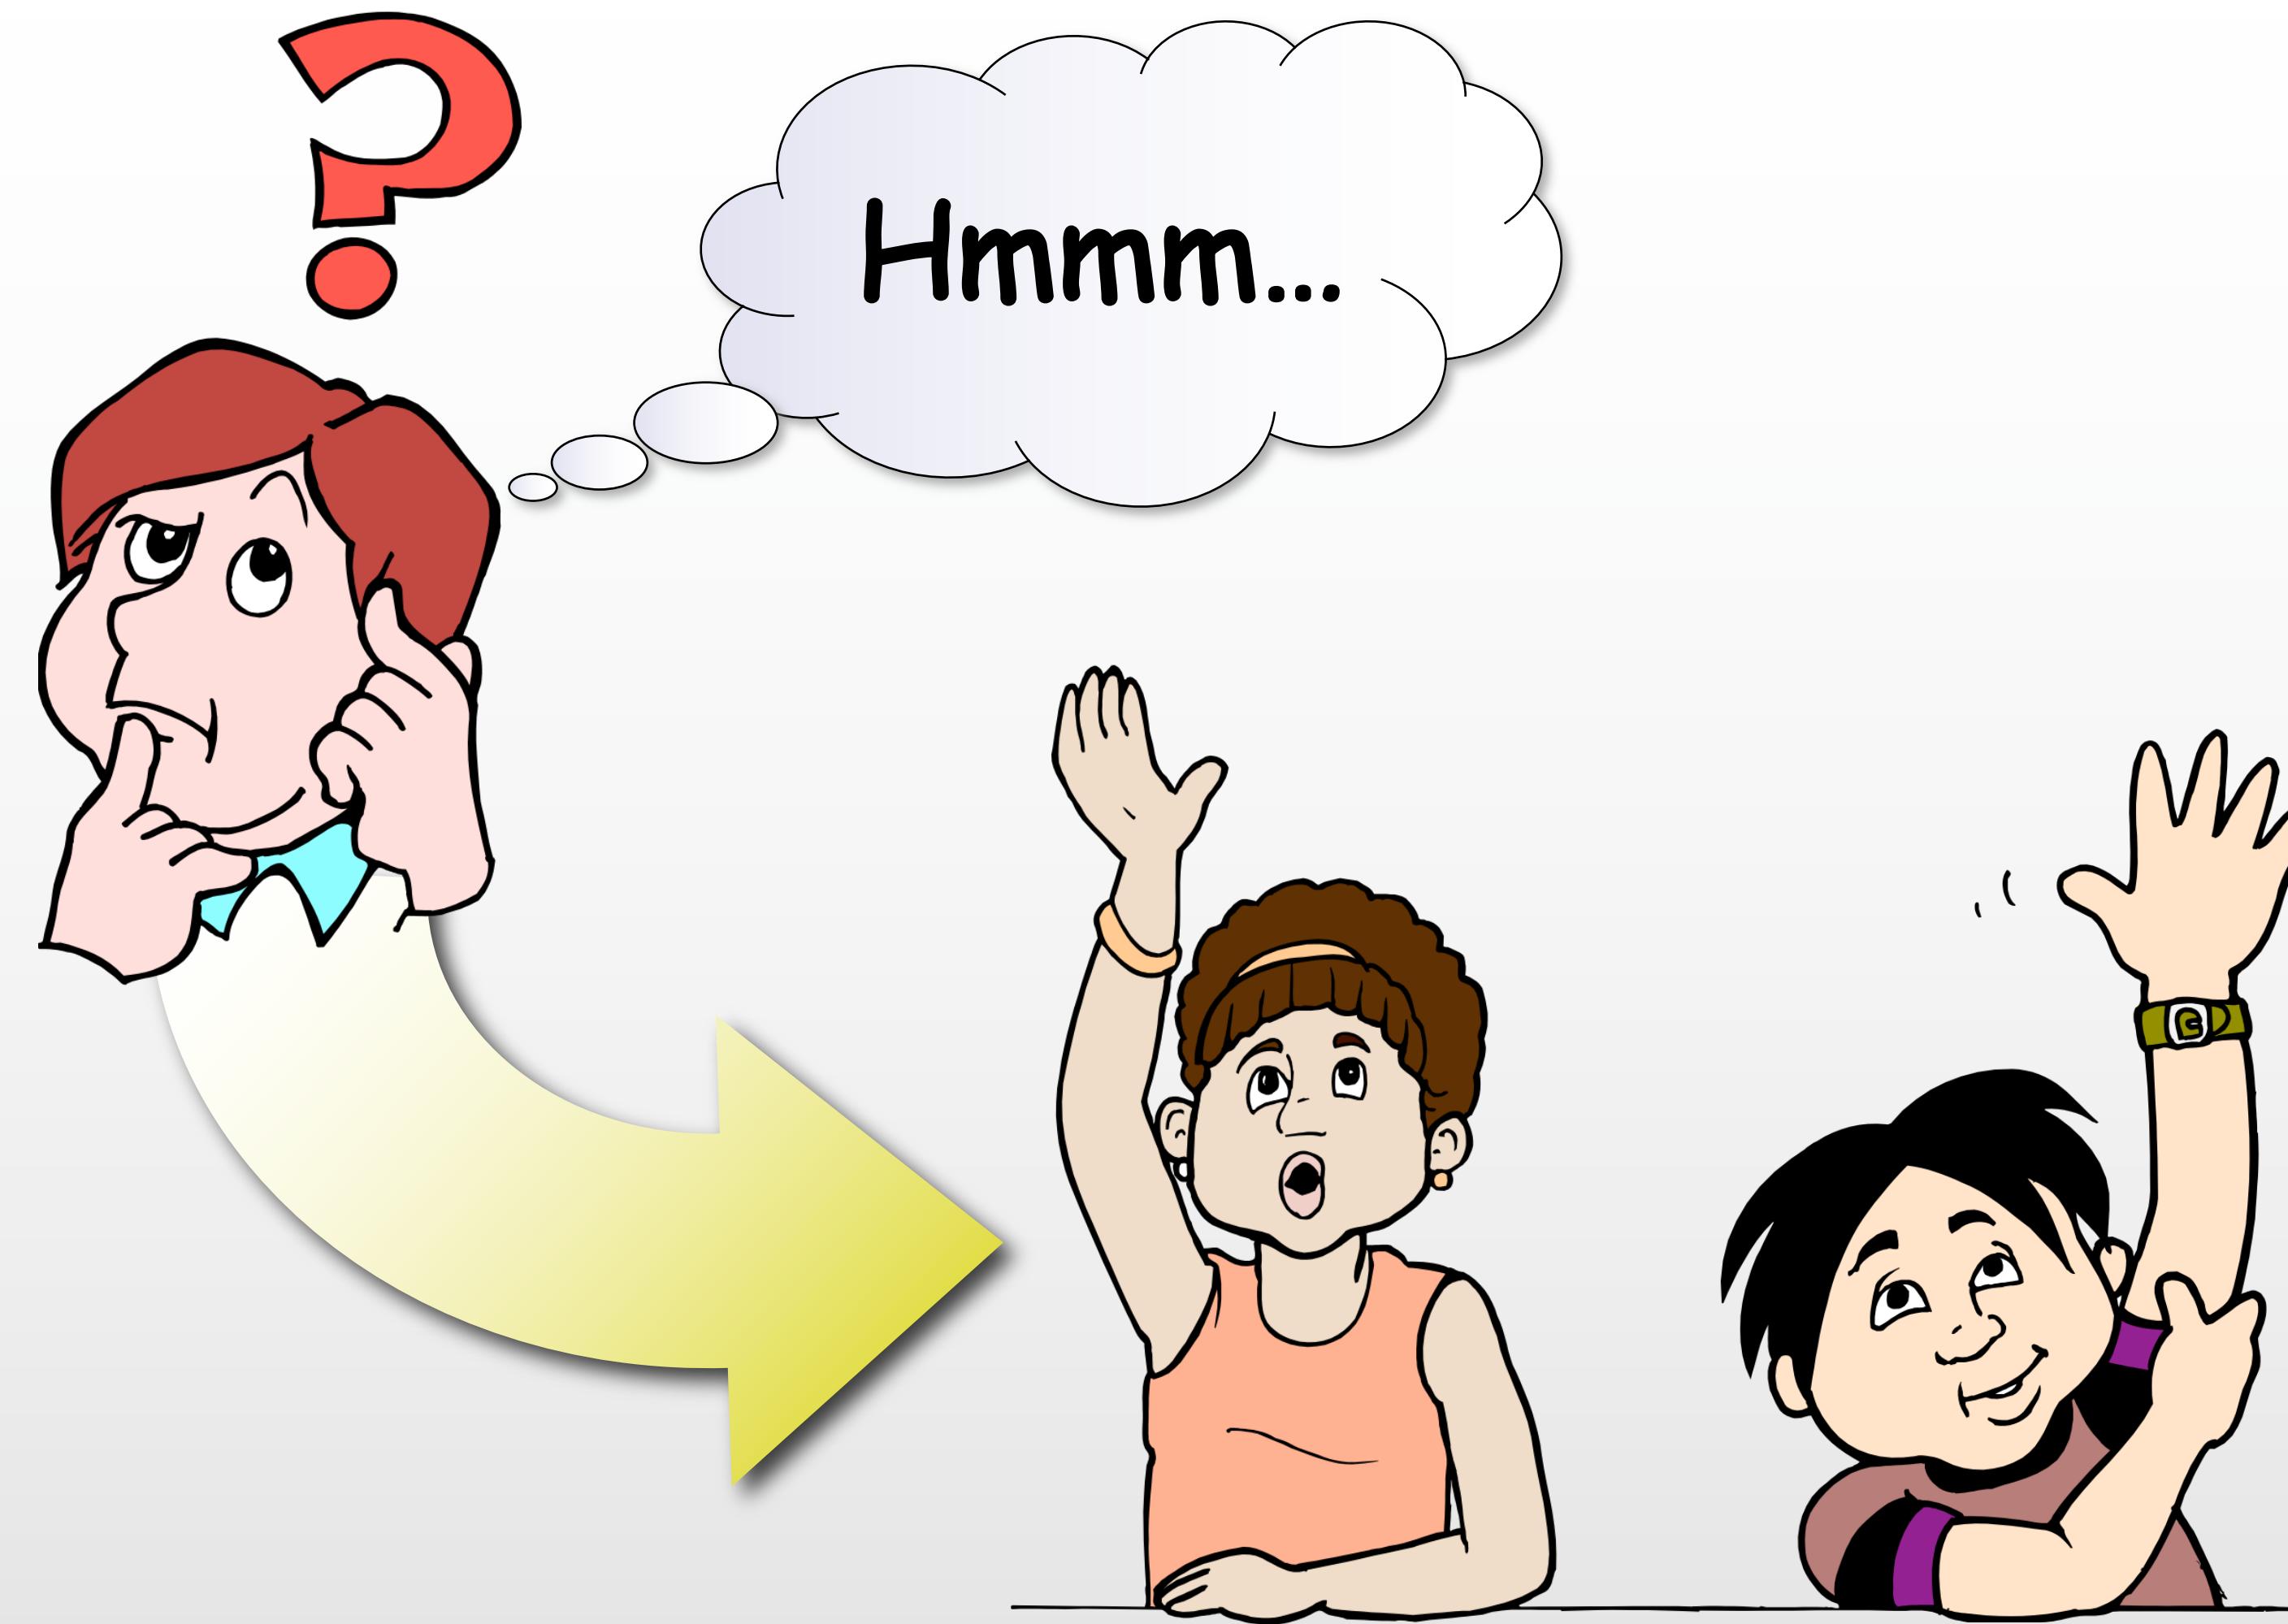


```
$ man gdb  
  
$ man ps  
$ man top  
  
$ man vmstat  
  
$ sudo apt install net-tools  
$ man netstat  
  
$ sudo apt install sysstat  
$ man iostat  
  
$ man strace  
  
$ sudo apt install linux-tools-common  
$ sudo apt install linux-tools-generic  
$ man perf  
  
$ man tcpdump  
  
$ sudo apt install bpfcc-tools  
$ sudo syscount-bpfcc
```

## Linux bcc/BPF Tracing Tools



<https://github.com/iovisor/bcc#tools 2017>



**COMP362 Operating Systems**  
**Prof. AJ Biesczad**