



Chapter 3: Processes

COMP362 Operating Systems

Prof. AJ Bieszczad

Outline: Processes



- Process Concept
- Process Lifecycle
- Process Control Block (PCB)
- Process Scheduling
- Context Switch
- Process Creation
 - cloning with `fork()`
 - morphing with `exec()`
- Processes Termination

Process Concept



- An operating system executes a variety of programs:
 - Batch system – jobs
 - Time-shared systems – user programs or tasks
- Textbook uses the terms **job** and **process** almost interchangeably
- **Process** – a program in execution
 - process execution must progress in sequential fashion
- A process includes memory for:
 - text
 - code
 - data
 - static data of the program
 - stack
 - for function calls; fixed size; grows “down”
 - heap
 - for dynamic allocation; OS can resize it as needed
- Some addresses randomized to increase security
- Resources can be controlled by system calls

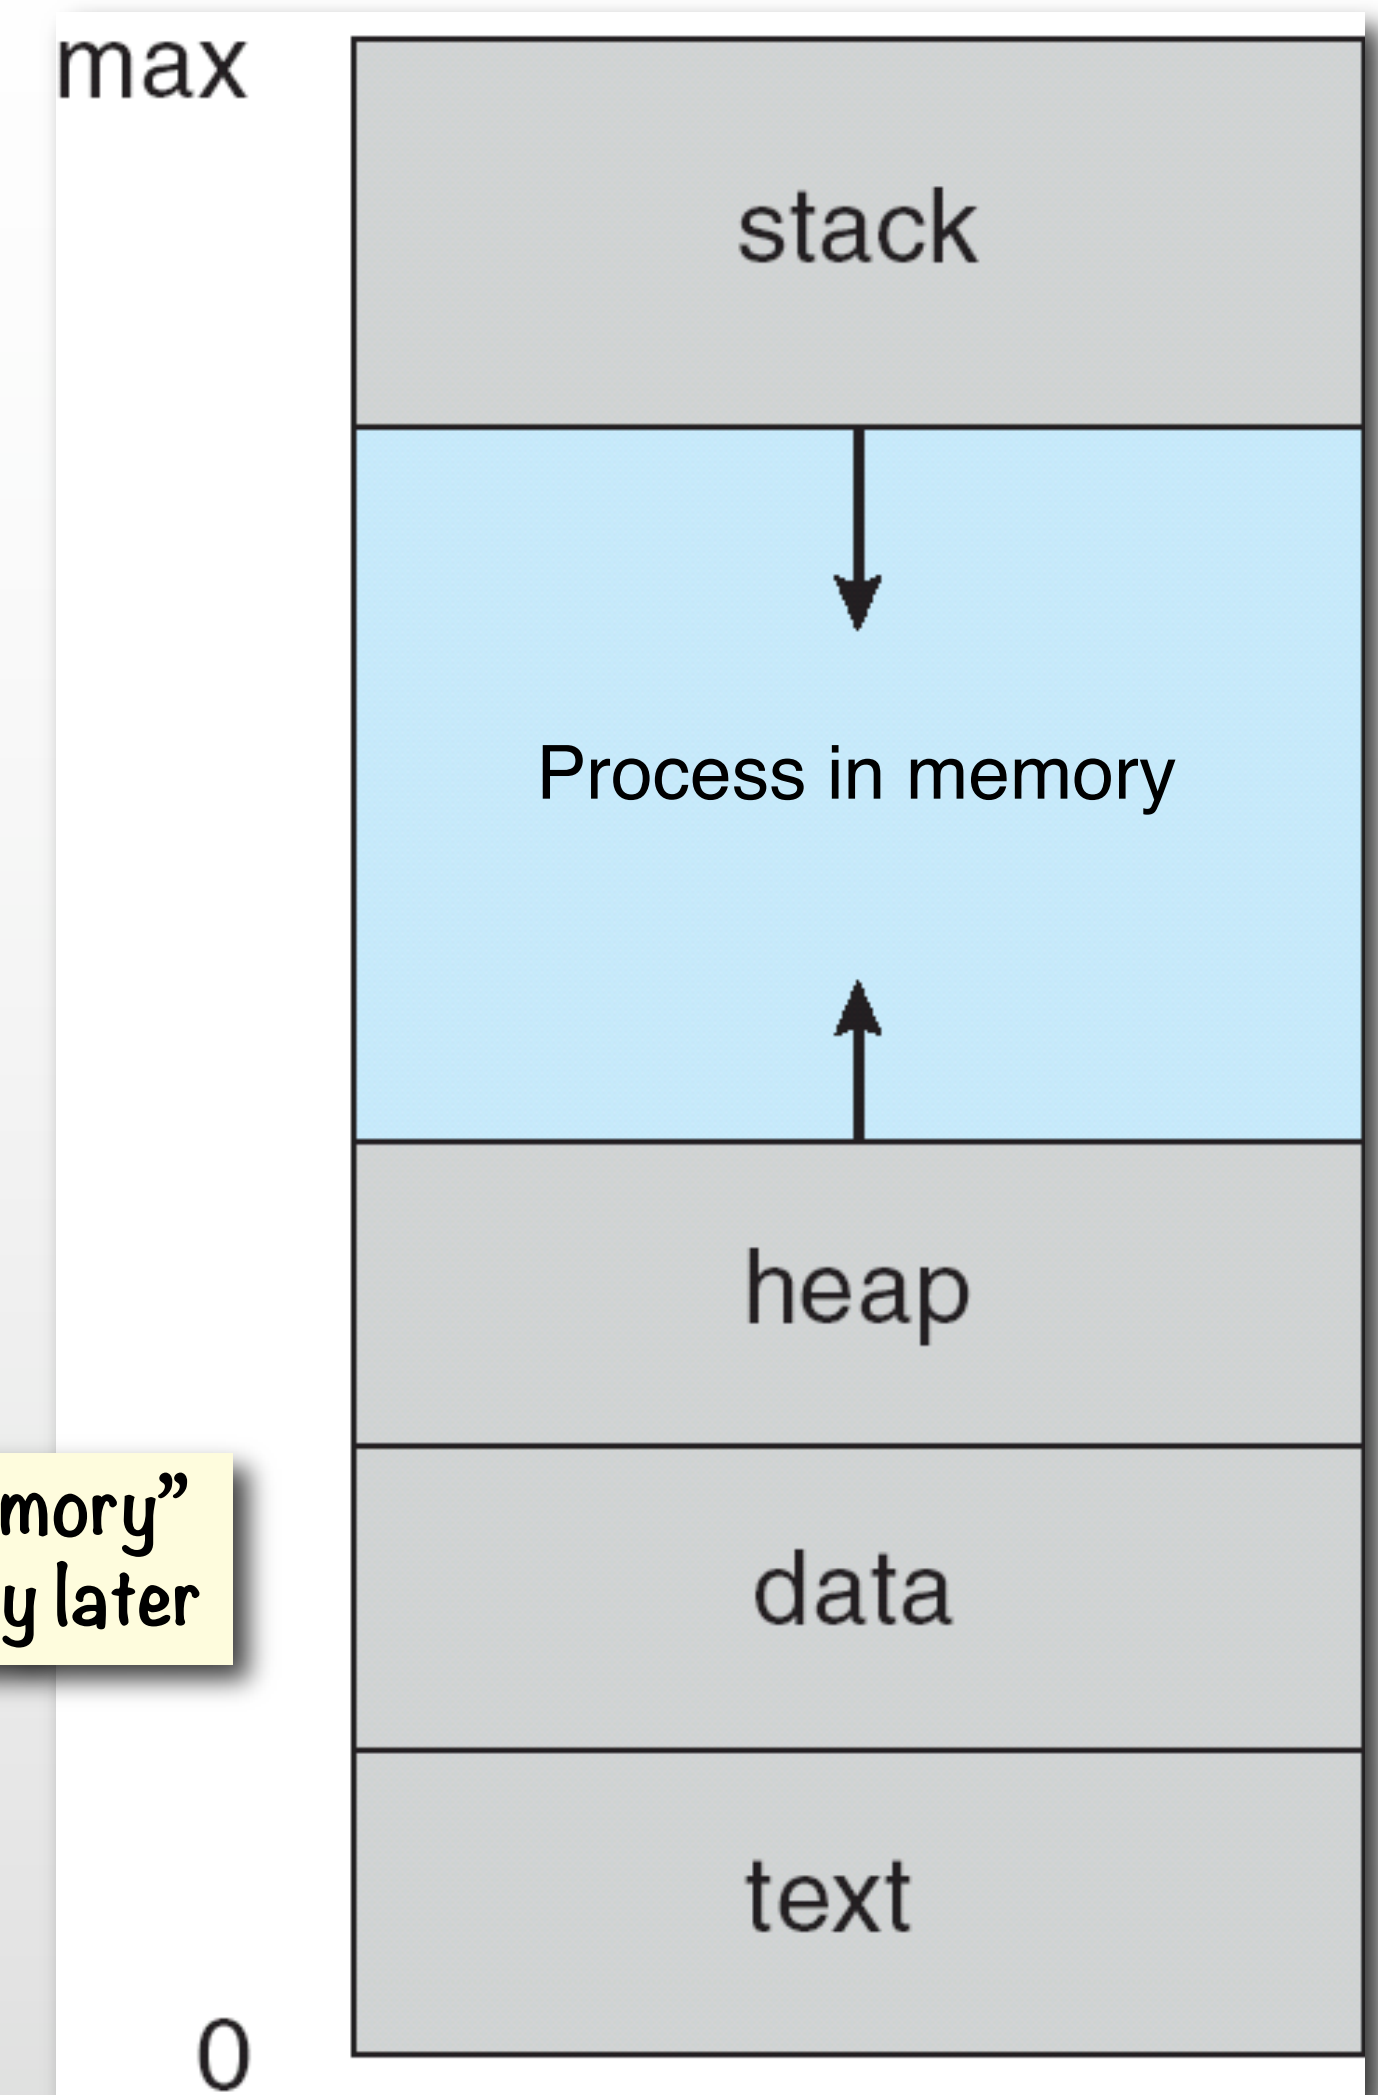
```
$ cat /proc/6514/maps  
$ pmap 6514
```

process ID

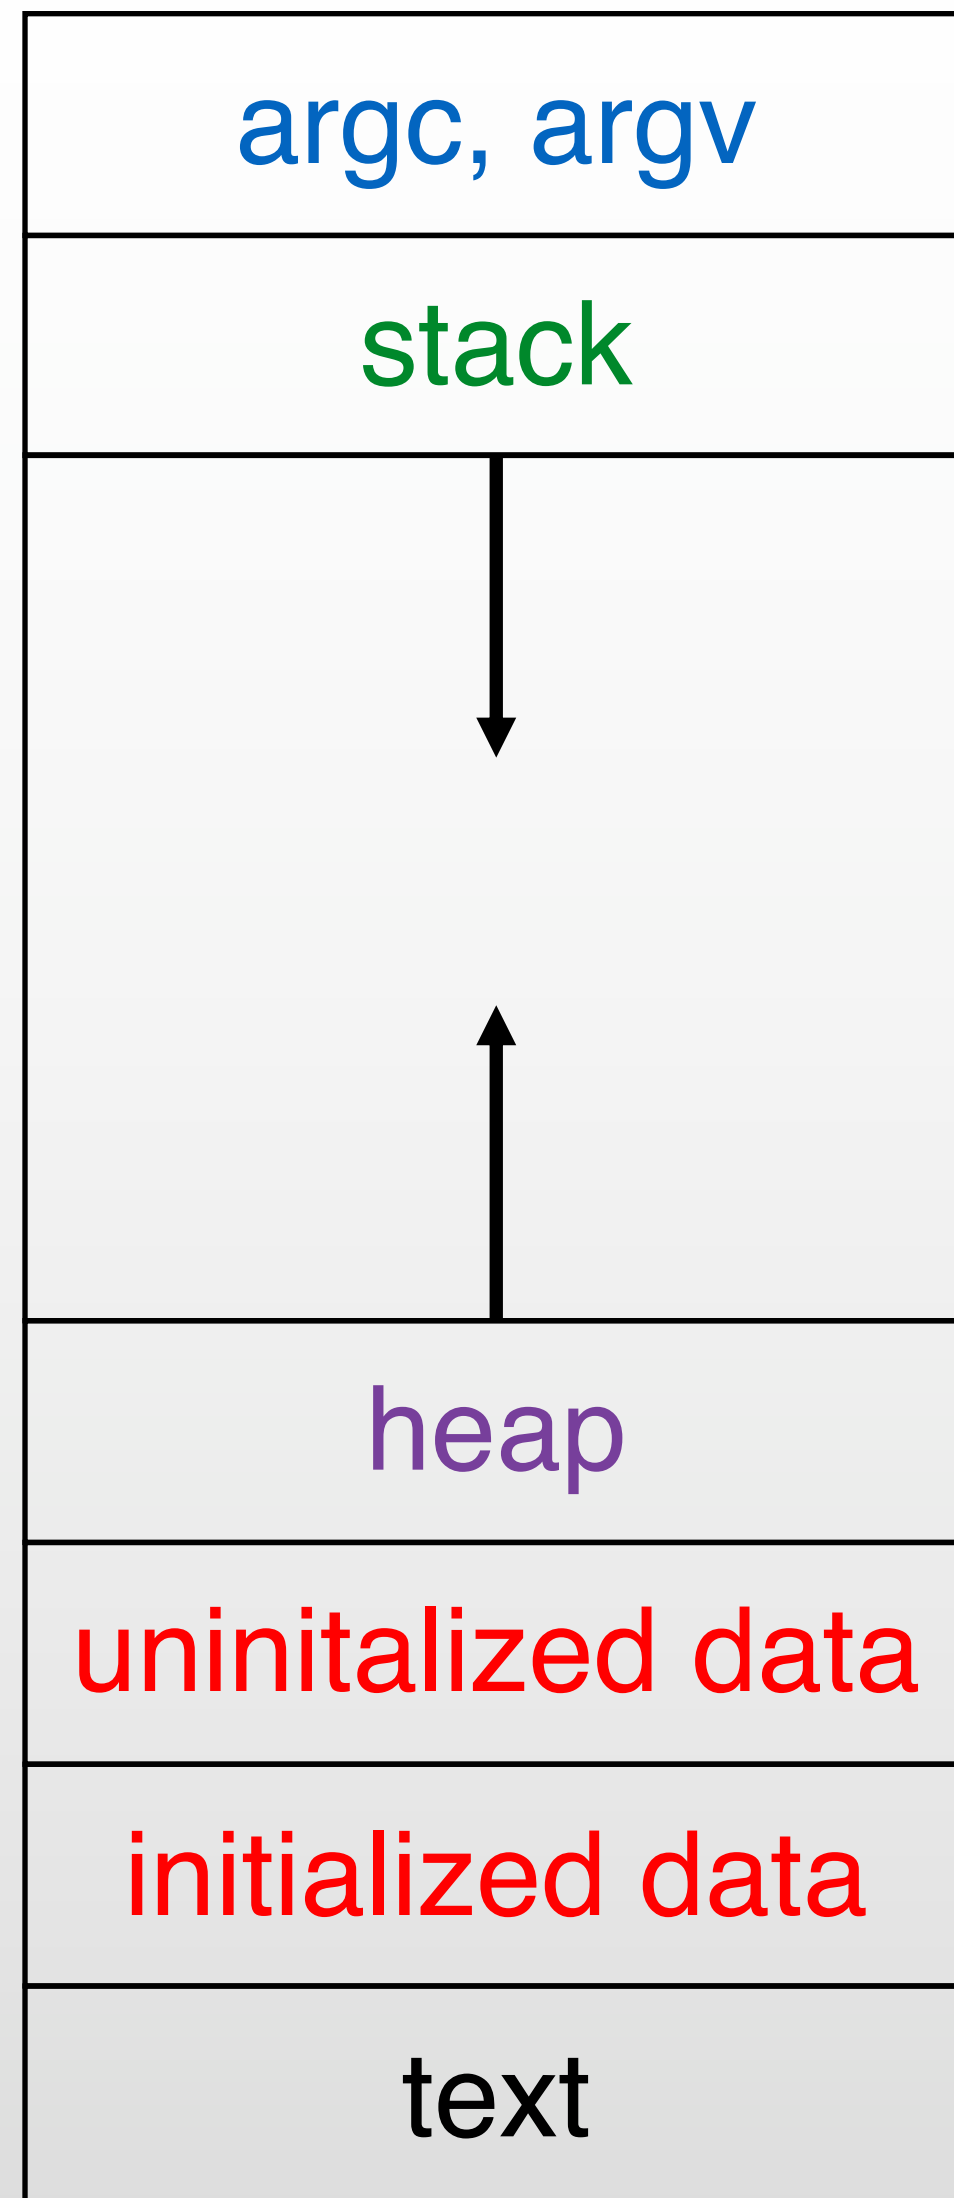
Linux provides procs
macOS does not

“virtual memory”
we will study later

```
$ man ulimit  
$ man setrlimit
```



Data Placement in a Process



```
#include <stdio.h>
#include <stdlib.h>
```

```
int error_type;
char *error_msg = "ERROR";
```

```
int main(int argc, char **argv)
{
    int *val;
```

```
    val = (int *) malloc(5 * sizeof(int));
```

```
    for (int i = 0; i < 5; i++)
        val[i] = i;
```

```
    return EXIT_SUCCESS;
```

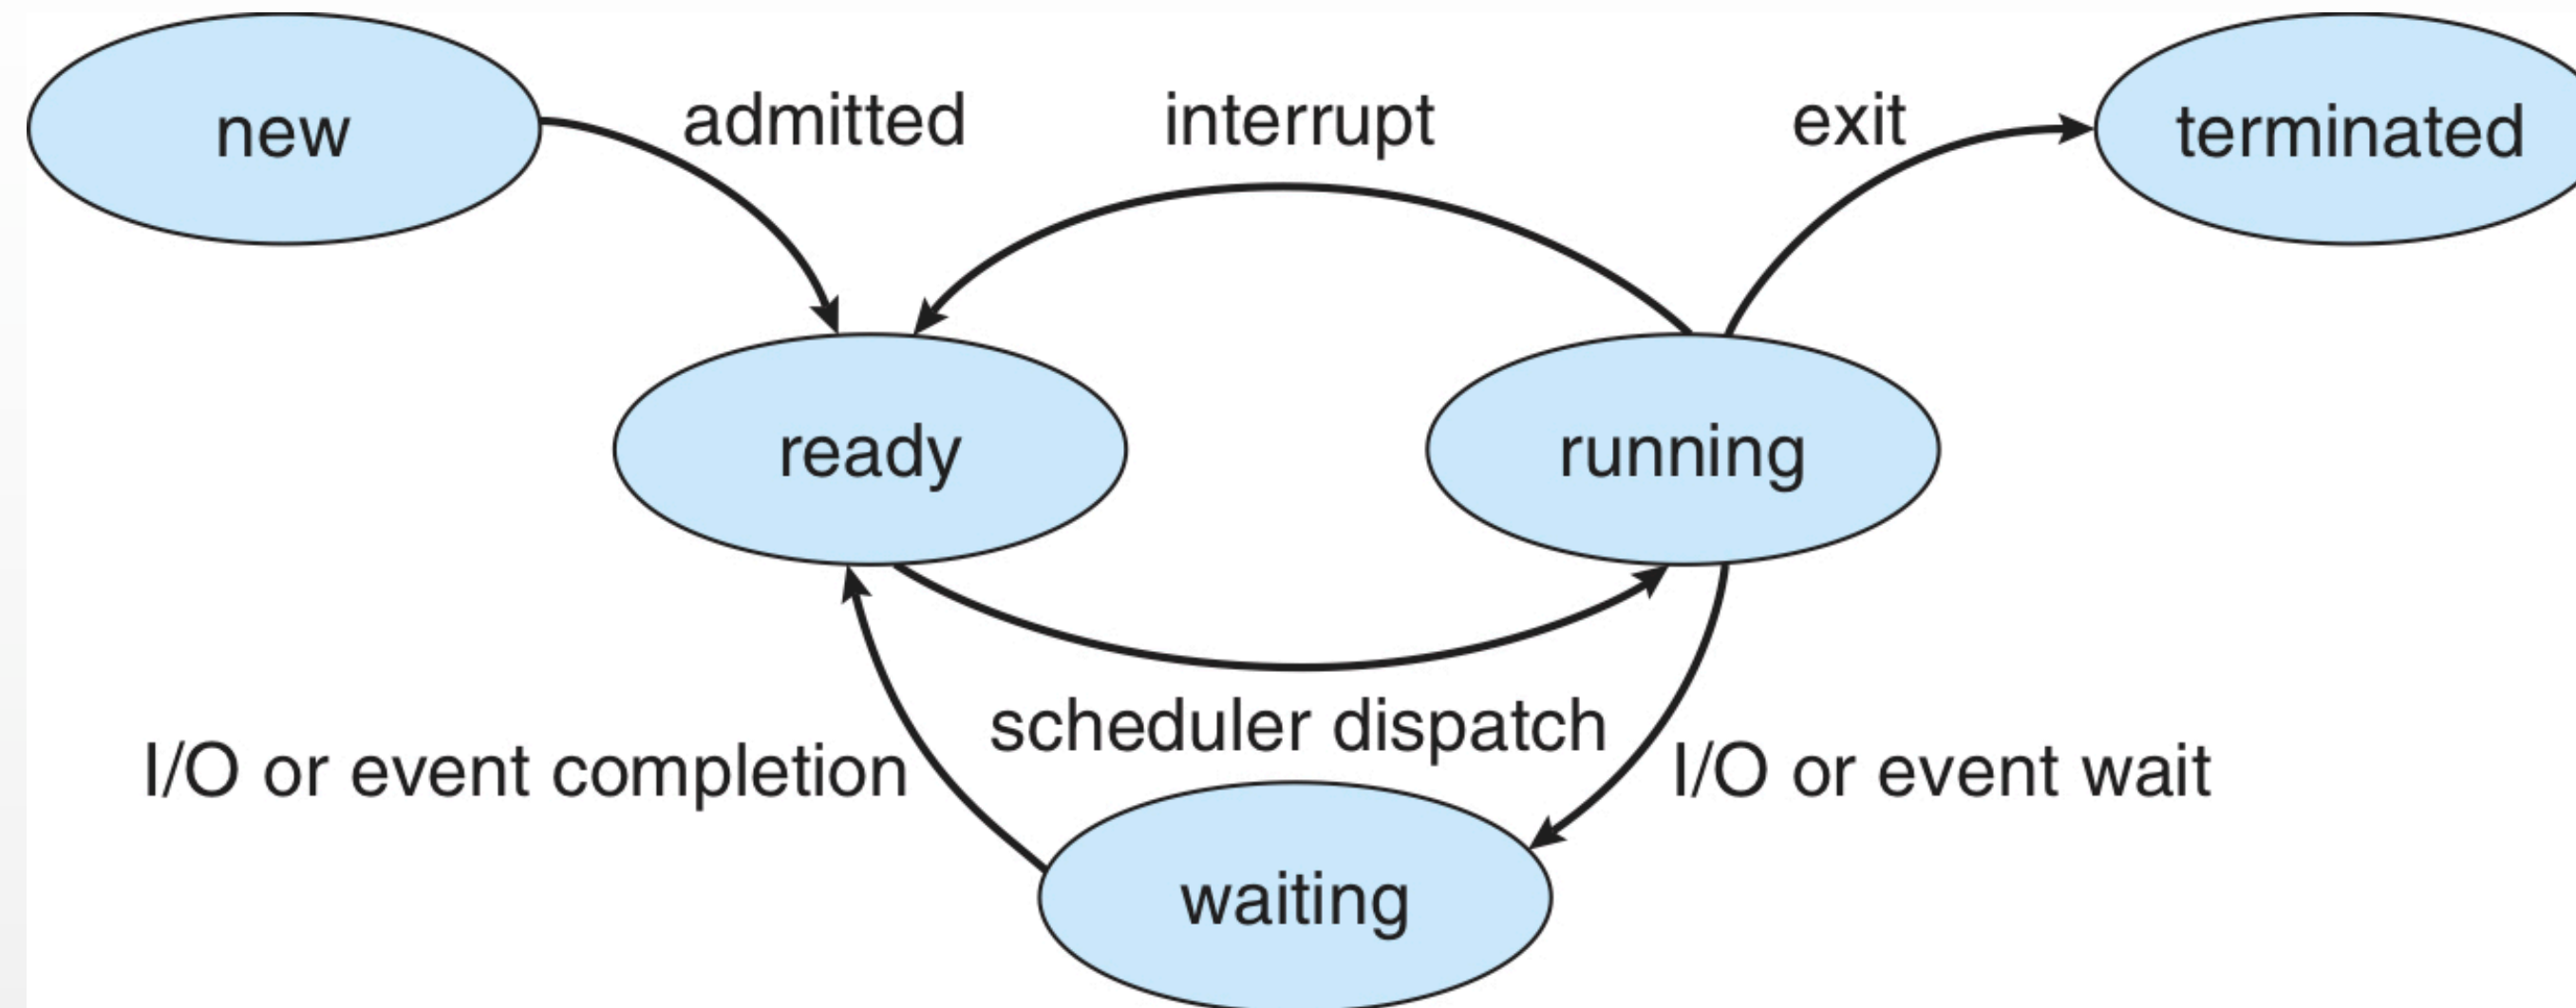
```
}
```

```
$ gcc memory_layout.c -o memory_layout
$ size memory_layout
```


Process Lifecycle



- Throughout its life in the system, a process changes states as follows

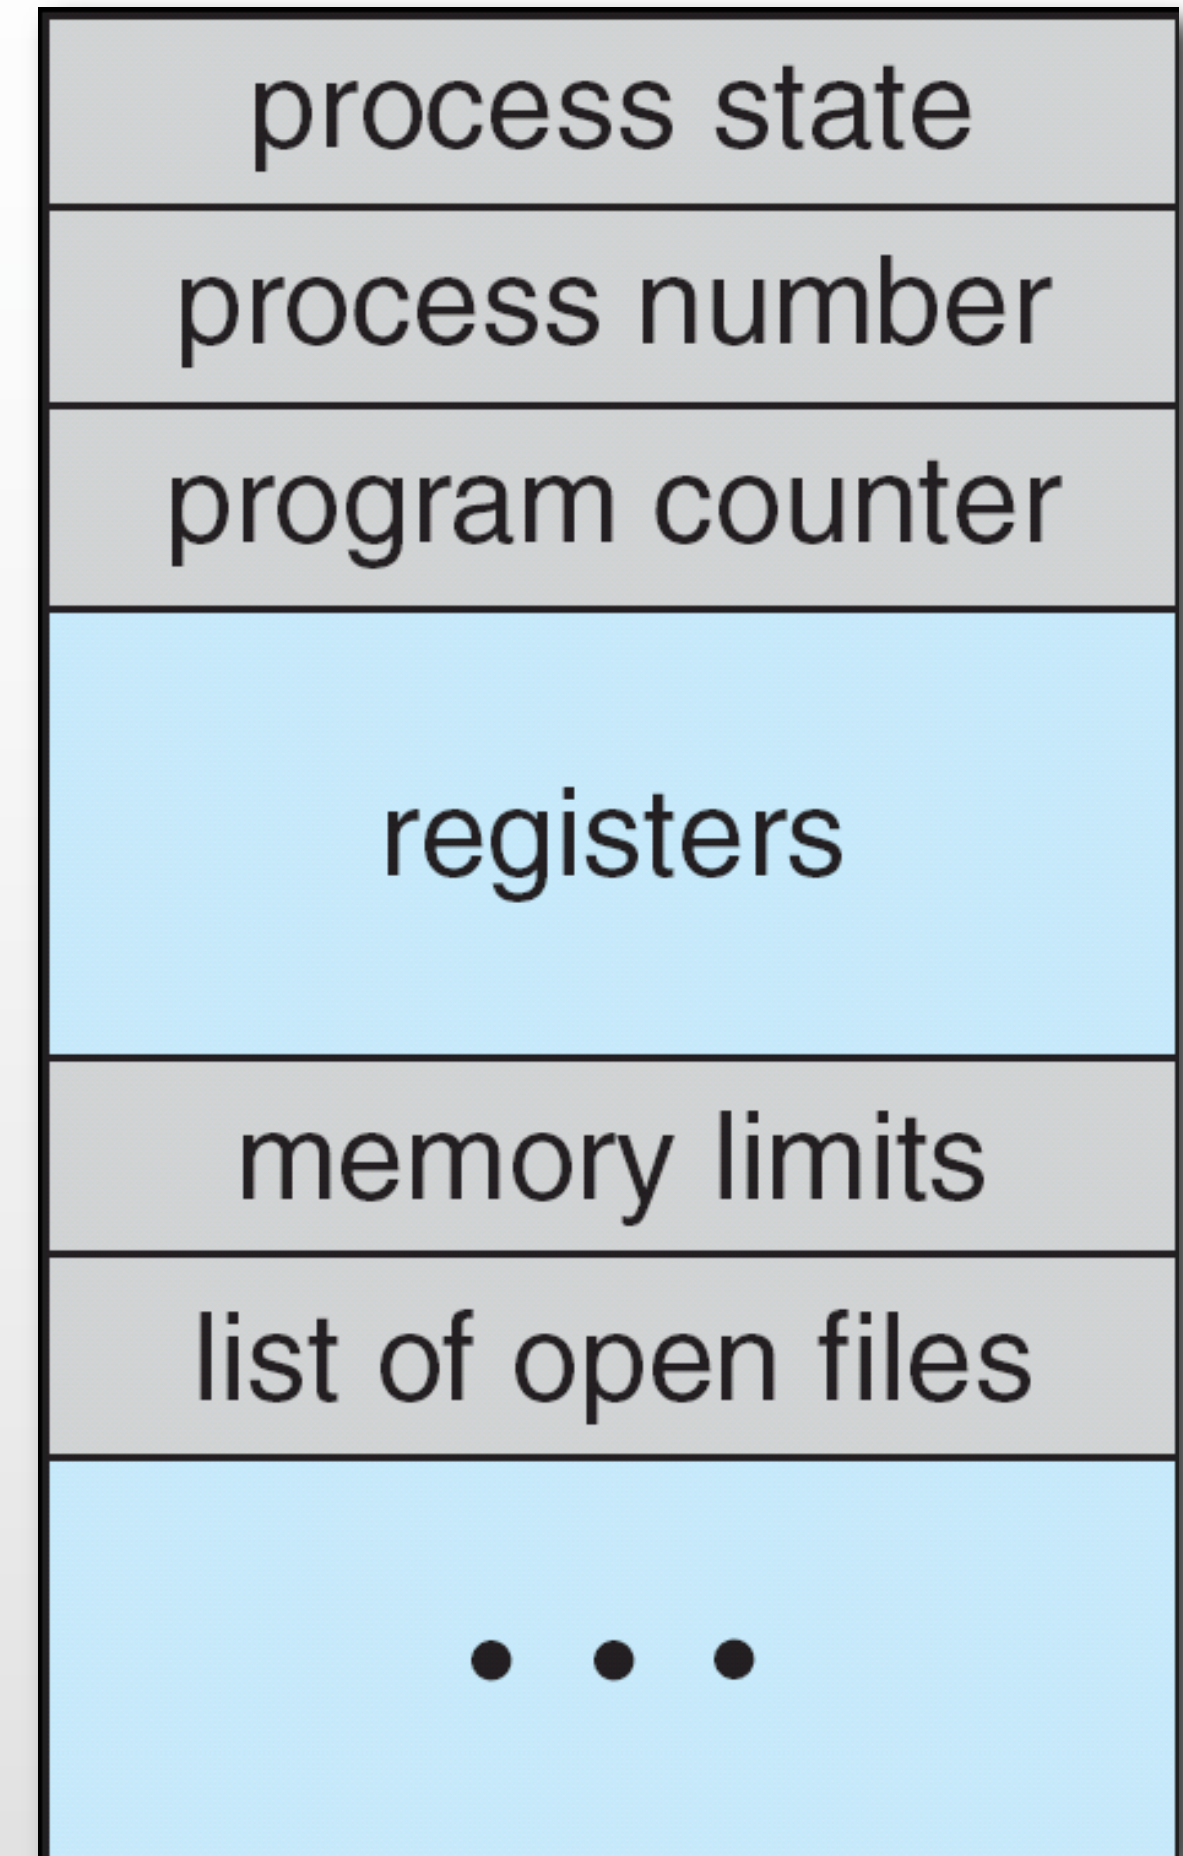


- **new**: The process is being created
- **ready**: The process is waiting to be assigned to a processor (CPU)
- **running**: Instructions are being executed
- **waiting**: The process is waiting for some event to occur
- **terminated**: The process has finished execution

Process Control Block (PCB)



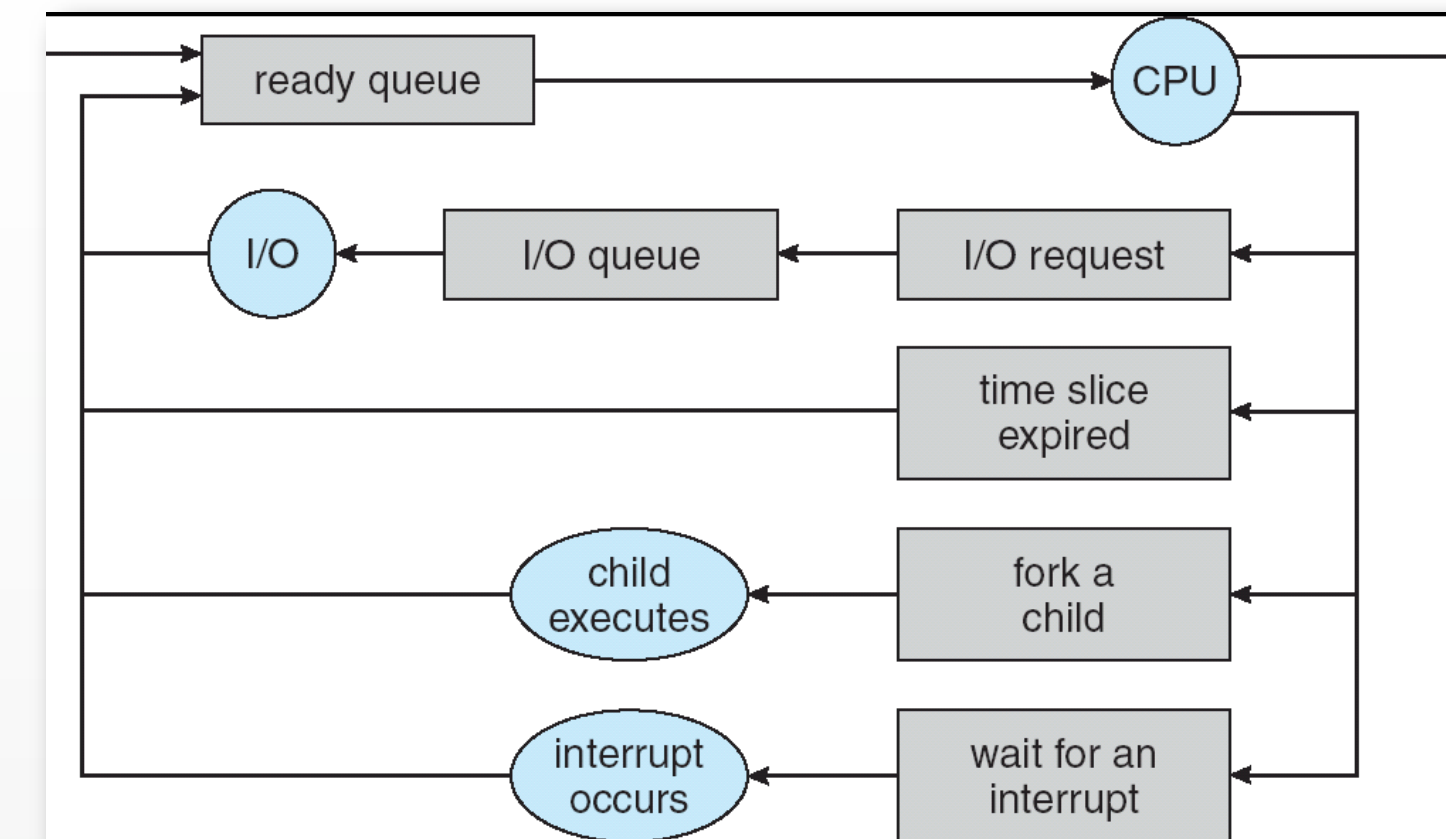
- Information associated with each process
 - Process state
 - Program counter
 - CPU registers
 - CPU scheduling information
 - Memory-management information
 - Accounting information
 - I/O status information
- OS maintains a list of PCBs for all processes; e.g.,
 - array
 - linked list
 - hybrid



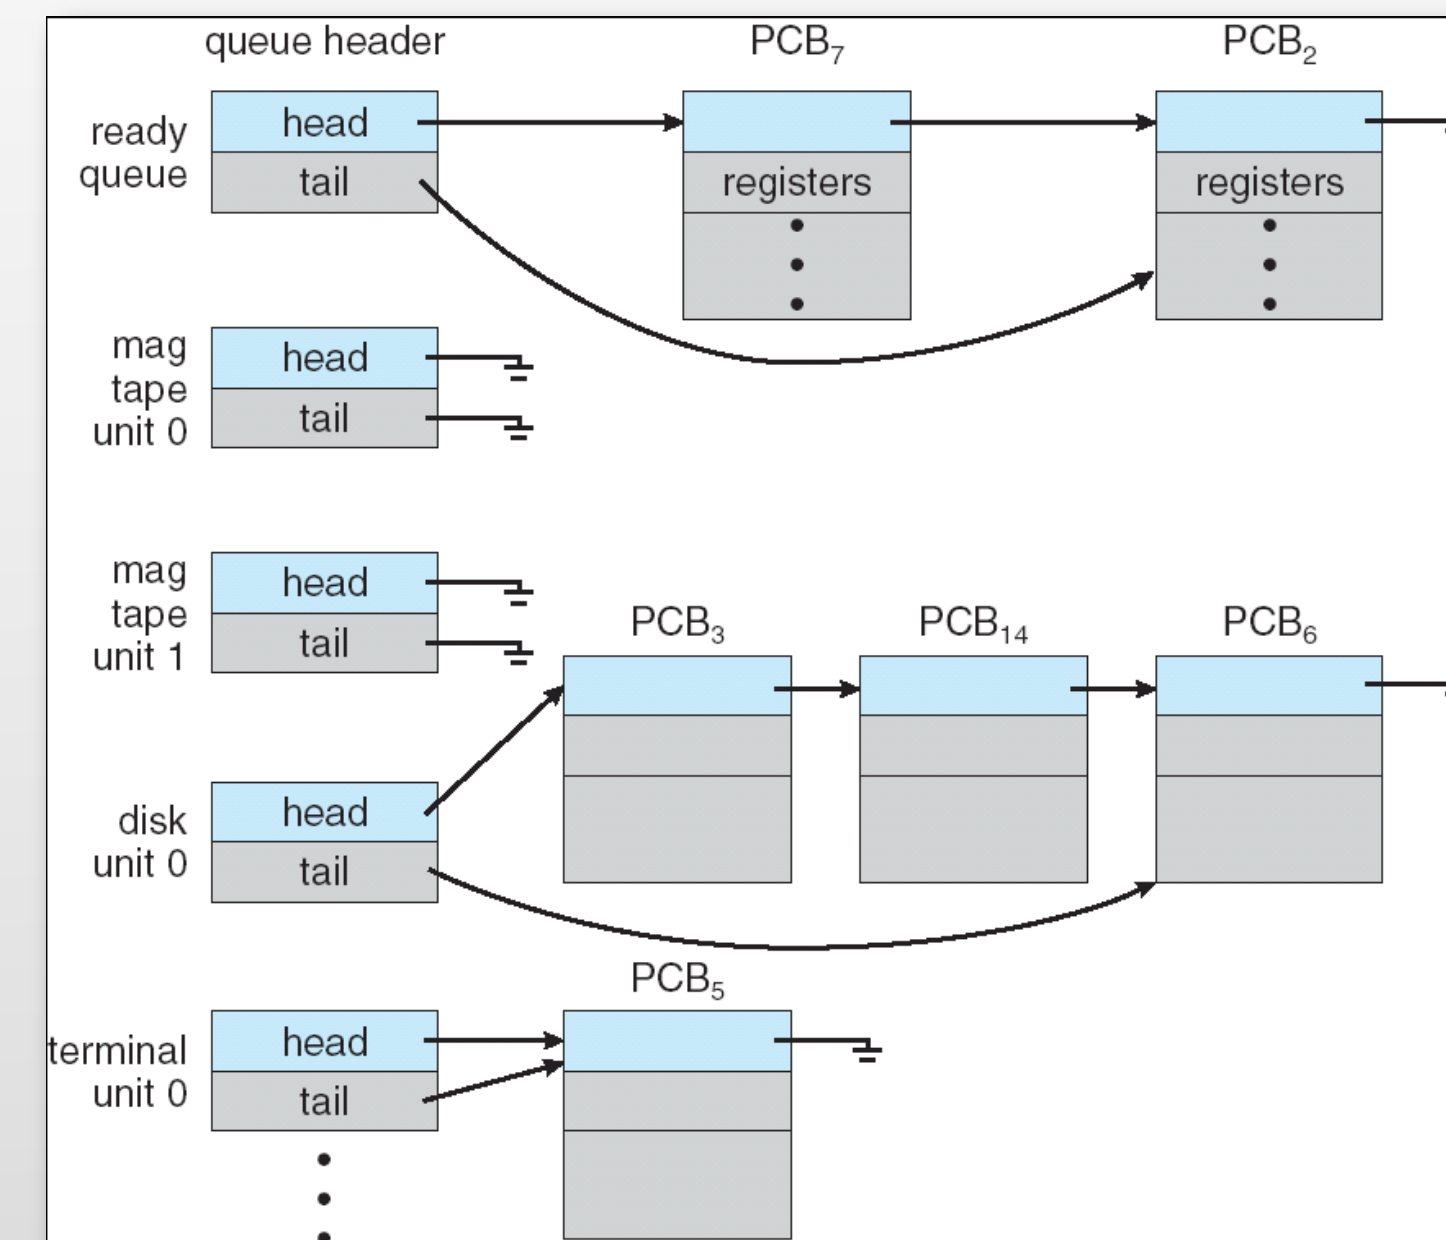
Process Scheduling Queues



- Processes migrate among various OS queues
 - actually: references to PCBs used
- Job queue**
 - set of all processes in the system
- Ready queue**
 - set of all processes residing in main memory, ready and waiting to execute
- Device queues**
 - set of processes waiting for an I/O device



some reasons
for suspending
a process



some device
queues

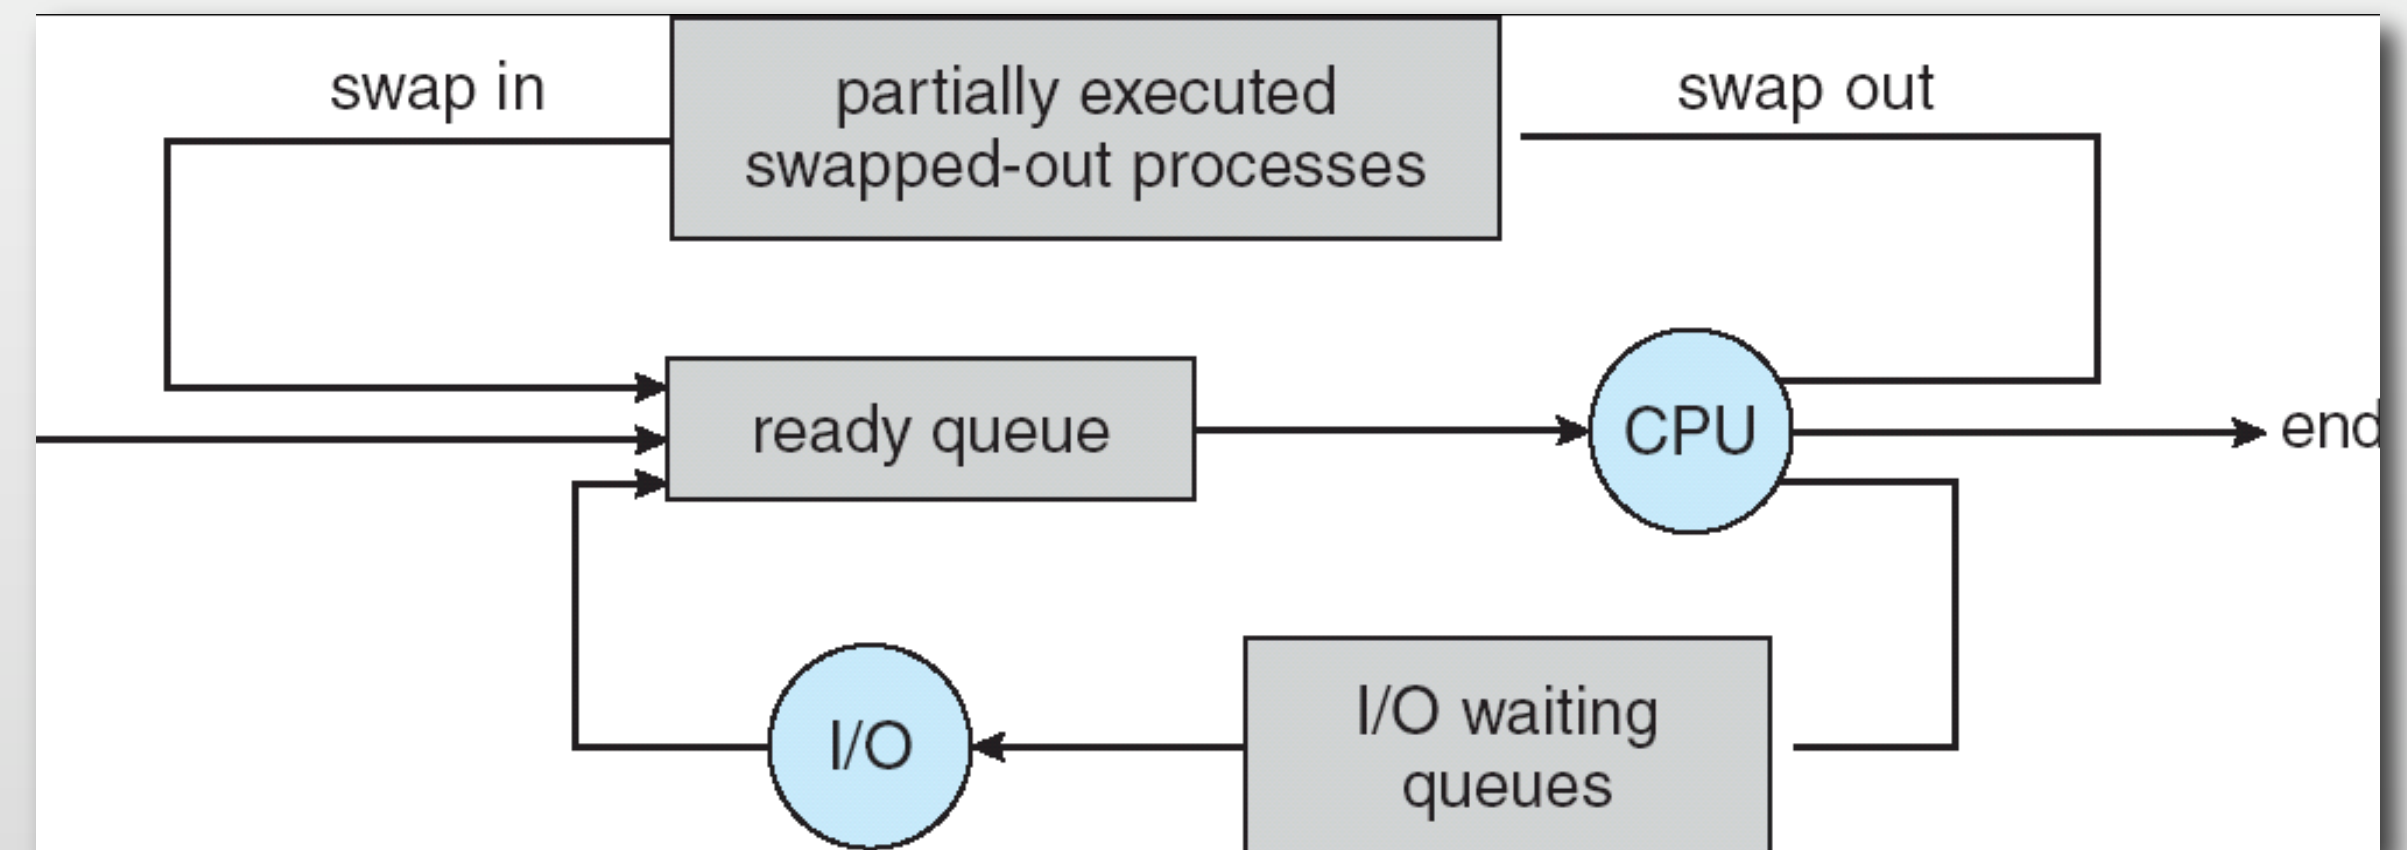
Schedulers



- **Long-term scheduler (or job scheduler)**
 - selects which processes should be brought into the ready queue
- **Short-term scheduler (or CPU scheduler)**
 - selects which process should be executed next and allocates CPU
- **Medium-term scheduler**
 - for swapping processes out and in
 - when inactive, consuming too much memory, etc.

Analogy: taking several courses

- ✓ how many can you take in a semester?
- ✓ which one should you attend now?
- ✓ which one can you put away for a week if you are swamped with work?



Schedulers (Cont.)

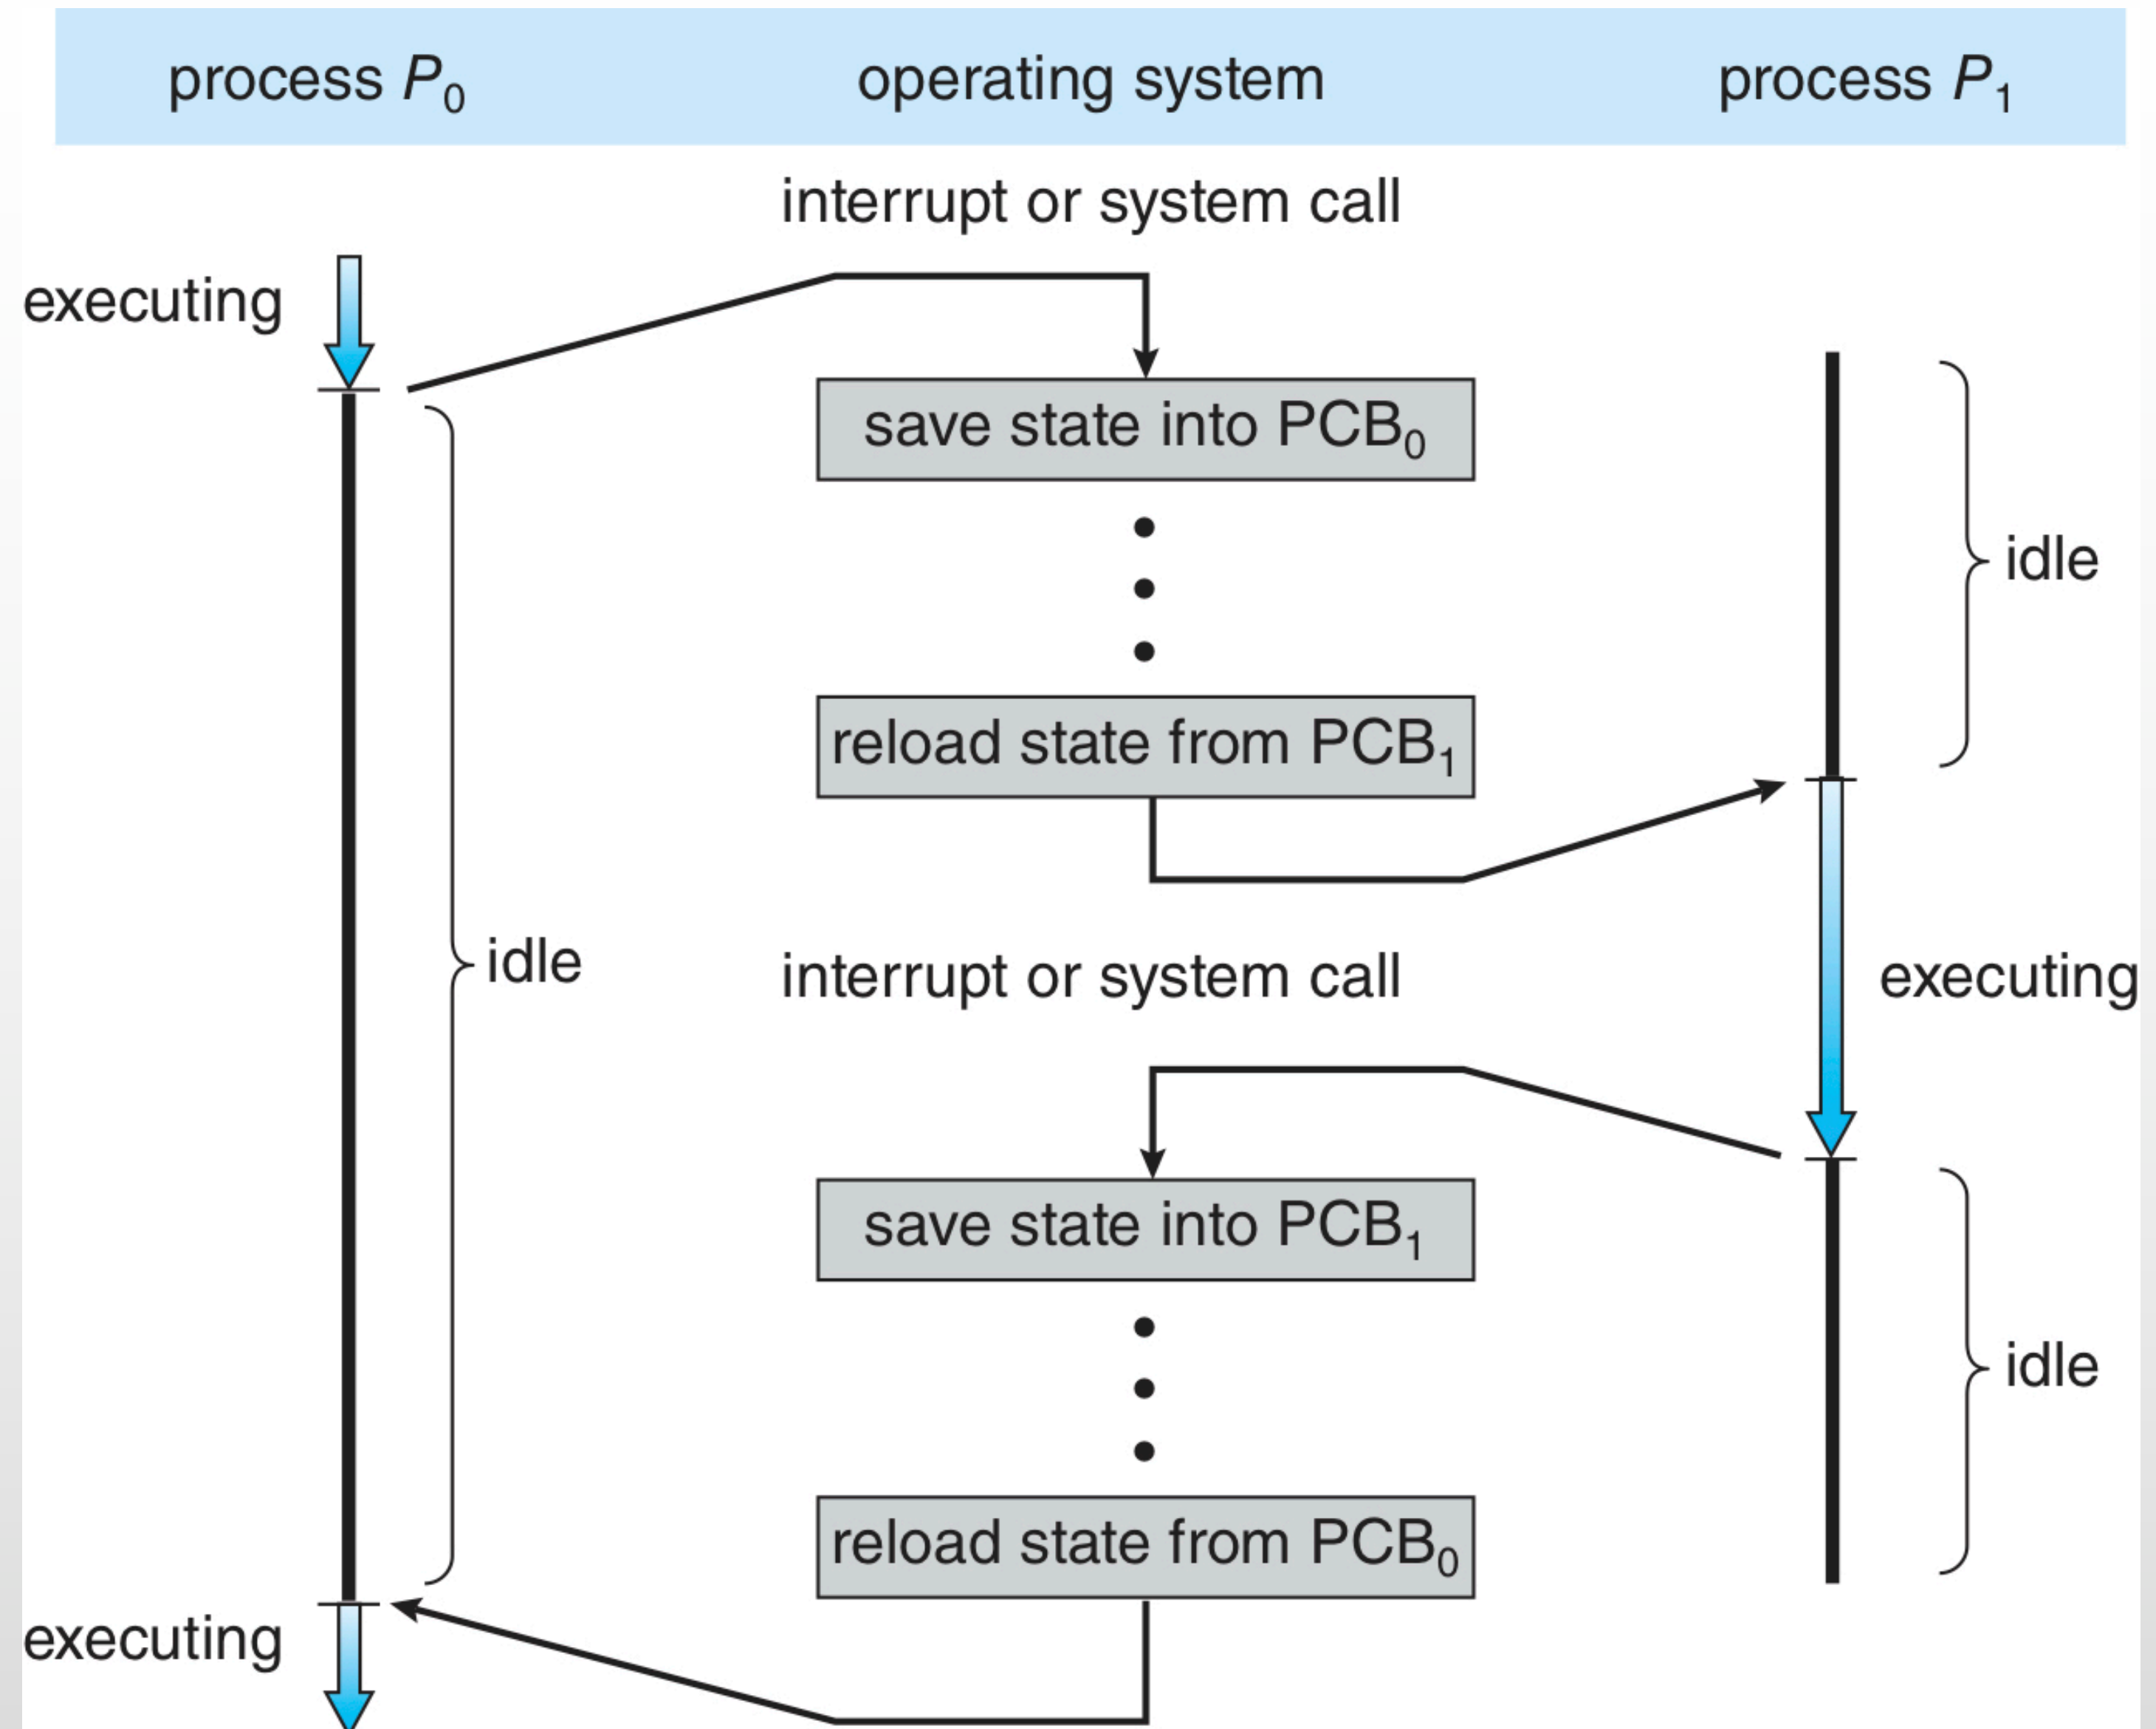


- Short-term scheduler is invoked very frequently (milliseconds)
 - ➡ must be fast
- Long-term scheduler is invoked very infrequently (seconds, minutes)
 - ➡ does not need to be very fast
- The long-term scheduler controls the degree of **multiprogramming**
 - long-term scheduler must take many things into account when admitting programs; that includes type of the process:
 - **I/O-bound process**
 - spends more time doing I/O than computations, many short CPU bursts
 - **CPU-bound process**
 - spends more time doing computations; few very long CPU bursts

Context Switch



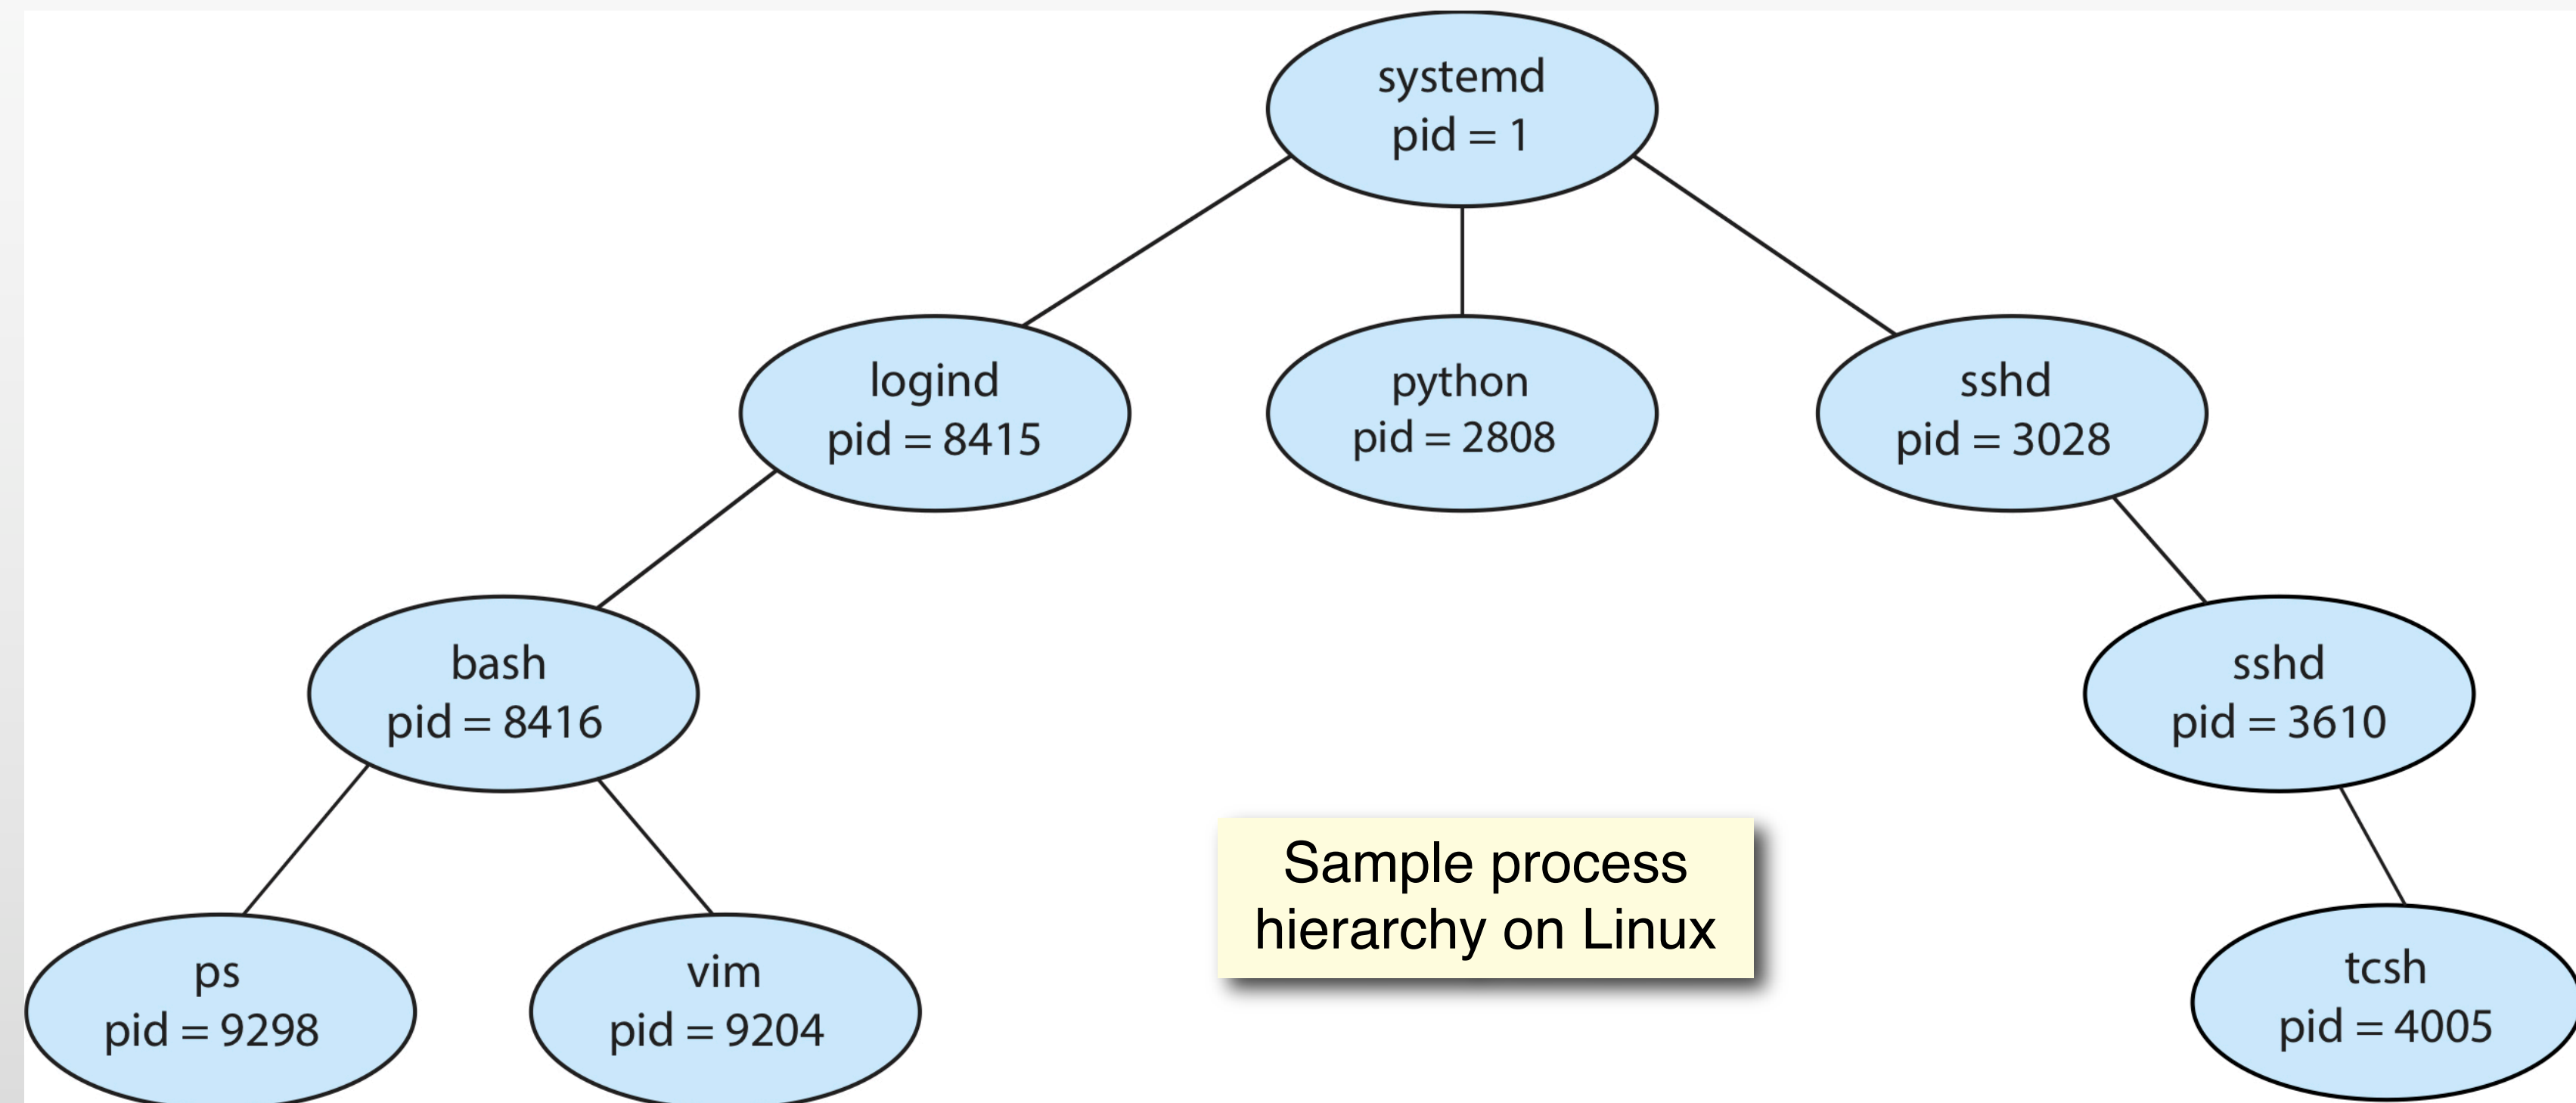
- When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process
- **Context-switch** time is an overhead
 - the system is not executing programs
- Time depends on hardware support



Process Creation



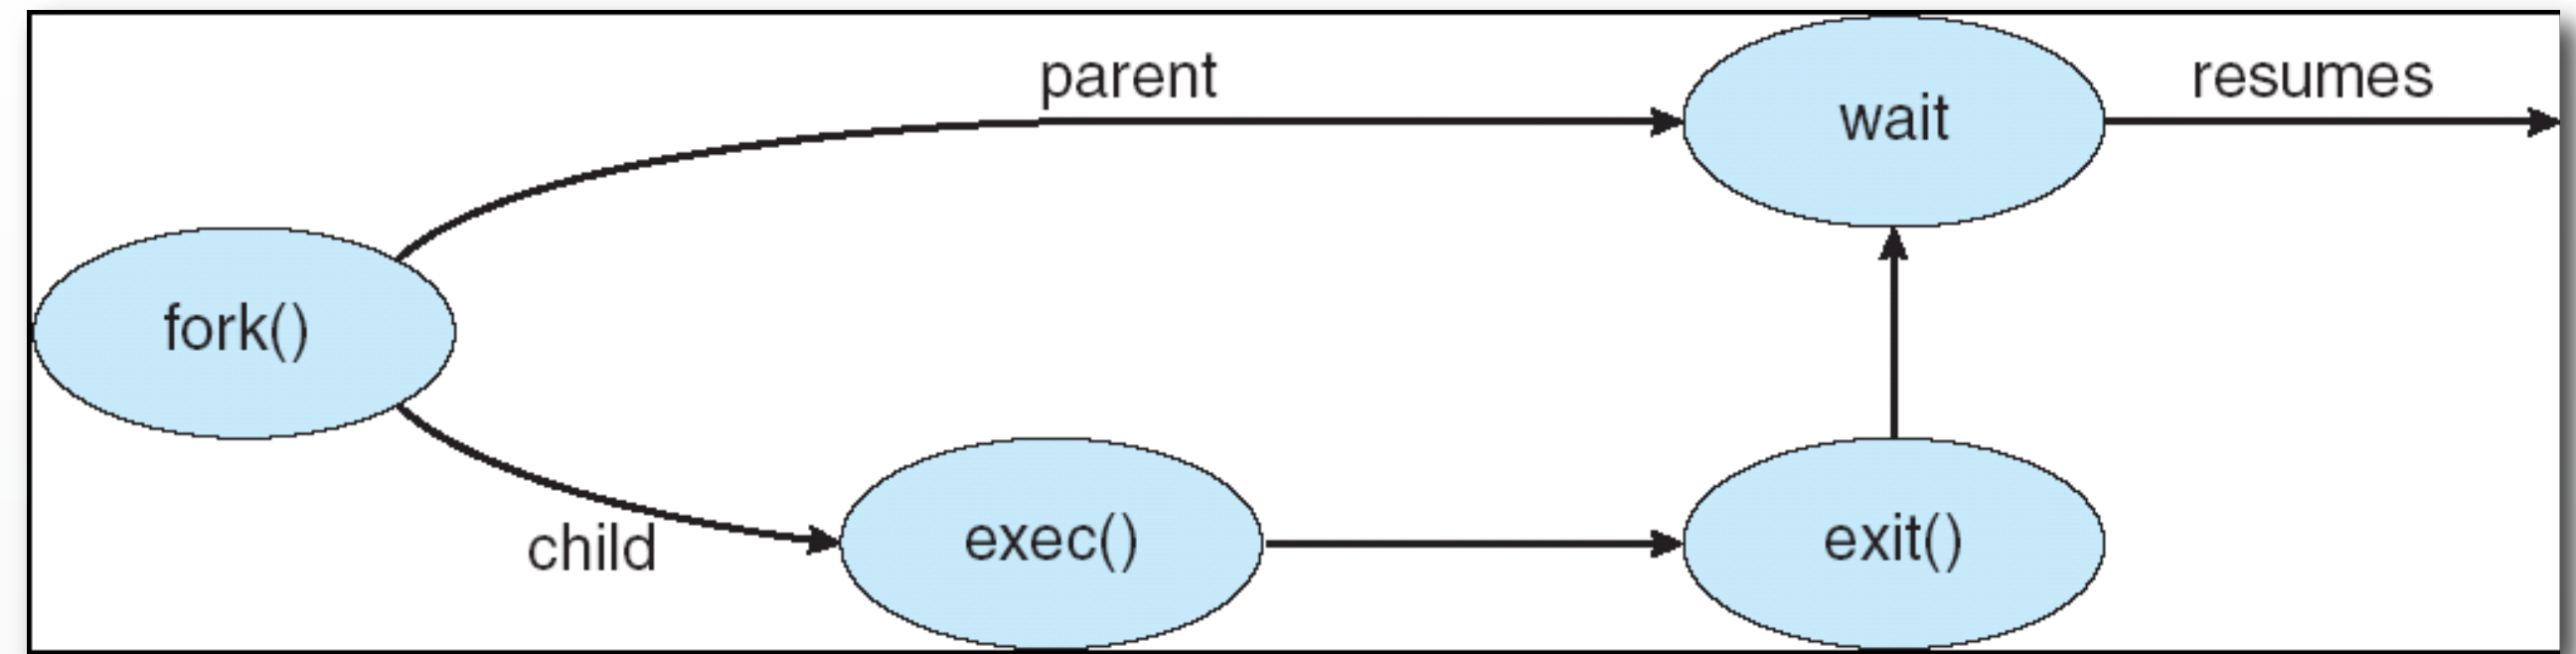
- **Parent process** creates **children processes**, which, in turn create other processes, forming a hierarchy of processes
- Resource sharing options
 - Parent and children may share all resources
 - Children may share subset of parent's resources
 - Parent and children may share no resources
- Execution options
 - Parent and children may execute concurrently
 - Parent may wait until children terminate



Process Creation



- Address space choices
 - Child may be a duplicate of the parent
 - Child has another program loaded into it
- UNIX examples
 - **fork** system call creates new process
 - **exec** system call is used after a fork to replace the process' memory space with a new program



```
/* fork a child process */
pid = fork();
if (pid < 0)
{
    /* error occurred */
}
else if (pid == 0)
{
    /* child process */
}
else /* pid > 0 */
{
    /* parent process */
}
```

create_process.c

Process Creation in POSIX



POSIX standard API

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
int main()
{
    pid_t pid;
    /* fork a child process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed\n");
        exit(EXIT_FAILURE);
    }
    else if (pid == 0) { /* child process */
        printf("I am the child %d\n", getpid());
        execvp("/bin/ls", "ls", NULL); // run /bin/ls in the child
        exit(EXIT_FAILURE); // executed only if execvp() fails
    }
    else { /* parent process; pid is the pid of the child */
        /* parent will wait for the child to complete */
        printf("I am the parent %d\n", getpid());
        waitpid(pid, NULL, 1); // wait for the child with the given pid
        printf("Child Complete\n");
        exit(EXIT_SUCCESS);
    }
}
```

Process Termination



- Process executes last statement and asks the operating system to delete it (**exit**)
 - output data from child to parent (via **wait**)
 - process' resources are **deallocated** by operating system
- Parent may terminate execution of children processes (**abort**)
 - for example, when:
 - child has exceeded allocated resources
 - task assigned to child is no longer required
 - if parent is exiting
 - some operating system do not allow a child to continue if its parent terminates
 - **cascading termination** - if all children along with their children, and so on, are being terminated
- Terminated process may enter so-called a “zombie” state
 - while the operating system is performing termination-related cleanup

Process Management



```
$ man ps
$
$ man kill
$
$ man top
$ top
<exit by pressing Q>
$
$ sudo apt install htop
$ man htop
$ htop
<exit by pressing Q>
$
$ man procfs
$ cat /proc/12345/maps
$ cat /proc/12345/sched
```

htop is a better version of top

accessing process
information using process
pseudo file system (content
of /proc) – only on Linux!

Process Management (cont.)



```
$ ps -elf | grep proc_
0 S aj      29836  2122  0 80  0 -  620 do_wai 16:14 pts/0    00:00:00 ./proc_hier
0 S aj      30333 29860  0 80  0 - 4415 pipe_w 16:17 pts/2    00:00:00 grep --color=auto proc_


$ ps -elf | grep iam
0 R aj      29837 29836 30 80  0 -  620 -      16:14 pts/0    00:00:04 iam 1.1
0 R aj      29838 29836 30 80  0 -  620 -      16:14 pts/0    00:00:04 iam 1.2
0 R aj      29839 29837 30 80  0 -  620 -      16:14 pts/0    00:00:04 iam 1.1.1
0 R aj      29840 29837 31 80  0 -  620 -      16:14 pts/0    00:00:04 iam 1.1.2
0 R aj      29841 29838 31 80  0 -  620 -      16:14 pts/0    00:00:04 iam 1.2.1
0 R aj      29842 29838 31 80  0 -  620 -      16:14 pts/0    00:00:04 iam 1.2.2
0 S aj      29888 29860  0 80  0 - 4382 pipe_w 16:14 pts/2    00:00:00 grep --color=auto iam

$ kill -9 29839 29840
$ ps -elf | grep iam
0 R aj      29837 29836 31 80  0 -  620 -      16:14 pts/0    00:00:16 iam 1.1
0 R aj      29838 29836 31 80  0 -  620 -      16:14 pts/0    00:00:16 iam 1.2
0 Z aj      29839 29837 29 80  0 -    0 -      16:14 pts/0    00:00:15 [iam] <defunct>
0 Z aj      29840 29837 29 80  0 -    0 -      16:14 pts/0    00:00:15 [iam] <defunct>
0 R aj      29841 29838 31 80  0 -  620 -      16:14 pts/0    00:00:16 iam 1.2.1
0 R aj      29842 29838 31 80  0 -  620 -      16:14 pts/0    00:00:16 iam 1.2.2
0 S aj      29994 29860  0 80  0 - 4382 pipe_w 16:15 pts/2    00:00:00 grep --color=auto iam

$ kill -9 29837
$ ps -elf | grep iam
0 R aj      29838    913 45 80  0 -  620 -      16:14 pts/0    00:03:51 iam 1.2
0 R aj      29841 29838 45 80  0 -  620 -      16:14 pts/0    00:03:52 iam 1.2.1
0 R aj      29842 29838 46 80  0 -  620 -      16:14 pts/0    00:03:52 iam 1.2.2
0 S aj      31168 29860  0 80  0 - 4382 pipe_w 16:22 pts/2    00:00:00 grep --color=auto iam
```

zombies!

zombies gone after
the parent terminated



COMP362 Operating Systems
Prof. AJ Bieszczad