



# Chapter 4: Inter-Process Communication (IPC)

**COMP362 Operating Systems**  
**Prof. AJ Bieszczad**

# Outline: Processes



- Cooperating Processes
- Main IPC Models
- Client-Server Systems
- Shared Memory-based IPC
- Kernel-facilitated IPC
- Direct Communication
- Indirect Communication
- Management of IPC Resources

# Interprocess Communication (IPC)

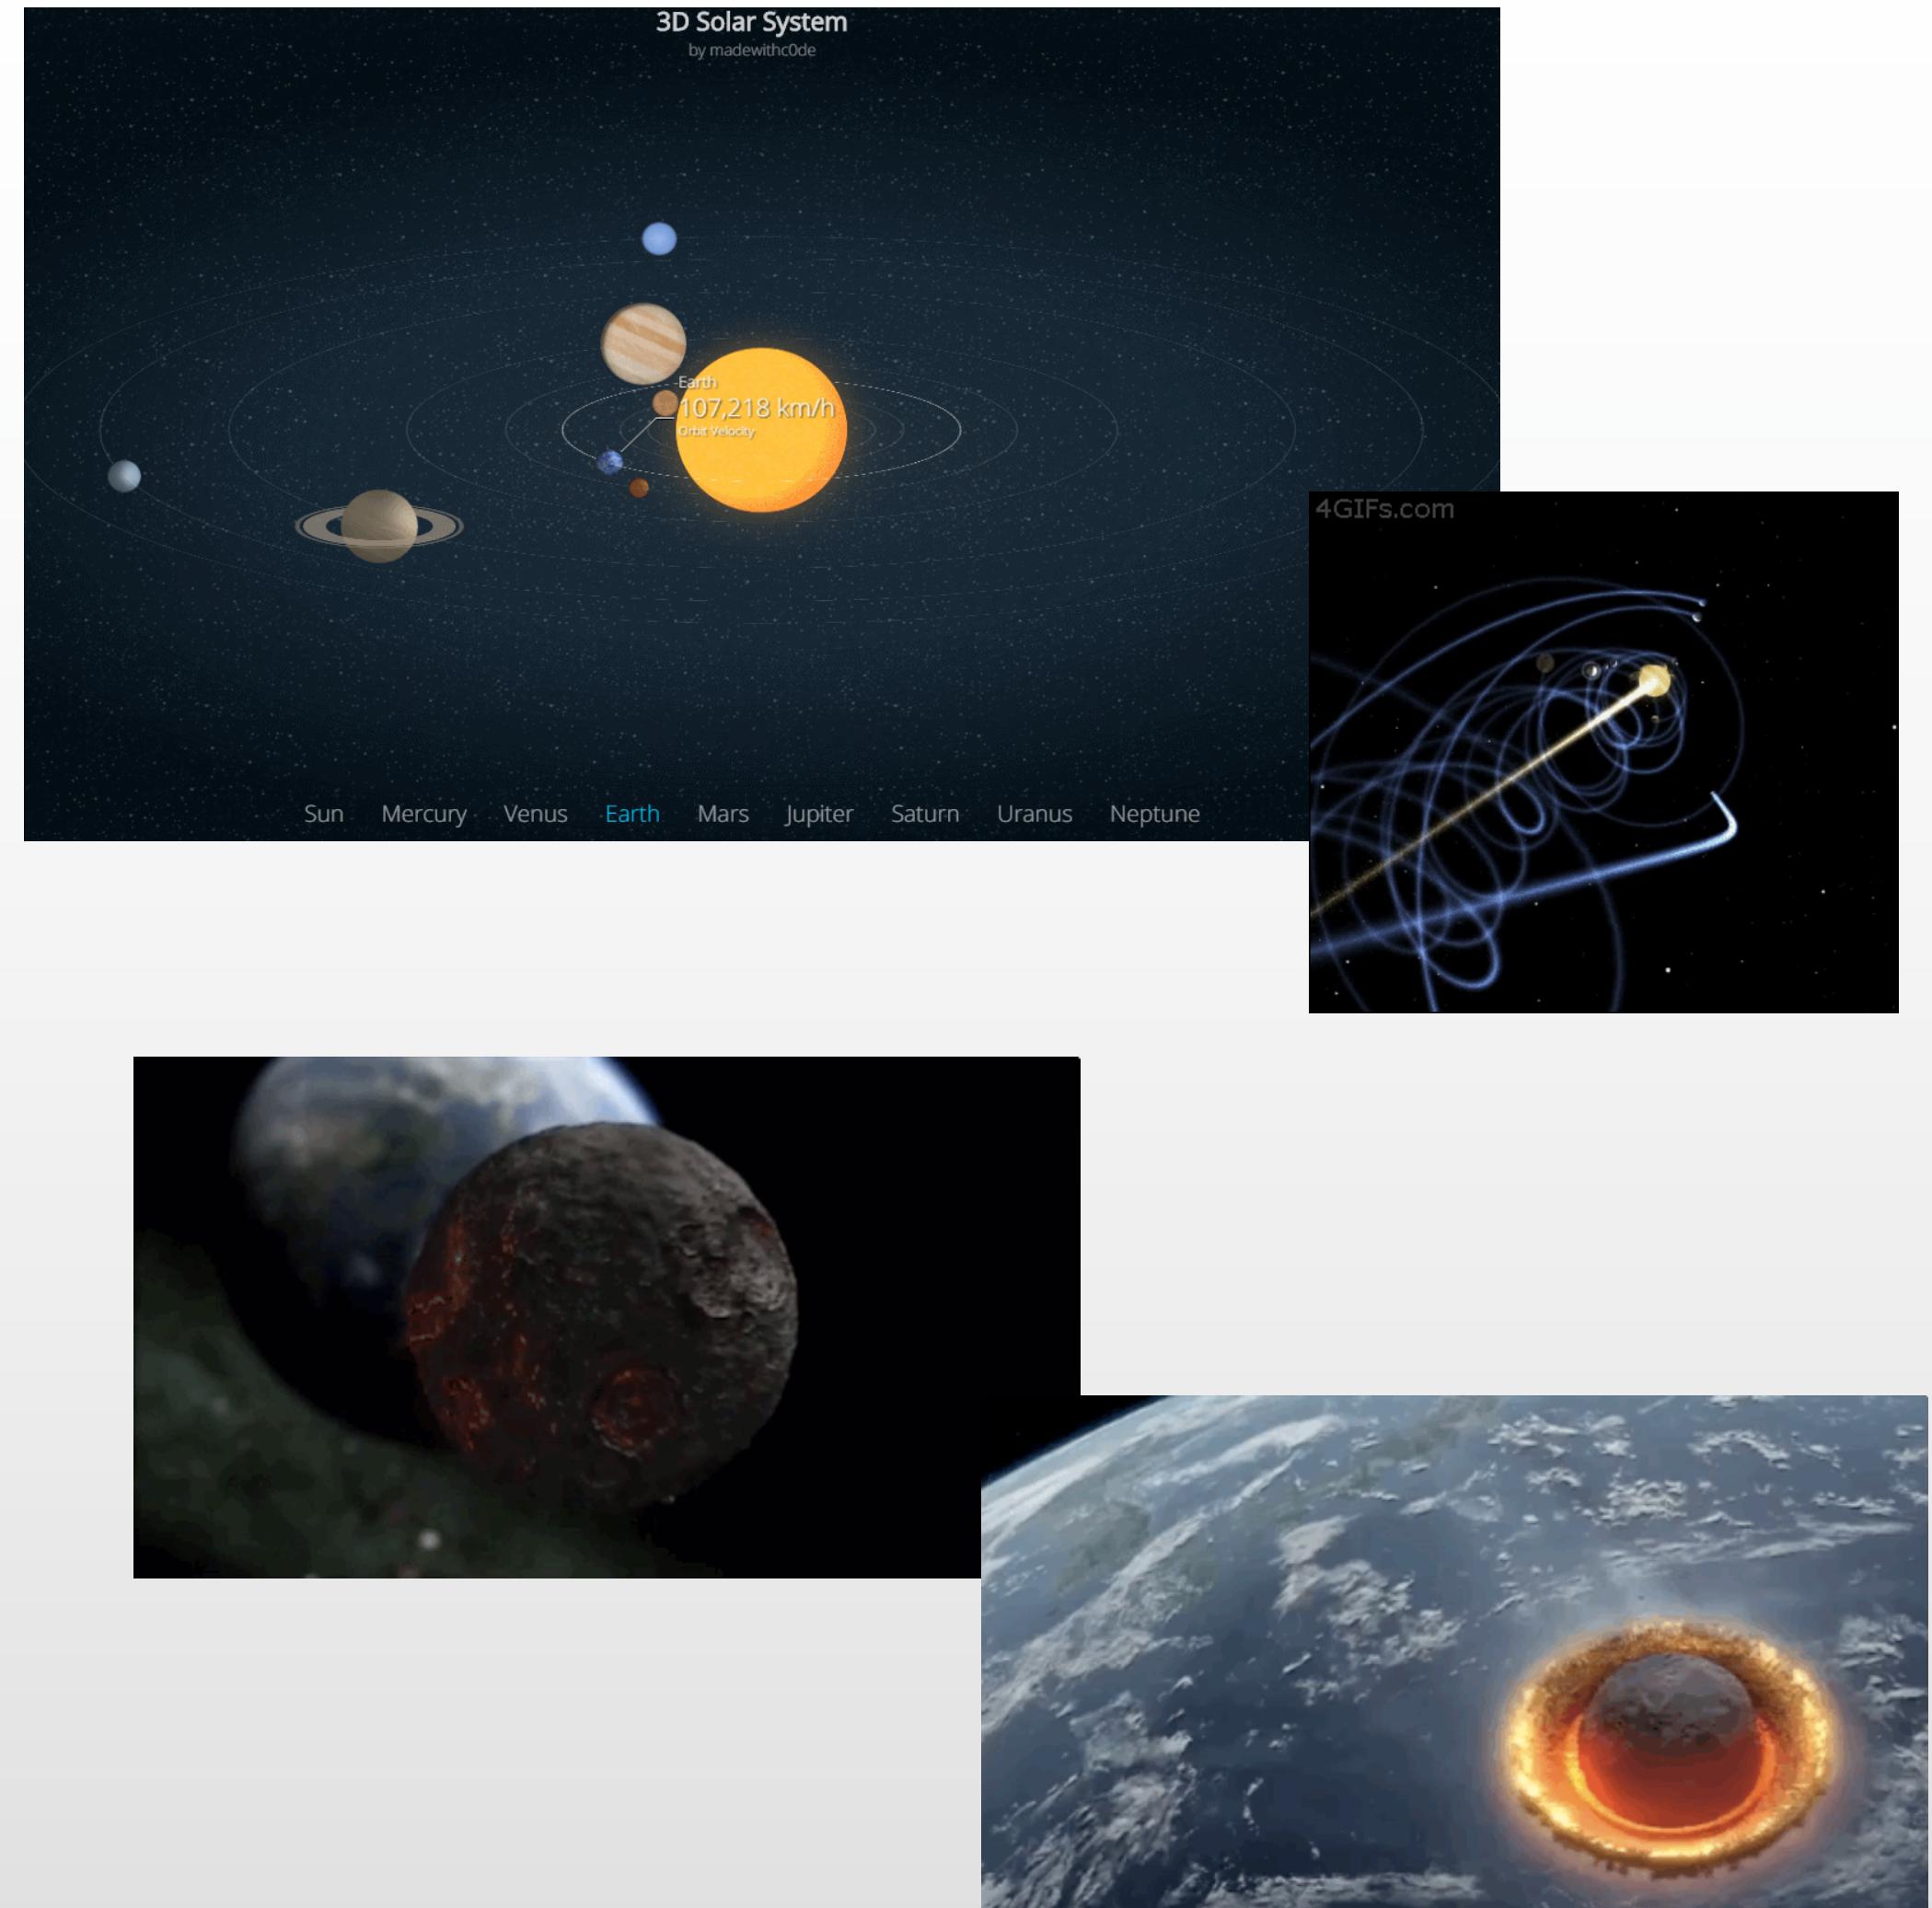


- We may want to use interprocess communication for:
  - **Information sharing**
    - files, shared memory
  - **Computation speedup**
    - e.g., on multi-core processors
  - **Modularity**
    - software engineering
  - **Convenience**
    - e.g., signaling events in parallel activities, in workflows, etc.
      - for example, printing, compiling, building, etc. -- user may stayed informed while doing something else

# Example: Dynamic Systems

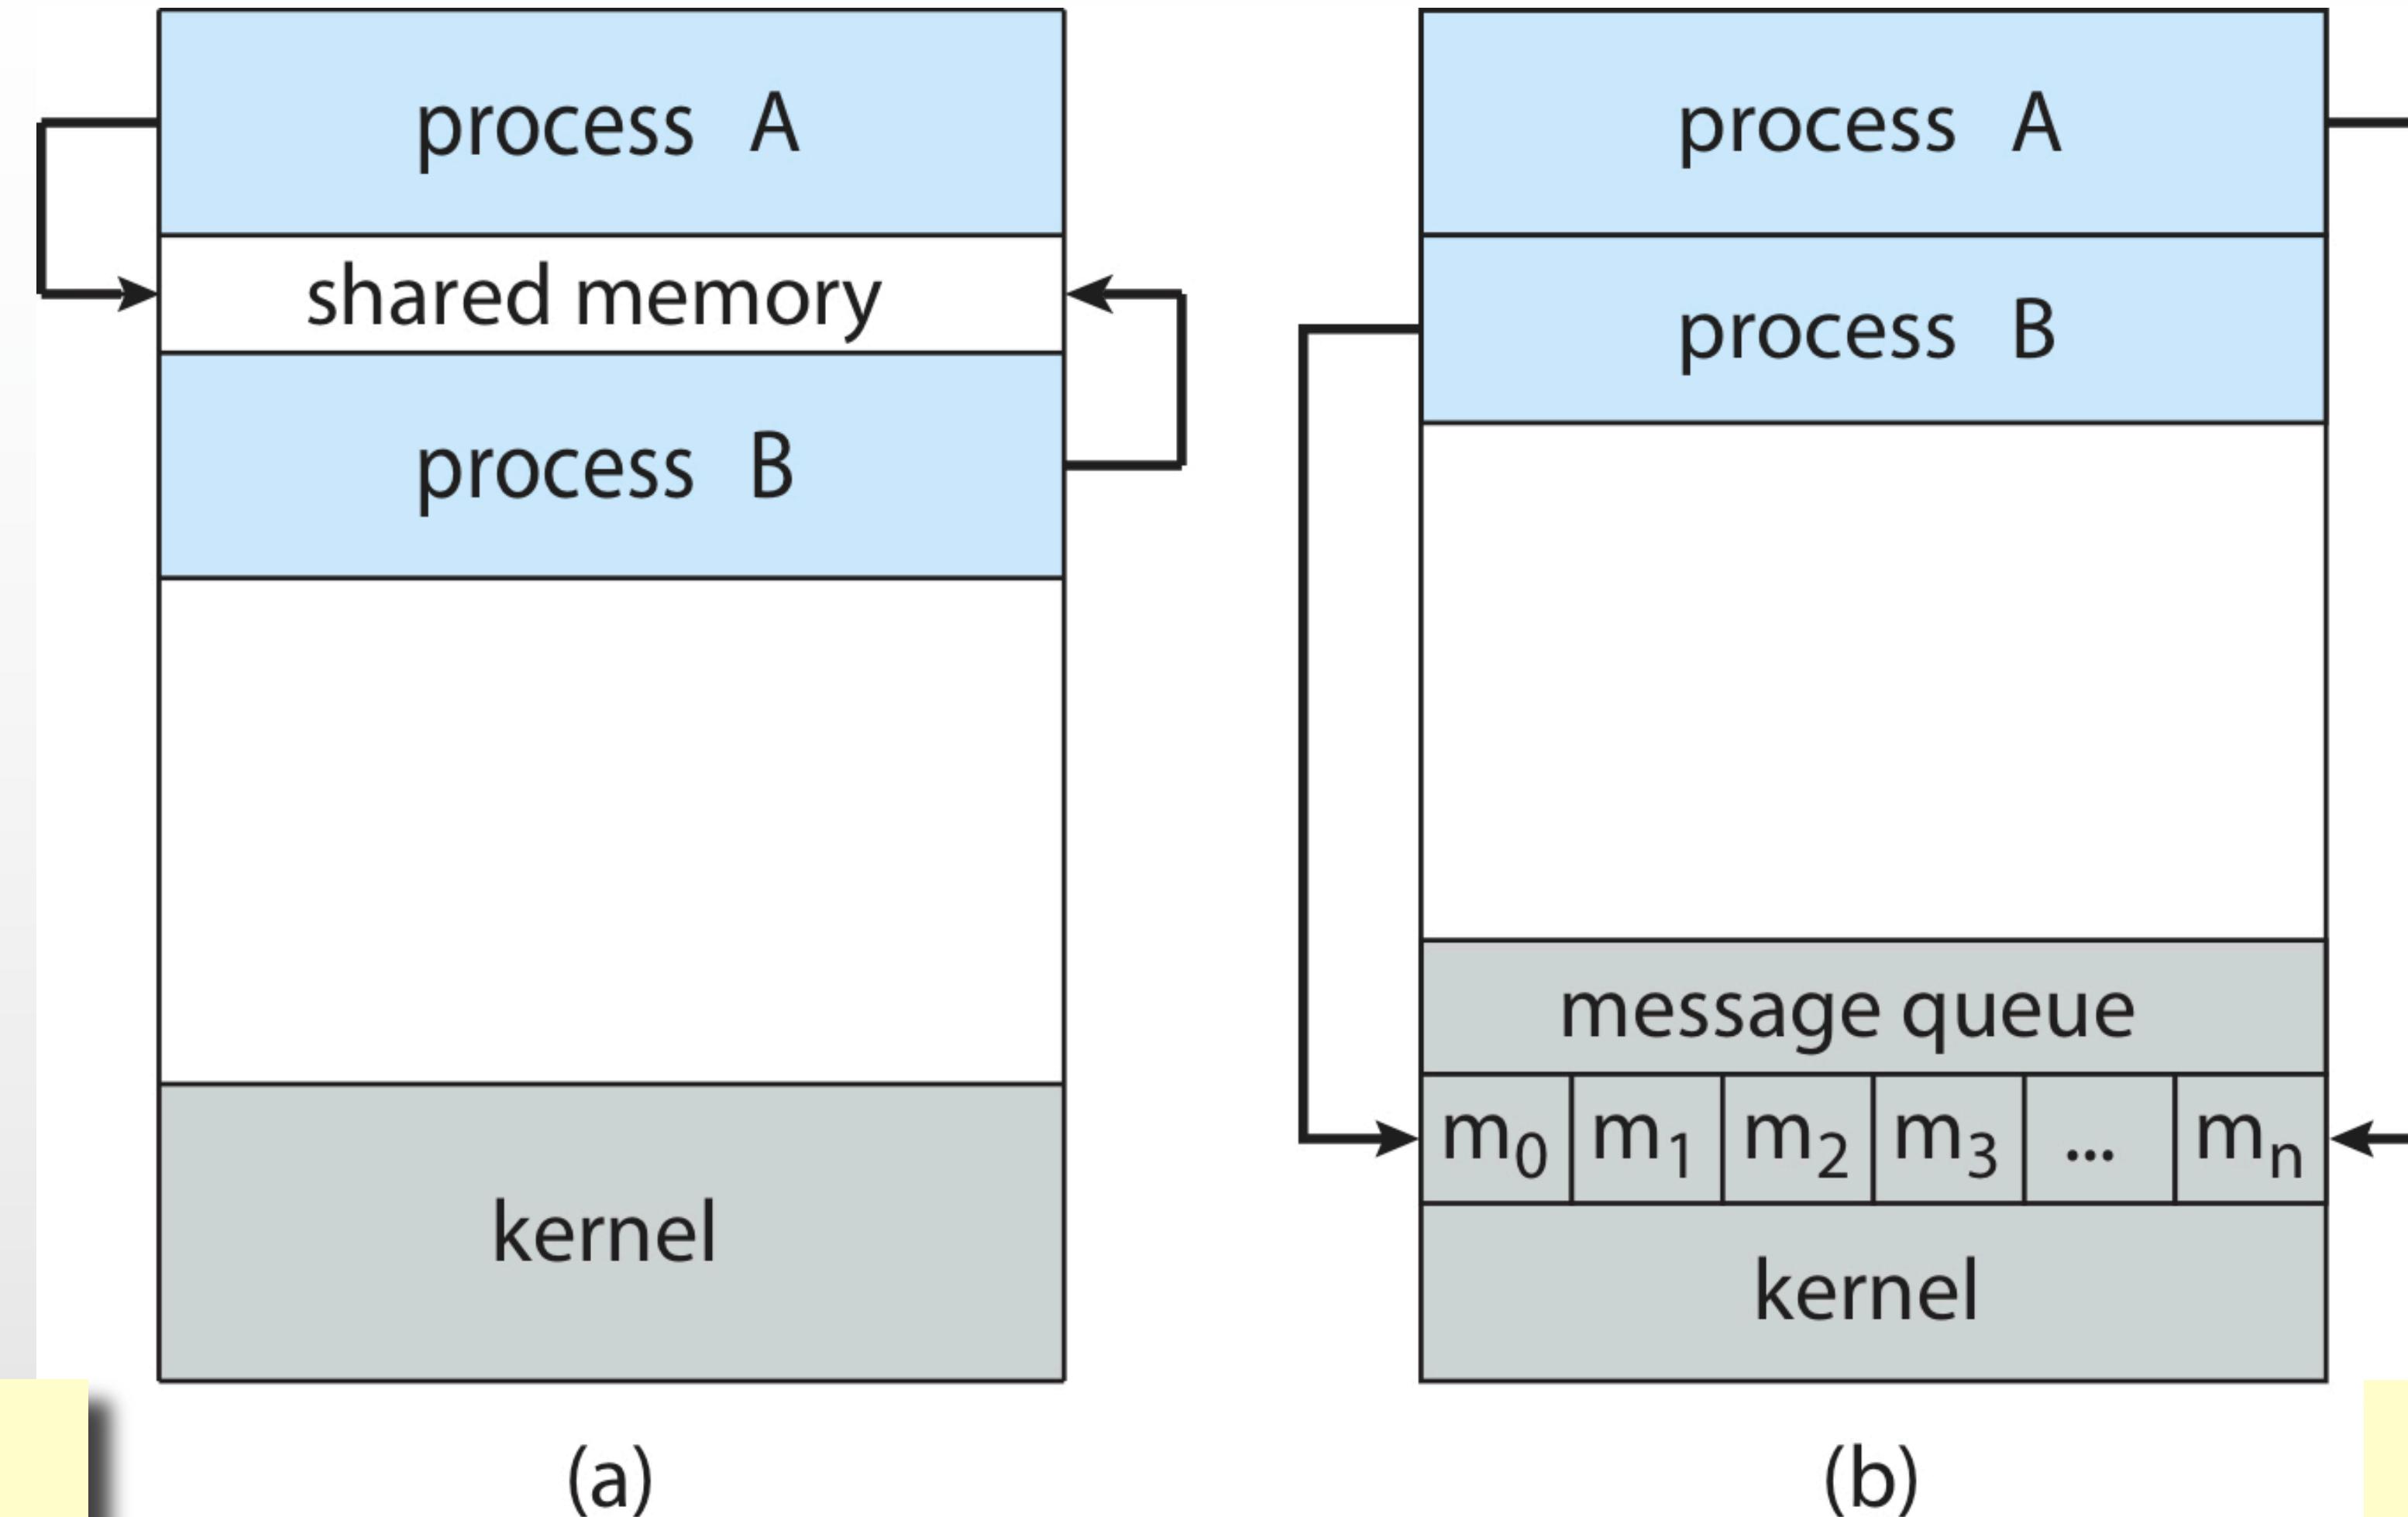


- Multiple processes collaborate to solve a certain problem
- For example, would we survive if a foreign celestial body enters the system and hits one of the native solar system objects?
  - the gravitational attraction between every pair of celestial bodies is proportional to the distance between them, and vice versa
  - therefore, changing the position of any object disturbs the whole system
- A process may represent a celestial body
  - it receives positions from other objects, calculates the gravitational forces, then its new position, and then sends it back to the others
  - the information is exchanged as long as the position of any object keeps changing
  - the system is considered stable when no object changes the position



This is a very simplified model!

# Two Main IPC Models



No kernel involvement after the shared space is allocated

(a)

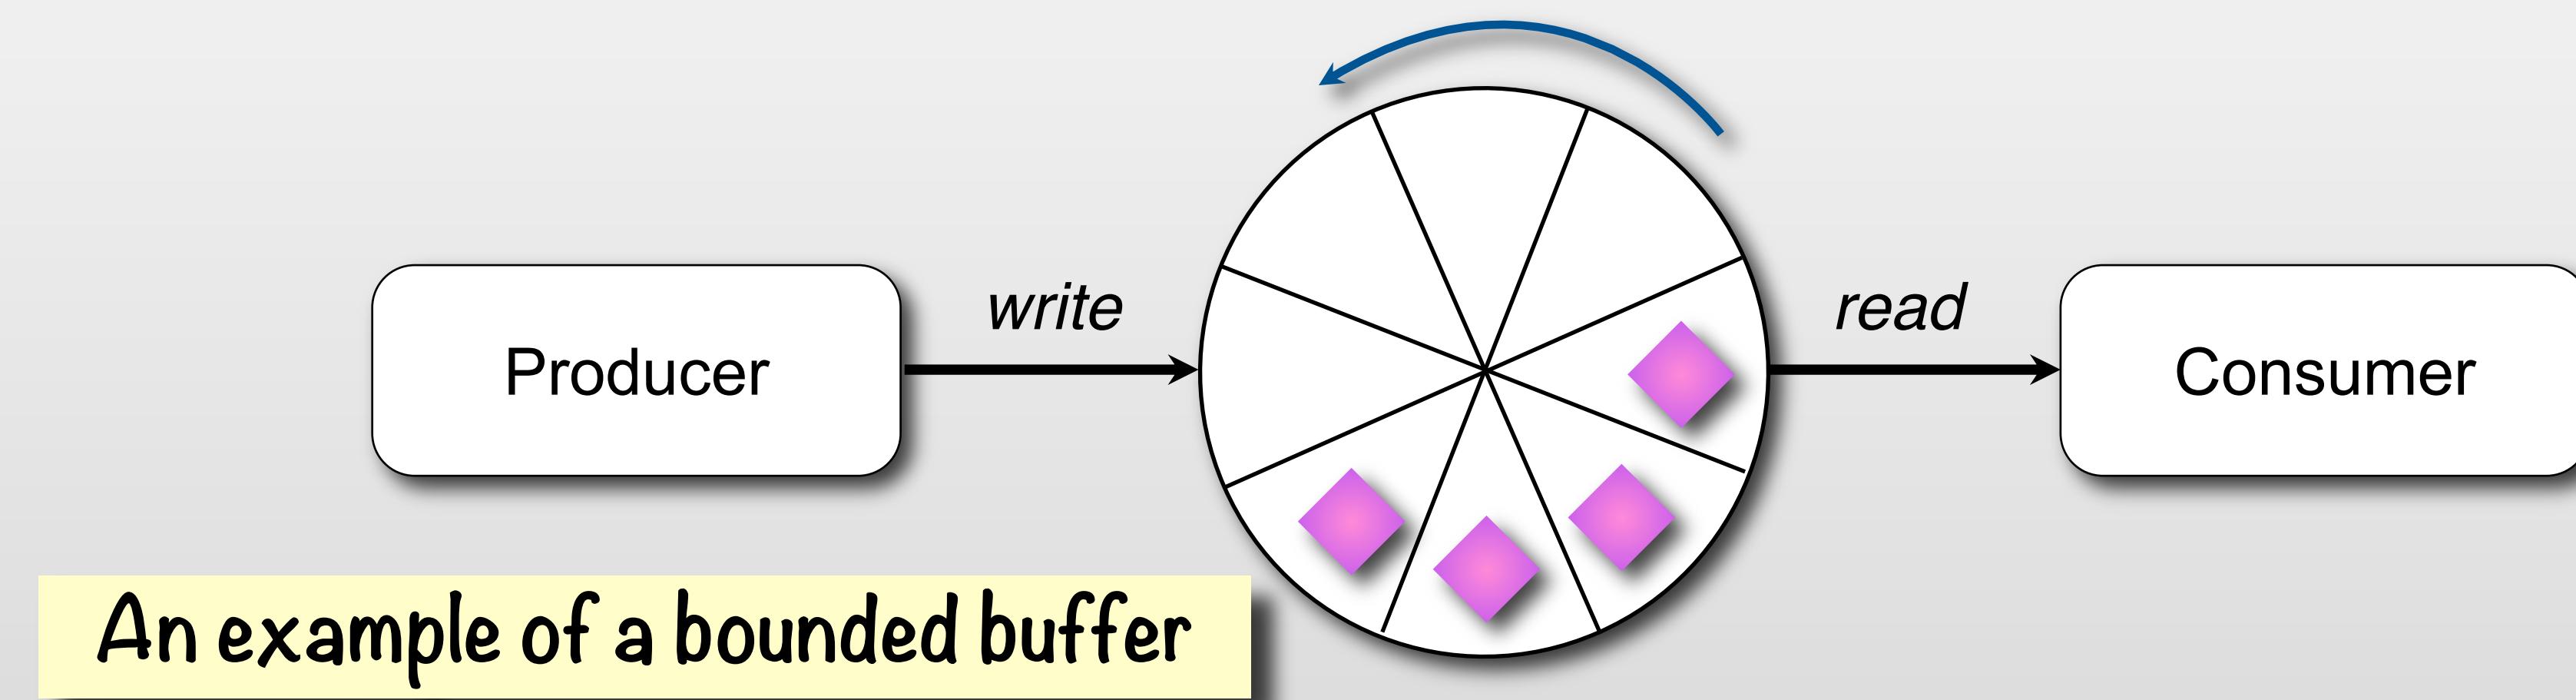
(b)

Kernel actively facilitates the communication

# Shared Memory: Producer-Consumer Paradigm



- Common paradigm for cooperating processes:
  - **producer process** produces information that is consumed by a **consumer process**
- The processes use a buffer:
  - **unbounded-buffer** places no practical limit on the size of the buffer
    - available memory is the only limit
  - **bounded-buffer** assumes that there is a fixed buffer size



# UNIX Shared Memory



## ● C Standard Library in UNIX:

```
int main() {
    long now; // time now
    int fd; // shared memory file descriptor
    char *mem_ptr; // pointer to shared memory object

    // create the shared memory object
    if ((fd = shm_open(SHM_NAME, O_CREAT | O_RDWR, S_IWUSR | S_IRUSR)) == -1)
        oops("CANNOT OPEN SHM", EXIT_FAILURE);

    ftruncate(fd, SHM_SIZE); // cut down to the desired size

    // memory map the shared memory object
    mem_ptr = (char *) mmap(0, SHM_SIZE, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
    // run a clock for a minute
    printf("Time ticks: ");
    for (int i = 0; i < 60; i++)
    {
        fflush(stdout);
        time(&now); // get the time
        strcpy(mem_ptr, ctime(&now)); // write to memory
        sleep(1); // wait a sec
        printf(".");
    }
    printf("\n");

    shm_unlink(SHM_NAME); // remove the shared memory object

    return EXIT_SUCCESS;
}
```

shm\_POSIX\_time\_server.c

server

Time is sent in this example from  
“time server” to “time client”

```
#define SHM_NAME "MY_SHM"
#define SHM_SIZE 64
#define oops(m,x) { perror(m); exit(x); }
```

shm\_POSIX.h

```
int main()
{
    int fd; // shared memory file descriptor
    char *mem_ptr; // pointer to shared memory object

    // open the shared memory object
    if ((fd = shm_open(SHM_NAME, O_RDONLY, 0400)) == -1)
        oops("CANNOT OPEN SHM", EXIT_FAILURE);

    // memory map the shared memory object
    mem_ptr = (char *) mmap(0, SHM_SIZE, PROT_READ, MAP_SHARED, fd, 0);

    // read from the shared memory object
    printf("TIME NOW: %s\n", mem_ptr);

    munmap(mem_ptr, SHM_SIZE);

    return EXIT_SUCCESS;
}
```

client

shm\_POSIX\_time\_client.c

# Message System: Overview



- Message system allows processes to communicate with each other without resorting to shared memory (shared variables)
- Kernel provides mechanisms to facilitate the communication
  - Message passing facility provides two operations:
    - **send**
    - **receive**
  - If processes P and Q wish to communicate, they need to **establish a communication link** between them before exchanging messages via send/receive

# Message System: Implementation Challenges

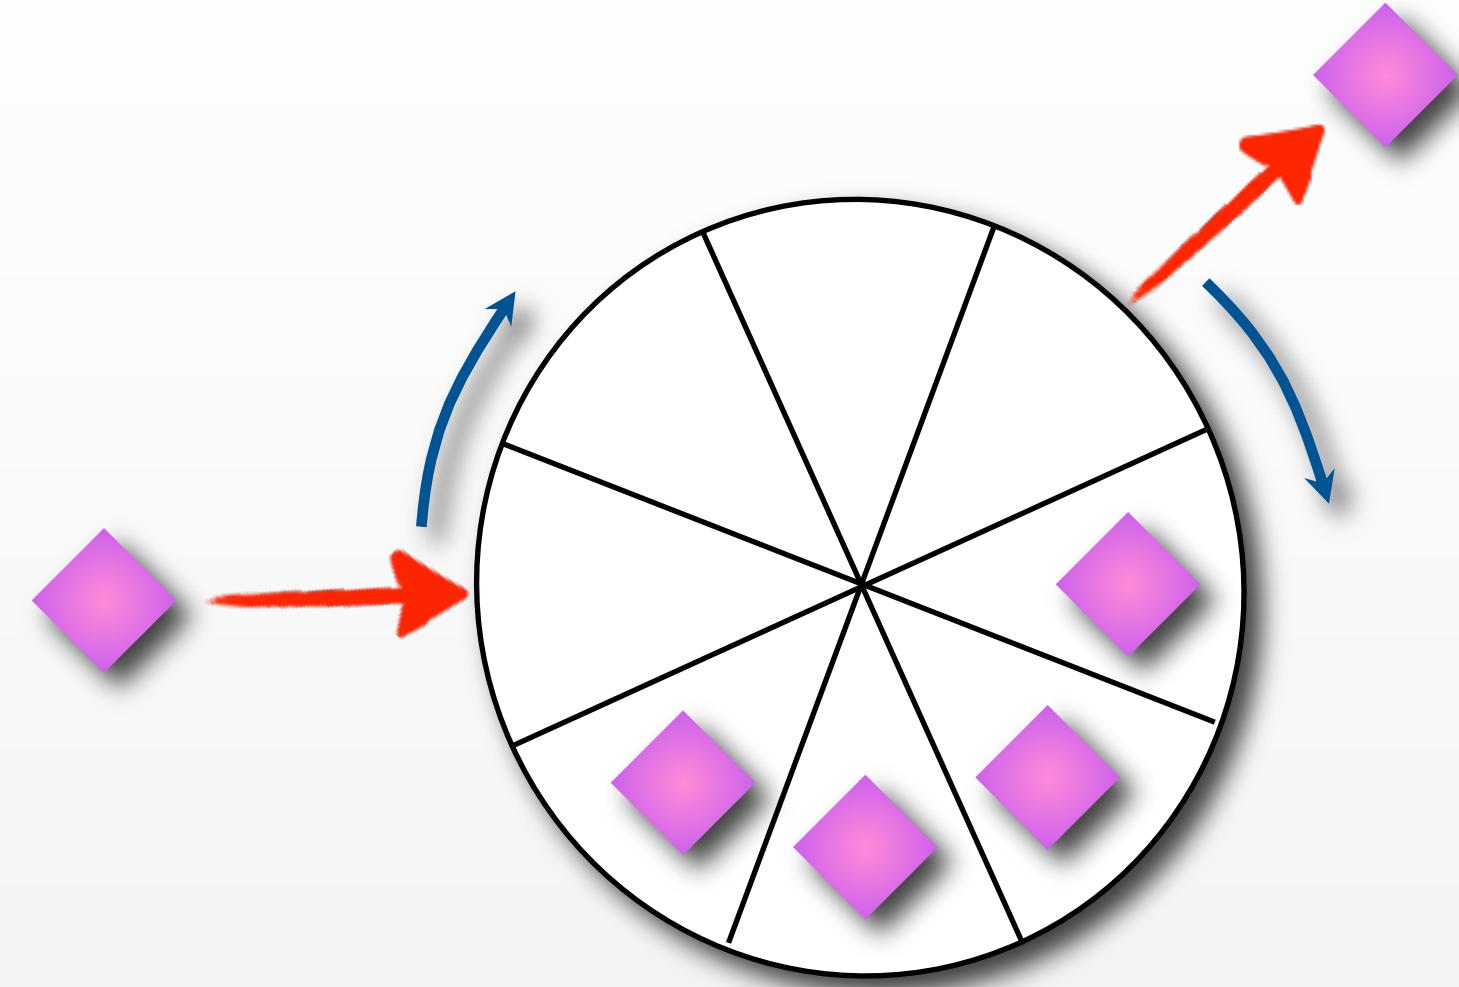


- Designers of a message system must decide:
  - how are links established
  - whether a link can be associated with more than two processes
  - the number of links that can there be between every pair of communicating processes
  - the capacity of a link
  - whether the size of a message that the link can accommodate is fixed or variable
  - whether the link is unidirectional or bi-directional
  - whether the link is direct or indirect
  - whether the link is synchronous or asynchronous
  - whether the link buffering is automatic or explicit

# Message System: Buffering



- Queue of messages attached to the link
- Implemented in one of three ways
  - **zero capacity** – 0 messages
    - sender must wait for receiver
  - **bounded capacity** – finite length of  $n$  messages
    - sender must wait if link full
  - **unbounded capacity** – infinite length
    - sender never waits
    - available memory is the practical limit however

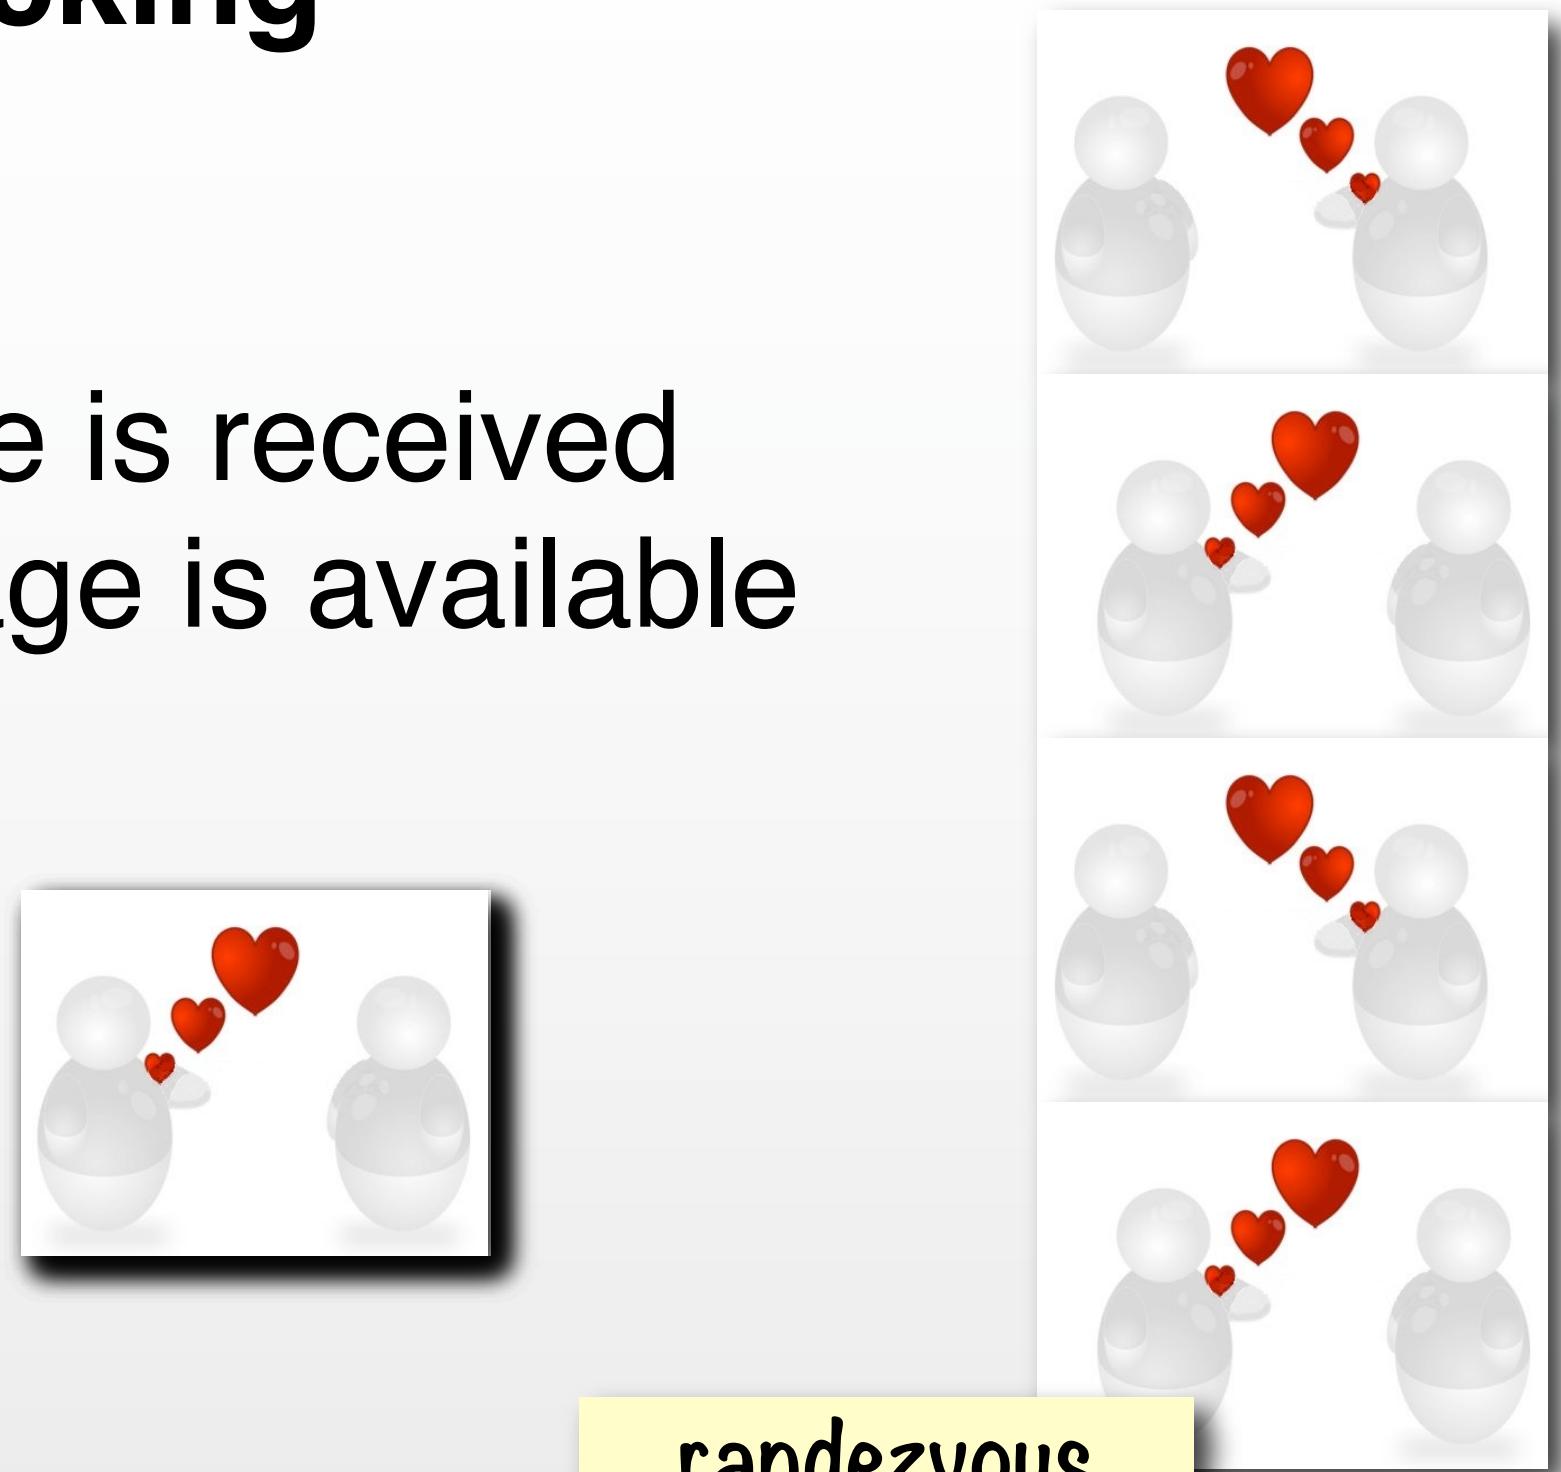


Sample illustration of a bounded buffer with two revolving indices for send (in; green) and receive (out; red)

# Message System: Synchronization



- Message passing may be either **blocking** or **non-blocking**
- **Blocking** is considered synchronous
  - **Blocking send** has the sender block until the message is received
  - **Blocking receive** has the receiver block until a message is available
    - Blocked processes waits in the waiting queue
- **Rendezvous**
  - communication between blocking sender and receiver
- **Non-blocking** is considered asynchronous
  - **Non-blocking send** has the sender send the message and continue
  - **Non-blocking receive** has the receiver receive a valid message or null (if there is no message)



rendezvous

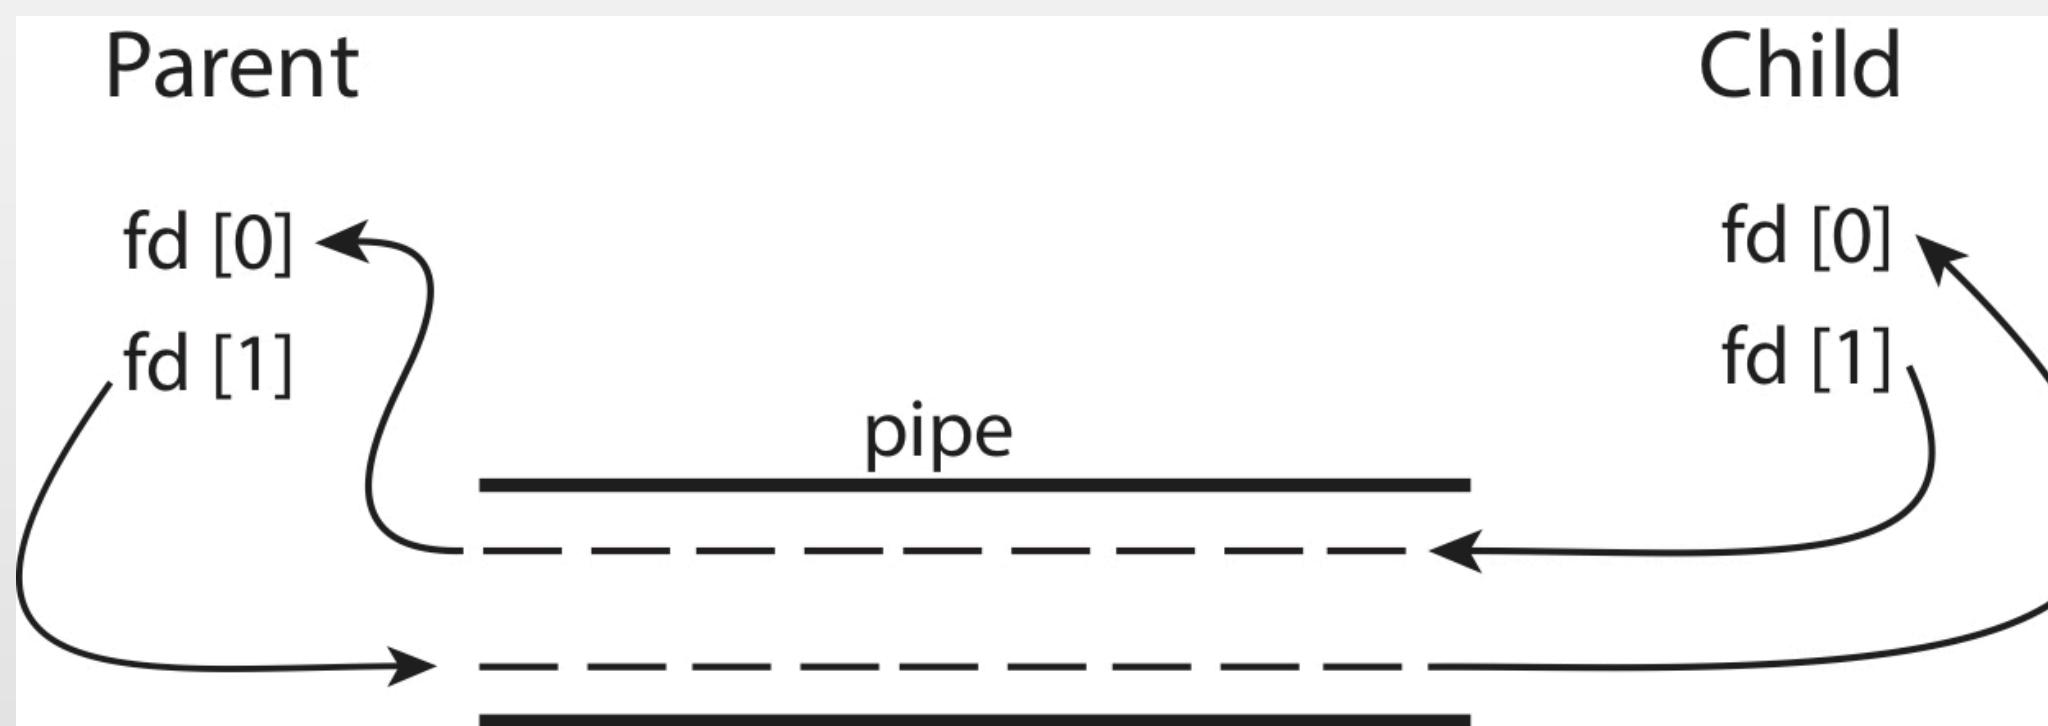
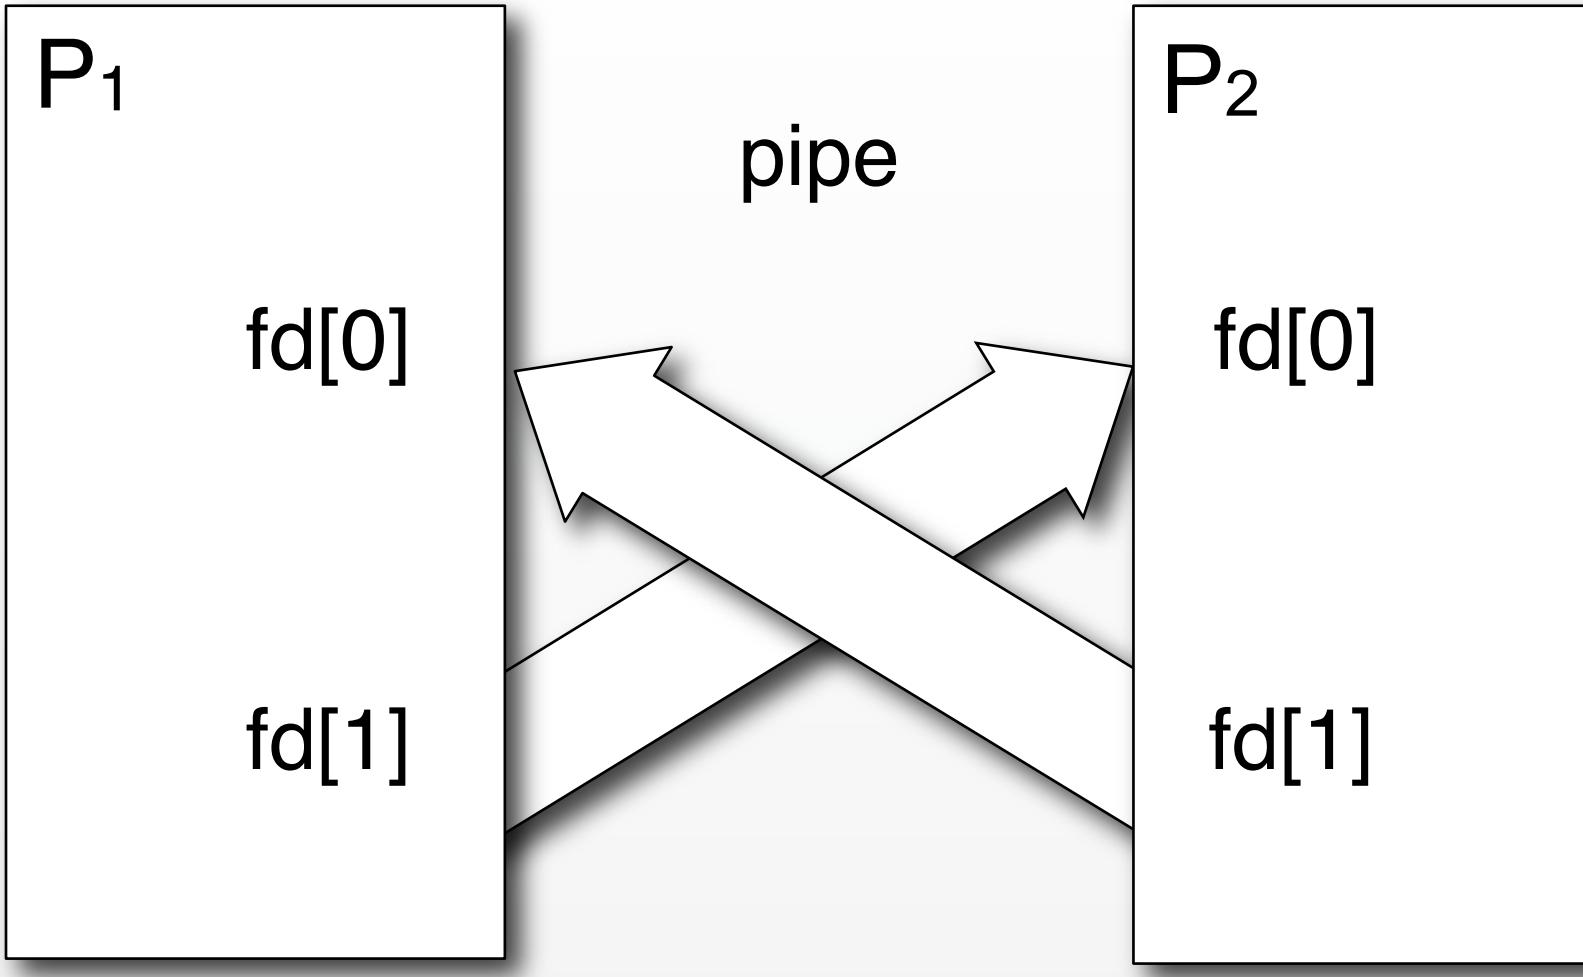
# Message System: Direct Communication



- Processes must name each other explicitly:
  - **send(P, message)** – send a message to process P
  - **receive(Q, message)** – receive a message from process Q
- Properties of the direct communication link
  - Links are established automatically
  - A link is associated with exactly one pair of communicating processes
  - Between each pair there exists exactly one link
  - The link may be unidirectional, but is usually bi-directional
  - The communication can be synchronous or asynchronous
    - “blocking” or “non-blocking”

It's like using an instant messaging system or a phone (with some “call-me-back” stretch for prompting messages from peers)

# UNIX (Unnamed) Pipes - Direct IPC

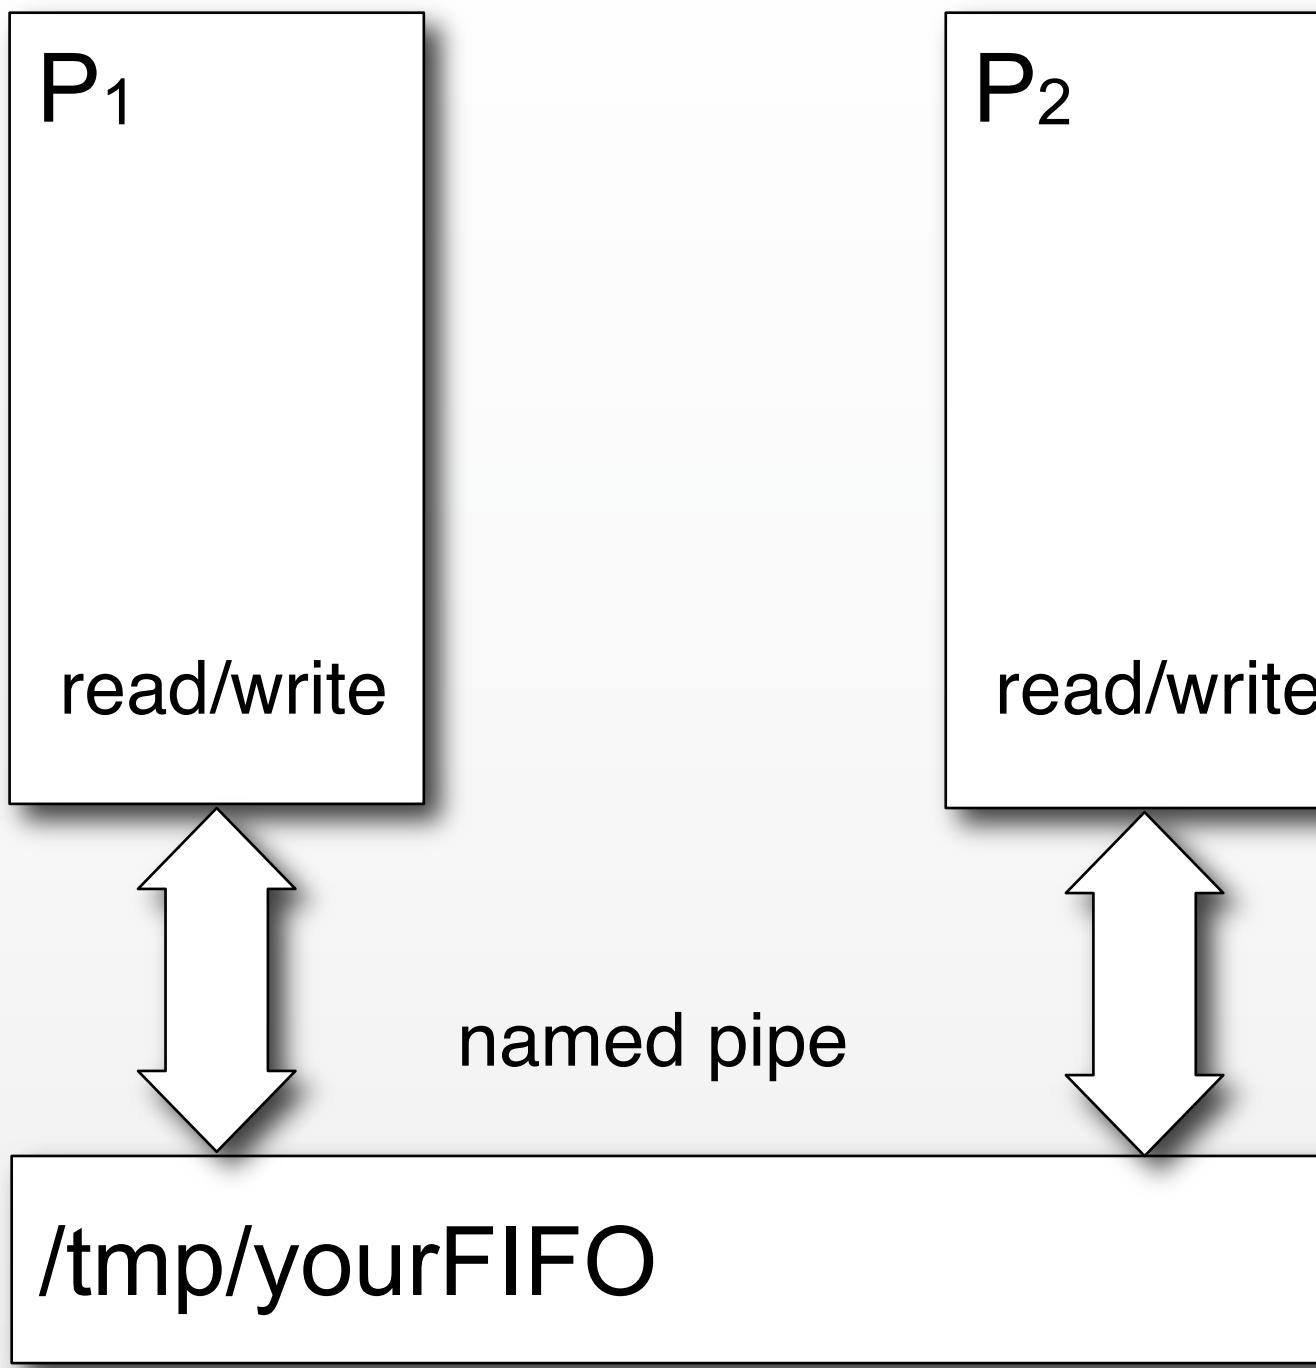


```
int      n, fd[2];
pid_t   pid;
char    line[MAXLINE];
char    line2[MAXLINE];

if (pipe(fd) < 0)
    printf("Cannot create a pipe.");
else if ( (pid = fork()) < 0)
    printf("Cannot fork.");
else if (pid == 0) /* child */
{
    write(fd[1], "Message for the parent\n", 23);
    n = read(fd[0], line, MAXLINE);
    printf("%s", line);
    close(fd[1]);
}
else /* parent */
{
    write(fd[1], "Message for the child\n", 22);
    n = read(fd[0], line2, MAXLINE);
    printf("%s", line2);
    close(fd[0]);
    wait(NULL);
}
```

See also pipe1.c and pipe2.c

# UNIX Named Pipes (FIFOs)



## Setting / clearing flags

NOTE: if flag == 00100000  
then ~flag == 11011111

SET: 11001100 | flag ==> 11101100

CLEAR: 11101100 & ~flag ==> 11001100

```
if (stat = mkfifo("/tmp/ajsFIFO", 0777) < 0)
    oops("Cannot make FIFO", stat);
// open a named pipe
pipe = open("/tmp/ajsFIFO", O_WRONLY); O_RDWR may be needed
while(1) {
    // get a line to send
    printf("Enter line: ");
    fgets(line, MAX_LINE, stdin);
    *strchr(line, '\n') = '\0';
    if (strcmp(line, "quit") == 0)
        break;
    // actually write out the data and close the pipe
    write(pipe, line, strlen(line) + 1);
}
// close the pipe
close(pipe);
```

npipe\_send.c

```
// open a named pipe
pipe = open("/tmp/ajsFIFO", O_RDONLY);
// set the mode to nonblocking
int flags;
flags |= O_NONBLOCK;
// flags &= ~O_NONBLOCK; // default
fcntl(pipe, F_SETFL, flags);
// read the data from the pipe
read(pipe, line, MAX_LINE);
printf("Received line: %s\n", line);
// close the pipe
close(pipe);
```

npipe\_receive.c

# Message System: Indirect Communication



- Messages are directed to and received from **mailboxes**
  - also referred to as **ports**
    - each mailbox has a unique id
    - processes can communicate only if they share a mailbox
- Properties of the indirect communication link
  - link established only if processes share a common mailbox
  - a link may be associated with many processes
  - each pair of processes may share several communication links
  - link may be unidirectional or bi-directional
- Operations
  - **create** a new mailbox
    - ...and tell others what it is
  - **send** and **receive** messages through mailbox
  - **destroy** a mailbox
- Primitives are defined as:
  - **send(A, message)** – send a message to mailbox A
  - **receive(A, message)** – receive a message from mailbox A

➤ It's like having multiple email accounts and reading them at your own leisure  
➤ No one can send anything to you without knowing your address

# Message System: Indirect Communication (cont.)



- Mailbox sharing dilemma
  - $P_1$ ,  $P_2$ , and  $P_3$  share mailbox A
  - $P_1$ , sends;  $P_2$  and  $P_3$  receive
  - Who gets the message?
- Possible solutions
  - constrain links to be associated with at most two processes
    - every message has one sender and one receiver
  - allow only one process at a time to execute a receive operation
  - allow the system to select arbitrarily the receiver
    - sender is notified who the receiver was
  - more complex schema could be used

# UNIX POSIX Message Queues: Sender



```
#include "ipc_POSIX_messg.h"

int main() {
    mqd_t msqid;
    unsigned int type;

    char *msg = calloc(MSG_SIZE, sizeof(char)); // put zeros in the allocated space
    strcpy(msg, "Hello receiver!");

    if ((msqid = mq_open(QUEUE_NAME, O_RDWR | O_CREAT, S_IRWXU | S_IRWXG, NULL)) < 0)
        oops("Error opening a queue.", errno);

    if (mq_send(msqid, msg, strlen(msg) + 1, MSG_TYPE) < 0) // type 1 (arbitrarily)
        oops("Error opening a queue.", errno);

    printf("Sent: \"%s\"\nRun the receiver.\n", msg);

    printf("To continue, press Enter on the keyboard...");
    getchar(); // just to stop and wait for the receiver to respond

    if (mq_receive(msqid, msg, MSG_SIZE, &type) < 0)
        oops("Error receiving data.", errno);

    printf("Received: \"%s\" of type %d\n", msg, type);

    mq_close(msqid);
    mq_unlink(QUEUE_NAME);

    exit(EXIT_SUCCESS);
}
```

ipc\_POSIX\_send.c

```
#ifndef _IPC_MESSG_H
#define _IPC_MESSG_H

#include <sys/types.h>
#include <sys/types.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <errno.h>
#include <fcntl.h>
#include <sys/stat.h>
#include <mqueue.h>

#define oops(ermmsg,errno)
    {perror(ermmsg); exit(errno); }

#define QUEUE_NAME "/myQueue"
#define MSG_SIZE 8192
#define MSG_TYPE 1

#endif
```

ipc\_POSIX\_messg.h

# POSIX Message Queues: Receiver



```
#include "ipc_POSIX_messg.h"

int main()
{
    mqd_t msqid;
    char msg_rcvd[MSG_SIZE], msg_send[MSG_SIZE];
    unsigned int type;

    if ((msqid = mq_open(QUEUE_NAME, O_RDWR)) < 0) // assuming it already exists
        oops("Error opening a queue.", errno);

    if (mq_receive(msqid, msg_rcvd, MSG_SIZE, &type) < 0)
        oops("Error receiving data.", errno);

    printf("Received: \"%s\" of type %d\n", msg_rcvd, type);

    sprintf(msg_send, "Sender, thank you for your message \"%s\"", msg_rcvd);

    if (mq_send(msqid, msg_send, MSG_SIZE, MSG_TYPE) < 0)
        oops("Error responding.", errno);

    printf("Sent back: \"%s\"\n", msg_send);

    mq_close(msqid);

    exit(0);
}
```

ipc\_POSIX\_receive.c

```
#ifndef _IPC_MESSG_H
#define _IPC_MESSG_H

#include <sys/types.h>
#include <sys/types.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <errno.h>
#include <fcntl.h>
#include <sys/stat.h>
#include <mqueue.h>

#define oops(ermmsg,errno)
    {perror(ermmsg); exit(errno); }

#define QUEUE_NAME "/myQueue"
#define MSG_SIZE 8192
#define MSG_TYPE 1

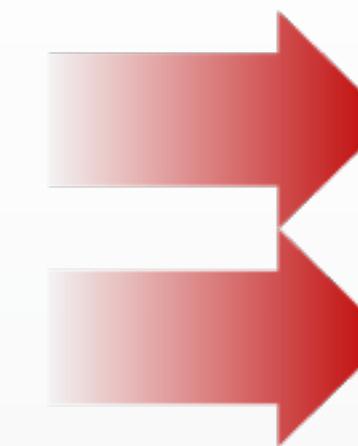
#endif
```

ipc\_POSIX\_messg.h

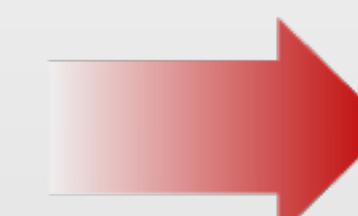
# System Management of IPC Resources



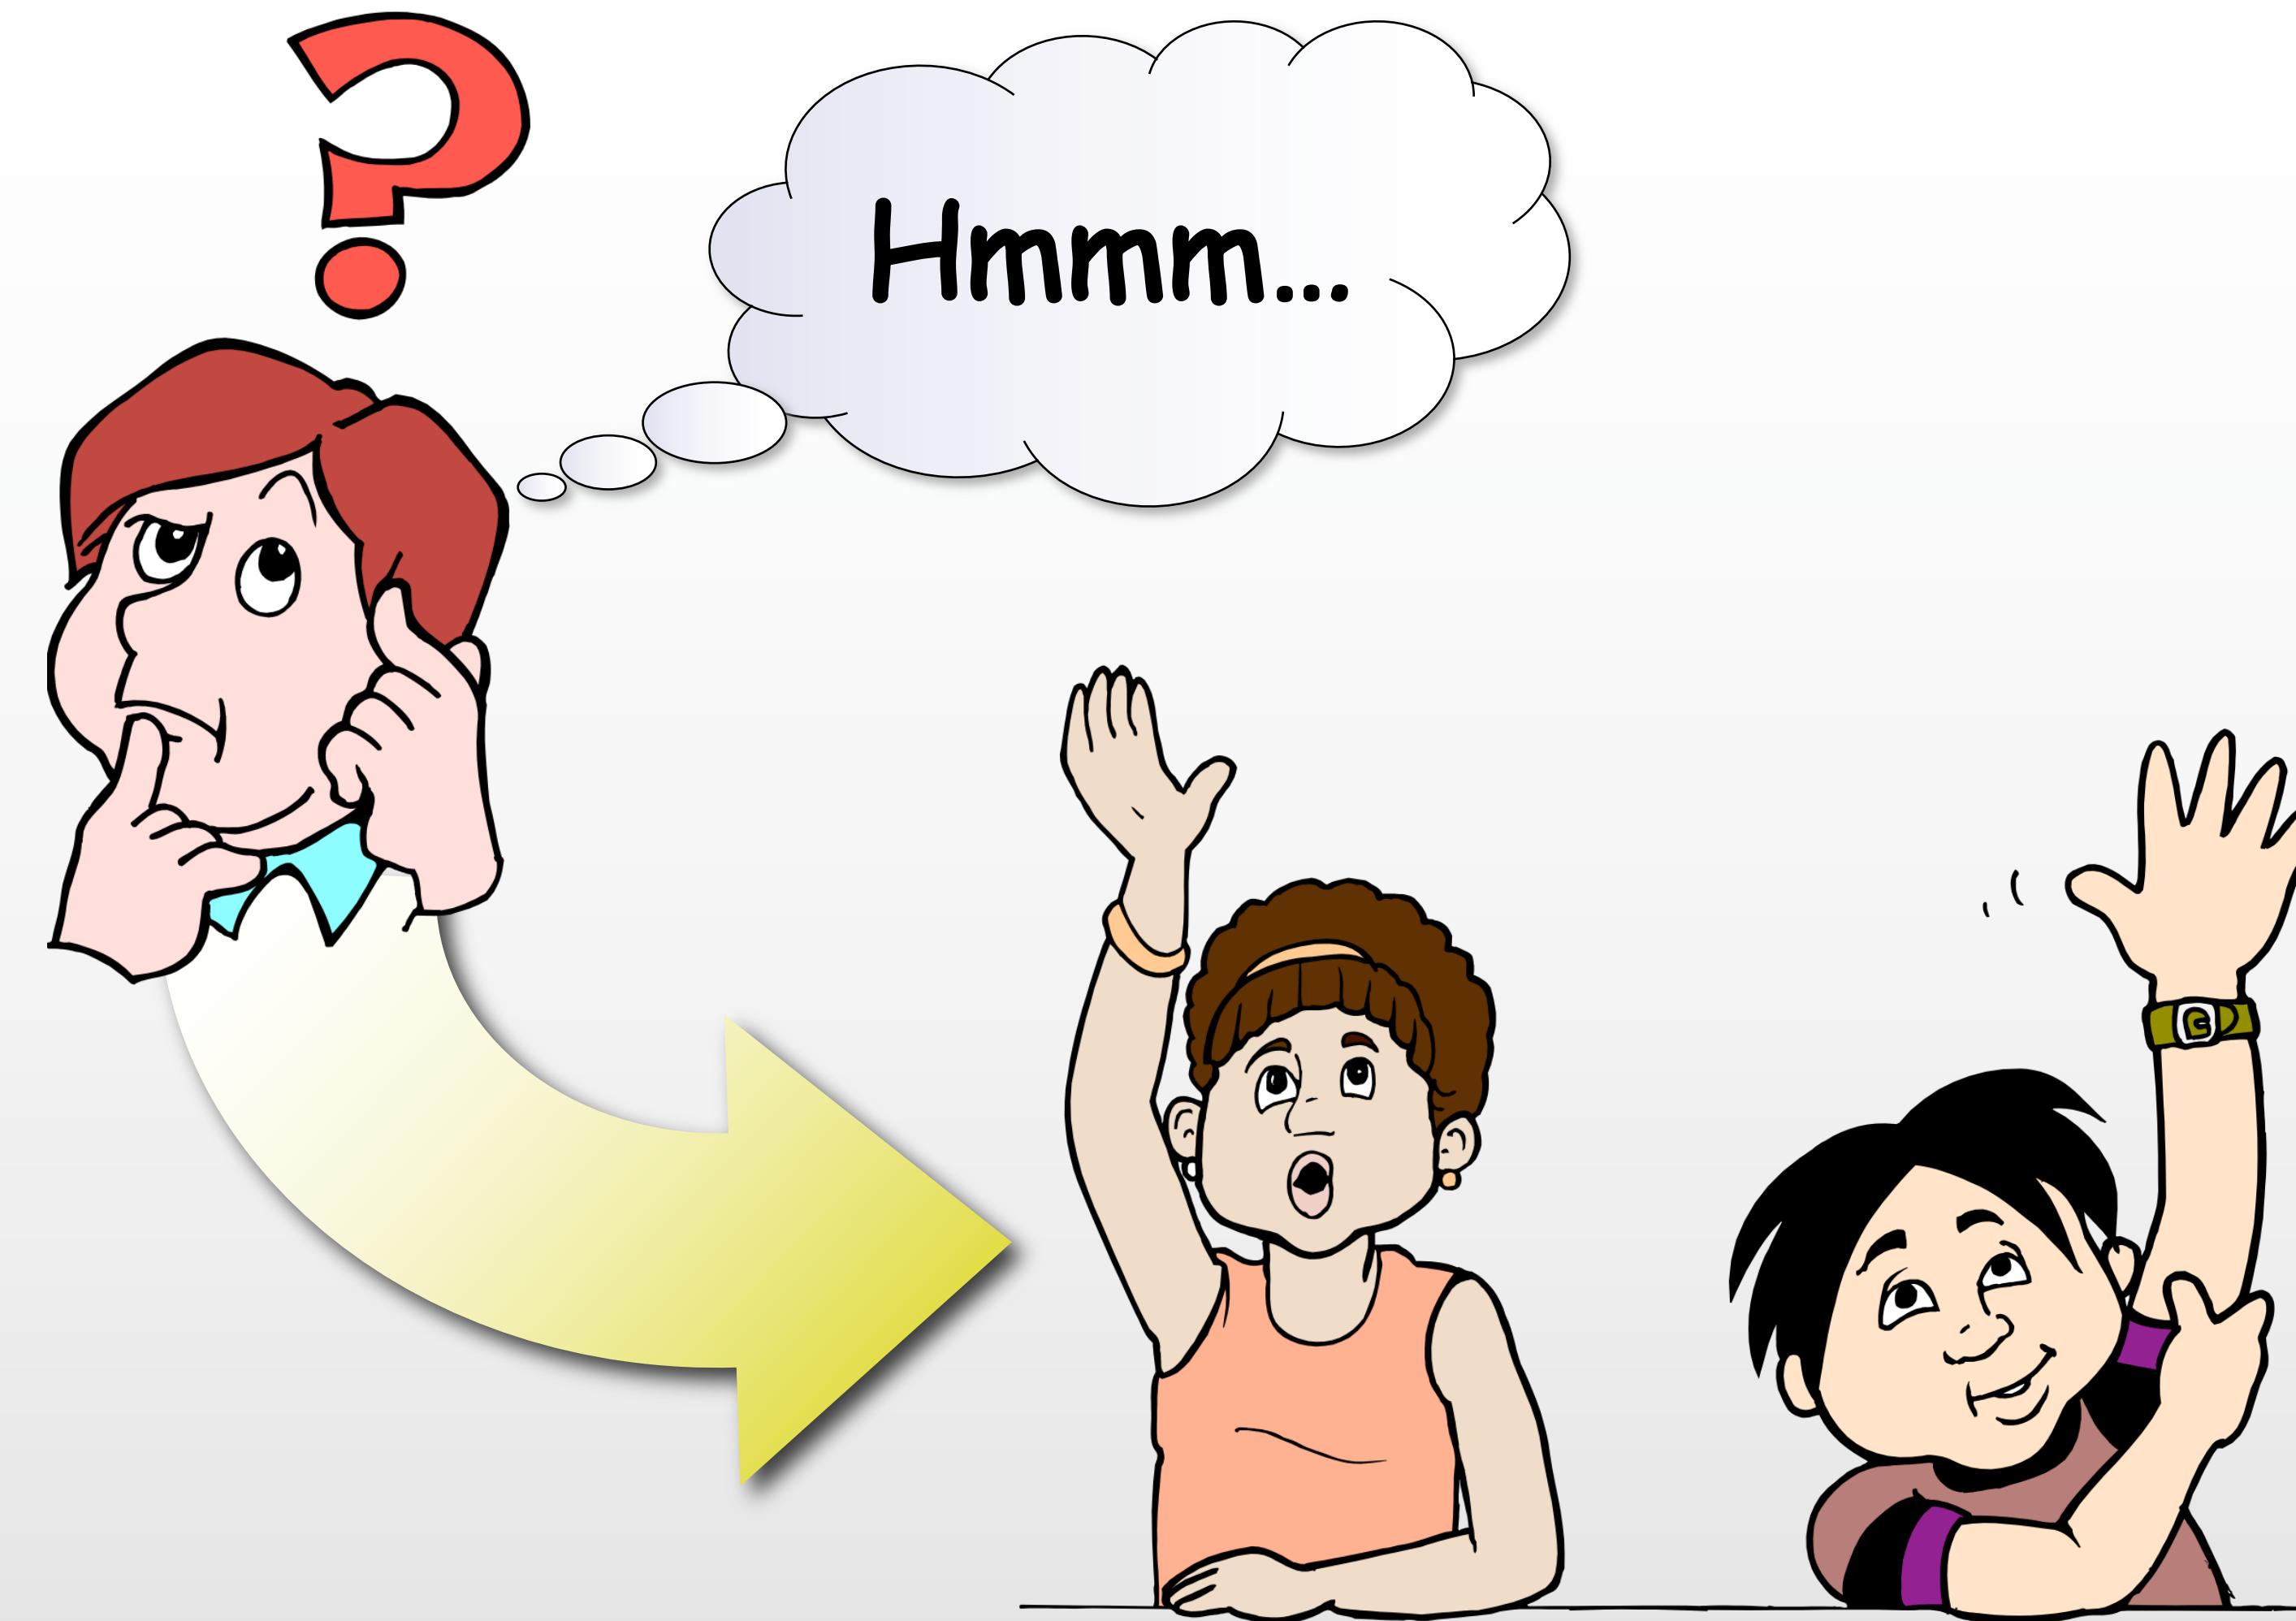
- IPC resources are assigned to the user, and not a process, so they are not returned automatically



```
$ man ipcs  
$ man ipcrm  
$  
$ ipcs  
----- Message Queues -----  
key msqid owner perms used-bytes messages  
----- Shared Memory Segments -----  
key shmid owner perms bytes nattch status  
0x00000000 65536 aj 600 524288 2 dest  
0x00000000 524289 aj 600 524288 2 dest  
0x00000000 425986 aj 600 524288 2 dest  
0x00000000 294915 aj 600 16777216 2 dest  
0x00000000 851972 aj 600 134217728 2 dest  
0x00000000 950277 aj 600 524288 2 dest  
0x00000000 655366 aj 600 524288 2 dest  
0x00000000 819207 aj 600 524288 2 dest  
0x00000000 1048584 aj 600 524288 2 dest  
0x00000000 1146889 aj 600 524288 2 dest  
0x00000000 1179658 aj 600 536870912 2 dest  
0x00000000 4128779 aj 600 393216 2 dest  
0x00000000 5570572 aj 600 393216 2 dest  
0x00000000 6848525 aj 600 67108864 2 dest  
----- Semaphore Arrays -----  
key semid owner perms nsems  
0xd20042e8 65536 aj 666 1  
0xd20042ea 98305 aj 666 1  
0xd20042e9 131074 aj 666 1  
0xd200430f 163843 aj 666 1  
0xd20042eb 196612 aj 666 1
```



- ipcs** and **ipcrm** are command that can be used to list and remove IPC resources



**COMP362 Operating Systems**  
**Prof. AJ Biesczad**