



# Lecture 8: Deadlocks

# COMP362 Operating Systems

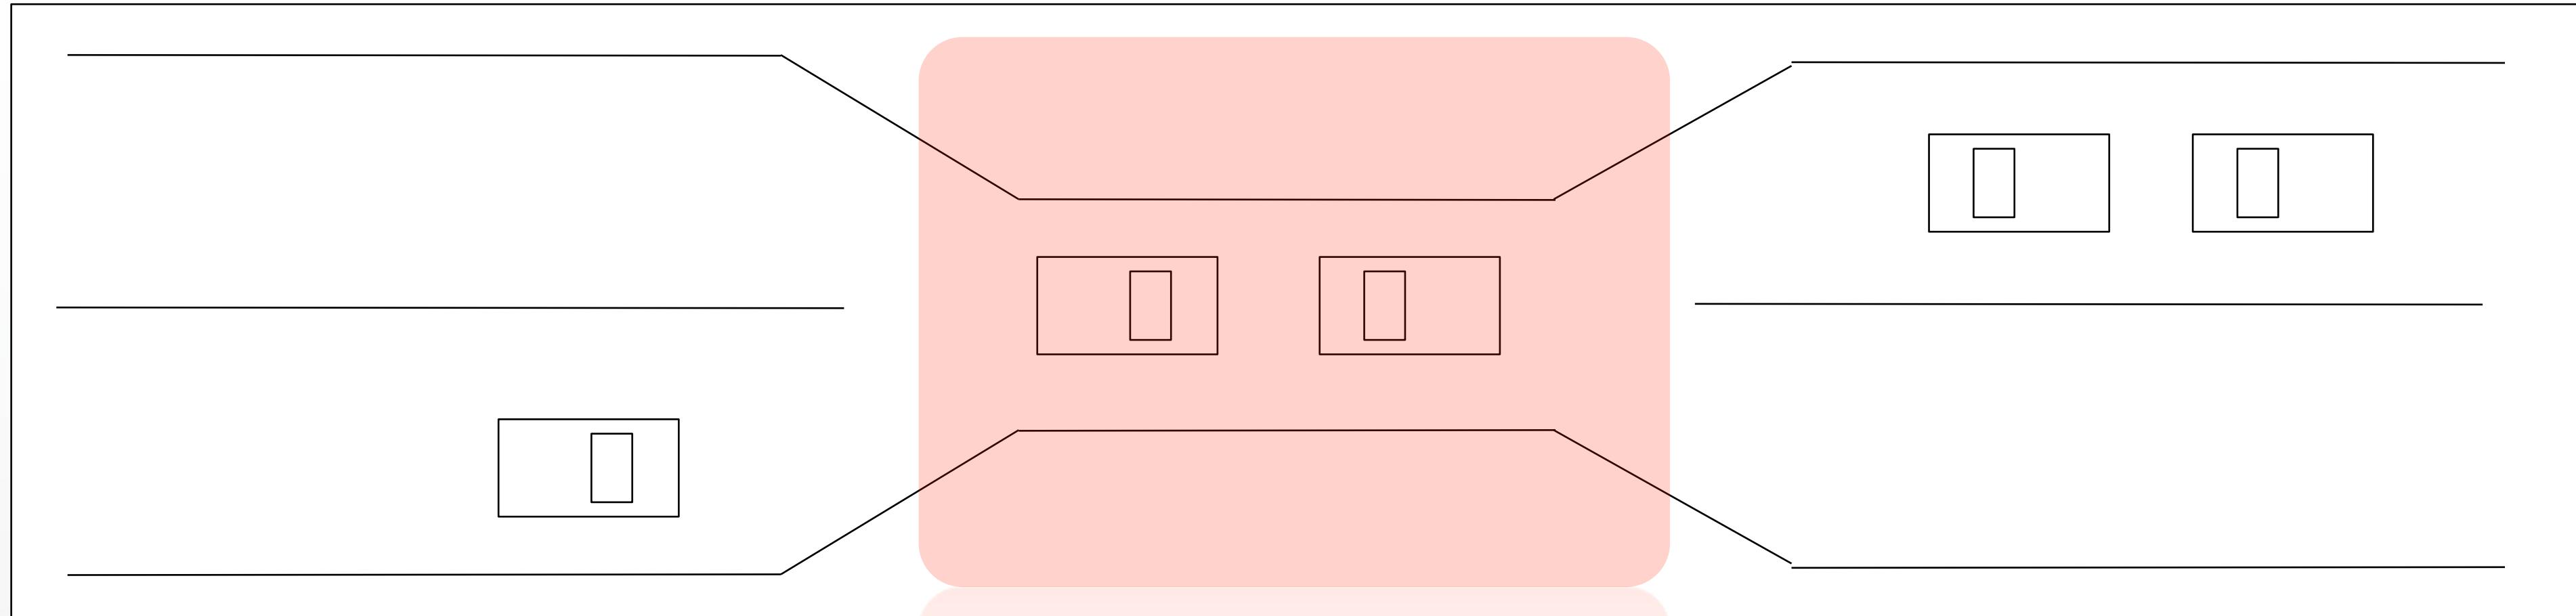
## Prof. AJ Bieszczad

# Outline



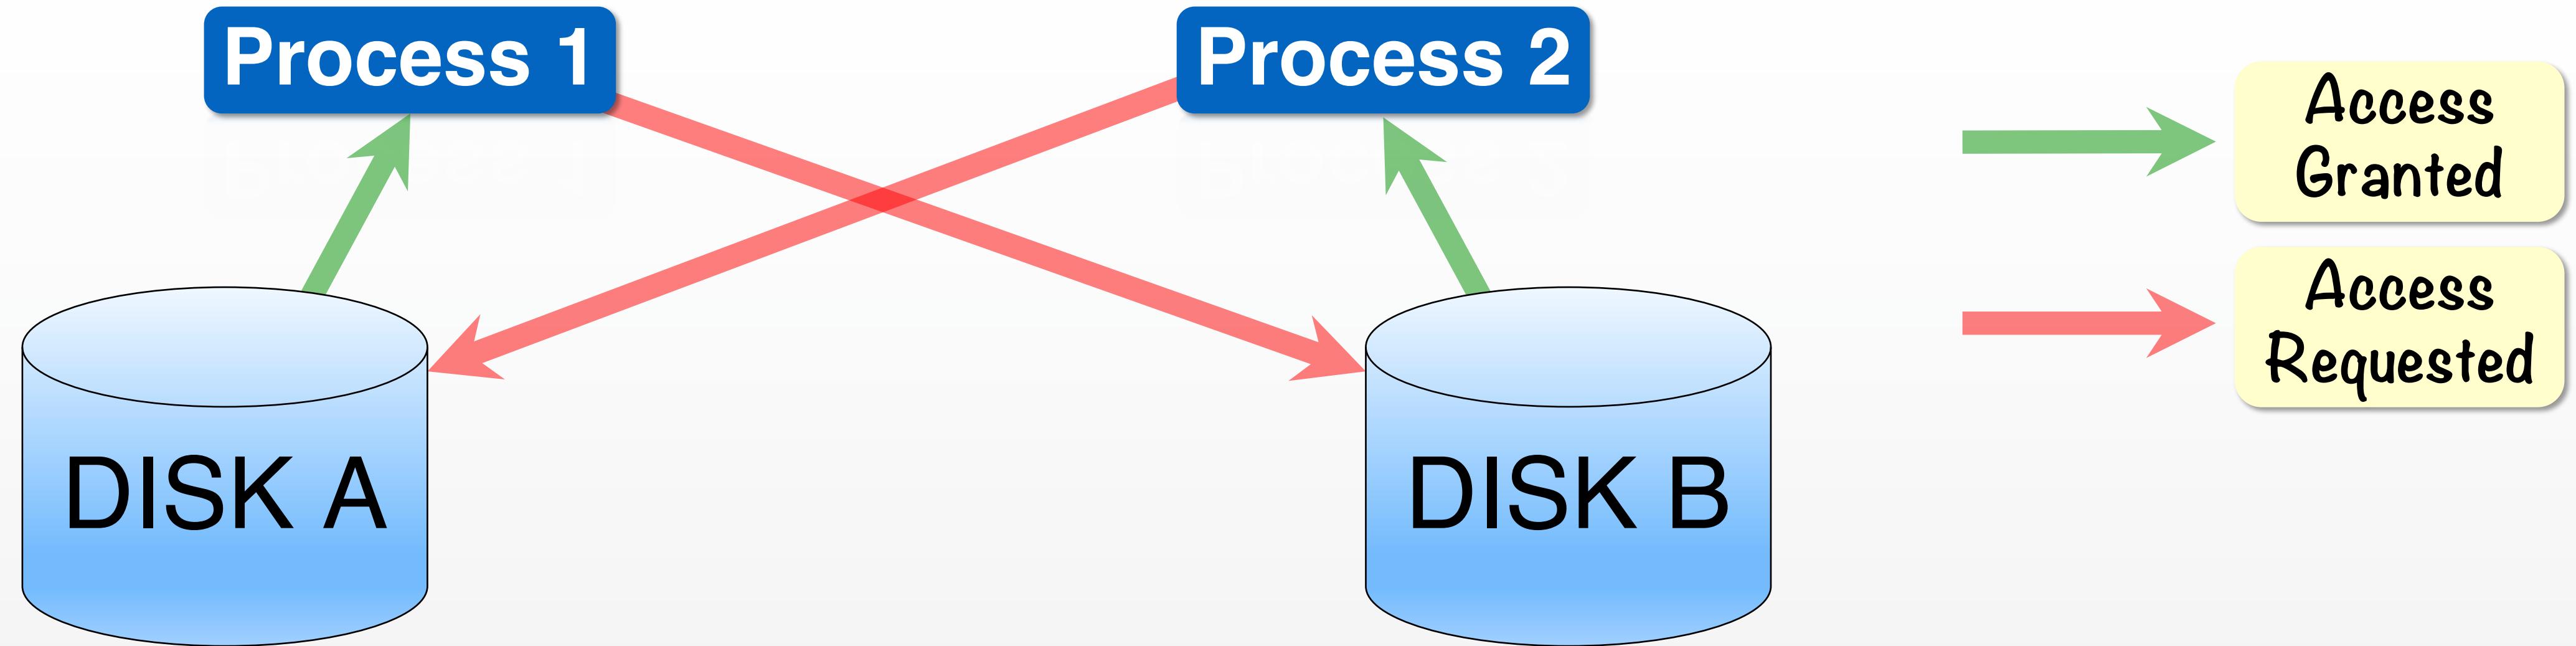
- The Deadlock Problem
- System Modeling
- Conditions for a Deadlock
- Methods for Handling Deadlocks
  - Deadlock Prevention
  - Deadlock Avoidance
  - Deadlock Detection
  - Recovery from Deadlock

# Problems with Sharing a Narrow Bridge



- Imagine that a section of a bridge is too narrow for two cars
  - i.e., the section is a resource shared by cars coming from both sides.
- If two cars start entering it from opposite directions, the bridge is deadlocked
  - i.e., no traffic in any direction is possible.
- The deadlock can be resolved if one or more cars back up.
- If only one side yields all the time, then the cars on the other side of the river will never cross.

# Problems with Sharing Computer Resources



- Say, we have a system with two disk drives, A and B.
- Imagine that Process 1 needs to copy a file from disk A to disk B, and Process 2 needs to copy a file from disk B to disk A.
- In this (somewhat primitive) system, there are no buffers, so a process needs to secure access to both disks to perform the operation.
- One possible scenario is that the Process 1 acquires access to disk A, and Process 2 acquires access to disk B.
- Each process needs two disks to carry out the copy operation, so to then they will both wait.
- Processes are deadlocked waiting for each other to release the other disk.

# Problem with Pthreads Locks



## Thread 1

```
void *do_work_one(void *param)
{
    pthread_mutex_lock(&first_mutex);
    pthread_mutex_lock(&second_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&second_mutex);
    pthread_mutex_unlock(&first_mutex);

    pthread_exit(0);
}
```

## Thread 2

```
void *do_work_two(void *param)
{
    pthread_mutex_lock(&second_mutex);
    pthread_mutex_lock(&first_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&first_mutex);
    pthread_mutex_unlock(&second_mutex);

    pthread_exit(0);
}
```

- Threads are deadlocked waiting for each other to release the mutex that they do not hold.

# Deadlock



- Deadlock can arise if the four following conditions hold simultaneously.
- **Mutual exclusion**
  - only one process at a time can use a resource.
- **Hold and wait**
  - a process holding at least one resource is waiting to acquire additional resources held by other processes.
- **No preemption**
  - a resource released only voluntarily by the process holding it after that process has completed its task
- **Circular wait:**
  - there exists a set  $\{p_0, p_1, \dots, p_n\}$  of waiting processes such that  $p_0$  is waiting for a resource that is held by  $p_1$ ,  $p_1$  is waiting for a resource that is held by  $p_2$ , ...,  $p_{n-1}$  is waiting for a resource that is held by  $p_n$ , and  $p_n$  is waiting for a resource that is held by  $p_0$ .



# Livelock

```
void *do_work_one(void *param){  
    bool done = false;  
    while (!done){  
        pthread_mutex_lock(&first_mutex);  
        if (pthread_mutex_trylock(&second_mutex)){  
            /**  
             * Do some work  
             **/  
            pthread_mutex_unlock(&second_mutex);  
            pthread_mutex_unlock(&first_mutex);  
            done = true;  
        }  
        else  
            pthread_mutex_unlock(&first_mutex);  
    }  
    pthread_exit(0);  
}
```

Thread 1

```
void *do_work_two(void *param){  
    bool done = false;  
    while (!done){  
        pthread_mutex_lock(&second_mutex);  
        if (pthread_mutex_trylock(&first_mutex)){  
            /**  
             * Do some work  
             **/  
            pthread_mutex_unlock(&first_mutex);  
            pthread_mutex_unlock(&second_mutex);  
            done = true;  
        }  
        else  
            pthread_mutex_unlock(&second_mutex);  
    }  
    pthread_exit(0);  
}
```

Thread 2

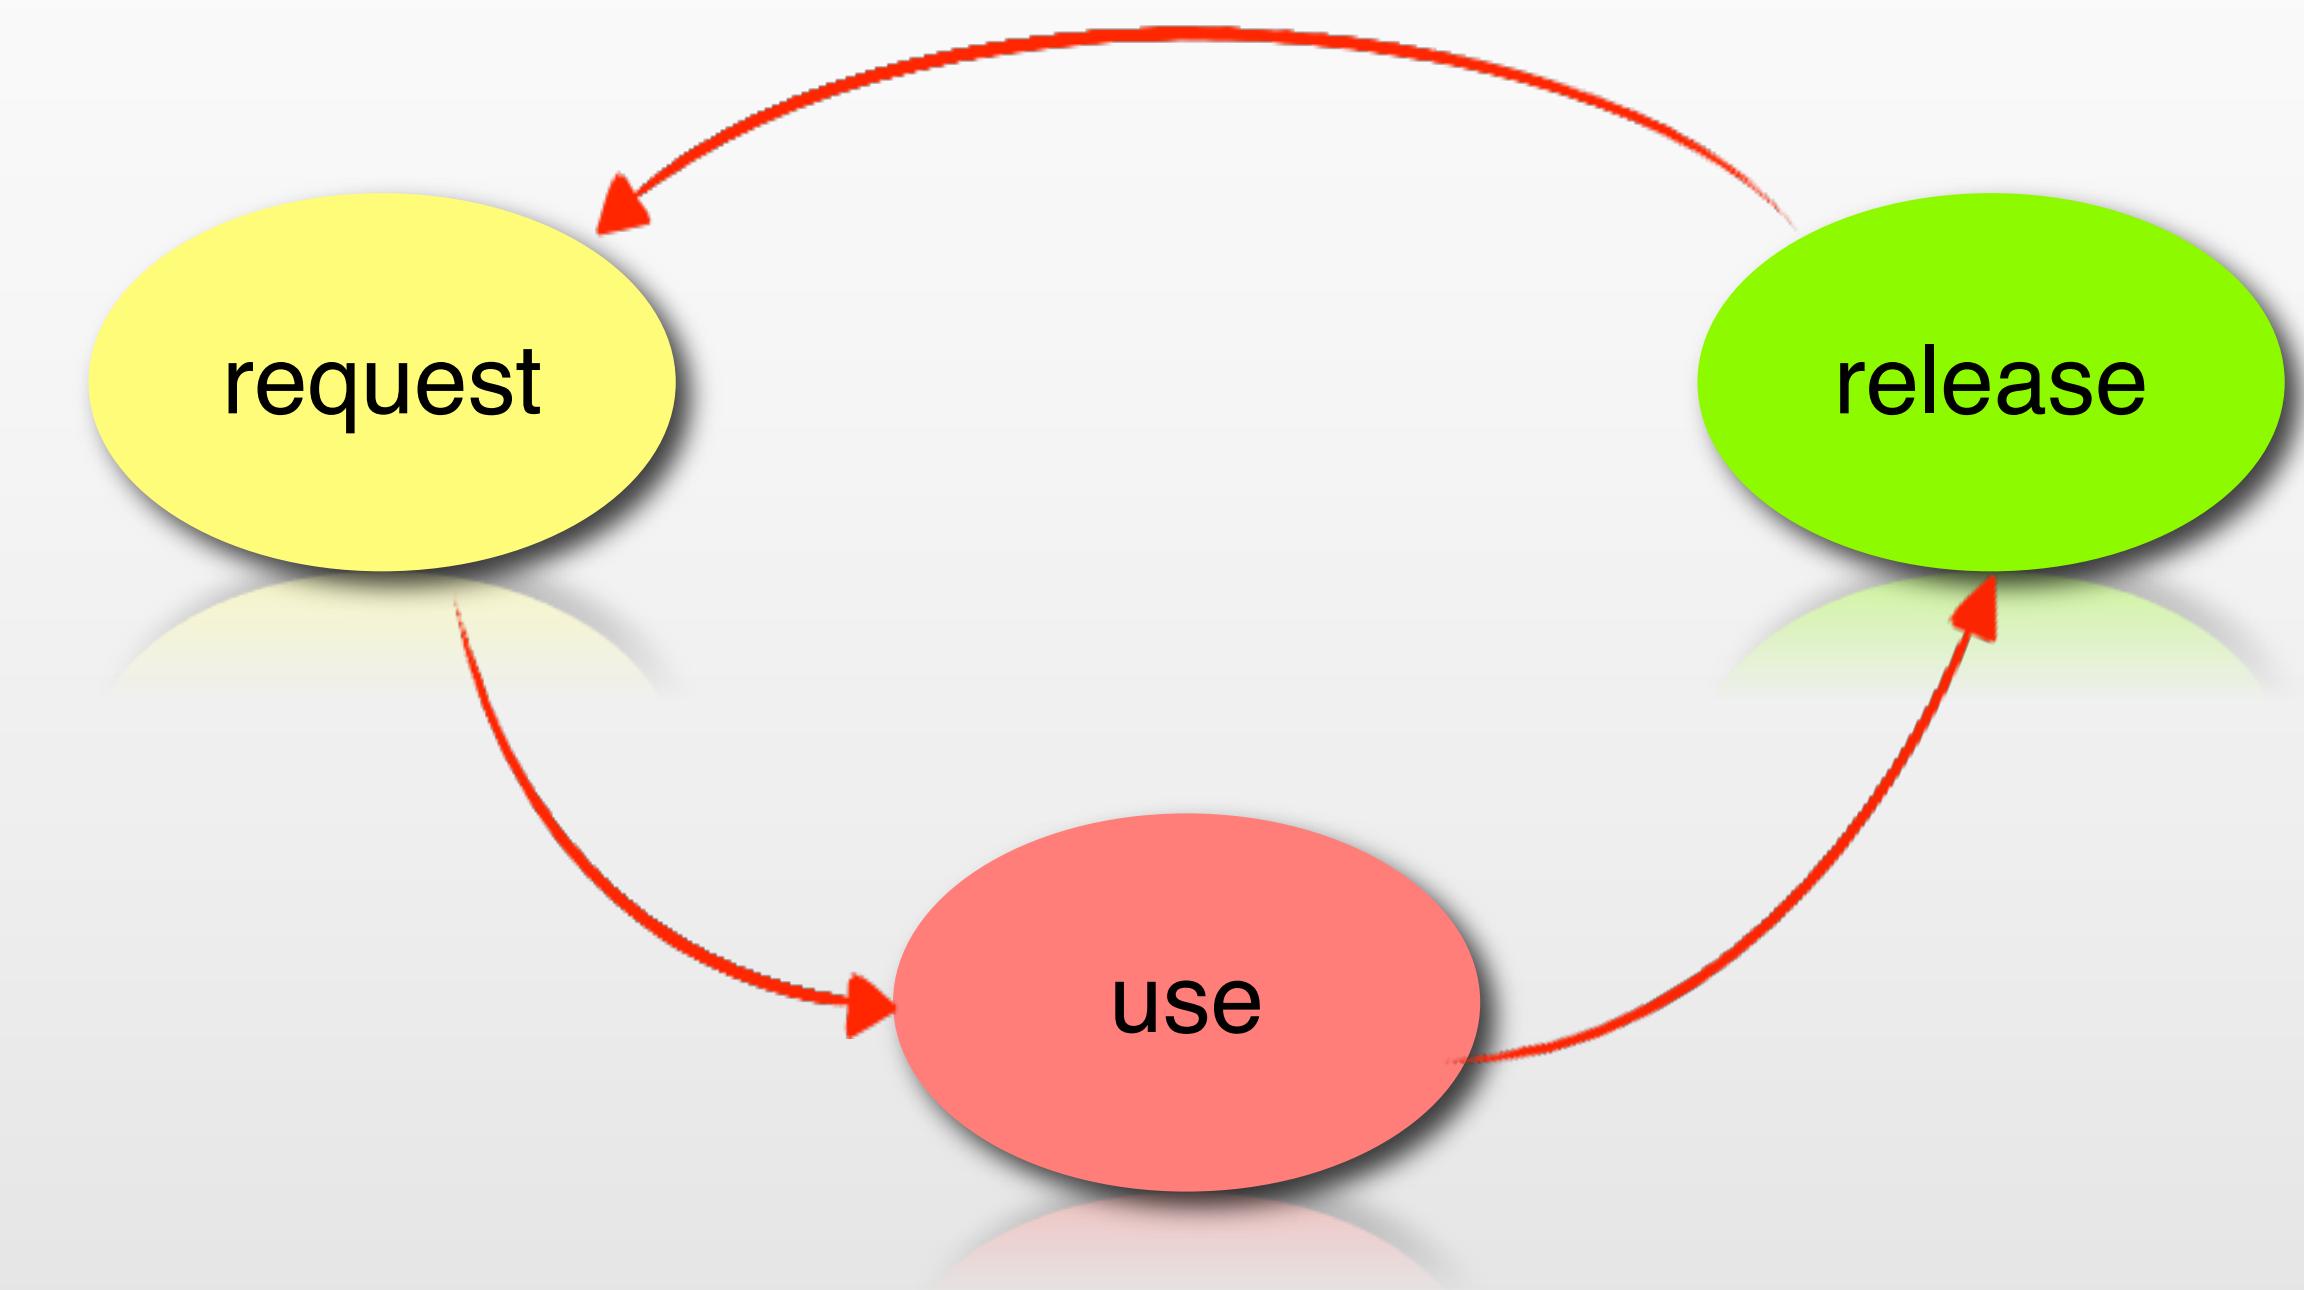
- Threads are not blocked, but they may no be able to do any work
- **Backoff** is a technique that randomizes access to resources to prevent contention
  - important concept used in other contexts

# System Modeling



- Let us institute the following formal model of a resource allocation system:
  - $n$  processes  $p_1, p_2, \dots, p_n$
  - $m$  resource types  $R_1, R_2, \dots, R_m$ 
    - such as CPU cycles, memory space, I/O devices, etc.
  - Each resource type  $R_i$  has  $w_i$  instances
    - e.g., five drives, two Ethernet ports, etc.

- Each process utilizes a resource as follows:

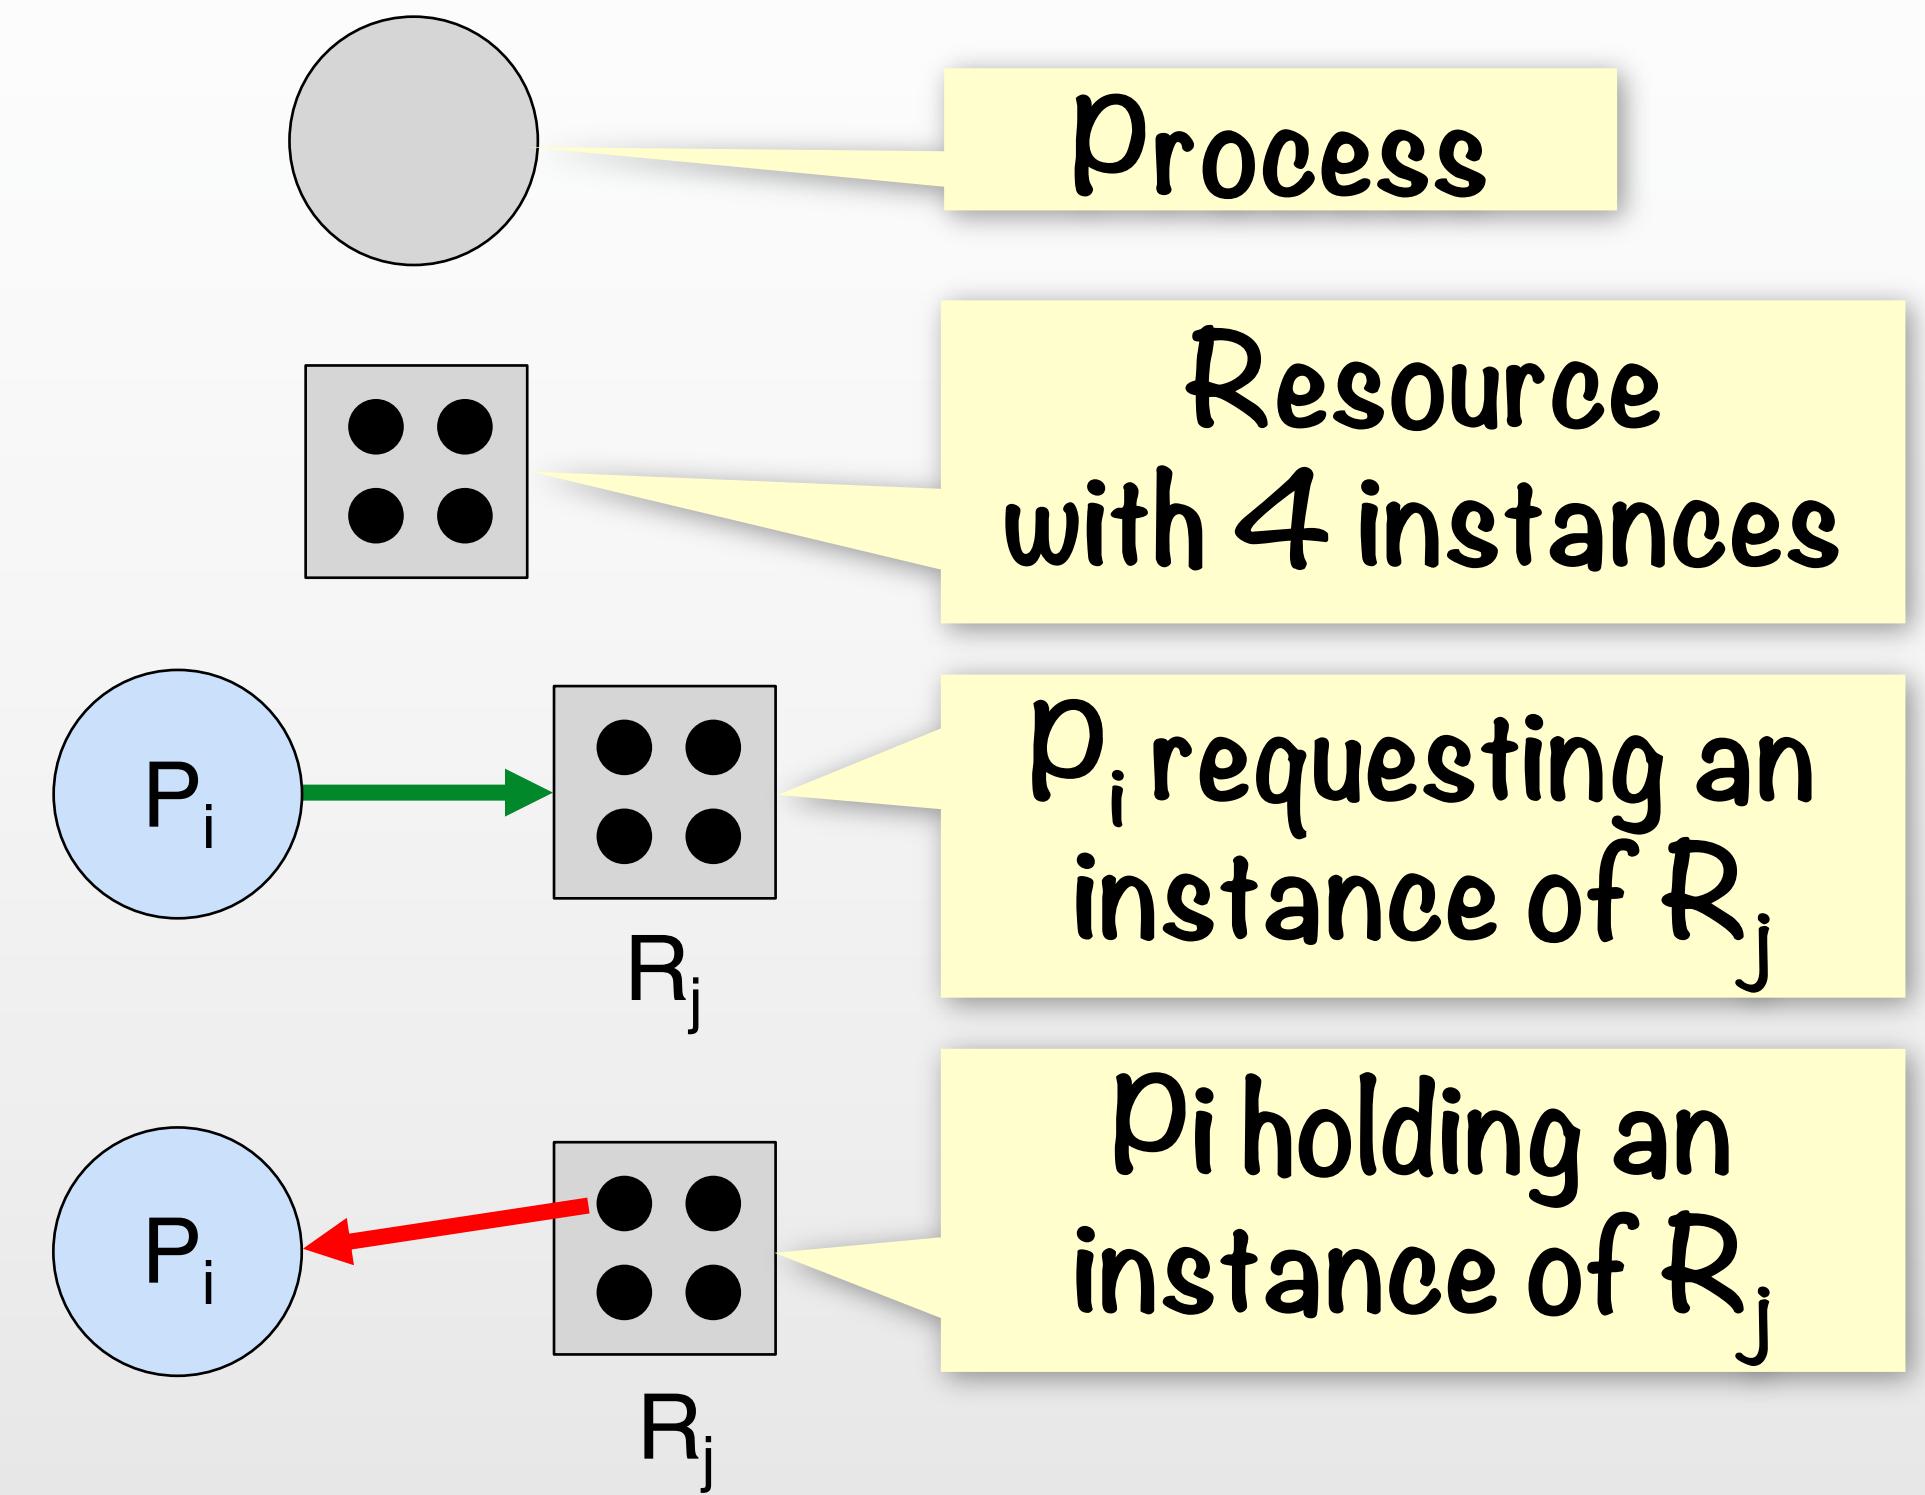




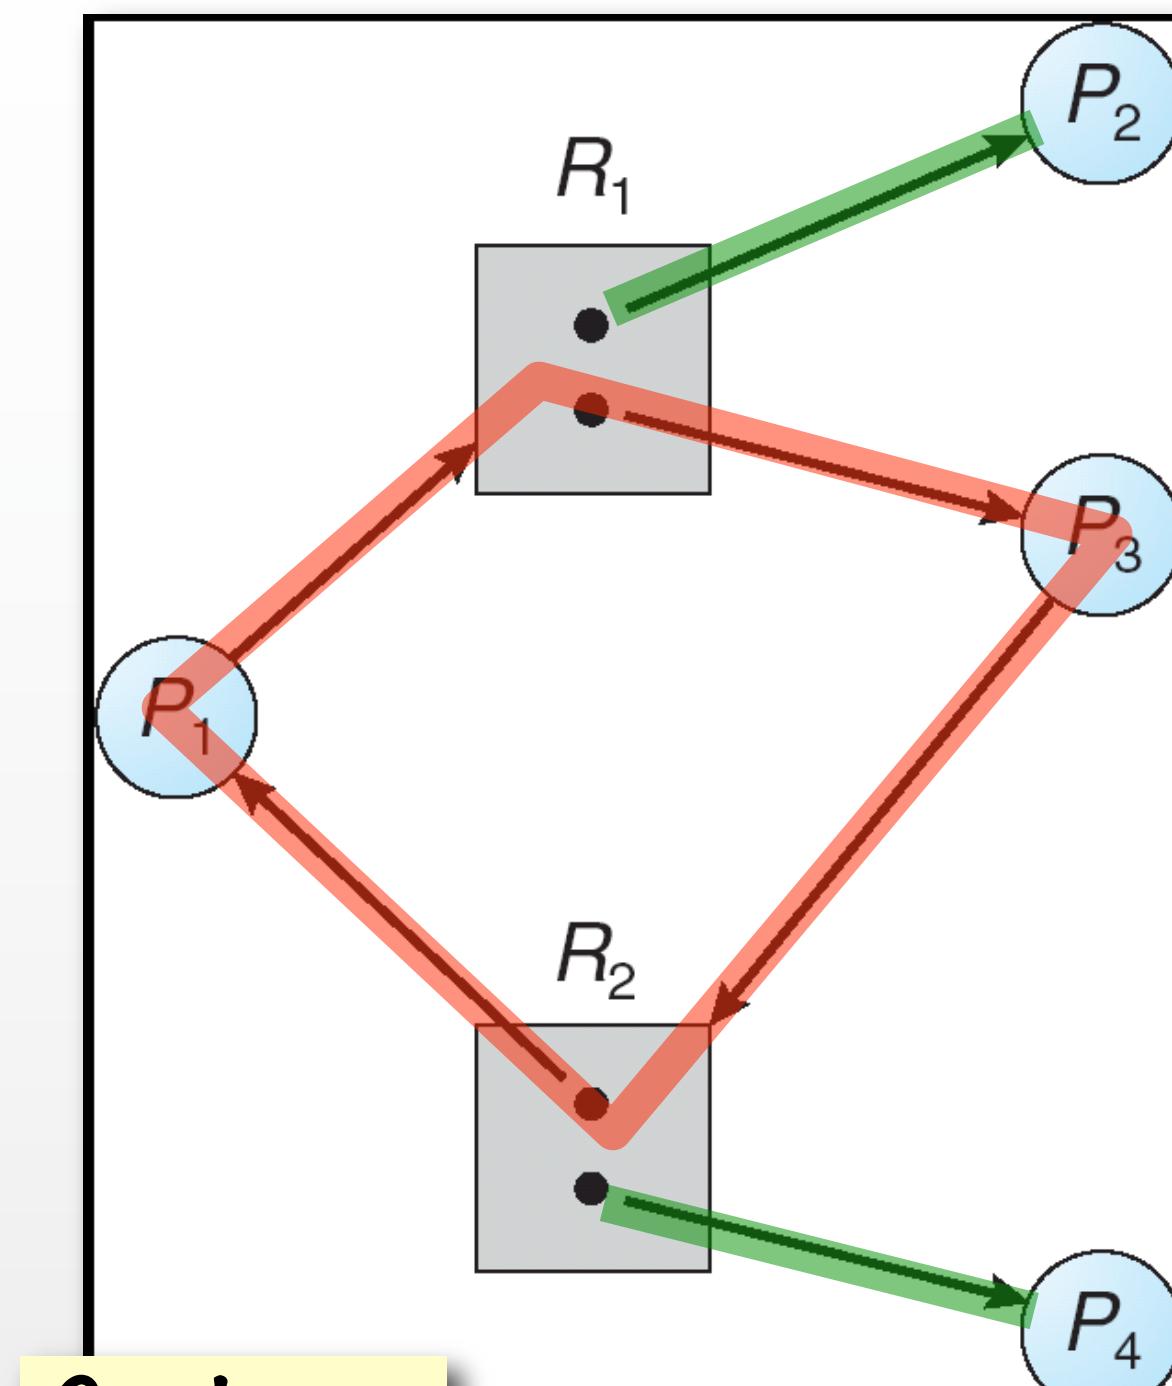
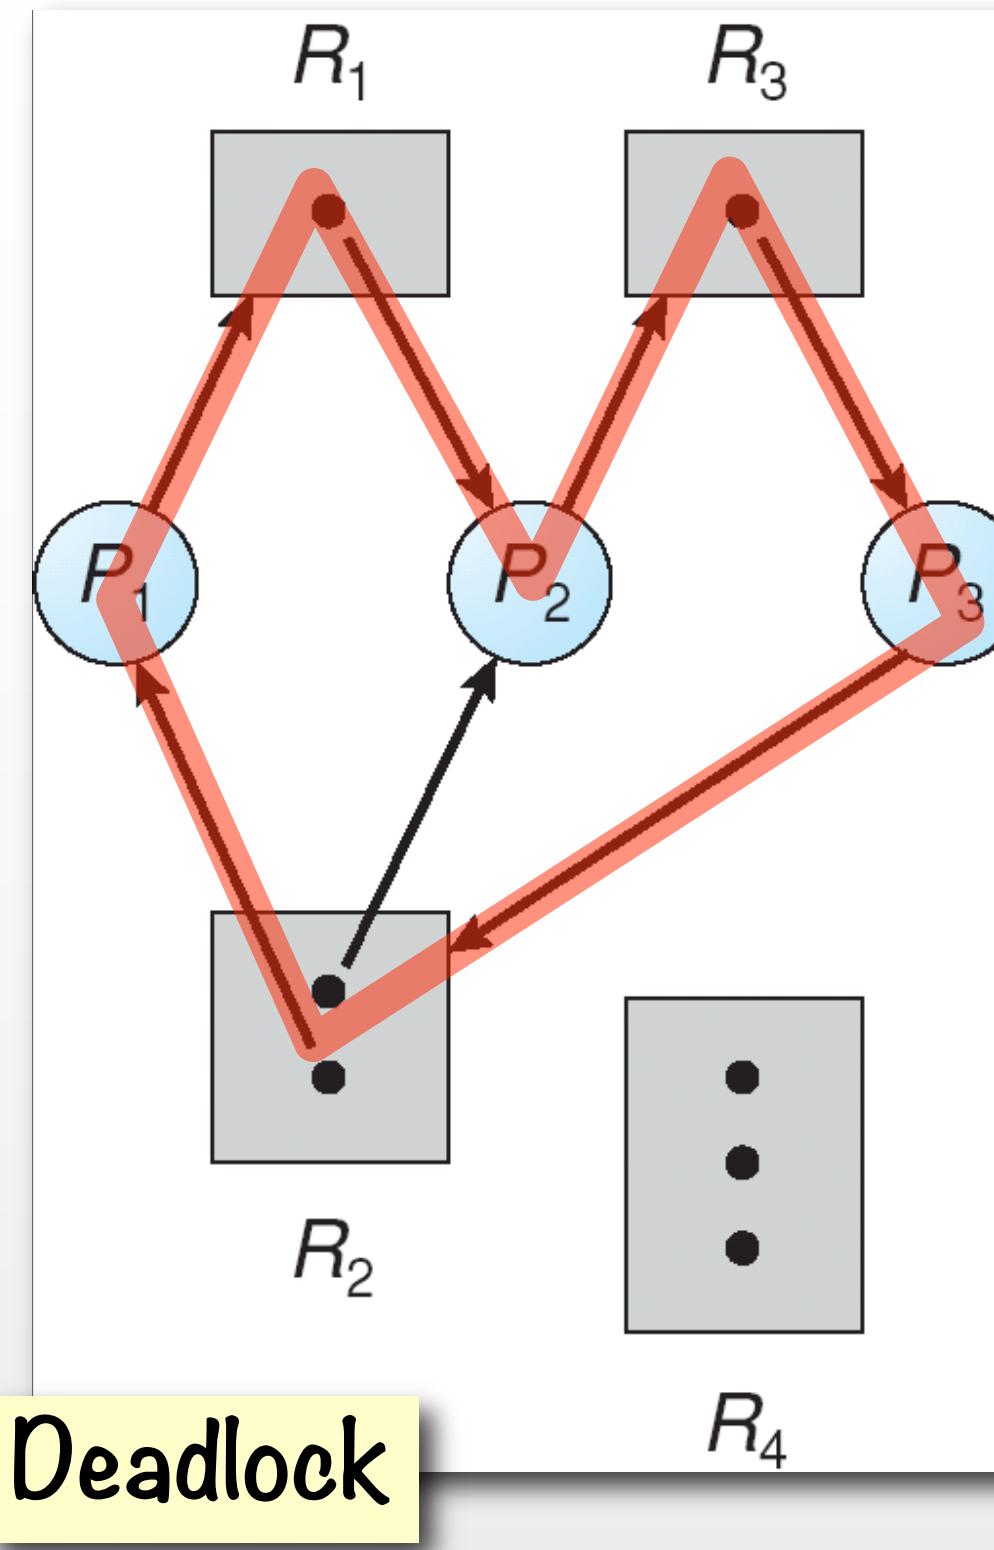
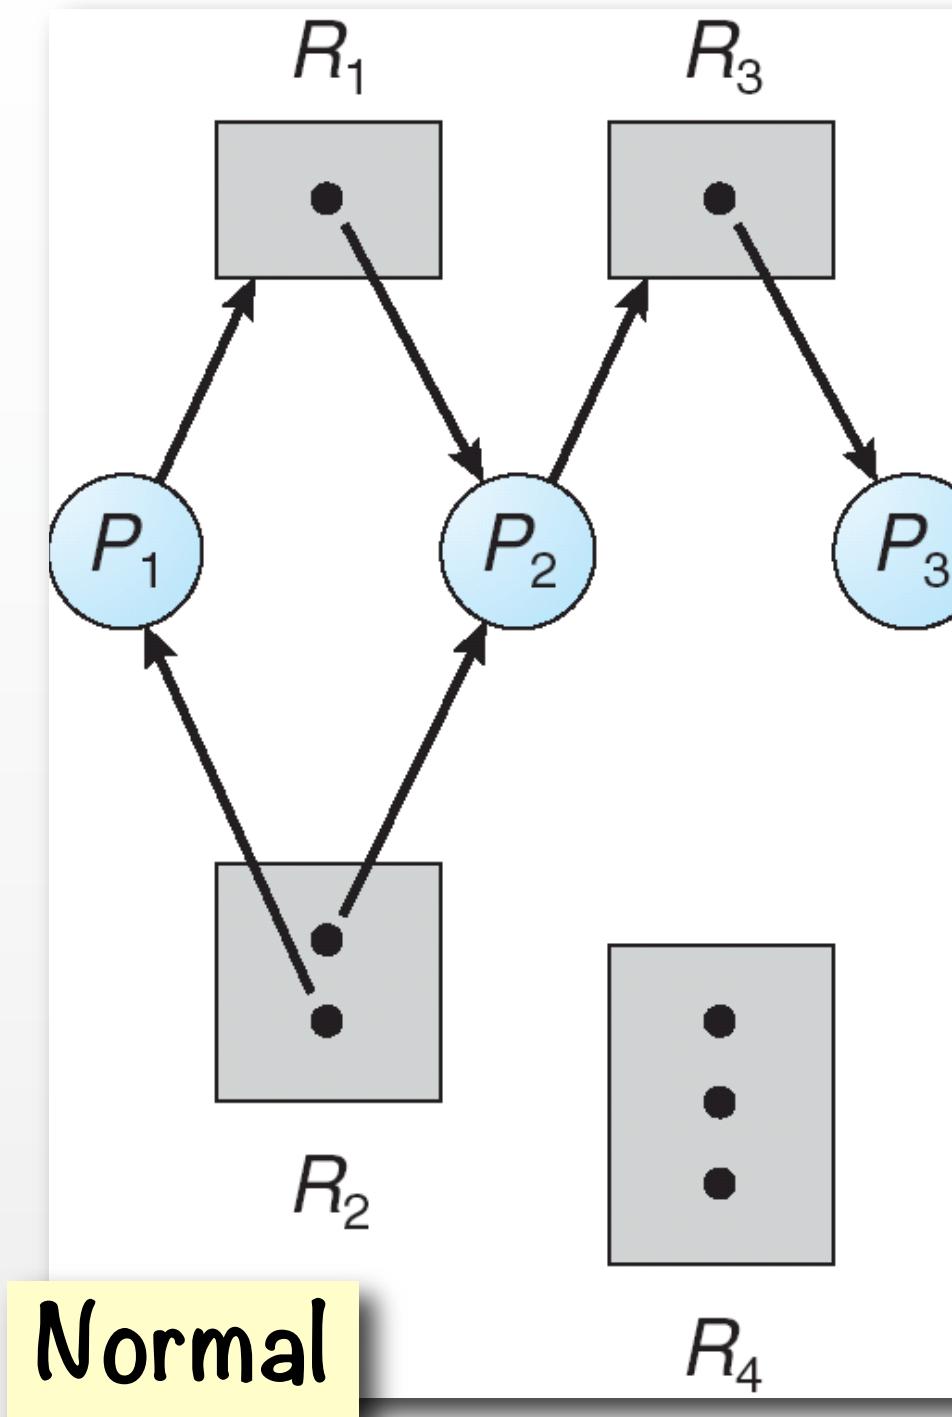
# Resource-Allocation Graph

- A set of vertices  $V = V_p \cup V_R$  and a set of edges  $E = E_R \cup E_A$  such that
  - $V$  is partitioned into two types:
    - $V_p = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the processes in the system.
    - $V_R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system.
  - **request edge**  $P_i \rightarrow R_j$
  - $E_R = \{ \dots, (P_i, R_j), \dots \}$
- **assignment edge**  $R_j \rightarrow P_i$
- $E_A = \{ \dots, (R_j, P_i), \dots \}$

- Graphically:



# Examples of a Resource Allocation Graph

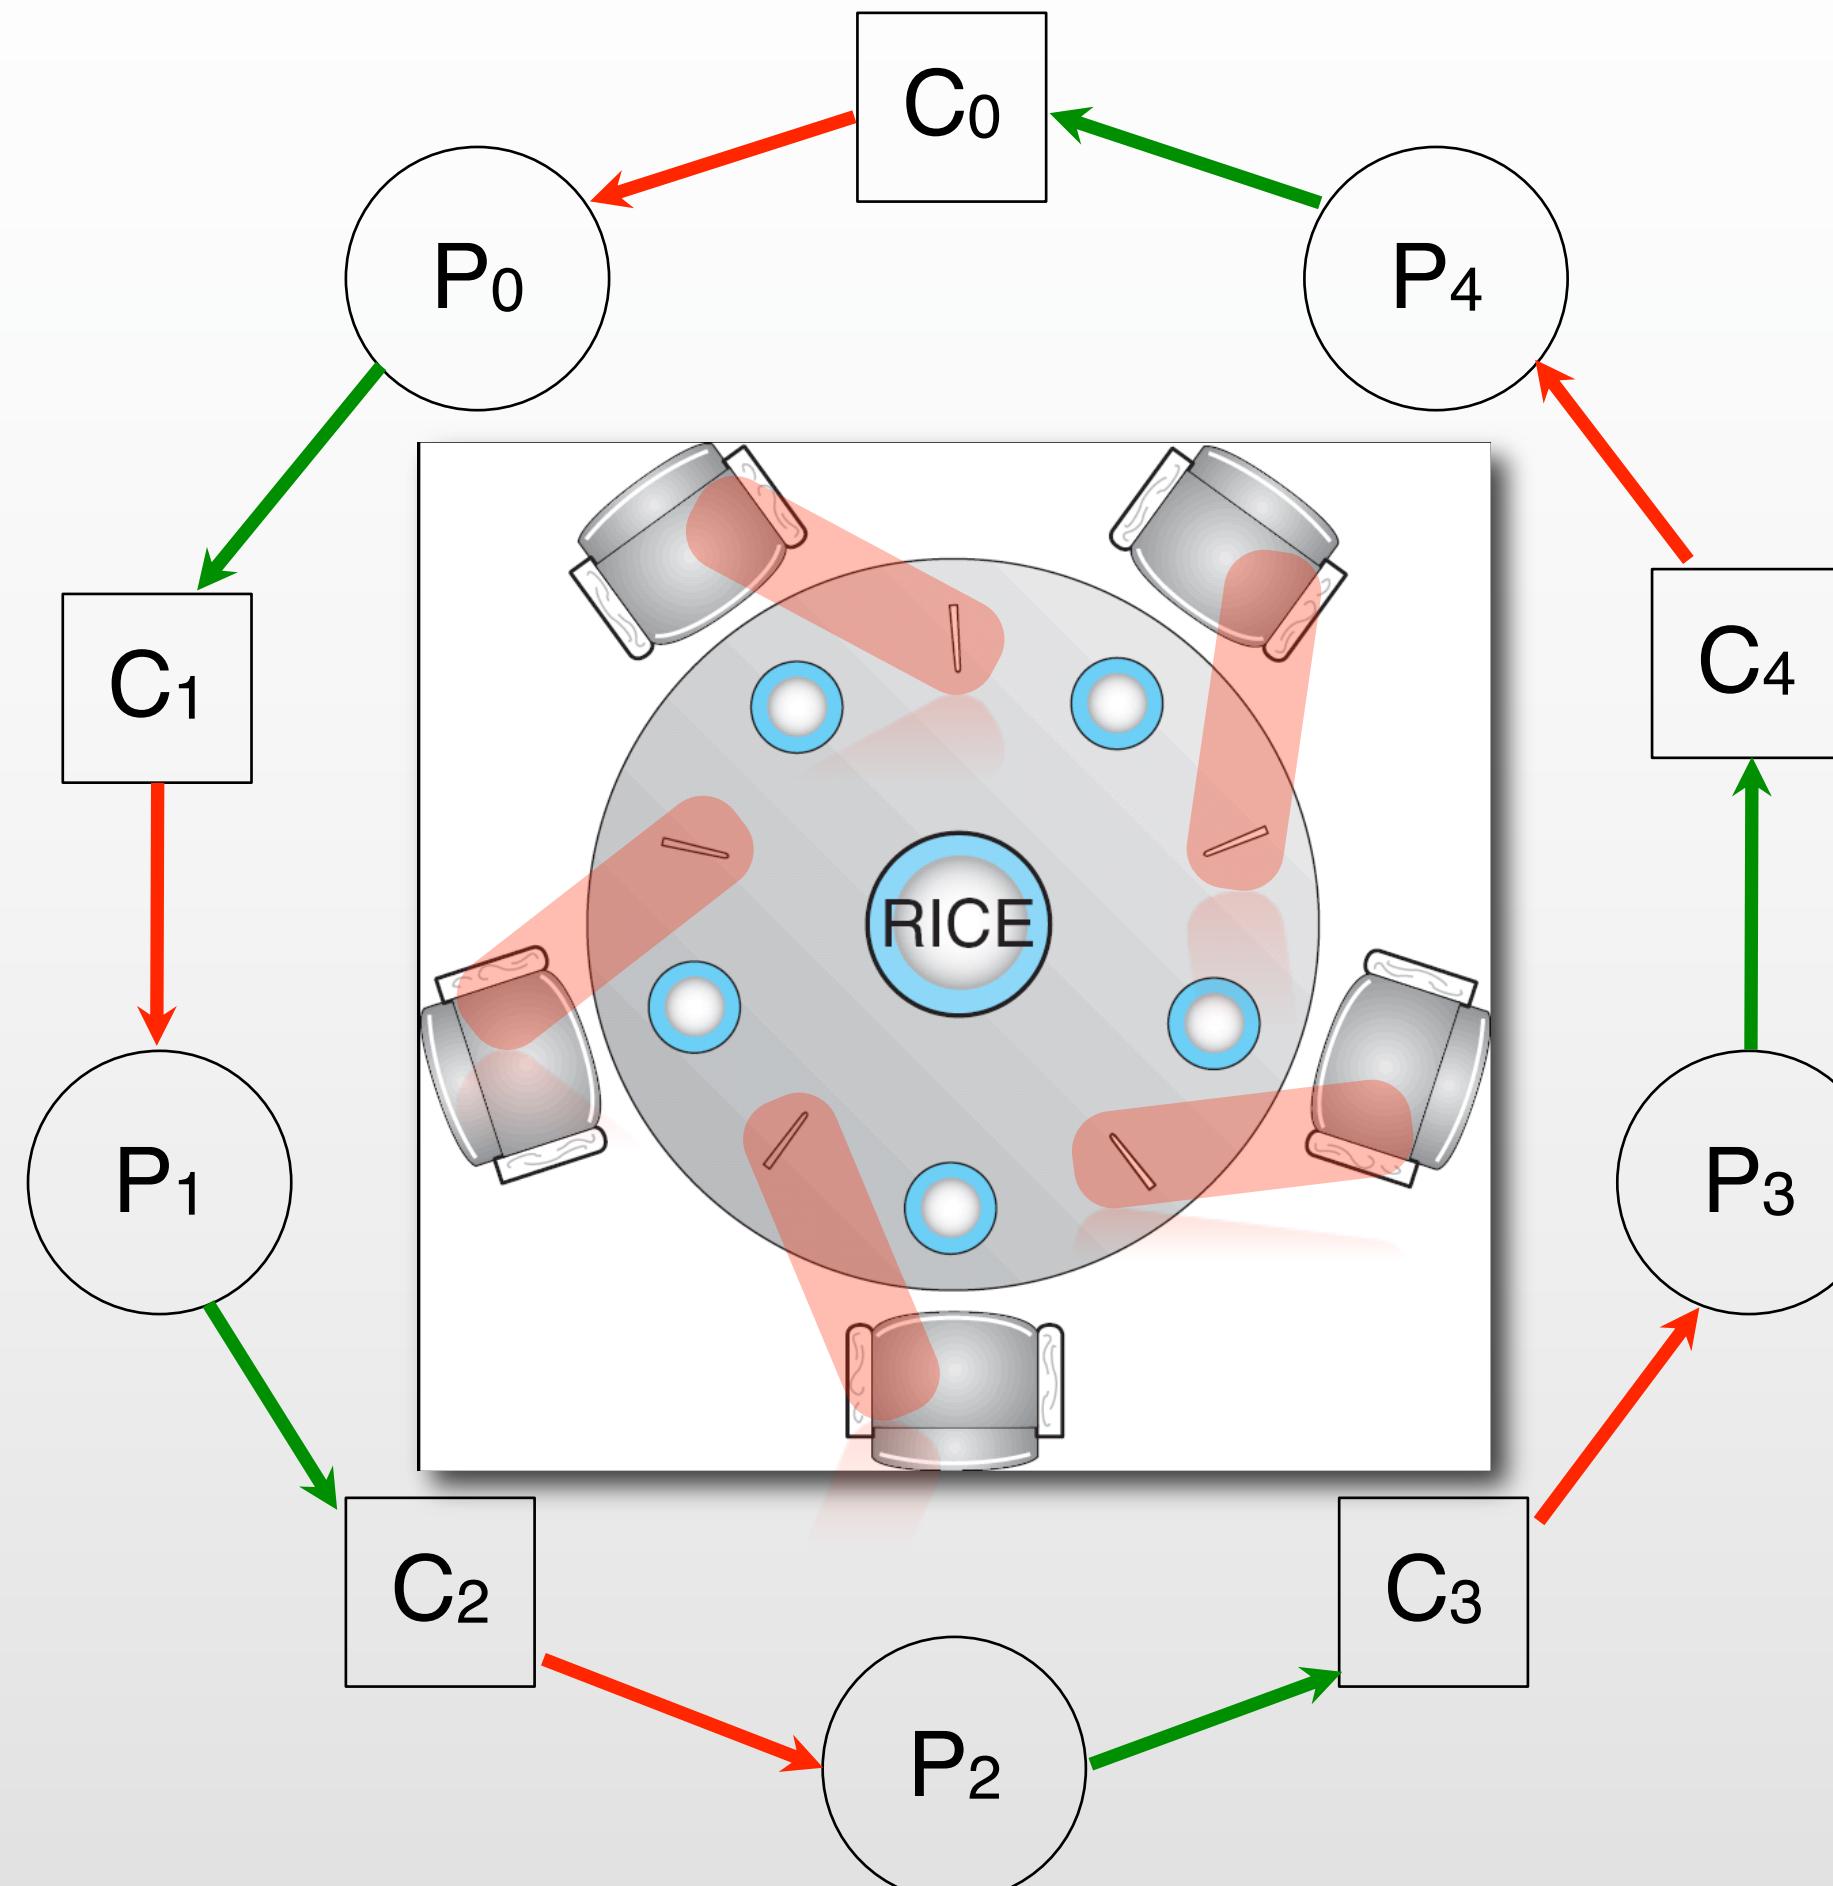


- If graph contains no cycles  $\Rightarrow$  no deadlock.
- If graph contains a cycle  $\Rightarrow$ 
  - if only one instance per resource type, then deadlock,
  - if several instances per resource type, possibility of deadlock.



# Deadlock with Dining Philosophers

- It was possible for a circle to form in the resource allocation graph.



```
do {  
    wait ( chopstick[i] );  
  
    // possible context switch  
  
    wait ( chopstick[ (i + 1) % 5] );  
  
    // eat (CRITICAL SECTION)  
  
    signal ( chopstick[i] );  
    signal ( chopstick[ (i + 1) % 5] );  
} while (TRUE);
```

- Note that we are using  $P_0, P_1, \dots, P_{n-1}$  to simplify the coding

# Methods for Handling Deadlocks



## • **Deadlock prevention**

- Engineer the system in a way that eliminates deadlocks completely.

## • **Deadlock avoidance**

- Manage system so it avoids “dangerous” states (i.e., ones that may lead to a deadlock)

## • **Deadlock recovery**

- Allow the system to enter a deadlock state but provide a method for deadlock detection and recovery.

## • Most common approach:

Ignore the problem and pretend that deadlocks never occur

- **used by most operating systems**, including UNIX
- leaves it up to the programmers to take care of the problem
- it's **unacceptable in certain situations**
  - e.g., real-time mission-critical systems
    - imagine a melting nuclear plant whose control system is in a deadlock!

# Deadlock Prevention



- Recall that all of the following conditions are **necessary** for a deadlock to occur :

- (1) mutual exclusion,
- (2) hold and wait,
- (3) no preemption, and
- (4) circular wait



To prevent a deadlock, restrain the ways request can be made to remove at least one of the conditions necessary for a deadlock.

- Let's try then to design a system with one of the conditions removed!

# Deadlock Prevention



- **1. Breaking the mutual exclusion condition**
  - cannot control this, because some resources are intrinsically non-sharable and require exclusive access
  - examples?
  
- **2. Breaking the hold and wait condition**
  - the system may be designed to demand that whenever a process requests a resource, it does not hold any other resources
    - EITHER require a process to request and be allocated all its resources before it begins execution,
    - OR, allow a process to request resources only when it holds none
      - the process must release held resources and then re-acquire them together with any desired additional resources
  - low resource utilization
  - starvation possible

recall the DP solution  
with `try_lock()`

# Deadlock Prevention (Cont.)



- 3. **Breaking the no preemption condition**
- the system is allowed to take away resources from a process forcefully
  - unlike (2), this is involuntary system action; a process may run without some resources that it does not immediately needs
  - if a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held by that process are forcefully released
    - preempted resources are added to the list of resources for which the process is waiting
  - process will be restarted only when both the old resources as well as the new ones that it requested are available
- low resource utilization
- starvation possible

# Deadlock Prevention (Cont.)



## 4. Breaking the circular wait condition

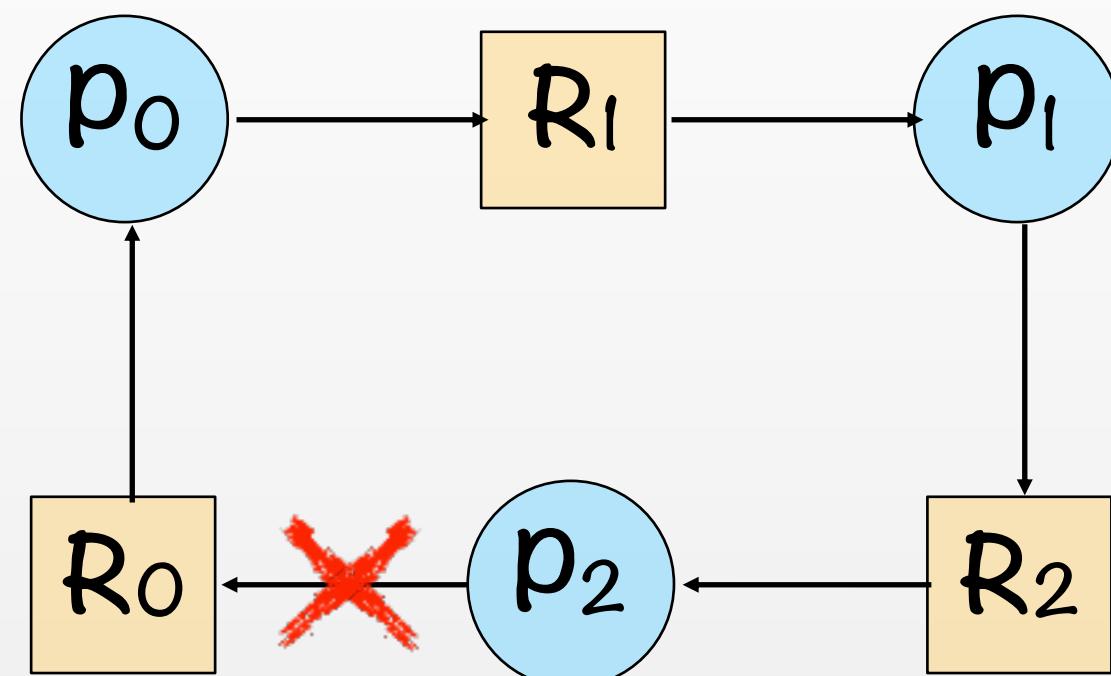
- impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration
  - we create tiers of resources

$$F: \mathcal{R} \rightarrow N$$

- where  $F$  is a **mapping function** from  $\mathcal{R}$  to  $N$  (**enumeration**)
  - $\mathcal{R} = \{R_1, R_2, \dots, R_m\}$  is a domain of resources
  - $N$  is the set of natural numbers
- if a process holds  $R_i$  then it can request  $R_j$  only if  $F(R_i) < F(R_j)$ 
  - alternatively,  
we can insist that to gain access to  $R_j$  process must release all  $R_k$  such that  $F(R_k) \geq F(R_j)$

# Resource Ordering: Prevention Proof

- To prove by contradiction, assume that we followed the rules but got a deadlock
  - therefore, we have a cycle, so for any  $i$ ,  $p_i$  waits for some  $R_{i+1}$  that is allocated to  $p_{i+1}$ ,



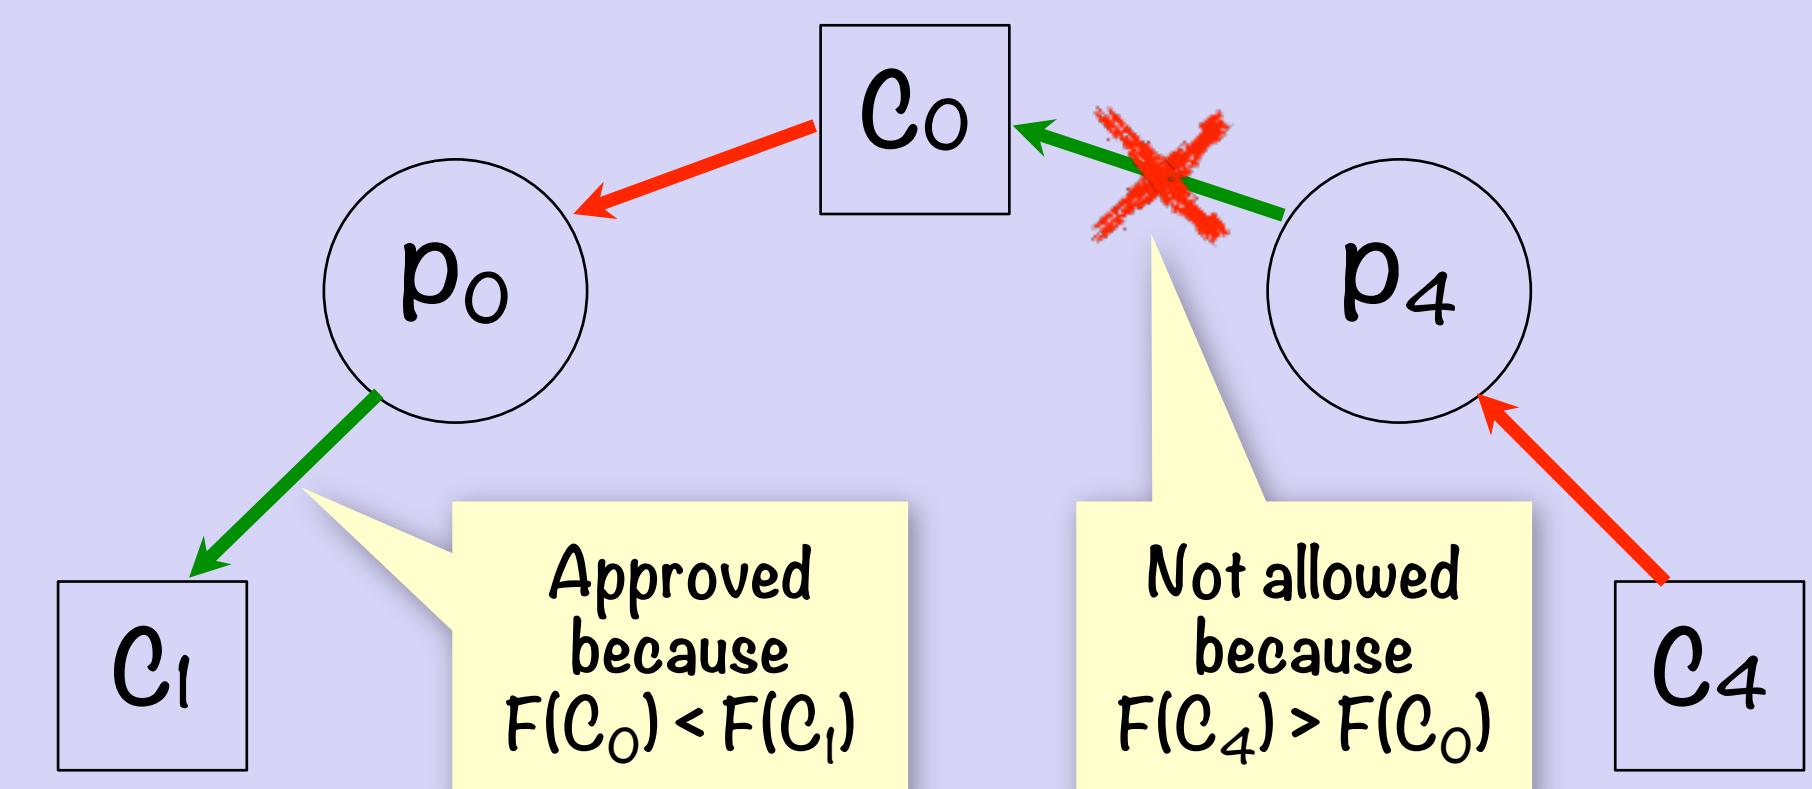
- however, for  $p_{i+1}$  to be allowed to request  $R_{i+1}$  it must have held that  $F(R_i) < F(R_{i+1})$ , so by Mathematical induction and using modulo- $n$  increment it holds that  $F(R_0) < F(R_1) < \dots < F(R_{n-1}) < F(R_0)$
- and that's not possible; hence, the proof

## Example

If we imposed the following order in our Dining Philosophers example

$$\begin{aligned}F(C_0) &= 1 \\F(C_1) &= 2 \\F(C_2) &= 3 \\F(C_3) &= 4 \\F(C_4) &= 5\end{aligned}$$

the deadlock would not happen because  $p_4$  has to give up  $C_4$  if it wants to request  $C_0$





# Deadlock Avoidance

- Deadlock prevention scheme is very expensive
  - the OS would be very restrictive, and therefore inefficient
    - however, it may be necessary in certain mission critical systems
- **Deadlock avoidance** takes a bit less restrictive approach
  - it requires that the system is provided with some additional a priori information about a process available
  - the system uses this information to make decision about resource allocation each time any process requests a resource
- Simplest and most useful model requires that each process declares the maximum number of resources of each type that it may need.
  - The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition.
  - **Resource-allocation state** is defined by the number of available and allocated resources, and the maximum demands of the processes.

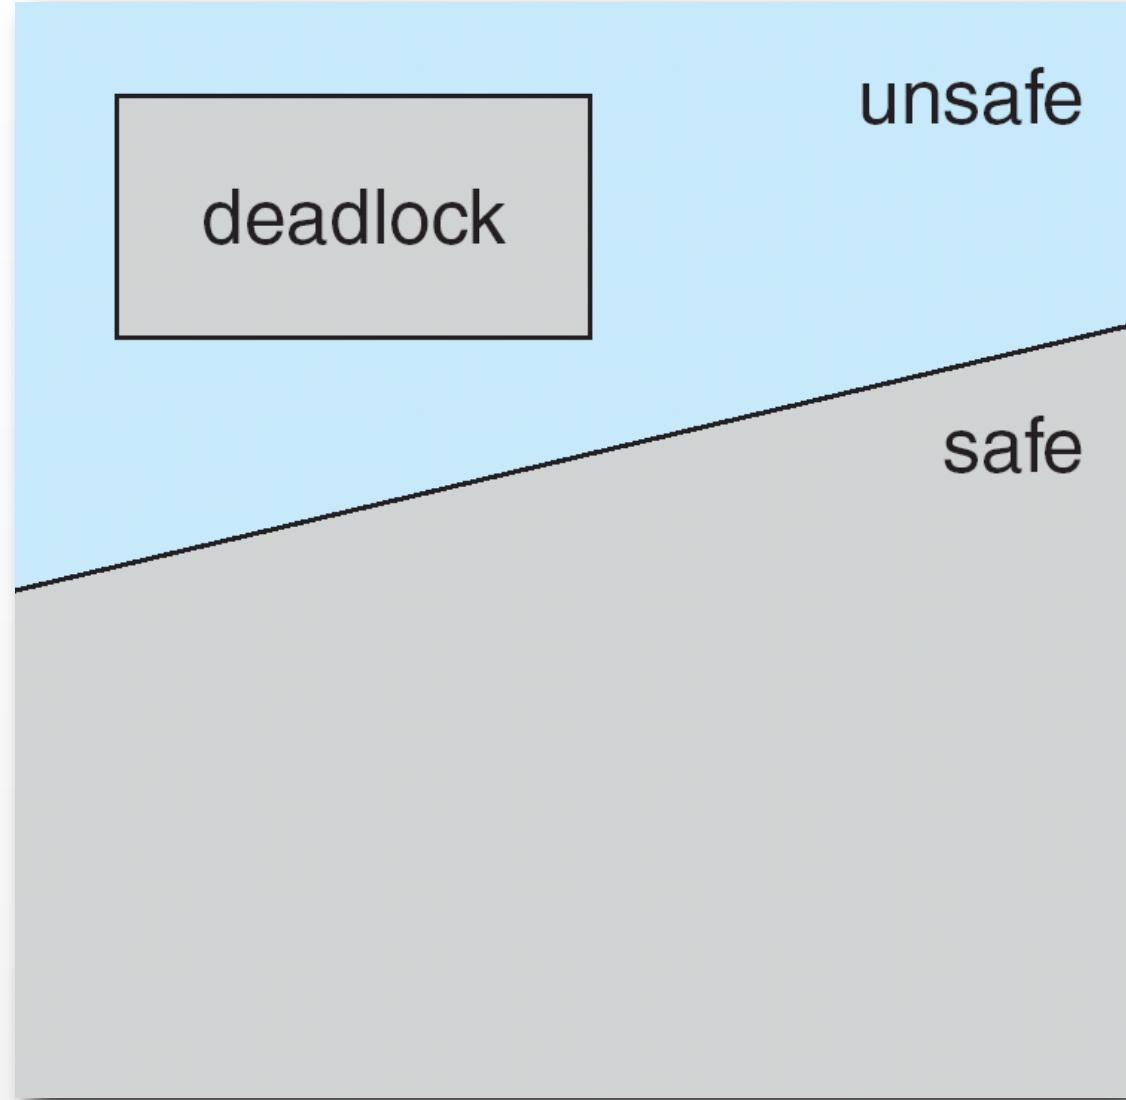
# Safe State



- When a process requests an available resource, OS must decide if immediate allocation leaves the system in a safe state

- If a system is in safe state  $\Rightarrow$  no deadlocks.
- If a system is in unsafe state  $\Rightarrow$  possibility of deadlock.
- Avoidance  $\Rightarrow$  ensure that a system will never enter an unsafe state.

- System is in **safe state** if there exists a scheduling sequence  $\langle p_1, p_2, \dots, p_n \rangle$  of all the processes in the system such that for each  $p_i$ , the resources that  $p_i$  may still request can be satisfied by currently available resources and resources held by all the  $p_j$  for  $j < i$  (i.e., the processes that precede  $p_i$  in the scheduling sequence; i.e., the ones that will terminate earlier releasing the resources that they hold).
  - When the system is in a safe state:
    - If  $p_i$  resource needs are not immediately available, then  $p_i$  can wait until sufficient number of  $p_j$  (for  $j < i$ ) have finished. The safe state definition ensures that at some point  $p_i$  will be able to get all that it needs to continue
    - When  $p_j$  is finished,  $p_i$  can obtain needed resources, execute, return allocated resources, and terminate.
    - In turn, when  $p_i$  terminates,  $p_k$  for some  $k > i$  can obtain its needed resources, and so on.





# Example of Safe State and Unsafe State

- Let the system with 13 devices have  $P_0$ ,  $P_1$ , and  $P_2$  processes with the shown needs

Process	Claims	Current allocation
$P_0$	10	5
$P_1$	4	2
$P_2$	9	2

note that there are  
4 devices free

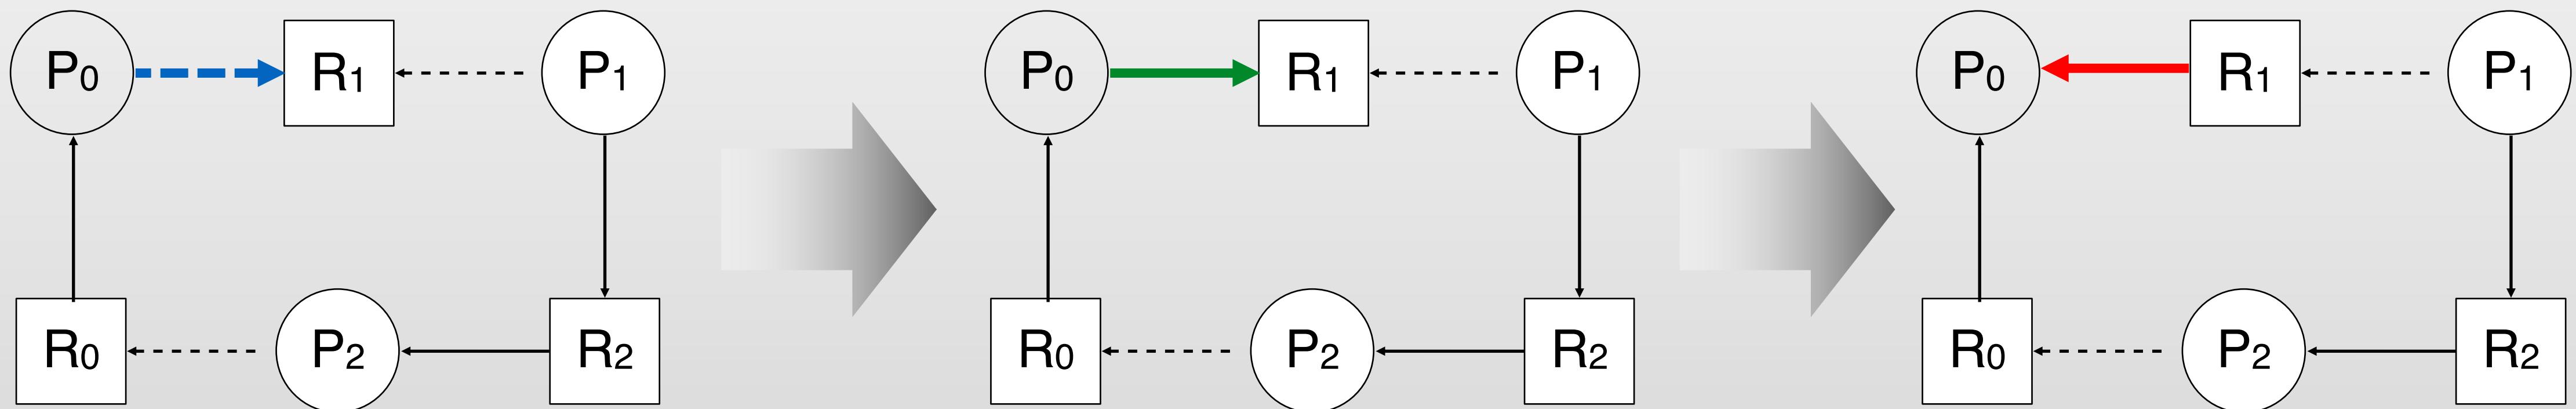
- $\langle P_1, P_0, P_2 \rangle$  sequence satisfies the safe state requirements
  - $P_1$  gets extra 2 devices (while 2 are left free) and then terminates releasing 4 devices; after that, we have 6 devices free in the pool
  - $P_0$  can now execute with 10 devices (getting 5 from the free pool leaving 1 still available), then it terminates and releases the resources; so we have 11 free
  - $P_2$  is now able to acquire 7 additional devices (up to the needed 9), execute and terminate
- Now, let's consider that  $P_2$  is allocated two more devices first, the system **may not be able** to fulfill future demands, so this is unsafe state; here is one bad case:
  - with  $P_2$  having 2 more resources (so it has 4 now), only 2 are still free
  - only  $P_1$  can be allocated all claimed resources ( $2+2=4$ ); none is left free
  - when  $P_1$  terminates and returns the resources, there are 4 free
  - now neither  $P_0$  (10, 5) nor  $P_2$  (9, 4) can continue if they need the maximum number of claimed resources; both  $P_0$  and  $P_2$  may need 5 extra

Note that unsafe state **may or may not** lead to a deadlock. Nevertheless, to prevent the worst case, we need to avoid such states.



# Avoidance with Resource-Allocation Graph

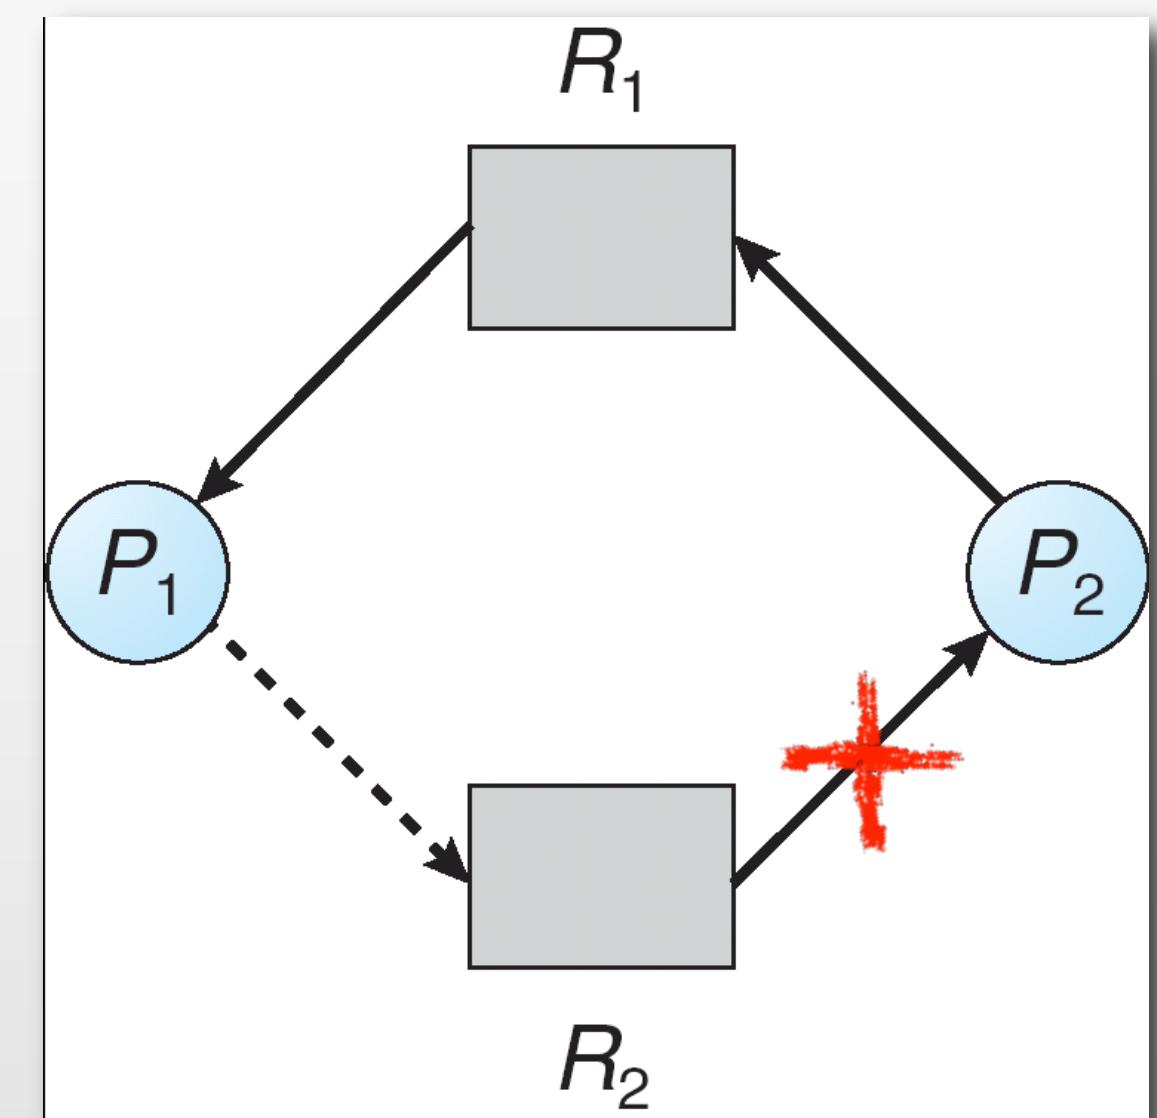
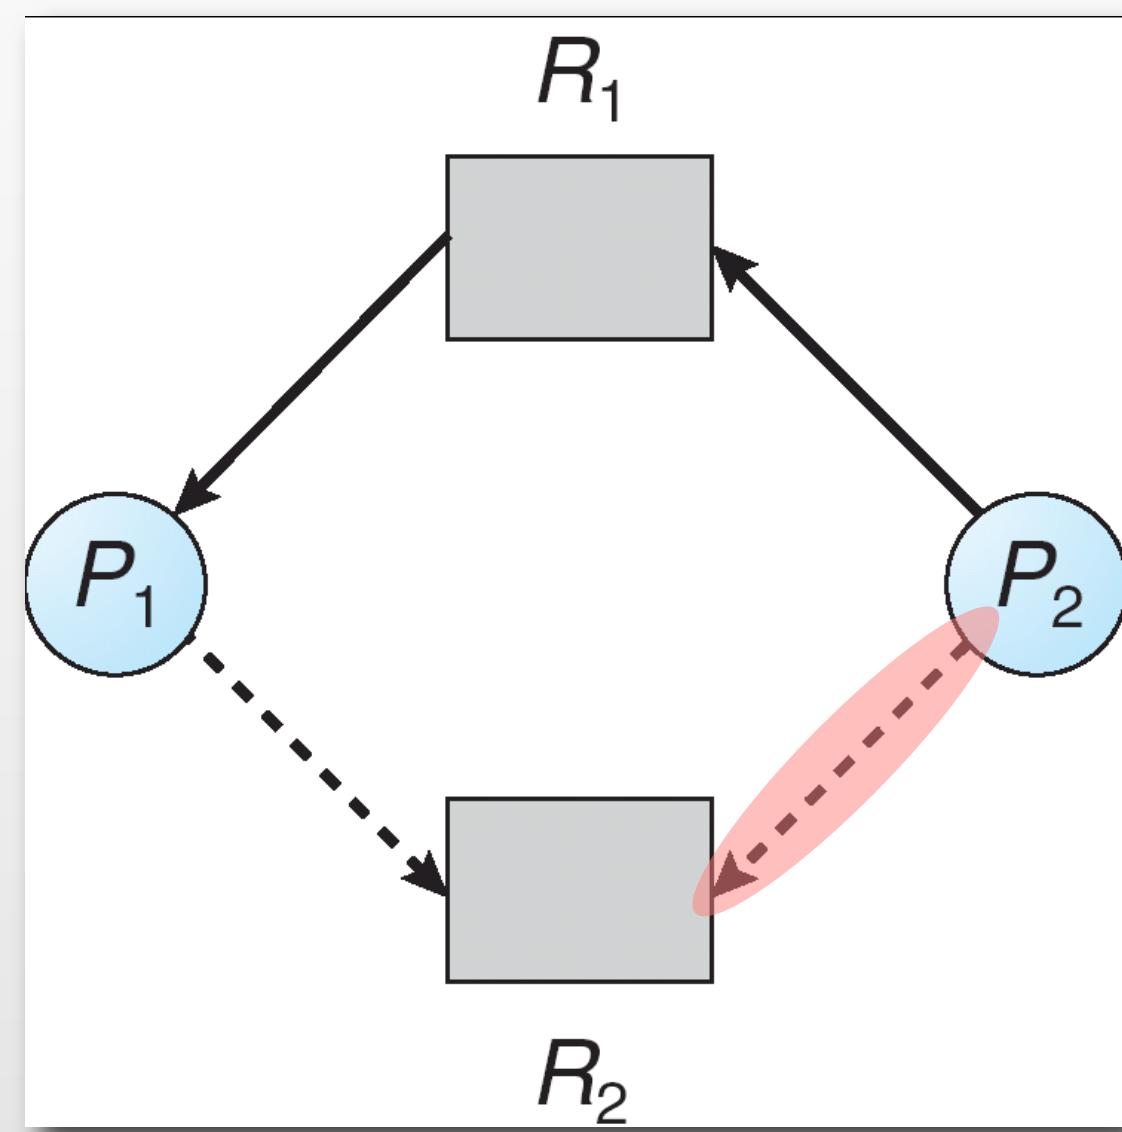
- First, we augment resource allocation graph by adding a **claim edge**  $P_i \rightarrow R_j$  indicating that process  $P_i$  may request resource  $R_j$  represented by a dashed line.
- **Claim edge** converts to a **request edge** when a process indeed requests a resource.
- **Request edge** is converted to an **assignment edge** when the resource is allocated to the process.
- When a resource is released by a process, **assignment edge** reconverts to a **claim edge**.
- Resources must be claimed a priori in the system.



# Avoidance Algorithm Example



- Suppose that process  $P_i$  requests a resource  $R_j$
- The request can be granted only if converting the request edge to an assignment edge does not result in the formation of a cycle in the resource allocation graph
- In the following example, allocating  $R_2$  to  $P_2$  is therefore not allowed!



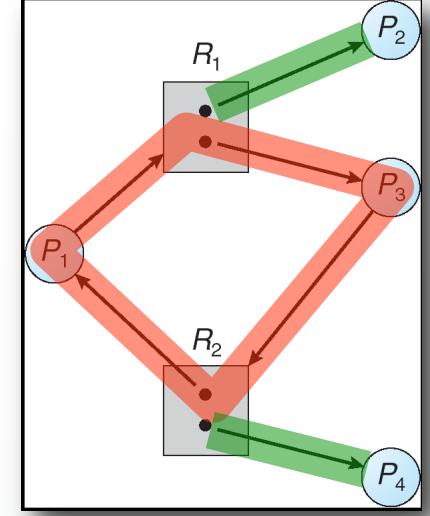
Note that there is no deadlock yet; just the possibility!

- Detecting a cycle in a graph is of  $O(n^2)$  complexity

# Avoidance with Banker's Algorithm



- Resource allocation graph can be used only with single instance resources
  - since – as we saw earlier - detecting a cycle with a multi-instance resource does not necessarily constitute a deadlock danger
- The **banker's algorithm** (invented by Dijkstra) is suitable for cases with multiple instances of resources.
  - the name comes from the fact that banks should use a similar algorithm to allocate funds (e.g., through loans) in such a way that they always can satisfy the current needs of their customers (be solvent)
- Assumptions:
  - each process must **claim maximum need for resources a priori**,
    - including the resources that may not be used at all in some cases
  - when a process requests a resource it **may have to wait**,
  - when a process gets all its resources it **must return them in a finite amount of time**.
- The banker's algorithm is a more complex deadlock avoidance scheme than the resource-allocation graph method
  - more overhead for the system



# Data Structures for the Banker's Algorithm



- Let
  - $n$  = number of processes  $P_i$  with  $0 \leq i < n$ , and
  - $m$  = number of resources  $R_j$  with  $0 \leq j < m$ .
- Define the following:
  - **Available:**
    - Vector of length  $m$ . If  $\text{Available}[j] = k$ , there are  $k$  instances of resource type  $R_j$  available.
  - **Max:**
    - $n \times m$  matrix. If  $\text{Max}[i,j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$ .
  - **Allocation:**
    - $n \times m$  matrix. If  $\text{Allocation}[i,j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$ .
  - **Need:**
    - $n \times m$  matrix. If  $\text{Need}[i,j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task.
    - Note that  $\text{Need}[i,j] = \text{Max}[i,j] - \text{Allocation}[i,j]$ .
- We can also use a shorthand notation (using **vector arithmetic**):
  - e.g.,  $\text{Need}[i] = \text{Max}[i] - \text{Allocation}[i]$ 
    - where  $i$  is a process number, so these vectors are rows in the matrices corresponding to  $P_i$
- Define a relation  $\leq$  as:  $X \leq Y$  if for every  $i$ :  $X[i] \leq Y[i]$ 
  - e.g.,  $(3, 2, 4) < (4, 3, 9)$ , and  $(3, 2, 4) \overset{\text{?}}{\leq} (2, 5, 7)$
  - furthermore,  $X < Y$  if  $X \leq Y$  and  $X \neq Y$

Assumption:  
processes do not lie!

Undetermined  
“ $\leq$ ” imposes only a partial order

# Resource Request Algorithm for Process $P_i$



- Determines if a resource request from process  $P_i$  is acceptable
- Let **Request**[ $i$ ] be a request vector for process  $P_i$ .
  - If **Request**[ $i, j$ ] =  $k$  then process  $P_i$  wants  $k$  instances of resource type  $R_j$ .
    - as before, we drop  $j$  and work with vectors
- Step 1
  - If **Request**[ $i$ ]  $\leq$  **Need**[ $i$ ] go to step 2. Otherwise, raise error condition, since process  $P_i$  has exceeded its maximum claim.
- Step 2
  - If **Request**[ $i$ ]  $\leq$  **Available**, go to step 3. Otherwise process  $P_i$  must wait, since resources are not available.
- Step 3
  - Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:
    - save **Available**, **Allocation**[ $i$ ], and **Need**[ $i$ ]
    - **Available** = **Available** – **Request**[ $i$ ]
    - **Allocation**[ $i$ ] = **Allocation**[ $i$ ] + **Request**[ $i$ ]
    - **Need**[ $i$ ] = **Need**[ $i$ ] – **Request**[ $i$ ]
  - Now run the safety algorithm (see next slide)
    - If safe state  $\Rightarrow$  the requested resources can be allocated to process  $P_i$ .
    - If state unsafe  $\Rightarrow$   $P_i$  must wait, and the old resource-allocation state is restored from the saved values

this contradicts the assumption that processes don't lie



# Safety Algorithm in Banker's Algorithm

- Recall: to determine if a certain state is safe we need to find a sequence in which processes can execute; the following algorithm will find a sequence if one exists
- Step 1.
  - Let **Work** and **Finish** be vectors of length **m** and **n**, respectively. Initialize:
    - Work** = **Available**
    - Finish**[*i*] = **false** for *i* = 0, 1, ..., *n*-1.
- number of resources
- number of processes
- Step 2.
  - Find an *i* such that both:
    - (a) **Finish**[*i*] = **false**
    - (b) **Need**[*i*] ≤ **Work**
  - If no such *i* exists, go to step 4.
- Step 3.
  - we pretend that the process *i* is done
  - Work** = **Work** + **Allocation**[*i*]
  - Finish**[*i*] = **true**
  - go to step 2.
- Step 4.
  - If  $\forall i \text{ } \text{Finish } [i] == \text{true}$  → the system is in a safe state.
- The complexity is  $O(m^*n^2)$

This is not good...

## NOTES

- This is just detection of a safe state. See the next slide for a complete algorithm for resource allocation.
- All operations are pretended; neither actual allocation nor termination is carried out.
- The order is not important! The amount of available resources is increased by any process that “terminates”, so we always improve the situation.

# State Safety Example



- $n = 5$  processes  $P_0$  through  $P_4$
- $m = 3$  resource types:
  - A (10 instances)
  - B (5 instances), and
  - C (7 instances).

	Allocation			Max		
	A	B	C	A	B	C
$P_0$	0	1	0	7	5	3
$P_1$	2	0	0	3	2	2
$P_2$	3	0	2	9	0	2
$P_3$	2	1	1	2	2	2
$P_4$	0	0	2	4	3	3

Available	A	B	C
	3	3	2

	Need		
	A	B	C
$P_0$	7	4	3
$P_1$	1	2	2
$P_2$	6	0	0
$P_3$	0	1	1
$P_4$	4	3	1

Work	A	B	C
	3	3	2

- The content of the matrix **Need** is defined to be **Max – Allocation**.

- The system is in a safe state since the sequences such as  $\langle P_1, P_3, P_4, P_2, P_0 \rangle$  and  $\langle P_3, P_4, P_1, P_2, P_0 \rangle$  satisfy safety criteria.



# Resource Request Example

- Given the state shown on the previous slide, can  $p_1$  request  $(1,0,2)$  be granted?

- Check the  $\text{Request} \leq \text{Need}$ 
  - that is,  $(1,0,2) \leq (1,2,2) \Rightarrow \text{true}$

- Check that  $\text{Request} \leq \text{Available}$ 
  - that is,  $(1,0,2) \leq (3,3,2) \Rightarrow \text{true}$
  - so if the Request was granted,  
we would have:

	Allocation			Need			Work		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	7	4	3			
P1	3	0	2	0	2	0			
P2	3	0	2	6	0	0			
P3	2	1	1	0	1	1			
P4	0	0	2	4	3	1			

- Executing safety algorithm shows that sequence  $\langle p_1, p_3, p_4, p_0, p_2 \rangle$  satisfies safety requirement  $\Rightarrow$  Request can be granted.



# More Examples

- Can request for  $(3,3,0)$  by  $P_4$  be granted given the current state as:

	Allocation			Max			Need		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	7	5	3	7	4	3
P1	2	0	0	3	2	2	1	2	2
P2	3	0	2	9	0	2	6	0	0
P3	2	1	1	2	2	2	0	1	1
P4	0	0	2	4	3	3	4	3	1

Work		
A	B	C
3	3	2

- ?
- Can request for  $(0,2,0)$  by  $P_0$  be granted assuming the same current state?

# Deadlock Detection

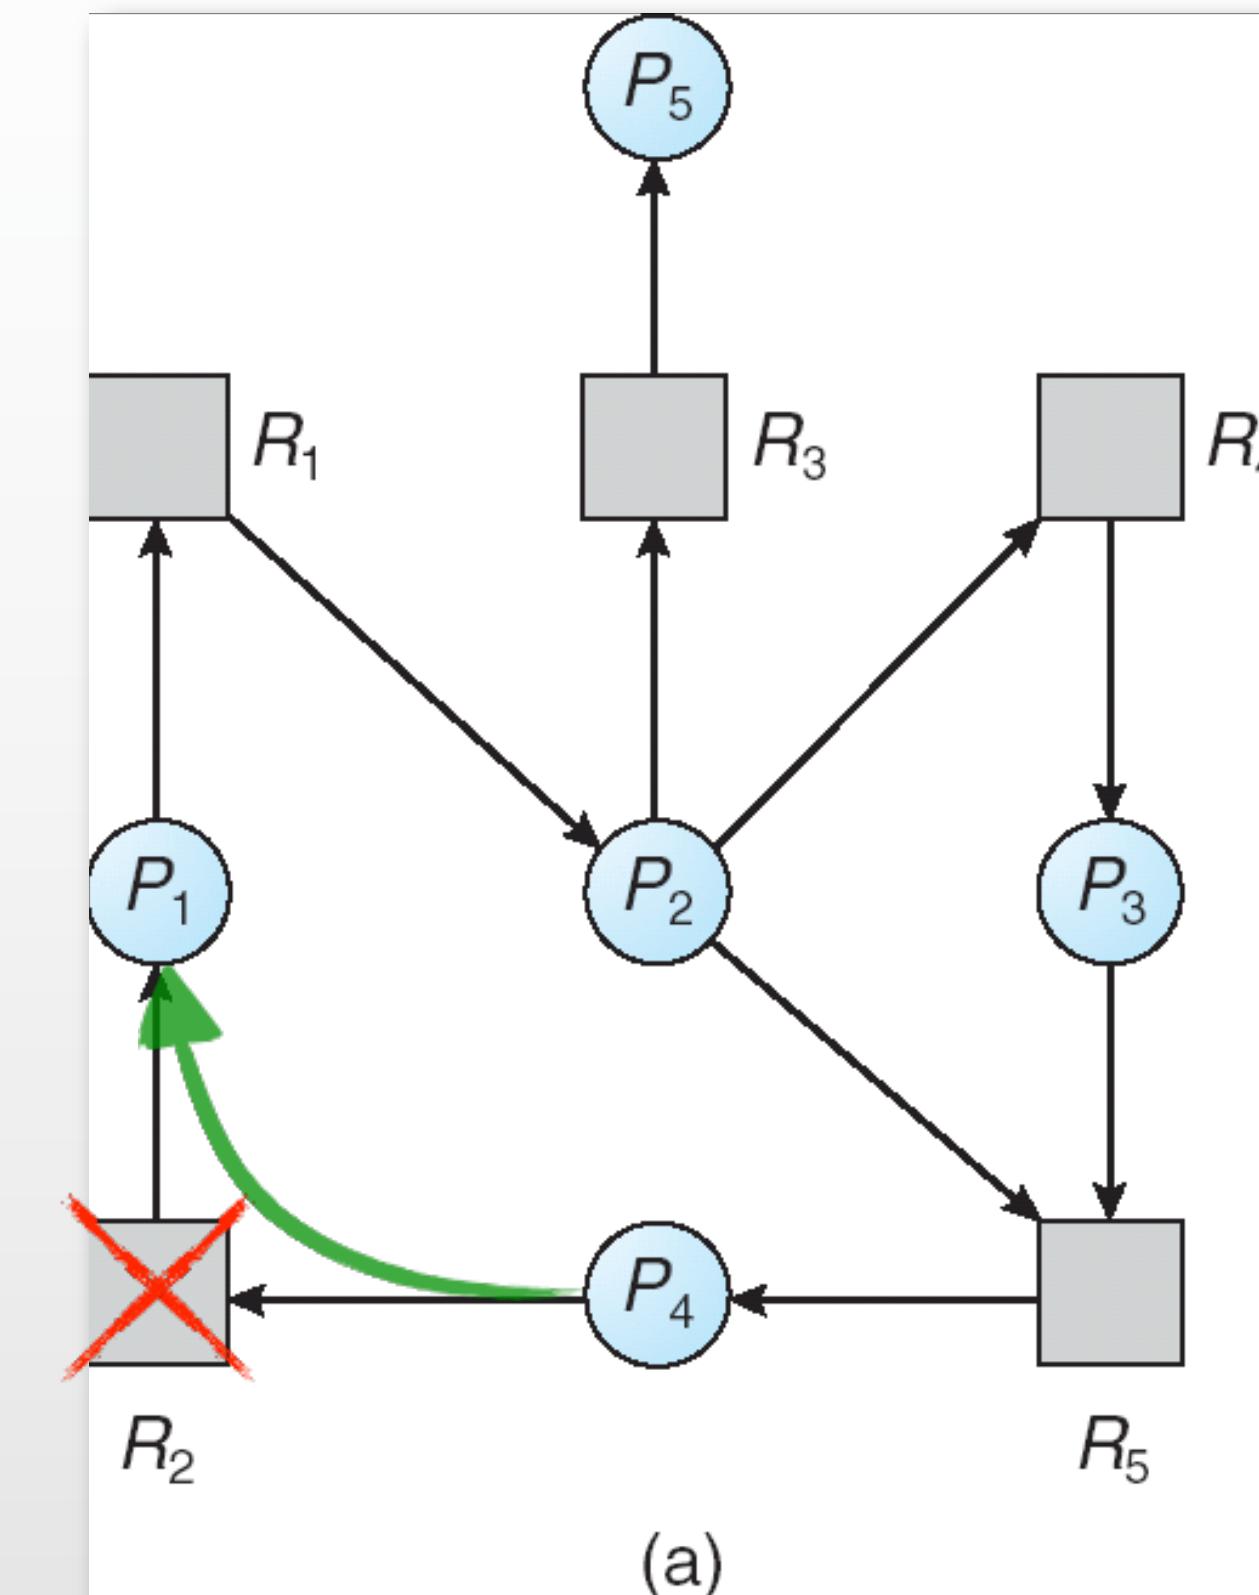


- Prevention and avoidance algorithms are time consuming, so they are used only in most critical systems
- Therefore, usually it is assumed that deadlocks may happen
- If a deadlock occurs, the operating system may be engineered to detect and resolve it
- OS that implements a deadlock detection scheme does not prevent the system from entering a deadlock state
- Instead, it tries to detect potential deadlocks, and if discovered, the OS tries to recover from the problem
  
- For that, the OS needs:
  - **Detection algorithm**
  - **Recovery scheme**

# Single Instance of Each Resource Type

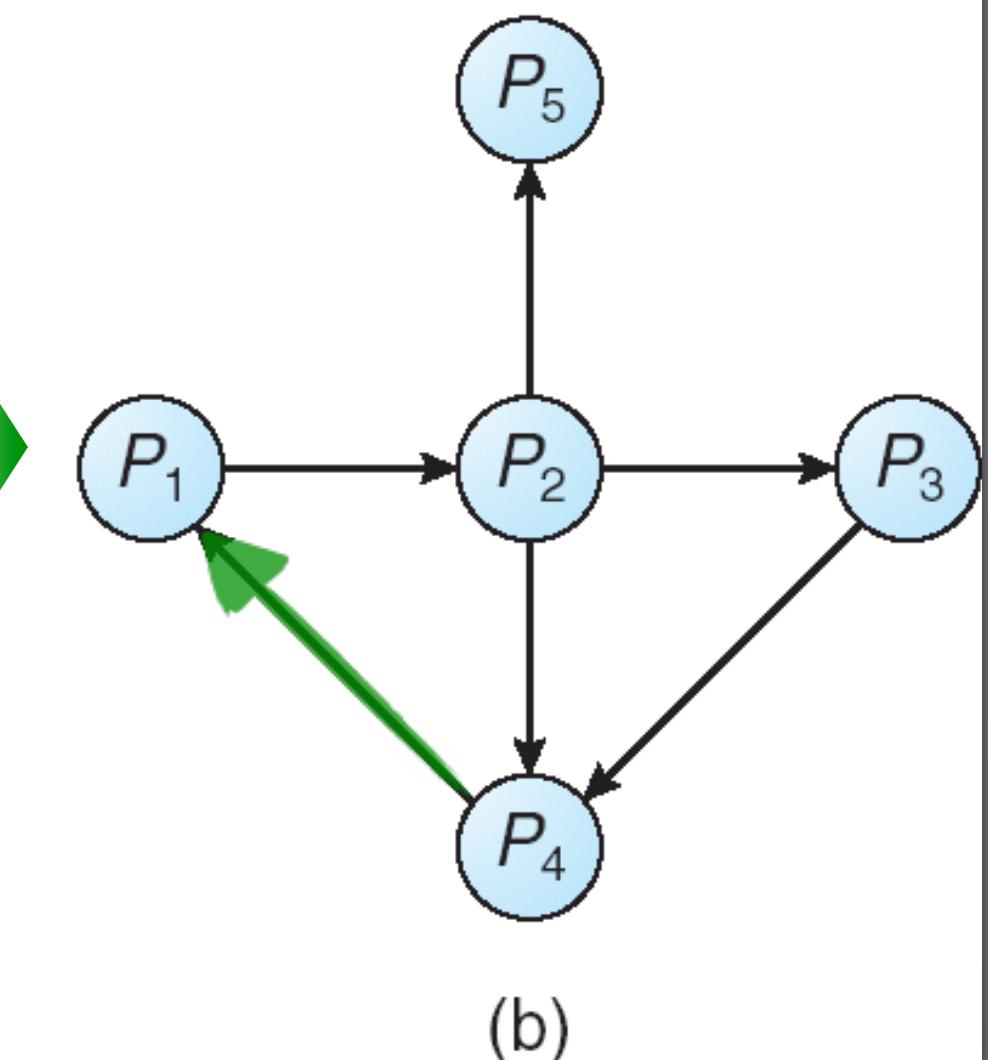
- Maintain a **wait-for graph**
  - Nodes are processes.
  - There is an edge  $p_i \rightarrow p_j$  if  $p_i$  is waiting for  $p_j$ 
    - i.e., if there was  $R_k$  such that
      - $p_i$  requested  $R_k$
      - $R_k$  was allocated to  $p_j$
- Periodically invoke an algorithm that searches for a cycle in the graph.
  - If there is a cycle, there exists a deadlock.
- An algorithm to detect a cycle in a graph requires an order of  $n^2$  operations, where  $n$  is the number of vertices in the graph.

Resource-Allocation Graph



This is not good either... but at least we control the frequency

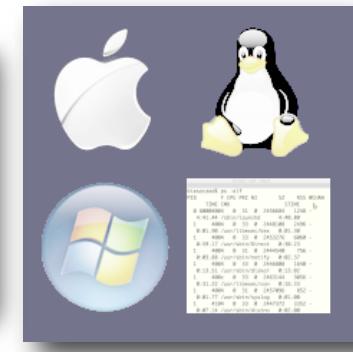
Corresponding  
wait-for graph





# Several Instances of a Resource Type

- We cannot use the wait-for graph to manage multiple instances of resources, so we need another algorithm (*again: because a cycle is not necessarily a deadlock*)
- Based on the banker's algorithm (see earlier slides)
- Like earlier, let
  - $n$  = number of processes  $P_i$  with  $0 \leq i < n$ , and
  - $m$  = number of resources types.
- Similarly to the deadlock avoidance algorithm, we use a number of data structures:
  - **Available**
    - A vector of length  $m$  indicates the number of available resources of each type.
  - **Allocation**
    - An  $n \times m$  matrix defines the number of resources of each type currently allocated to each process.
  - **Request**
    - An  $n \times m$  matrix indicates the current request of each process.  
If  $\text{Request}[i,j] = k$ , then process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$ .



# Detection Algorithm

## • Step 1

- Let **Work** and **Finish** be vectors of length  $m$  and  $n$ , respectively
- Initialize:
  - **Work** = **Available**
  - for  $i = 0, 1, 2, \dots, n$ , if **Allocation**[ $i$ ]  $\neq 0$ , then **Finish**[ $i$ ] = **false**; otherwise, **Finish**[ $i$ ] = **true**.

## • Step 2

- Find an index  $i$  such that both:
  - **Finish**[ $i$ ] == **false**
  - **Request**[ $i$ ]  $\leq$  **Work**
  - if no such  $i$  exists, go to step 4.

only processes with some resources allocated may be deadlocked

## • Step 3

- **Work** = **Work** + **Allocation**[ $i$ ]
- **Finish**[ $i$ ] = **true**
- go to step 2.

simulate process termination

## • Step 4

- If  $\exists i, 0 \leq i < n$ , **Finish**[ $i$ ] == **false**  $\Rightarrow$  then the system is in deadlock state.
- Moreover, if **Finish**[ $i$ ] == **false**, then  $P_i$  is deadlocked.
  - So, we not only know that there is a deadlock, but also which processes are deadlocked
- Algorithm requires an order of  $O(m \times n^2)$  operations to detect whether the system is in deadlocked state.

As earlier, this is not good... but... at least we control the frequency here



# Deadlock Detection Example

- Five processes  $p_0$  through  $p_4$ 
  - three resource types
    - A (7 instances)
    - B (2 instances)
    - C (6 instances).
- Sequence  $\langle p_0, p_2, p_3, p_1, p_4 \rangle$  will result in  $\text{Finish}[i] = \text{true}$  for all  $i$ .

	Allocation			Request			Work		
	A	B	C	A	B	C	A	B	C
$p_0$	0	1	0	0	0	0	0	0	0
$p_1$	2	0	0	2	0	2	0	0	0
$p_2$	3	0	3	0	0	0	0	0	0
$p_3$	2	1	1	1	0	0	0	0	0
$p_4$	0	0	2	0	0	2	0	0	0

O.K.

- However, if - for example -  $p_2$  is allocated an additional instance of type C then there is a problem:
  - we can reclaim resources held by process  $p_0$ , but the processes  $p_1, p_2, p_3$ , and  $p_4$  are stuck with  $\text{Finish}[i] = \text{false}$ , so there is a deadlock, and all of these processes are deadlocked.

	Allocation			Request			Work		
	A	B	C	A	B	C	A	B	C
$p_0$	0	1	0	0	0	0	0	0	0
$p_1$	2	0	0	2	0	2	0	0	0
$p_2$	3	0	3	0	0	1	0	0	0
$p_3$	2	1	1	1	0	0	0	0	0
$p_4$	0	0	2	0	0	2	0	0	0

Houston!

# Detection-Algorithm Usage

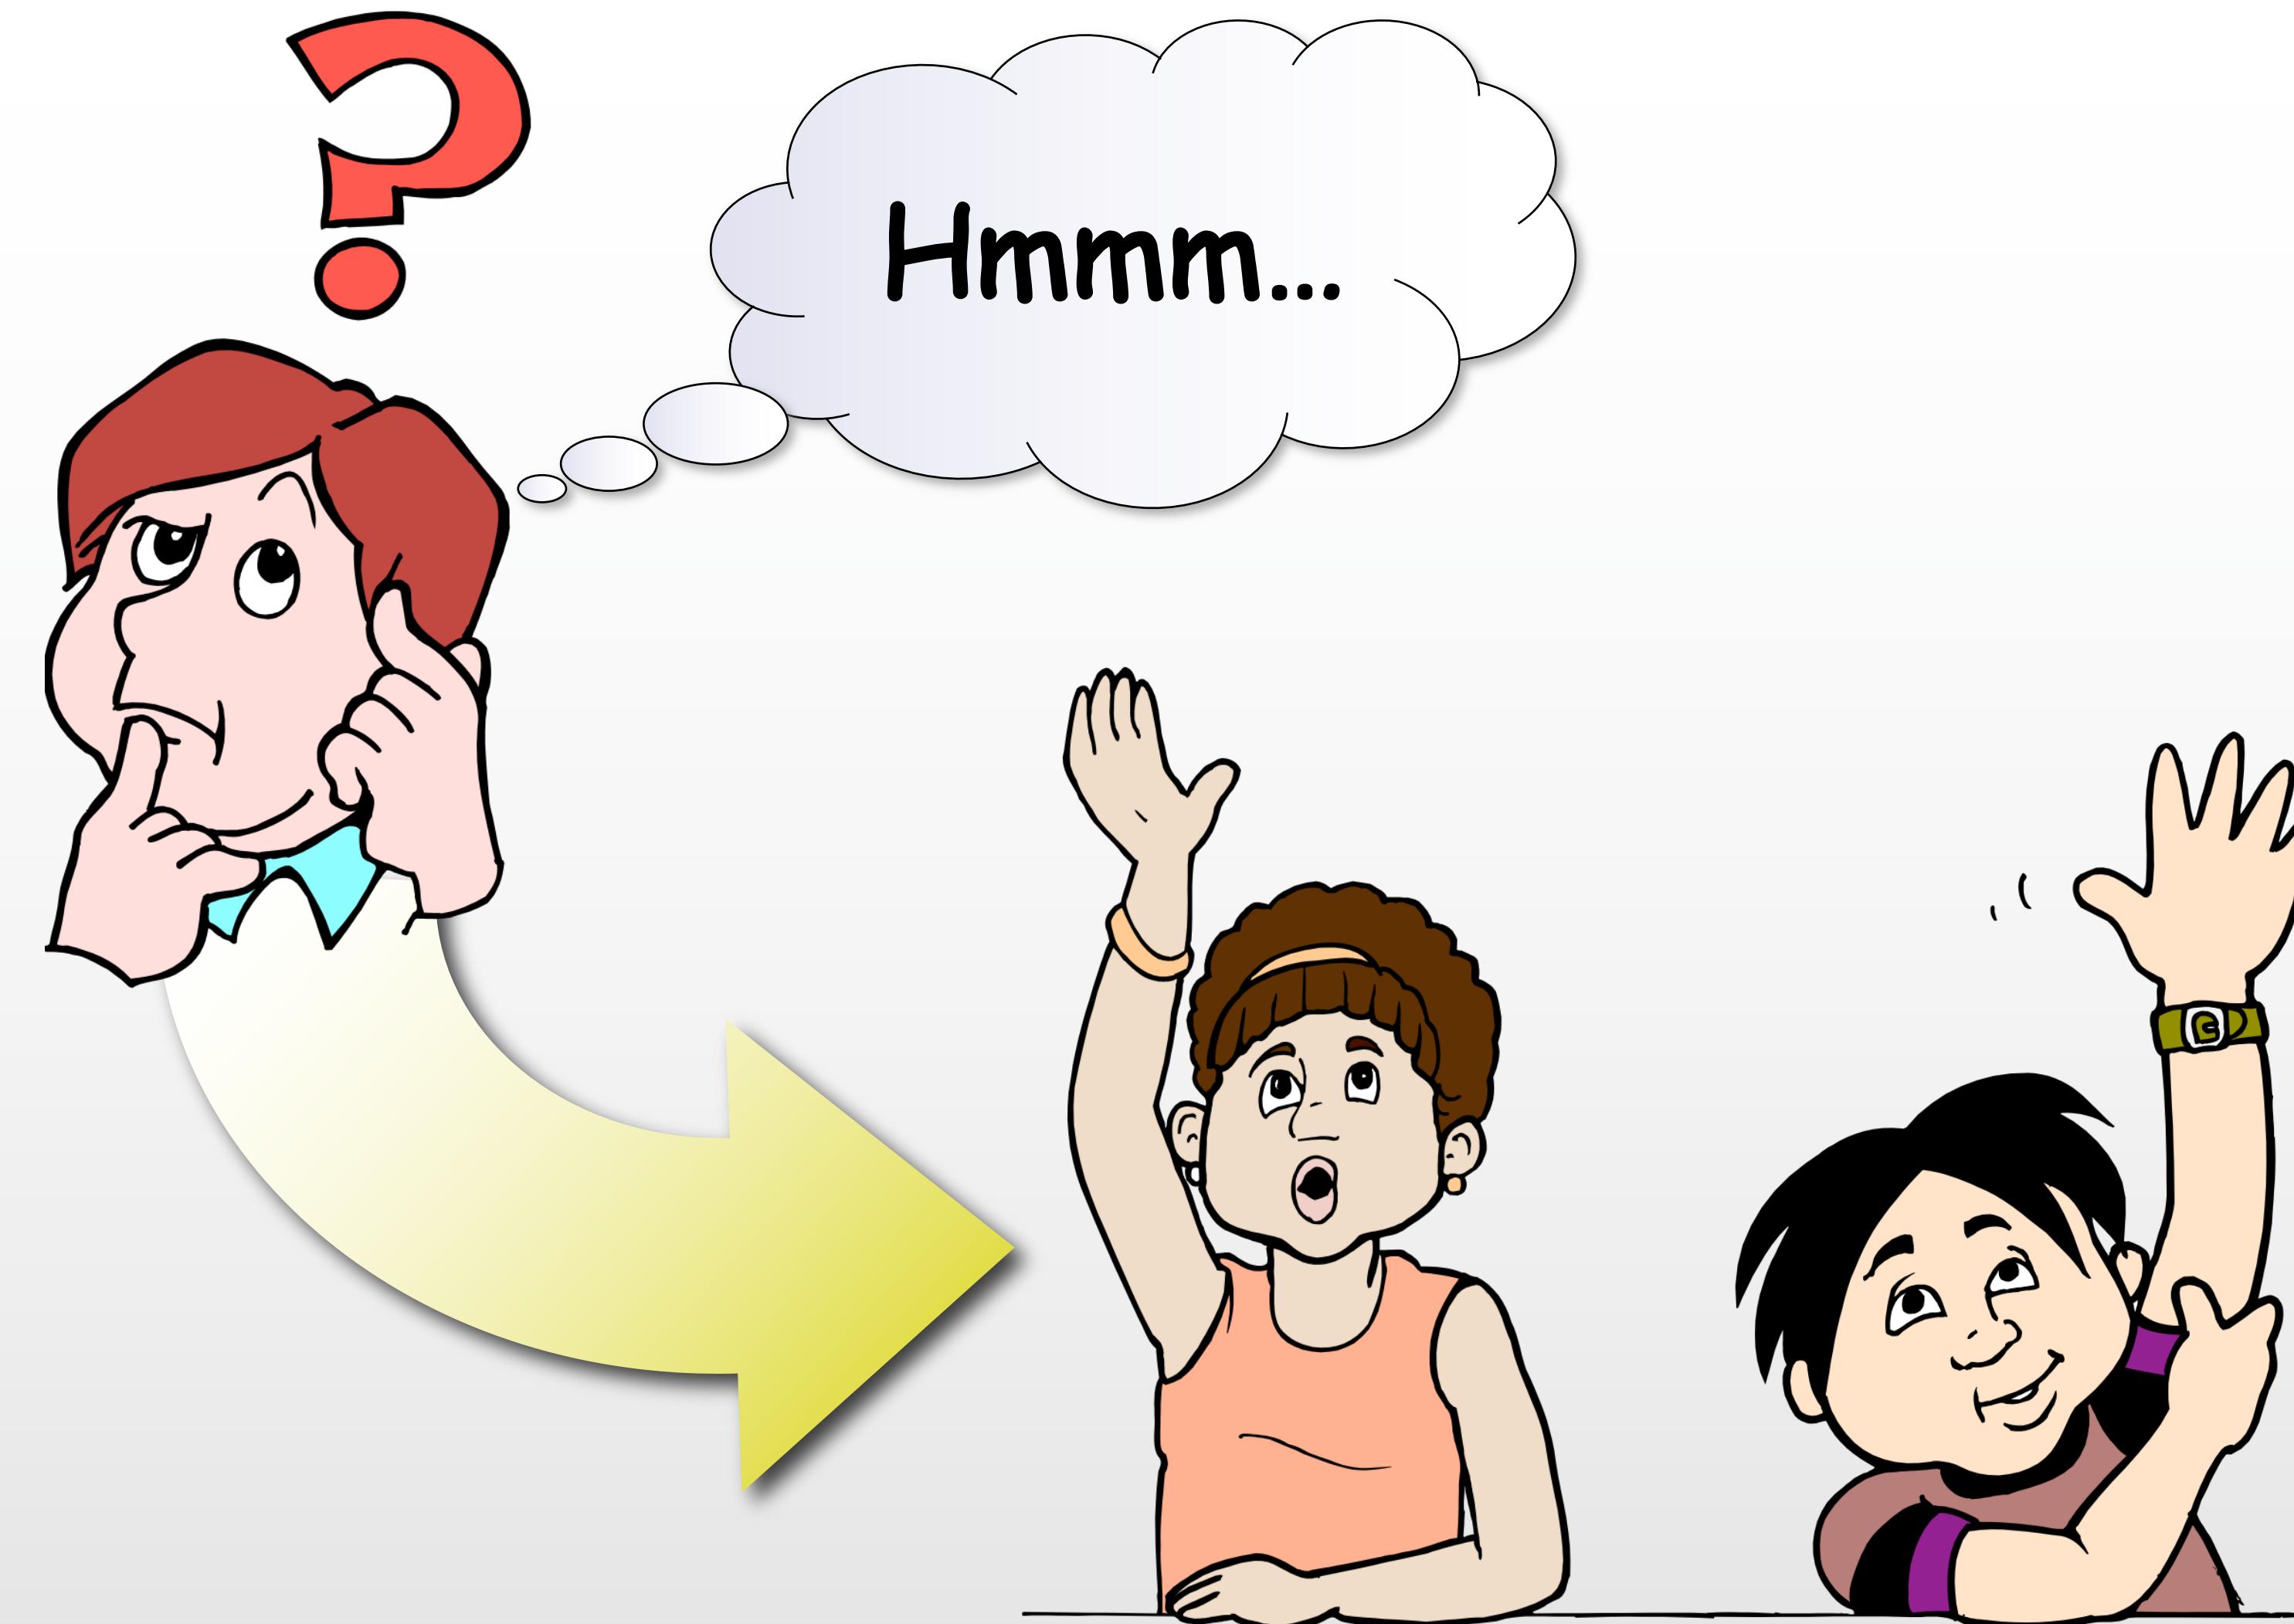


- Many "engineering" questions to answer
- When, and how often to invoke?
  - depends on:
    - how often a deadlock is likely to occur?
    - how many processes will need to be rolled back/terminated?
      - e.g., one for each disjoint cycle
- There may be many cycles in the wait-for graph, which one is “more important” to break?
- Which of the many deadlocked processes “causes” the deadlock?

# Recovery from Deadlock: Process Termination



- Recovery from a deadlock can be done by terminating all or some processes, or by reclaiming resources
  - leaving the process alive and just reclaiming is very difficult
- How to select a victim?
  - minimize cost - but what exactly is the cost?
  - Abort all deadlocked processes.
    - simple but dirty
  - Abort one process at a time until the deadlock cycle is eliminated.
    - Question: In which order should we choose to abort?
      - Priority of the process.
      - How long process has computed, and how much longer to completion.
      - Resources the process has used.
      - Resources process needs to complete.
      - How many processes will need to be terminated.
      - Is process interactive or batch?
- Problem: Starvation
  - same process may always be picked as victim
- May consider **rollbacks**
  - i.e., return a deadlocked process to some safe state, and restart the process from that state
  - extremely expensive, since the system must save safe states



**COMP362 Operating Systems**  
**Prof. AJ Biesczad**